



---

## Feature Engineering

Evan Jones

# Feature Engineering

---

Scale to large datasets

Find good features

Preprocess with Cloud MLE

# Objectives

---

**Turn raw data to features**

What raw data do we need to collect to predict the price of a house?





# Lot Size

## Number of Rooms



# Historical sale price





# Location, location, location




Can we use this raw data  
we've cleaned and collected?

```
0 : {  
  house_info : {  
    num_rooms: 6  
    num_bedrooms: 3  
    street_name: "Main Street"  
    num_basement_rooms: -1  
    ...  
  }  
}
```



Raw data must be mapped  
into numerical feature vectors

```
0 : {  
  house_info : {  
    num_rooms: 6  
    num_bedrooms: 3  
    street_name: "Main Street"  
    num_basement_rooms: -1  
    ...  
  }  
}
```



```
[  
  6.0,  
  1.0,  
  0.0,  
  0.0,  
  0.0,  
  9.321,  
  -2.20,  
  1.01,  
  0.0,  
  ...,  
]
```

# What makes a feature “good”?

```
0 : {  
  house_info : {  
    ✓ num_rooms: 6  
    ✓ num_bedrooms: 3  
    ✓ street_name: "Main Street"  
    ✓ num_basement_rooms: -1  
    ...  
  }  
}
```

# Objectives

---

Turn raw data to features

**Compare good vs bad features**



What makes a good feature?

# What makes a good feature?

- 1 Be related to the objective

# What makes a good feature?

1 Be related to the objective

2 Be known at  
prediction-time



# What makes a good feature?

1 Be related to the objective

2 Be known at prediction-time

3 Be numeric with meaningful magnitude

4 Have enough examples

5 Bring human insight to problem



---

Be related to the objective





# Choose the good features



A) Breed

B) Age

C) Eye Color

# Objective: Good racehorse



**A) Breed**

**B) Age**

**C) Eye Color**



# Objective: Eye disease



**A) Breed**

**B) Age**

**C) Eye Color**



Different problems in the  
same domain may need  
different features

## Quiz

Are these features related to the objective or not?

Predict total number of customers who will use a  
certain discount coupon



# Predict total number of customers who will use a certain discount coupon

1

Font of the text with which the discount is advertised on partner websites



# Predict total number of customers who will use a certain discount coupon

1

Font of the text with which the discount is advertised on partner websites

2

Price of the item the coupon applies to





# Predict total number of customers who will use a certain discount coupon

1

Font of the text with which the discount is advertised on partner websites

2

Price of the item the coupon applies to

3

Number of items in stock



Predict whether a credit card transaction is  
fraudulent



# Predict whether a credit card transaction is fraudulent

1

Whether cardholder has purchased these items at this store before



# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Credit card chip reader speed



# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Credit card chip reader speed
- 3 Category of item being purchased





# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Credit card chip reader speed
- 3 Category of item being purchased
- 4 Expiry date of credit card





---

Be known at prediction-time





Some data could be known immediately, and some other data is not known in real time.





**SOLD: \$300,000**

What's wrong with the second feature?

✔ `city_id: "br/sao_paulo"`

✘ `inferred_city_cluster_id: 219`

Feature definitions should  
not change over time

✔ `city_id: "br/sao_paulo"`

✘ `inferred_city_cluster_id: 219`

## Quiz

Is the value knowable at prediction time or not?



# Predict total number of customers who will use a certain discount coupon

1 Total number of discountable items sold



# Predict total number of customers who will use a certain discount coupon

- 1 Total number of discountable items sold
- 2 Number of discountable items sold the previous month



# Predict total number of customers who will use a certain discount coupon

- 1 Total number of discountable items sold
- 2 Number of discountable items sold the previous month
- 3 Number of customers who viewed ads about item



# Predict whether a credit card transaction is fraudulent

1

Whether cardholder has purchased these items at this store before



You cannot train with  
current data and predict  
with stale data





```
SELECT name, COUNT(trans_id) AS count
FROM sales_warehouse
WHERE
    # filter out last three days
    trans_time <
    TIMESTAMP_SUB(
        CURRENT_TIMESTAMP(), INTERVAL 3 DAY
    )
```

# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Whether item is new at store (and can not have been purchased before)



# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Whether item is new at store (and can not have been purchased before)
- 3 Category of item being purchased



# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Whether item is new at store (and can not have been purchased before)
- 3 Category of item being purchased
- 4 Online or in-person purchase?

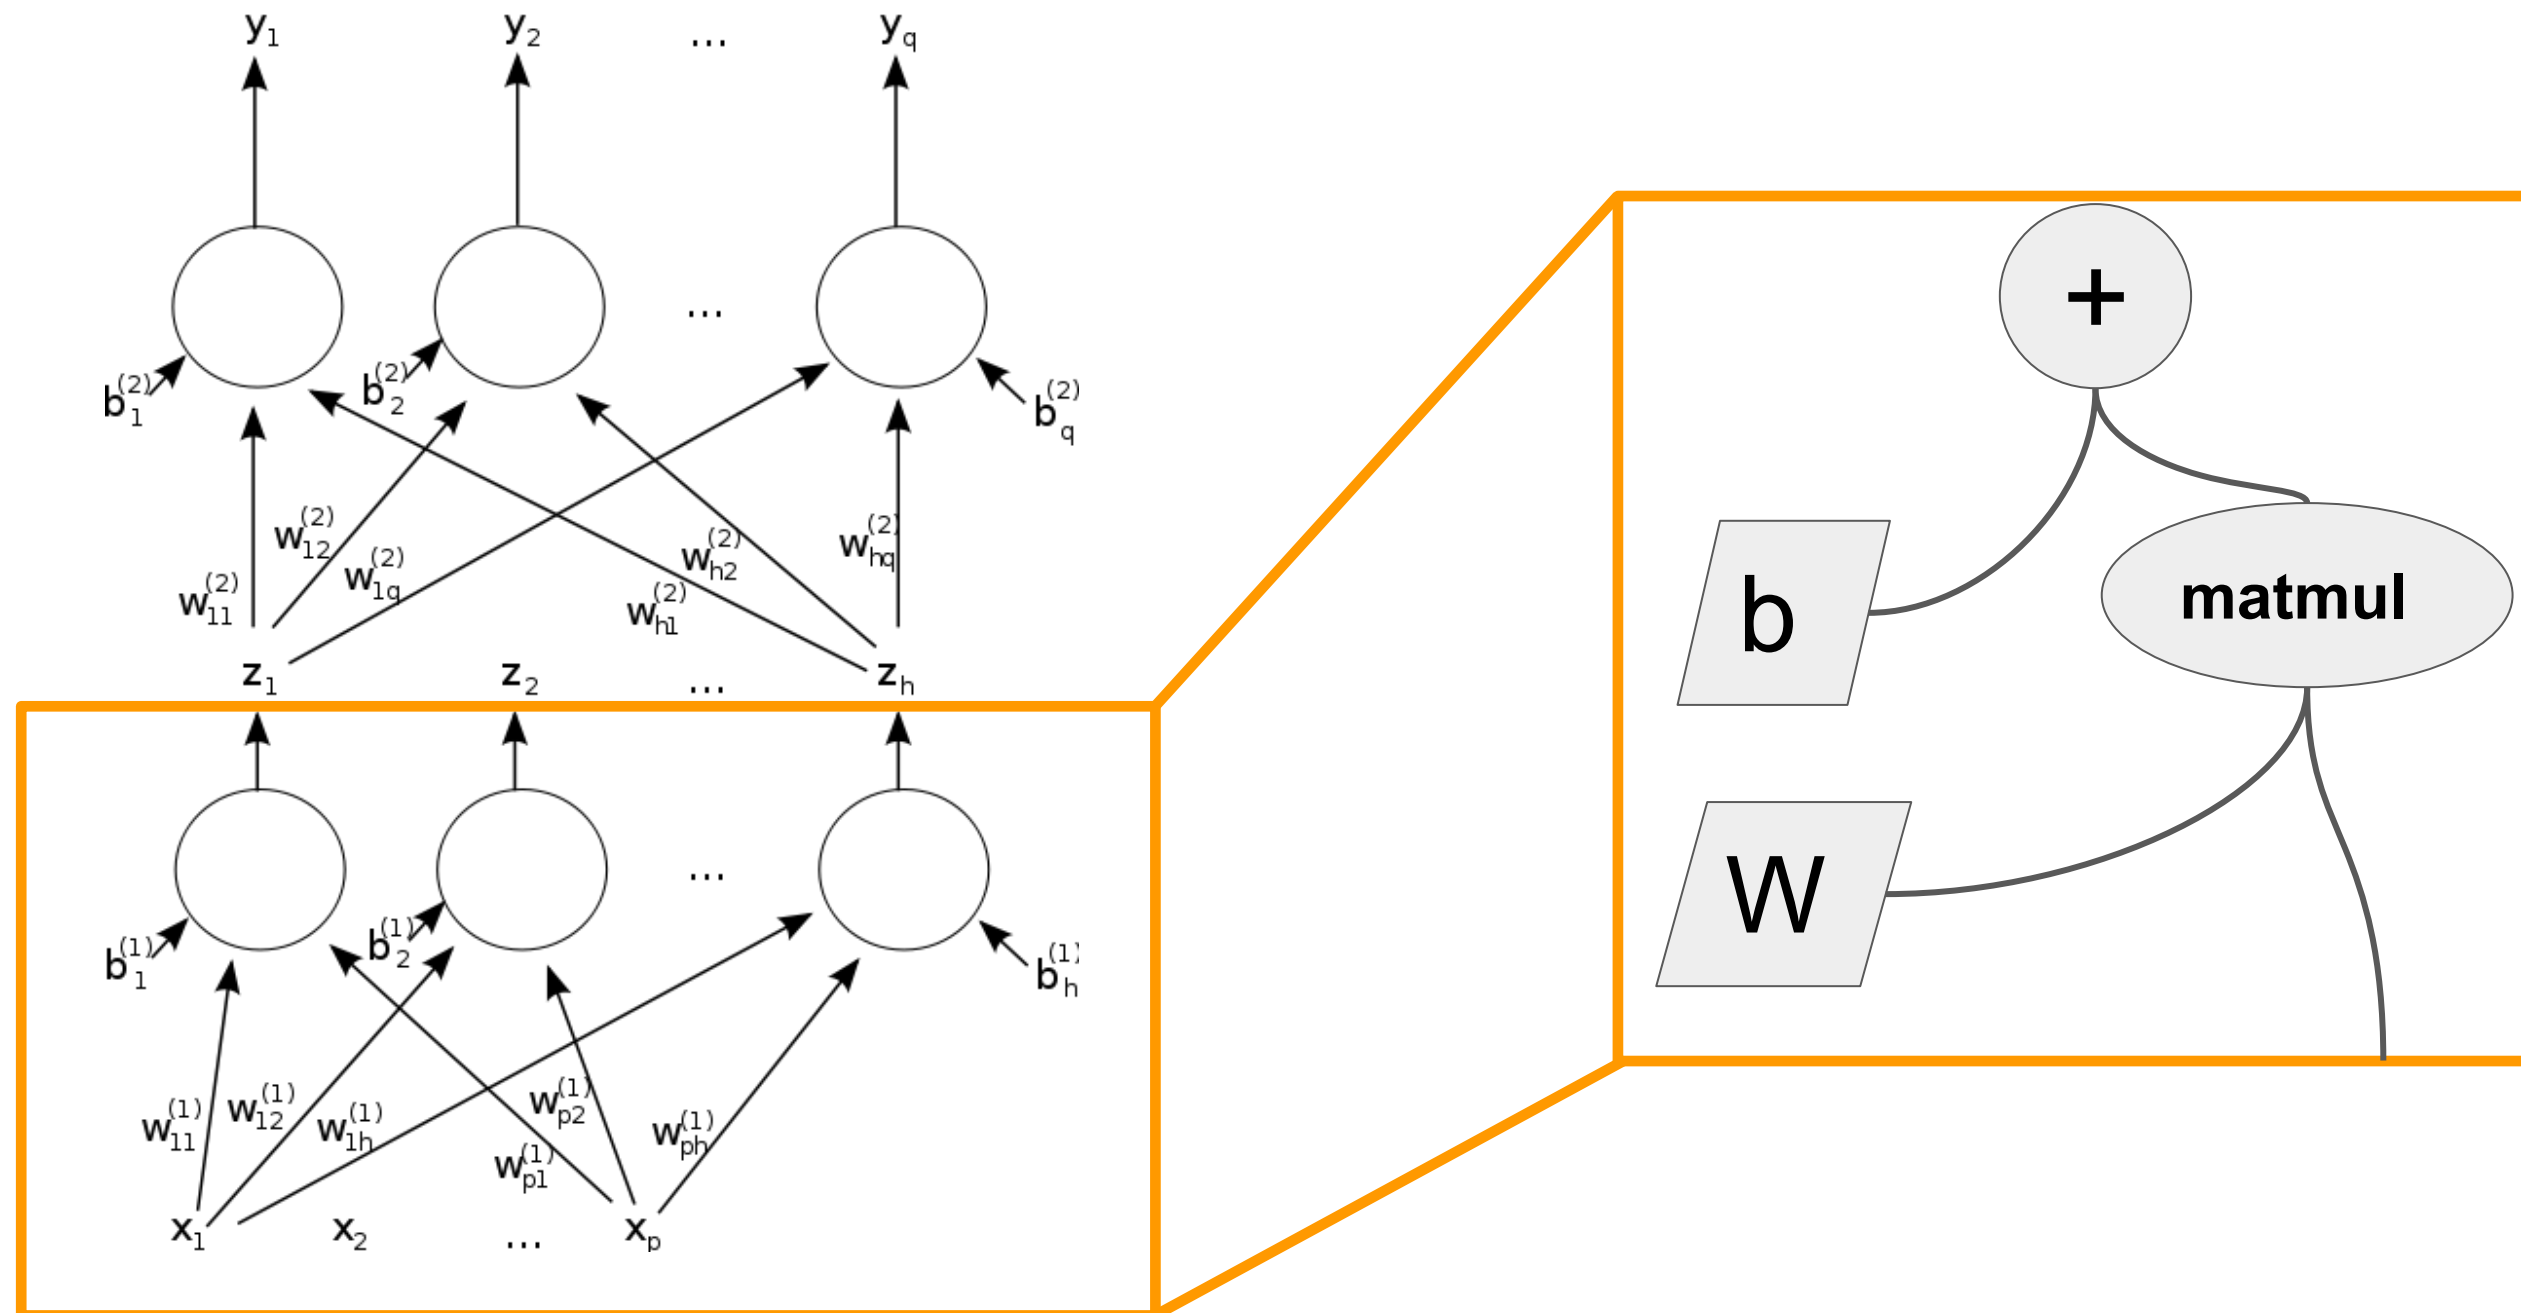




---

Be numeric with meaningful  
magnitude

# Neural networks are weighing and adding machines





## Quiz

Which of these are  
numeric?

# Predict total number of customers who will use a certain discount coupon

- 1 Percent value of the discount (e.g. 10% off, 20% off, etc.)



# Discount coupon usage

- 1 Percent value of the discount  
(e.g. 10% off, 20% off, etc.)
- 2 Size of the coupon  
(e.g. 4 cm<sup>2</sup>, 24 cm<sup>2</sup>, 48 cm<sup>2</sup>, etc.)

# Discount coupon usage

- 1 Percent value of the discount  
(e.g. 10% off, 20% off, etc.)
- 2 Size of the coupon  
(e.g. 4 cm<sup>2</sup>, 24 cm<sup>2</sup>, 48 cm<sup>2</sup>, etc.)
- 3 Font an advertisement is in  
(Arial, Times New Roman, etc.)

# Discount coupon usage


4 Color of coupon (red, black, blue, etc.)

# Discount coupon usage

- 4 Color of coupon (red, black, blue, etc.)
- 5 Item category (1 for dairy, 2 for deli, 3 for canned goods, etc.)



# Word2vec



Word  
Vectors

# Word2vec





---

Have enough examples

Evan Jones



Avoid having values of  
which you don't have  
enough examples







## Quiz

Which of these will it be difficult to get enough examples?

# Discount coupon usage

- 1 Percent discount of coupon  
(20%, 30%, etc.)

# Discount coupon usage

1 Percent discount of coupon  
(20%, 30%, etc.)

2 Date that promotional offer  
starts

# Discount coupon usage

1 Percent discount of coupon  
(20%, 30%, etc.)

2 Date that promotional offer  
starts

3 Number of customers who  
opened advertising email

# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before



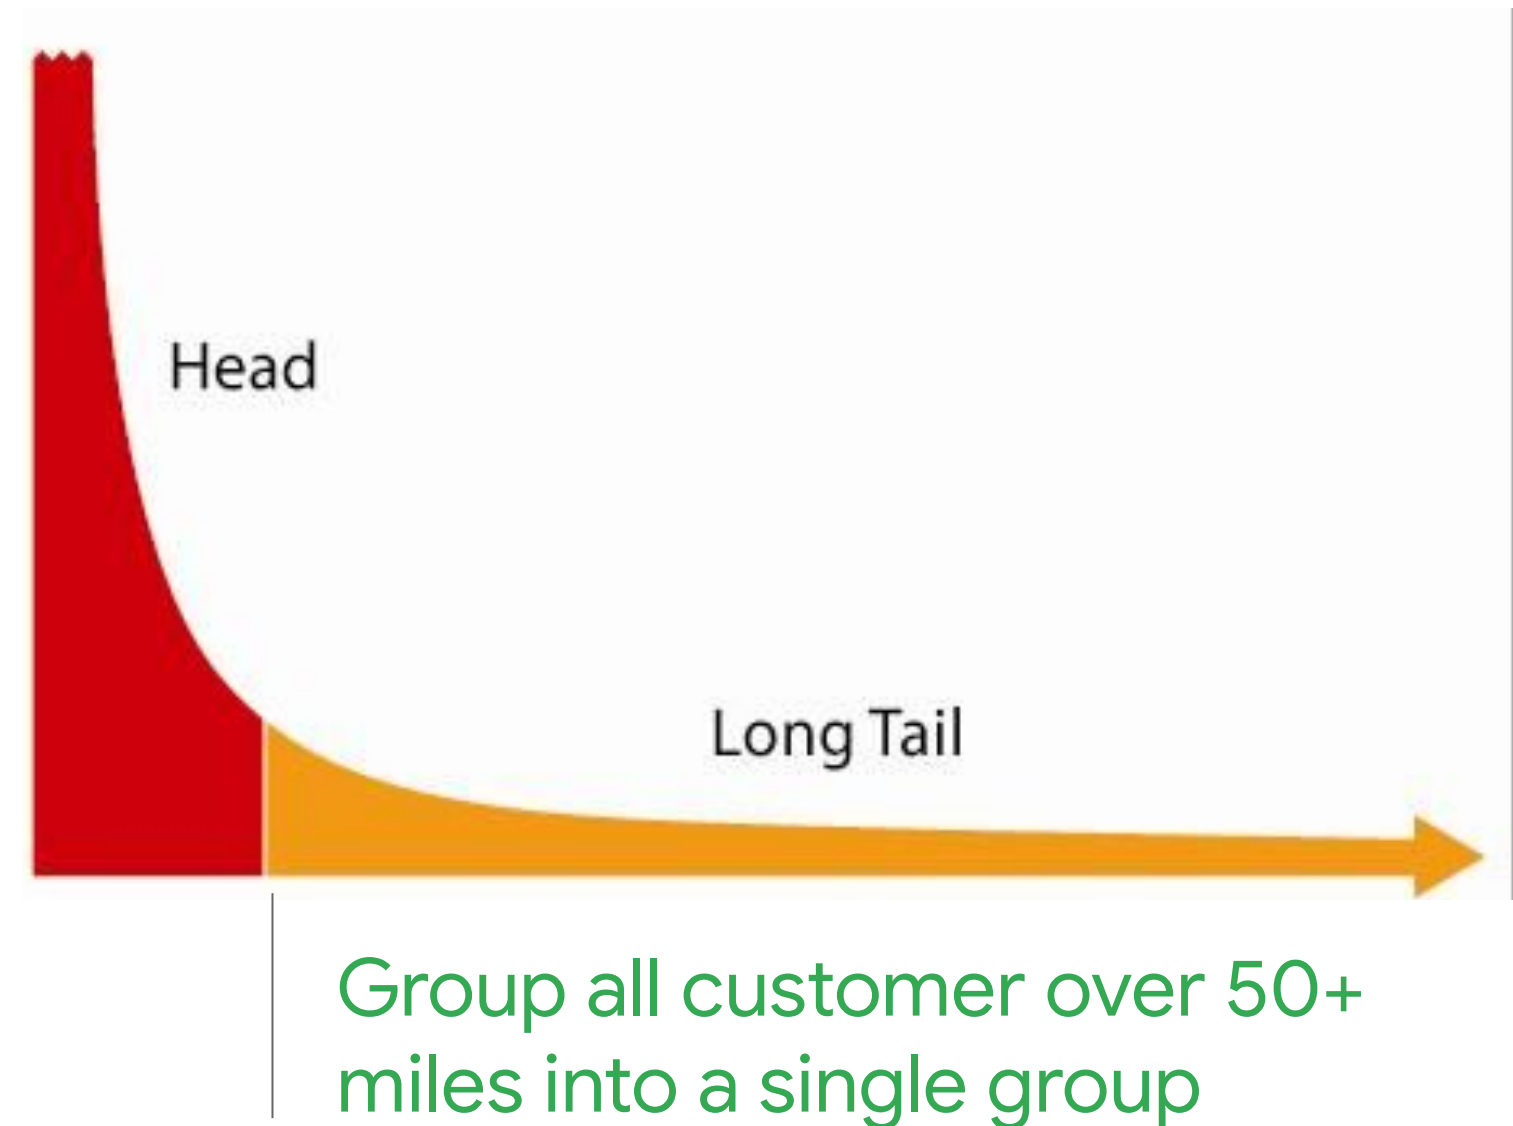
# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Distance between cardholder address and store



# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Distance between cardholder address and store




# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Distance between cardholder address and store

# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Distance between cardholder address and store



# Predict whether a credit card transaction is fraudulent

- 1 Whether cardholder has purchased these items at this store before
- 2 Distance between cardholder address and store
- 3 Category of item being purchased
- 4 Online or in-person purchase?



---

Bring human insight to problem

Evan Jones

# Objectives

---

Turn raw data to features

Compare good vs bad features

**Represent features**

# Raw data are converted to numeric features in different ways

```
{  
  "transactionId": 42,  
  "name": "Ice Cream",  
  "price": 2.50,  
  "tags": ["cold", "dessert"],  
  "servedBy": {  
    "employeeId": 72365,  
    "waitTime": 1.4,  
    "customerRating": 4  
  },  
  "storeLocation": {  
    "latitude": 35.3,  
    "longitude": -98.7  
  }  
},
```



# Raw data are converted to numeric features in different ways

```
{  
  "transactionId": 42,  
  "name": "Ice Cream",  
  "price": 2.50,  
  "tags": ["cold", "dessert"],  
  "servedBy": {  
    "employeeId": 72365,  
    "waitTime": 1.4,  
    "customerRating": 4  
  },  
  "storeLocation": {  
    "latitude": 35.3,  
    "longitude": -98.7  
  }  
},
```




```
[..., 1, 2.50, ..., ]  
[..., 0, 8.99, ..., ]  
[..., 0, 3.45, ..., ]  
...
```

In estimator API, This  
is a feature column

# Numeric values can be used as-is

```
{  
  "transactionId": 42,  
  "name": "Ice Cream",  
  "price": 2.50,  
  "tags": ["cold", "dessert"],  
  "servedBy": {  
    "employeeId": 72365,  
    "waitTime": 1.4,  
    "customerRating": 4  
  },  
  "storeLocation": {  
    "latitude": 35.3,  
    "longitude": -98.7  
  }  
},
```



```
[ , 2.50, ..., 1.4, ]  
...
```

```
INPUT_COLUMNS = [  
    ...,  
    tf.feature_column.numeric_col  
umn('price'),  
    ...  
]
```

numeric\_column is a  
type of feature column

# Overly specific attributes should be discarded

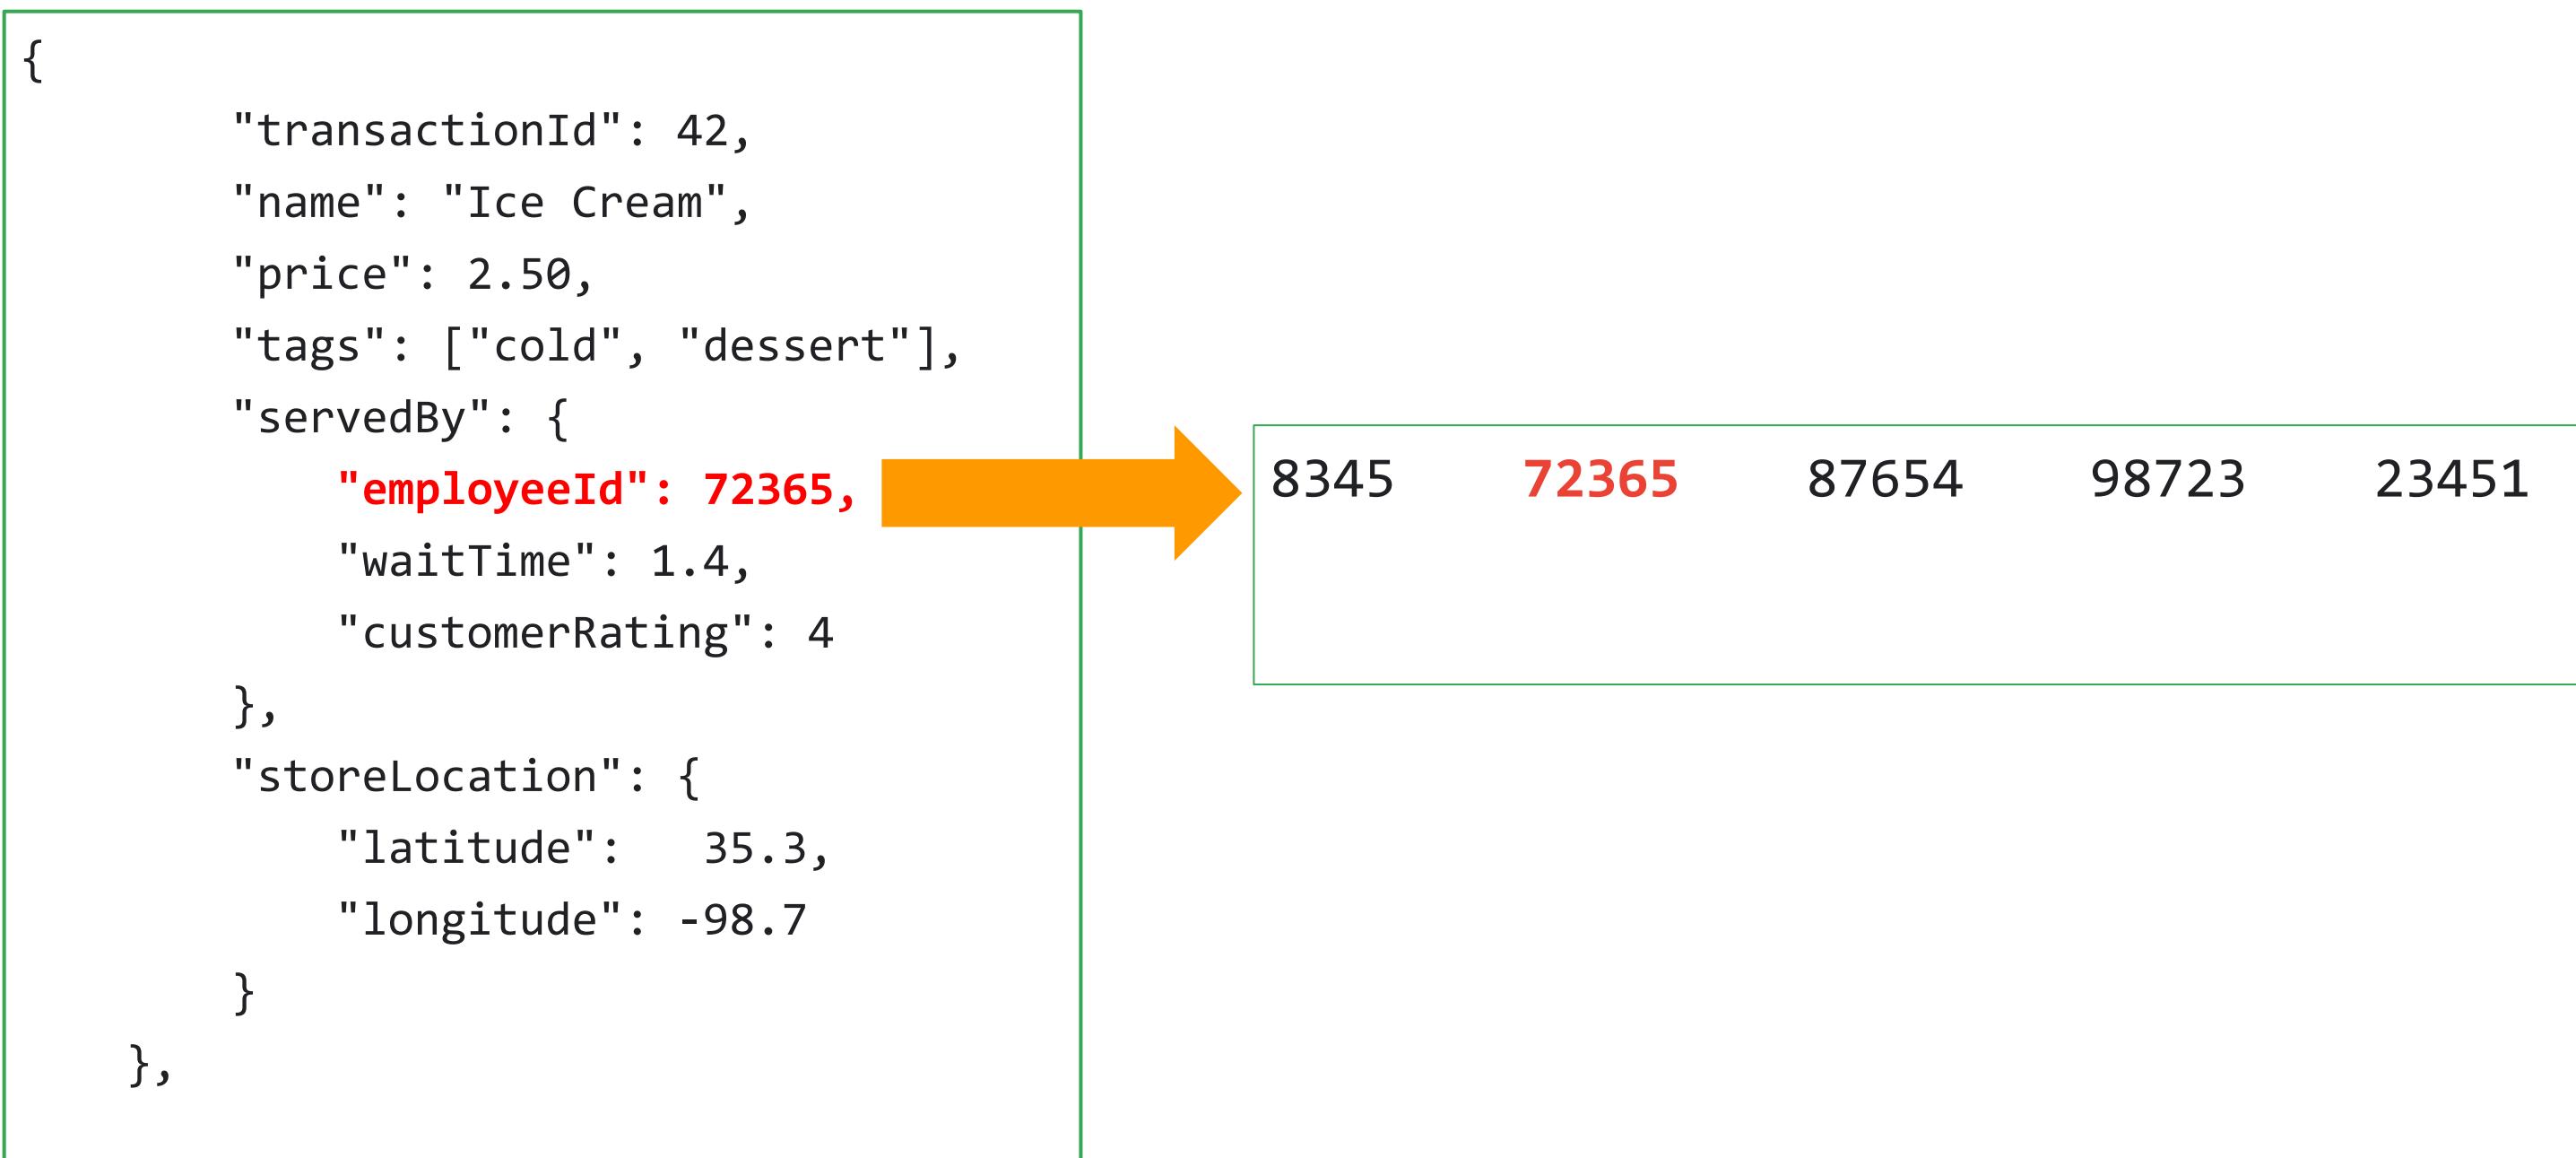
```
{  
  "transactionId": 42,  
  "name": "Ice Cream",  
  "price": 2.50,  
  "tags": ["cold", "dessert"],  
  "servedBy": {  
    "employeeId": 72365,  
    "waitTime": 1.4,  
    "customerRating": 4  
  },  
  "storeLocation": {  
    "latitude": 35.3,  
    "longitude": -98.7  
  }  
},
```



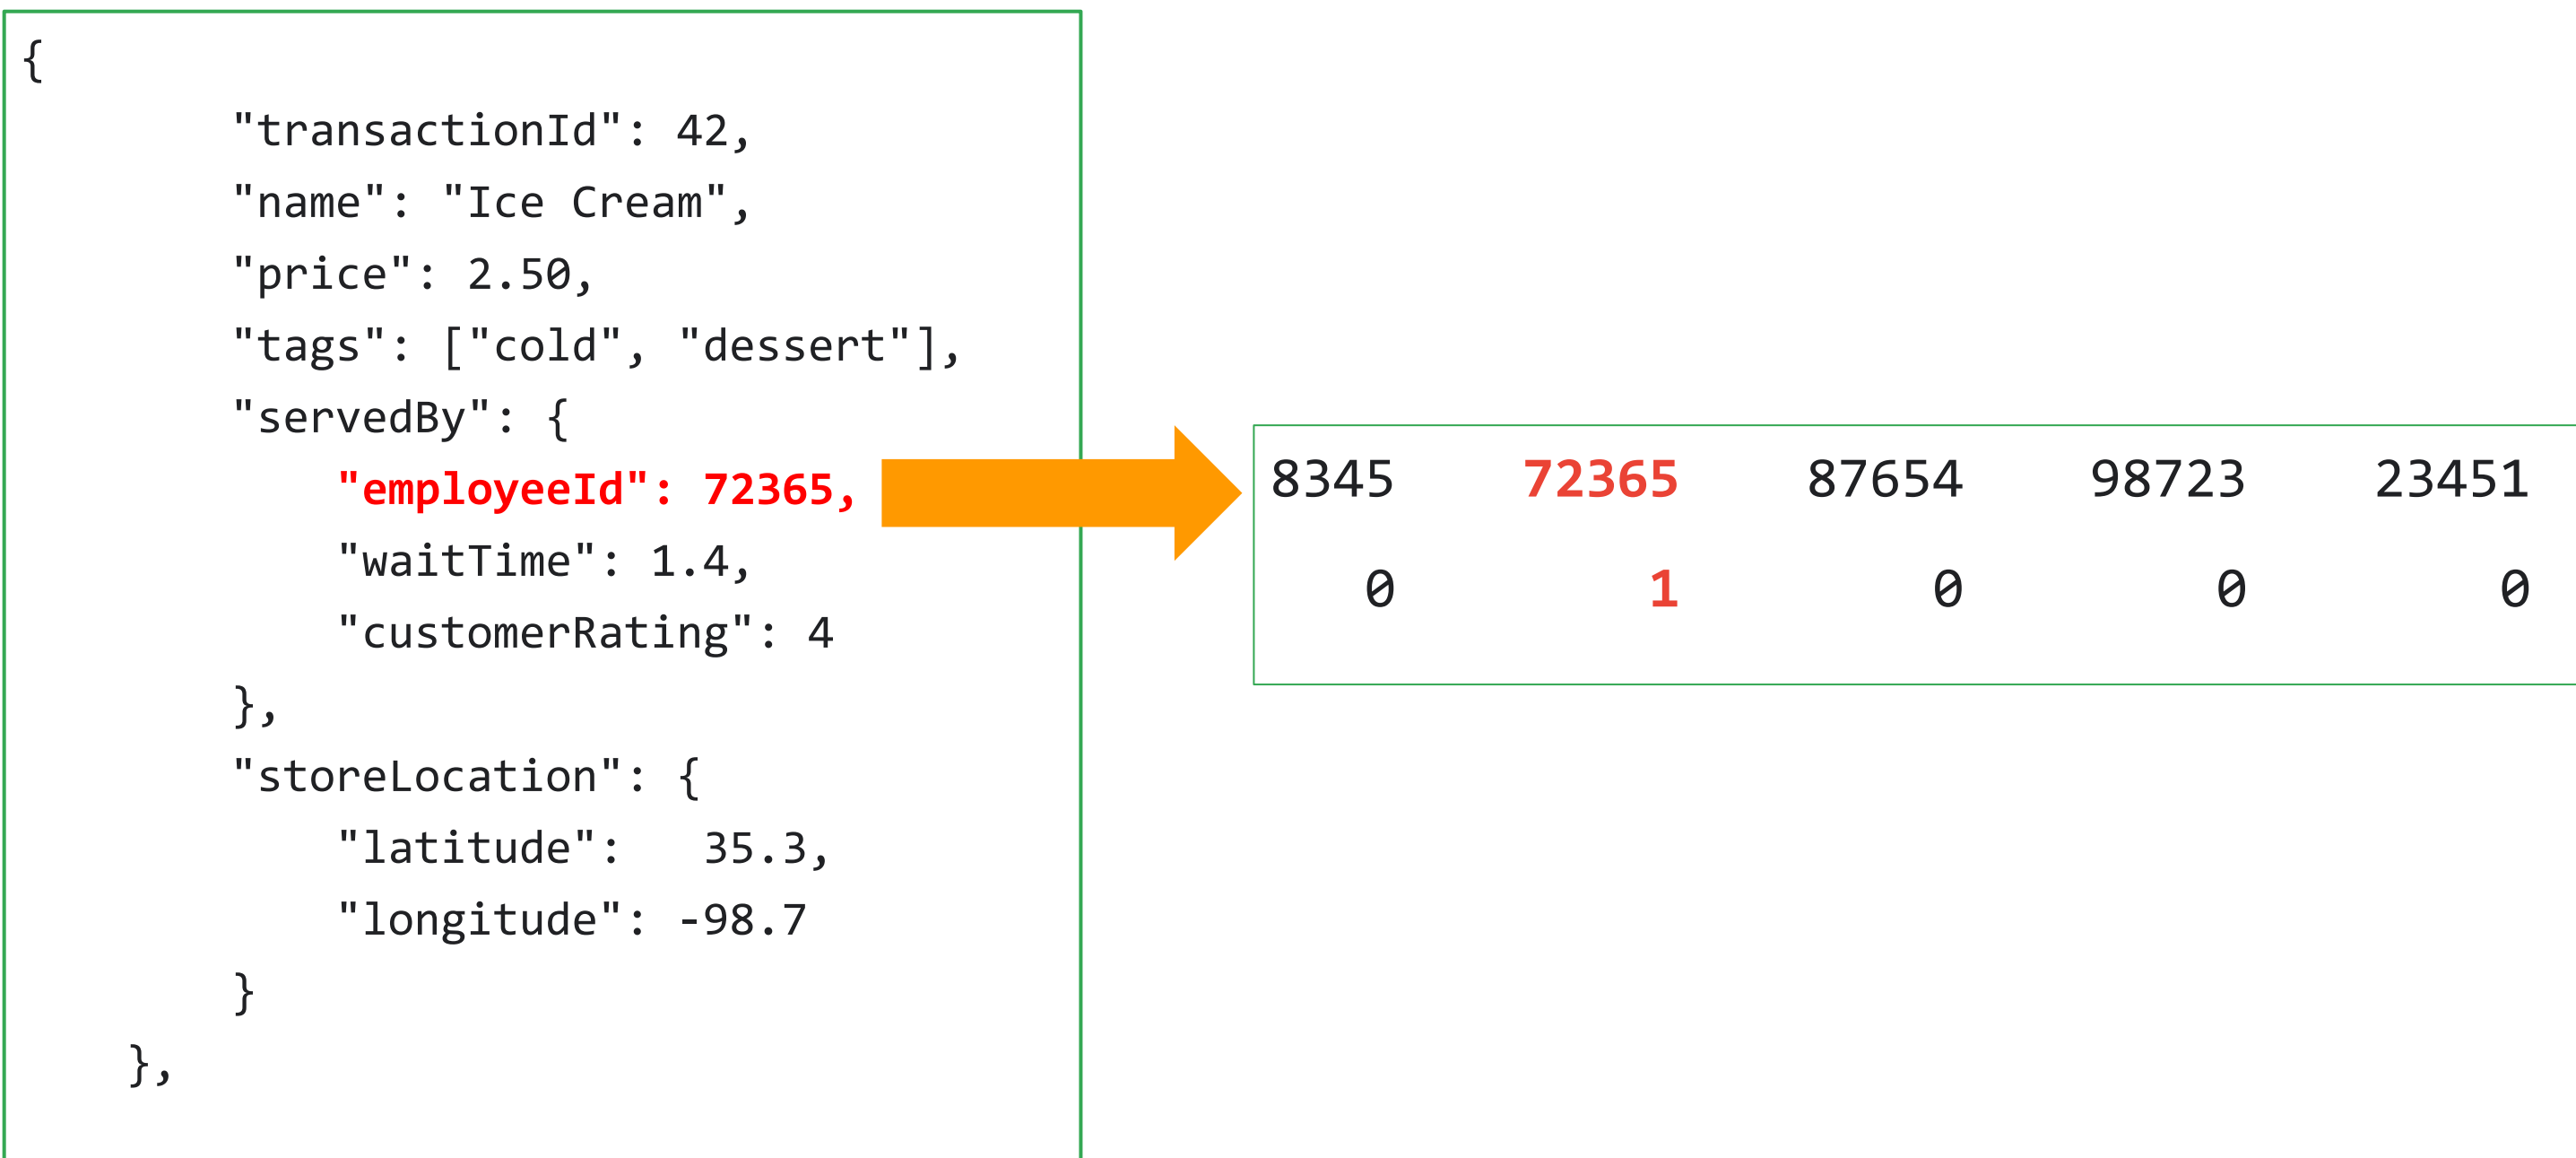
# Overly specific attributes should be discarded

```
{  
  "transactionId": 42,  
  "name": "Ice Cream",  
  "price": 2.50,  
  "tags": ["cold", "dessert"],  
  "servedBy": {  
    "employeeId": 72365,  
    "waitTime": 1.4,  
    "customerRating": 4  
  },  
  "storeLocation": {  
    "latitude": 35.3,  
    "longitude": -98.7  
  }  
},
```

# Categorical variables should be one-hot encoded



# Categorical variables should be one-hot encoded



# Categorical variables should be one-hot encoded

```
{  
  "transactionId": 42,  
  "name": "Ice Cream",  
  "price": 2.50,  
  "tags": ["cold", "dessert"],  
  "servedBy": {  
    "employeeId": 72365,  
    "waitTime": 1.4,  
    "customerRating": 4  
  },  
  "storeLocation": {  
    "latitude": 35.3,  
    "longitude": -98.7  
  }  
},
```



8345	72365	87654	98723	23451
0	1	0	0	0



# Categorical variables should be one-hot encoded

```
{  
  "transactionId": 42,  
  "name": "Ice Cream",  
  "price": 2.50,  
  "tags": ["cold", "dessert"],  
  "servedBy": {  
    "employeeId": 72365,  
    "waitTime": 1.4,  
    "customerRating": 4  
  },  
  "storeLocation": {  
    "latitude": 35.3,  
    "longitude": -98.7  
  }  
},
```

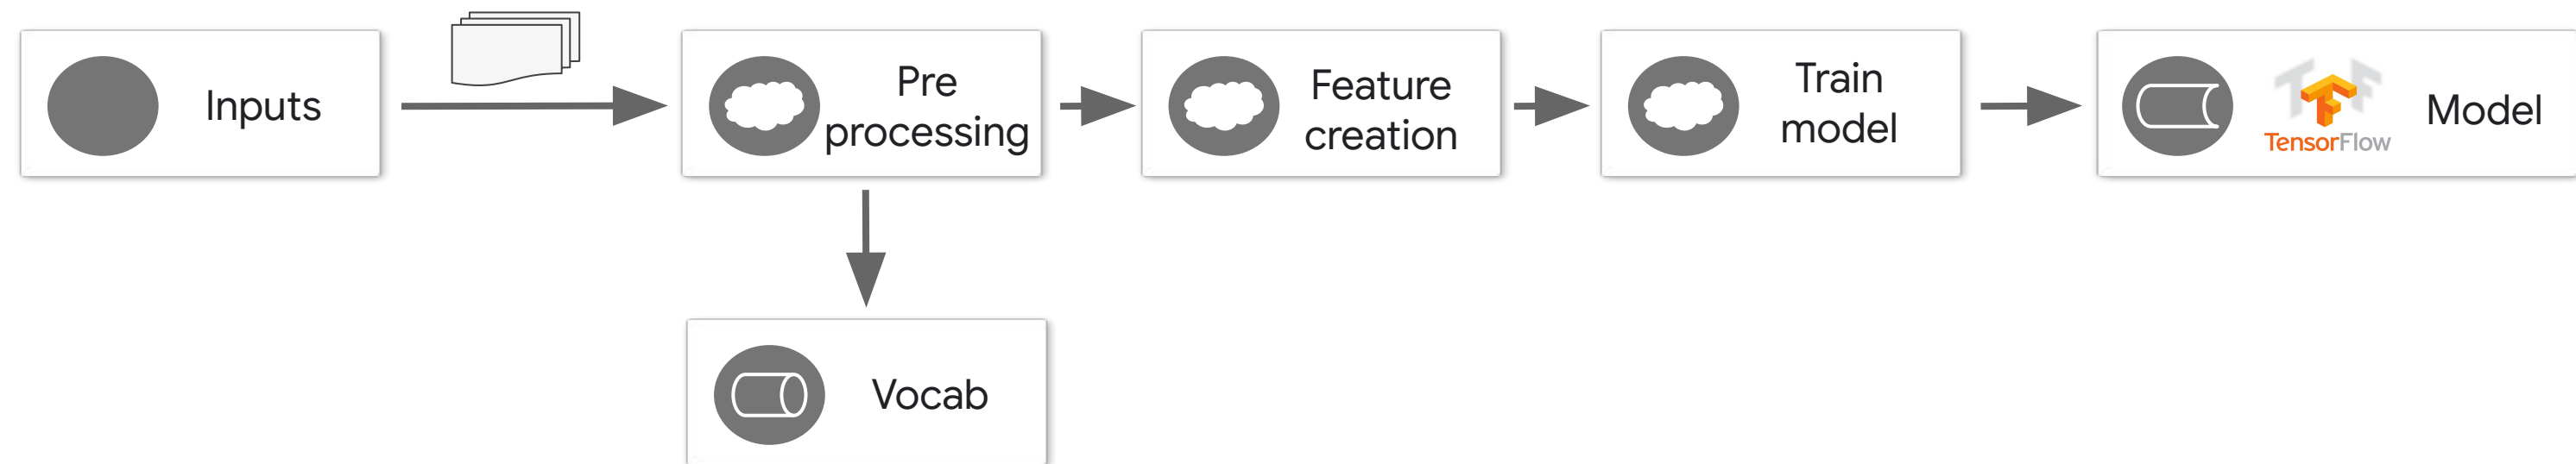


8345	72365	87654	98723	23451
0	1	0	0	0

```
tf.feature_column.categorical_column_with  
_vocabulary_list('employeeId',  
  Vocabulary_list = ['8345',  
    '72365', '87654', '98723', '23451']),
```

Don't know the list of  
keys? Create a  
vocabulary

# Preprocess data to create a vocabulary of keys



Vocabulary of keys:

8345	72365	87654	98723	23451
0	1	0	0	0

The vocabulary and the  
mapping of the vocabulary  
needs to be identical at  
prediction time

8345	72365	87654	98723	??????
0	0	0	0	0

# Options for encoding categorical data

If you know the keys beforehand:

```
tf.feature_column.categorical_column_with_vocabulary_list('employeeId',  
    vocabulary_list = ['8345', '72345', '87654', '98723', '23451']),
```



# Options for encoding categorical data

If you know the keys beforehand:

```
tf.feature_column.categorical_column_with_vocabulary_list('employeeId',  
    vocabulary_list = ['8345', '72345', '87654', '98723', '23451']),
```

If your data is already indexed; i.e., has integers in [0-N):

```
tf.feature_column.categorical_column_with_identity('employeeId',  
    num_buckets = 5)
```

# Options for encoding categorical data

If you know the keys beforehand:

```
tf.feature_column.categorical_column_with_vocabulary_list('employeeId',  
    vocabulary_list = ['8345', '72345', '87654', '98723', '23451']),
```

If your data is already indexed; i.e., has integers in [0-N):

```
tf.feature_column.categorical_column_with_identity('employeeId',  
    num_buckets = 5)
```

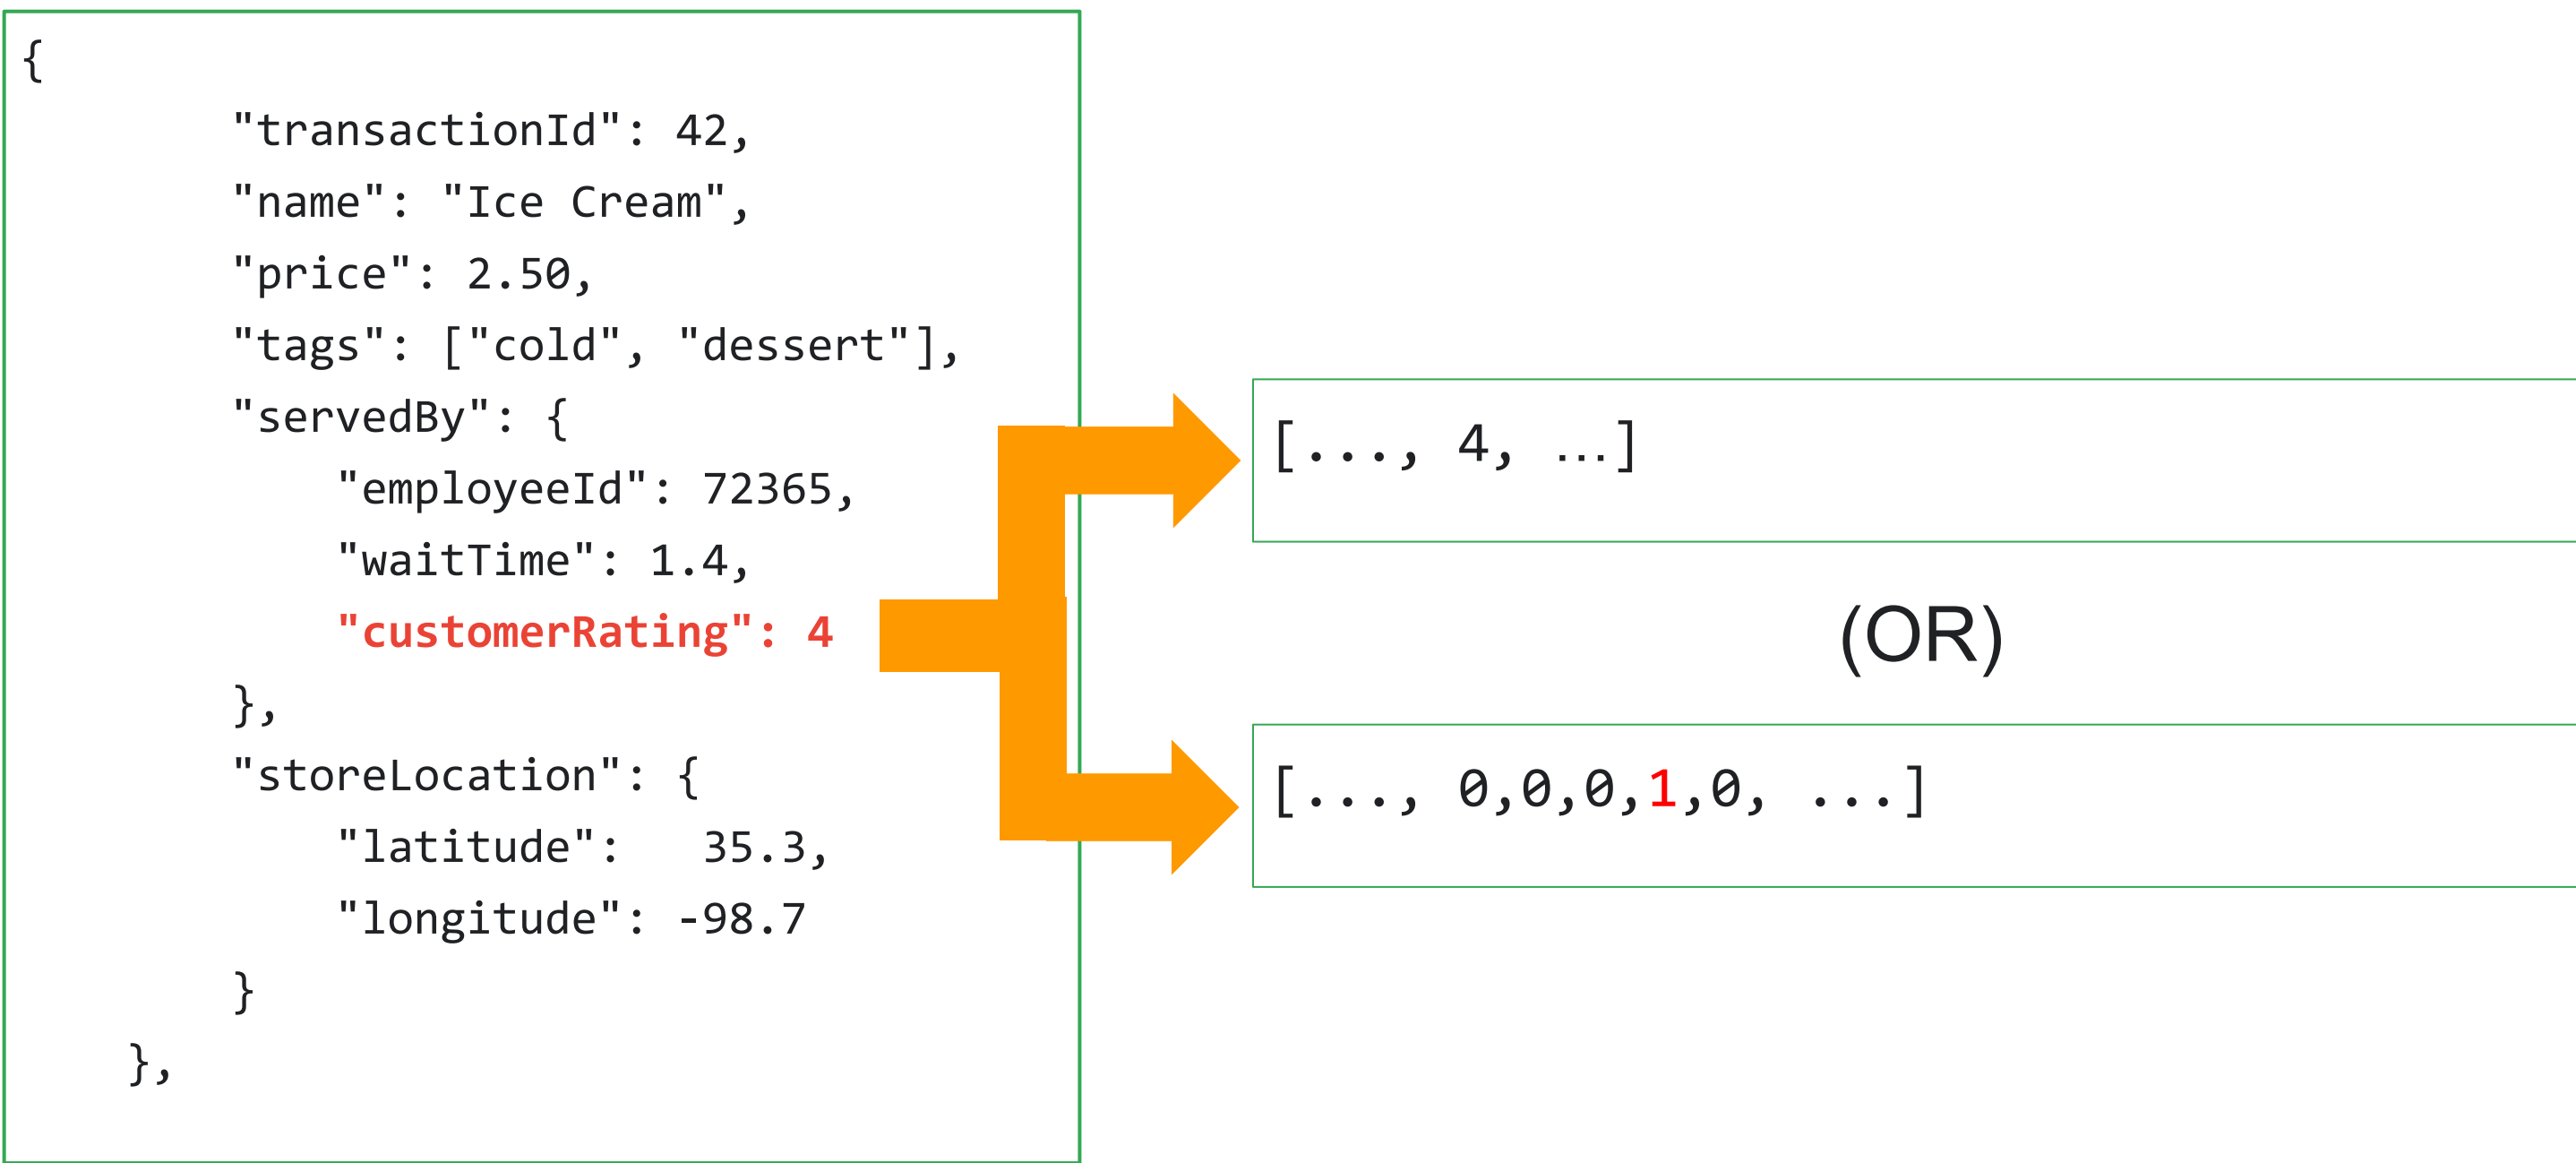
If you don't have a vocabulary of all possible values:

```
tf.feature_column.categorical_column_with_hash_bucket('employeeId',  
    hash_bucket_size = 500)
```

# Categorical variables should be one-hot encoded

```
{
  "transactionId": 42,
  "name": "Ice Cream",
  "price": 2.50,
  "tags": ["cold", "dessert"],
  "servedBy": {
    "employeeId": 72365,
    "waitTime": 1.4,
    "customerRating": 4
  },
  "storeLocation": {
    "latitude": 35.3,
    "longitude": -98.7
  }
},
```

# Categorical variables should be one-hot encoded



# Don't mix magic numbers with data

```
{
  "transactionId": 42,
  "name": "Ice Cream",
  "price": 2.50,
  "tags": ["cold", "dessert"],
  "servedBy": {
    "employeeId": 72365,
    "waitTime": 1.4,
    "customerRating": -1
  },
  "storeLocation": {
    "latitude": 35.3,
    "longitude": -98.7
  }
},
```

# Don't mix magic numbers with data

```
{  
  "transactionId": 42,  
  "name": "Ice Cream",  
  "price": 2.50,  
  "tags": ["cold", "dessert"],  
  "servedBy": {  
    "employeeId": 72365,  
    "waitTime": 1.4,  
    "customerRating": -1  
  },  
  "storeLocation": {  
    "latitude": 35.3,  
    "longitude": -98.7  
  }  
},
```

```
[..., 4, 1, ...] # 4  
[..., 0, 0, ...] # -1
```

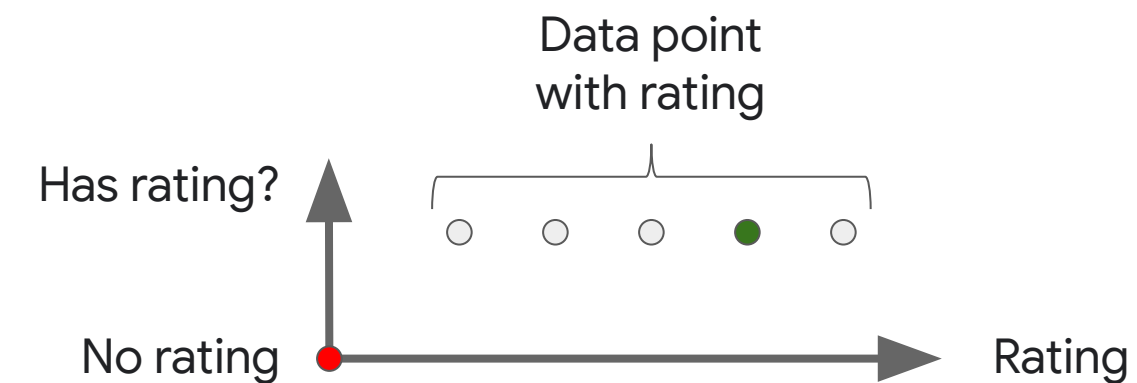
# Don't mix magic numbers with data

```
{
  "transactionId": 42,
  "name": "Ice Cream",
  "price": 2.50,
  "tags": ["cold", "dessert"],
  "servedBy": {
    "employeeId": 72365,
    "waitTime": 1.4,
    "customerRating": -1
  },
  "storeLocation": {
    "latitude": 35.3,
    "longitude": -98.7
  }
},
```

```
[..., 4, 1, ...] # 4
[..., 0, 0, ...] # -1
```

(OR)

```
[..., 0, 0, 0, 1, 1, ...] # 4
[..., 0, 0, 0, 0, 0, ...] # -1
```






# ML vs Statistics

ML = lots of data, keep  
outliers and build models  
for them

**ML** = lots of data, keep  
outliers and build models  
for them

**Statistics** = “I’ve got all  
the data I’ll ever get”,  
throw away outliers


# Exact floats are not meaningful



LOS ANGELES


SAN FRANCISCO

# Discretize floating point values into bins

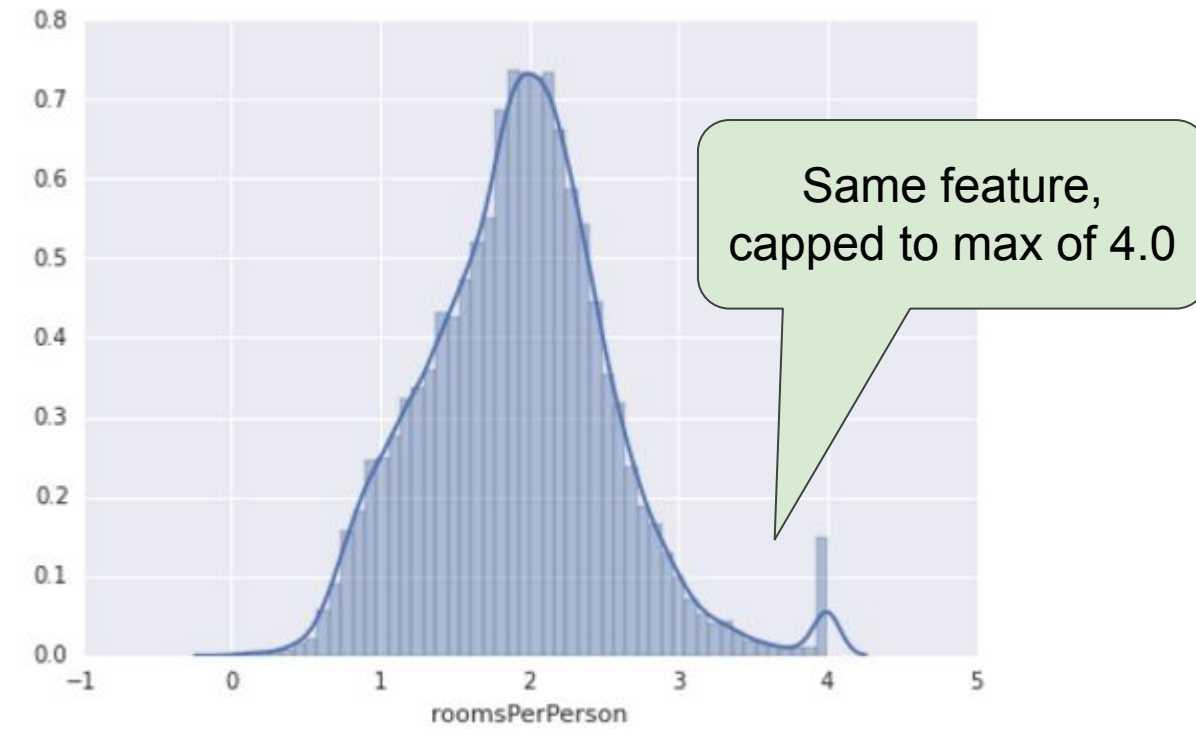


```
lat = tf.feature_column.numeric_column('latitude')
dlat = tf.feature_column.bucketized_column(
    lat, boundaries=np.arange(32,42,1).tolist()
)
```

# Crazy outliers will hurt trainability



Rooms Per Person



Capped Rooms Per Person

```
features['capped_rooms'] = tf.clip_by_value(  
    features['rooms'] ,  
    clip_value_min=0,  
    clip_value_max=4  
)
```

Ideally, features should have  
a similar range

Typically [0,1] or [-1,1]

```
features['scaled_price'] =  
    (features['price'] - min_price) /  
    (max_price - min_price)
```



# Lab

---

Improve the accuracy of a model by adding new features with the appropriate representation