



Preprocessing and feature
creation

Carl Osipov

Learn how to...

Get started with preprocessing
and feature creation

Learn how to...

Get started with preprocessing
and feature creation

Use Apache Beam and Cloud
Dataflow for feature engineering

Feature engineering often requires global statistics and vocabularies

```
features['scaled_price'] =  
    (features['price'] - min_price) / (max_price - min_price)
```

```
tf.feature_column.categorical_column_with_vocabulary_list('city',  
    keys=['San Diego', 'Los Angeles', 'San Francisco', 'Sacramento']),
```

How do you get full
vocabulary of cities in
the training dataset?

Feature engineering often requires global statistics and vocabularies

```
features['scaled_price'] =  
    (features['price'] - min_price) / (max_price - min_price)
```

```
tf.feature_column.categorical_column_with_vocabulary_list('city',  
    keys=['San Diego', 'Los Angeles', 'San Francisco', 'Sacramento']),
```

How do you get full
vocabulary of cities in
the training dataset?

Feature engineering often requires global statistics and vocabularies

```
features['scaled_price'] =  
    (features['price'] - min_price) / (max_price - min_price)
```

```
tf.feature_column.categorical_column_with_vocabulary_list('city',  
    keys=['San Diego', 'Los Angeles', 'San Francisco', 'Sacramento']),
```

How do you get full
vocabulary of cities in
the training dataset?

Feature engineering often requires global statistics and vocabularies

```
features['scaled_price'] =  
    (features['price'] - min_price) / (max_price - min_price)
```

```
tf.feature_column.categorical_column_with_vocabulary_list('city',  
    keys=['San Diego', 'Los Angeles', 'San Francisco', 'Sacramento']),
```

How do you get full
vocabulary of cities in
the training dataset?

Preprocess with...

1. BigQuery

2. Apache Beam

3. TensorFlow

Things that are commonly done in preprocessing

Remove examples that you don't want to train on



In BigQuery
or Beam

Things that are commonly done in preprocessing

Remove examples that you don't want to train on

Compute vocabularies for categorical columns

Compute aggregate statistics for numeric columns

In BigQuery
or Beam

Things that are commonly done in preprocessing

Remove examples that you don't want to train on

Compute vocabularies for categorical columns

Compute aggregate statistics for numeric columns

Compute time-windowed statistics (e.g. number of products sold in previous hour) for use as input features

In BigQuery
or Beam

In Beam only

Things that are commonly done in preprocessing

Scaling, discretization, etc. of numeric features

Splitting, lower-casing, etc. of textual features

Resizing of input images

Normalizing volume level of input audio

In TensorFlow
or Beam

Example of preprocessing in BigQuery

```
SELECT
  (tolls_amount + fare_amount)
    AS fare_amount,
  DAYOFWEEK(pickup_datetime)
    AS dayofweek,
  HOUR(pickup_datetime)
    AS hourofday,
  ...
FROM
  `nyc-tlc.yellow.trips`
WHERE
  trip_distance > 0
```

Example of preprocessing in BigQuery

```
SELECT
  (tolls_amount + fare_amount)
    AS fare_amount,
  DAYOFWEEK(pickup_datetime)
    AS dayofweek,
  HOUR(pickup_datetime)
    AS hourofday,
  ...
FROM
  `nyc-tlc.yellow.trips`
WHERE
  trip_distance > 0
```

Example of preprocessing in BigQuery

```
SELECT
  (tolls_amount + fare_amount)
  AS fare_amount,
  DAYOFWEEK(pickup_datetime)
  AS dayofweek,
  HOUR(pickup_datetime)
  AS hourofday,
  ...
FROM
  `nyc-tlc.yellow.trips`
WHERE
  trip_distance > 0
```

Example of preprocessing in BigQuery

```
SELECT
  (tolls_amount + fare_amount)
    AS fare_amount,
  DAYOFWEEK(pickup_datetime)
    AS dayofweek,
  HOUR(pickup_datetime)
    AS hourofday,
  ...
FROM
  `nyc-tlc.yellow.trips`
WHERE
  trip_distance > 0
```


There are two places for feature creation in TensorFlow

```
features['capped_rooms'] = tf.clip_by_value(
    features['rooms'] ,
    clip_value_min=0,
    clip_value_max=4
)
```

1. Features are preprocessed in input_FN (train, eval, serving)

```
lat = tf.feature_column.numeric_column('latitude')
dlat = tf.feature_column.bucketized_column(lat,
    boundaries=np.arange(32,42,1).tolist())
```

2. Feature columns are Passed into the estimator during construction

There are two places for feature creation in TensorFlow

```
features['capped_rooms'] = tf.clip_by_value(
    features['rooms'] ,
    clip_value_min=0,
    clip_value_max=4
)
```

1. Features are preprocessed in input_FN (train, eval, serving)

```
lat = tf.feature_column.numeric_column('latitude')
dlat = tf.feature_column.bucketized_column(lat,
    boundaries=np.arange(32,42,1).tolist())
```

2. Feature columns are Passed into the estimator during construction

1. Example of preprocessing in TensorFlow input_fn

```
def add_engineered(features):  
    lat1 = features['pickuplat']  
    ...  
    dist = tf.sqrt(latdiff*latdiff + londiff*londiff)  
    features['euclidean'] = dist  
    return features
```

How do we make sure this function gets called during both training and prediction?

1. Example of preprocessing in TensorFlow input_fn

```
def add_engineered(features):  
    lat1 = features['pickuplat']  
    ...  
    dist = tf.sqrt(latdiff*latdiff + londiff*londiff)  
    features['euclidean'] = dist  
    return features
```

How do we make sure this function gets called during both training and prediction?

Wrap features by call to the feature engineering to function

Wrap features in training/evaluation input function:

```
def input_fn():  
    features = ...  
    label = ...  
    return add_engineered(features), label
```

Wrap features in serving input function also:

```
def serving_input_fn():  
    feature_placeholders = ...  
    features = ...  
    return tf.estimator.export.ServingInputReceiver(  
        add_engineered(features), feature_placeholders)
```

Wrap features by call to the feature engineering to function

Wrap features in training/evaluation input function:

```
def input_fn():  
    features = ...  
    label = ...  
    return add_engineered(features), label
```

Wrap features in serving input function also:

```
def serving_input_fn():  
    feature_placeholders = ...  
    features = ...  
    return tf.estimator.export.ServingInputReceiver(  
        add_engineered(features), feature_placeholders)
```

Wrap features by call to the feature engineering to function

Wrap features in training/evaluation input function:

```
def input_fn():  
    features = ...  
    label = ...  
    return add_engineered(features), label
```

Wrap features in serving input function also:

```
def serving_input_fn():  
    feature_placeholders = ...  
    features = ...  
    return tf.estimator.export.ServingInputReceiver(  
        add_engineered(features), feature_placeholders)
```

2. Example of preprocessing via feature columns

```
def build_estimator(model_dir, nbuckets):  
    latbuckets = np.linspace(38.0, 42.0, nbuckets).tolist()  
    b_plat = tf.feature_column.bucketized_column(plat, latbuckets)  
    b_dlat = tf.feature_column.bucketized_column(dlat, latbuckets)  
  
    return tf.estimator.LinearRegressor(  
        model_dir=model_dir,  
        feature_columns=[..., b_plat, b_dlat, ...])
```


2. Example of preprocessing via feature columns

```
def build_estimator(model_dir, nbuckets):  
    latbuckets = np.linspace(38.0, 42.0, nbuckets).tolist()  
    b_plat = tf.feature_column.bucketized_column(plat, latbuckets)  
    b_dlat = tf.feature_column.bucketized_column(dlat, latbuckets)  
  
    return tf.estimator.LinearRegressor(  
        model_dir=model_dir,  
        feature_columns=[..., b_plat, b_dlat, ...])
```

2. Example of preprocessing via feature columns

```
def build_estimator(model_dir, nbuckets):  
    latbuckets = np.linspace(38.0, 42.0, nbuckets).tolist()  
    b_plat = tf.feature_column.bucketized_column(plat, latbuckets)  
    b_dlat = tf.feature_column.bucketized_column(dlat, latbuckets)  
  
    return tf.estimator.LinearRegressor(  
        model_dir=model_dir,  
        feature_columns=[..., b_plat, b_dlat, ...])
```

2. Normalization can be done in feature columns

```
def zscore(col):  
    mean = 3.04  
    std = 1.2  
    return (col - mean)/std  
  
feature_name = 'total_bedrooms'  
normalized_feature = tf.feature_column.numeric_column(  
    feature_name,  
    normalizer_fn=zscore)
```

Example of preprocessing in Beam (covered next)

```
def to_csv(rowdict):  
    if distance(rowdict['pickuplon'], ...) > 10: # only rides of more than 10km  
        CSV_COLUMNS = 'fare_amount,dayofweek,...,key'.split(',')  
        yield ','.join([str(rowdict[k]) for k in CSV_COLUMNS])  
  
def preprocess():  
    ...  
    for n, step in enumerate(['train', 'valid']):  
        (p | 'read_{}'.format(step) >>  
beam.io.Read(beam.io.BigQuerySource(query=query))  
        | 'tocsv_{}'.format(step) >> beam.FlatMap(to_csv)  
        | 'write_{}'.format(step) >> beam.io.Write(beam.io.WriteToText(outfile))  
        )  
    p.run()
```

Recap: things that are commonly done in preprocessing

Remove examples that you don't want to train on

Compute vocabularies for categorical columns

Compute aggregate statistics for numeric columns

Compute time-windowed statistics (e.g. number of products sold in previous hour) for use as input features

Scaling, discretization, etc. of numeric features

Splitting, lower-casing, etc. of textual features

Resizing of input images

Normalizing volume level of input audio

In BigQuery
or Beam

In Beam only


In TensorFlow
or Beam



Apache Beam/Cloud Dataflow

Carl Osipov


Beam is a way to write elastic data
processing pipelines



Cloud
Dataflow

Beam is a way to write elastic data processing pipelines

```
def packageHelp(record, keyword):  
    count=0  
    package_name=''  
    if record is not None:  
        lines=record.split('\n')  
        for line in lines:  
            if line.startswith(keyword):  
                package_name=line  
            if 'FIXME' in line or 'TODO' in line:  
                count+=1  
    packages = (getPackages(package_name))  
    for p in packages:  
        yield (p, count)
```



Cloud
Dataflow



Beam is a way to write elastic data
processing pipelines



BigQuery

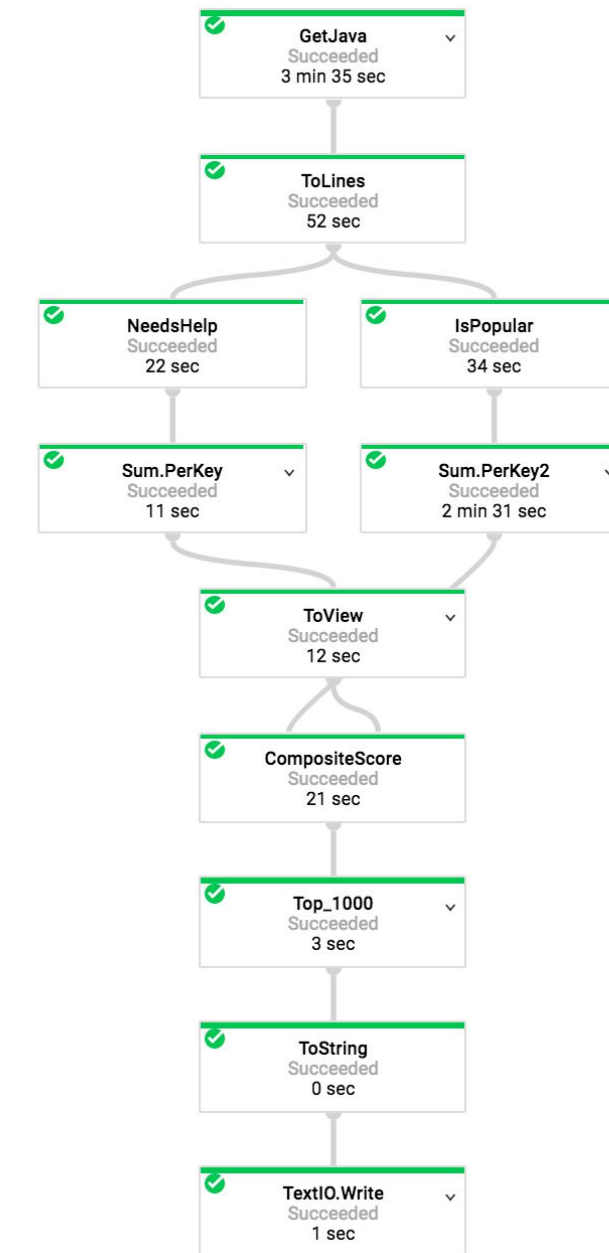
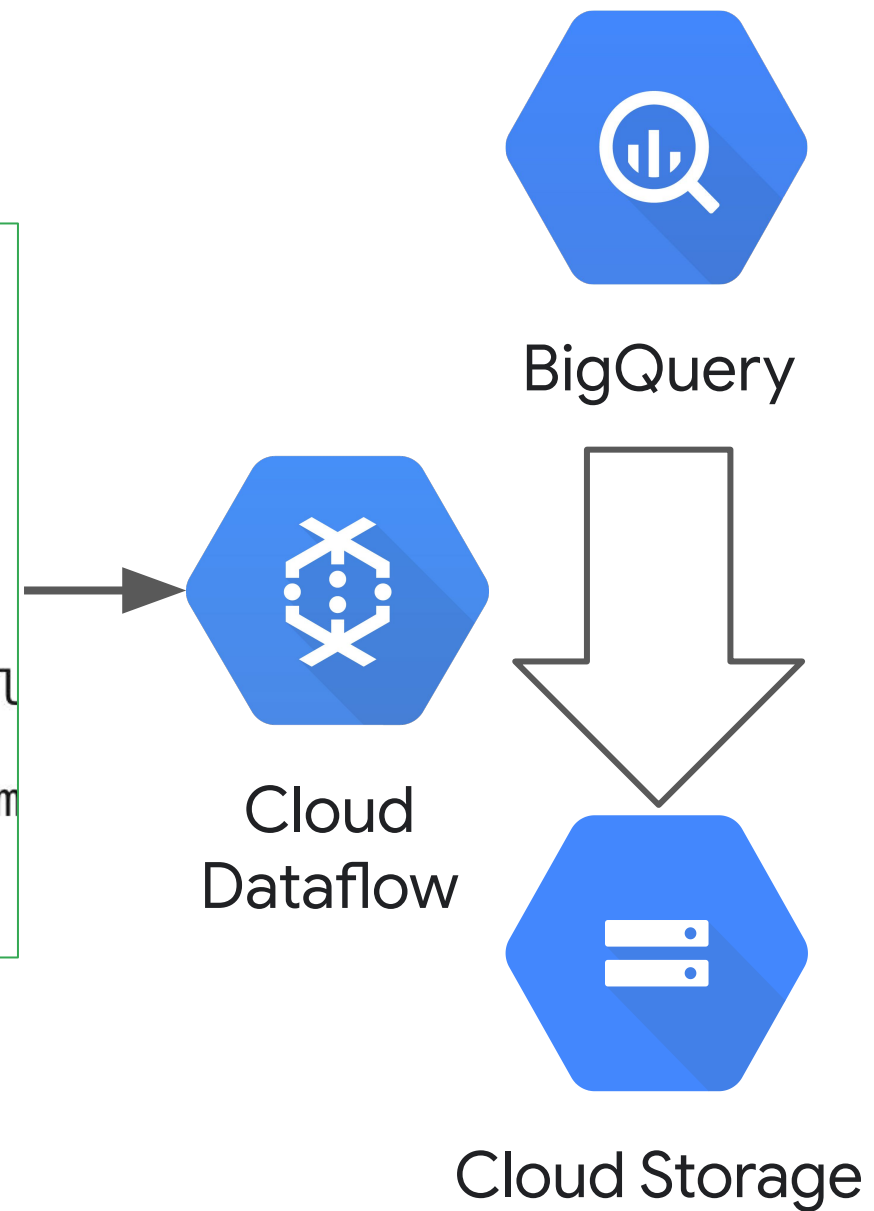
Beam is a way to write elastic data processing pipelines

```
def packageHelp(record, keyword):  
    count=0  
    package_name=''  
    if record is not None:  
        lines=record.split('\n')  
        for line in lines:  
            if line.startswith(keyword):  
                package_name=line  
            if 'FIXME' in line or 'TODO' in line:  
                count+=1  
    packages = (getPackages(package_name))  
    for p in packages:  
        yield (p, count)
```

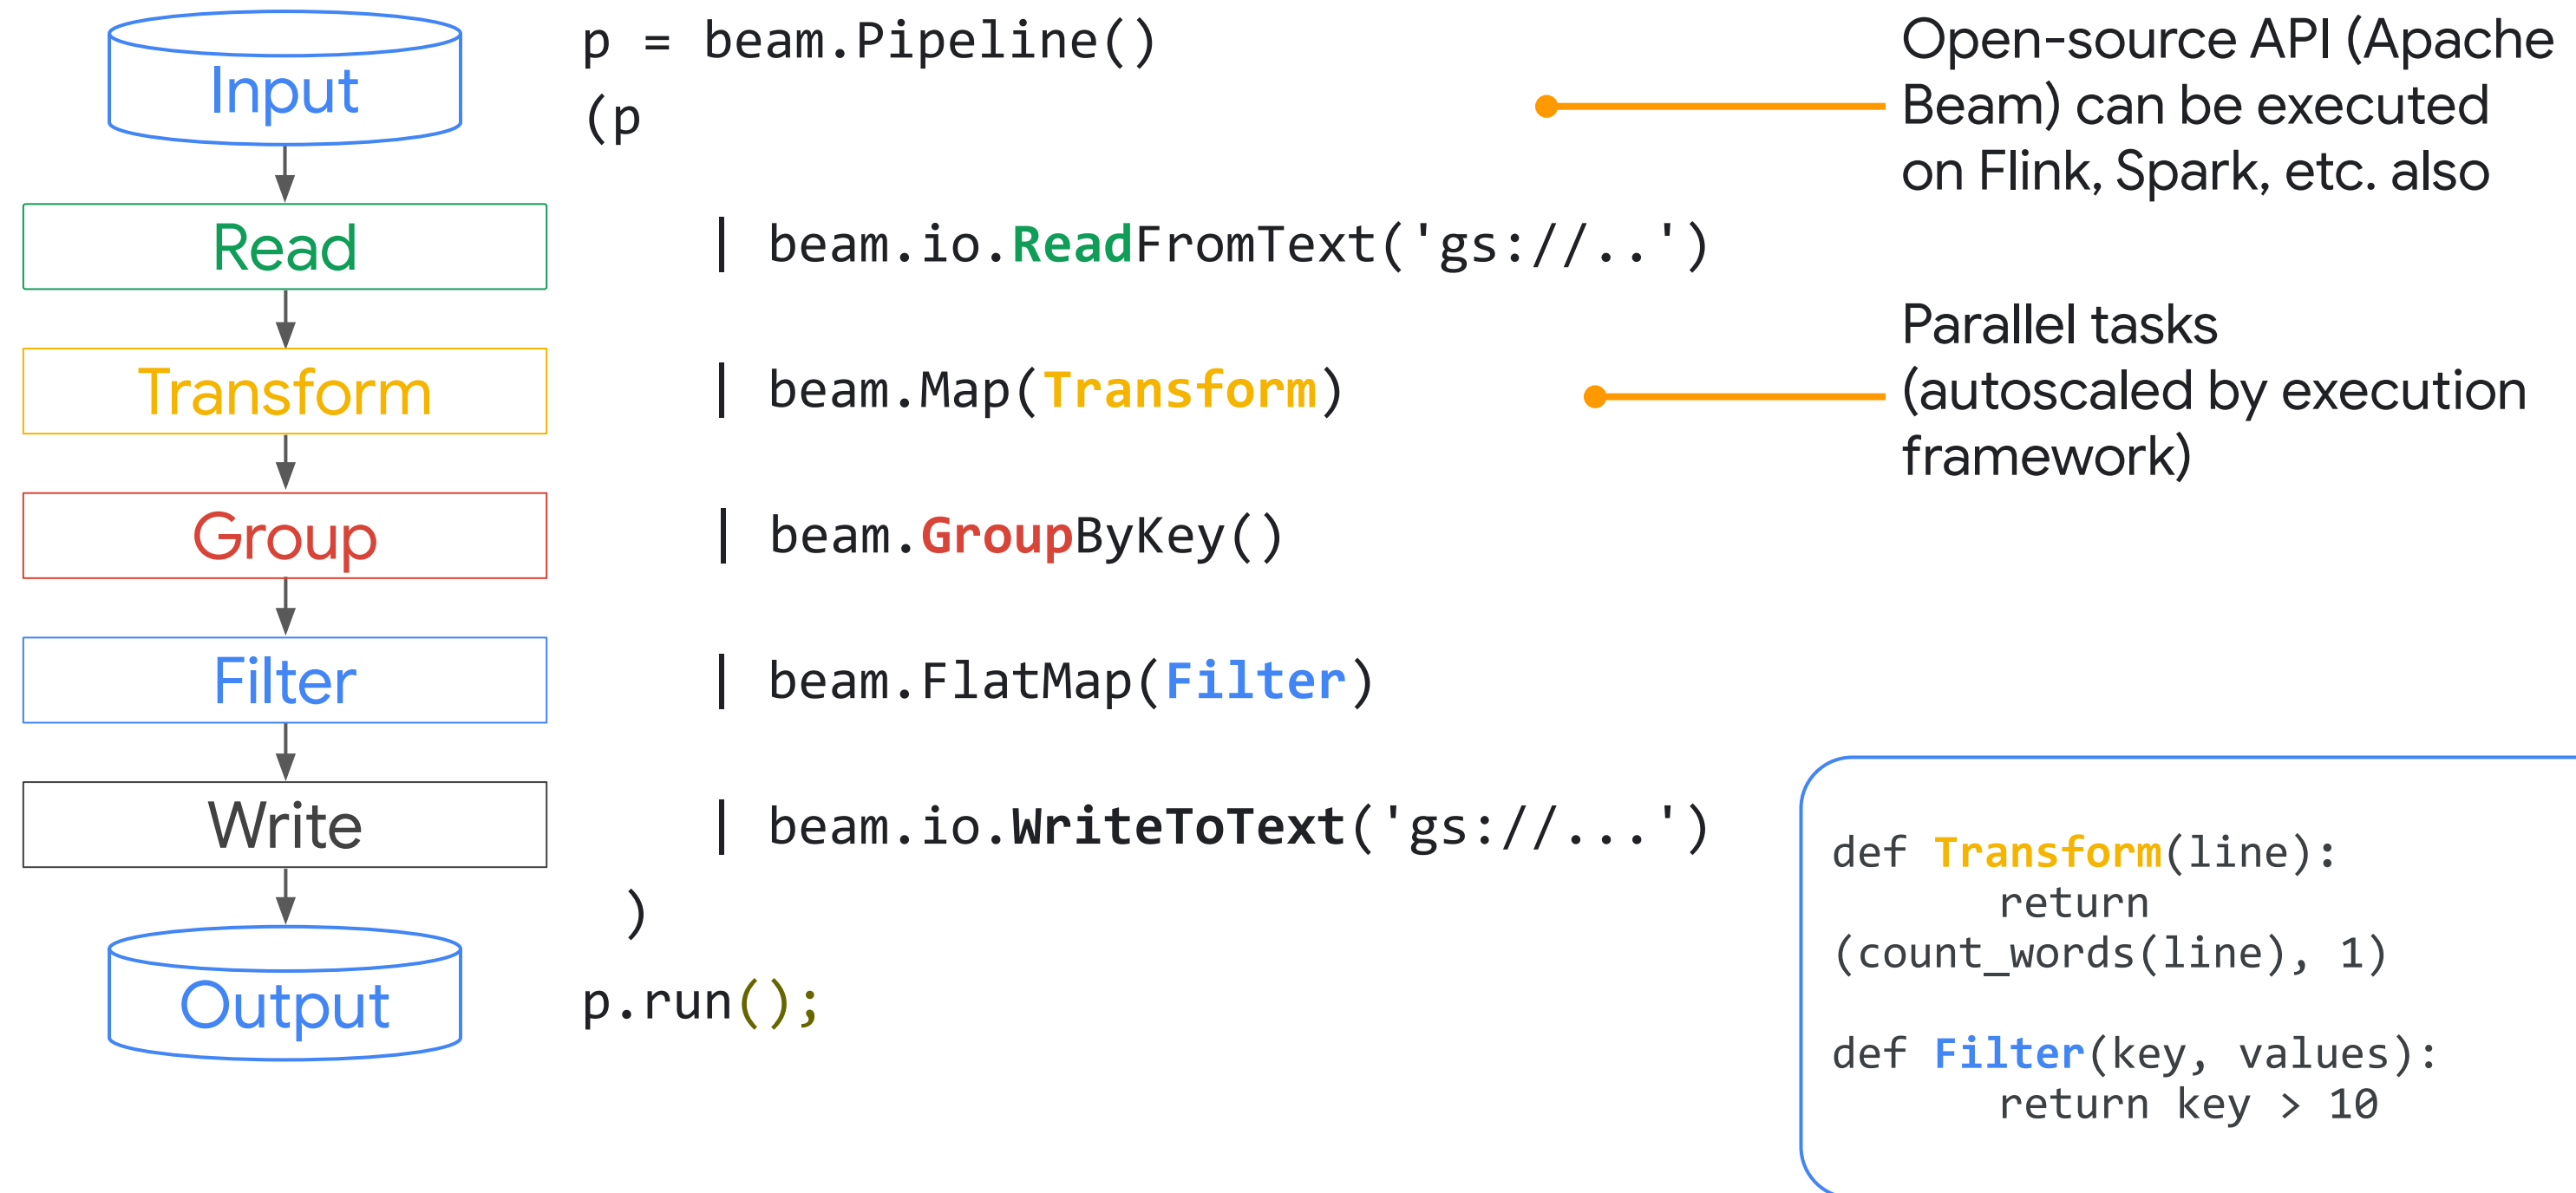


Beam is a way to write elastic data processing pipelines

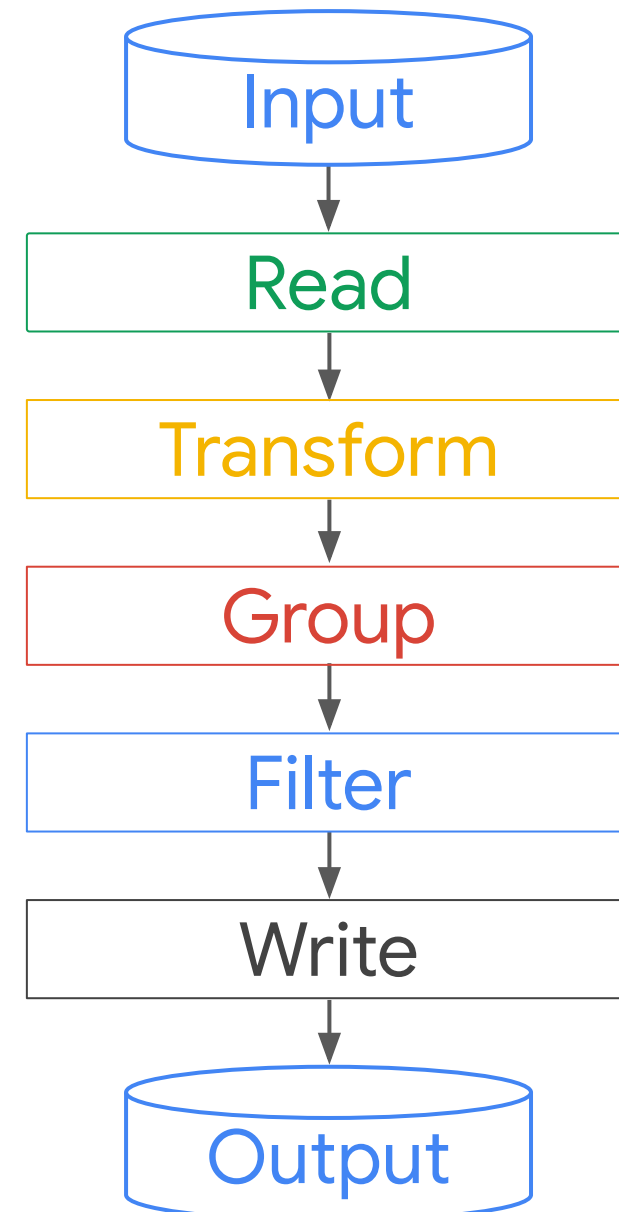
```
def packageHelp(record, keyword):  
    count=0  
    package_name=''  
    if record is not None:  
        lines=record.split('\n')  
        for line in lines:  
            if line.startswith(keyword):  
                package_name=line  
            if 'FIXME' in line or 'TODO' in line:  
                count+=1  
    packages = (getPackages(package_name))  
    for p in packages:  
        yield (p, count)
```



Open-source API, Google infrastructure



Open-source API, Google infrastructure



```
p = beam.Pipeline()  
(p
```

```
| beam.io.ReadFromText('gs://..')  
| beam.Map(Transform)  
| beam.GroupByKey()  
| beam.FlatMap(Filter)  
| beam.io.WriteToText('gs://...')
```

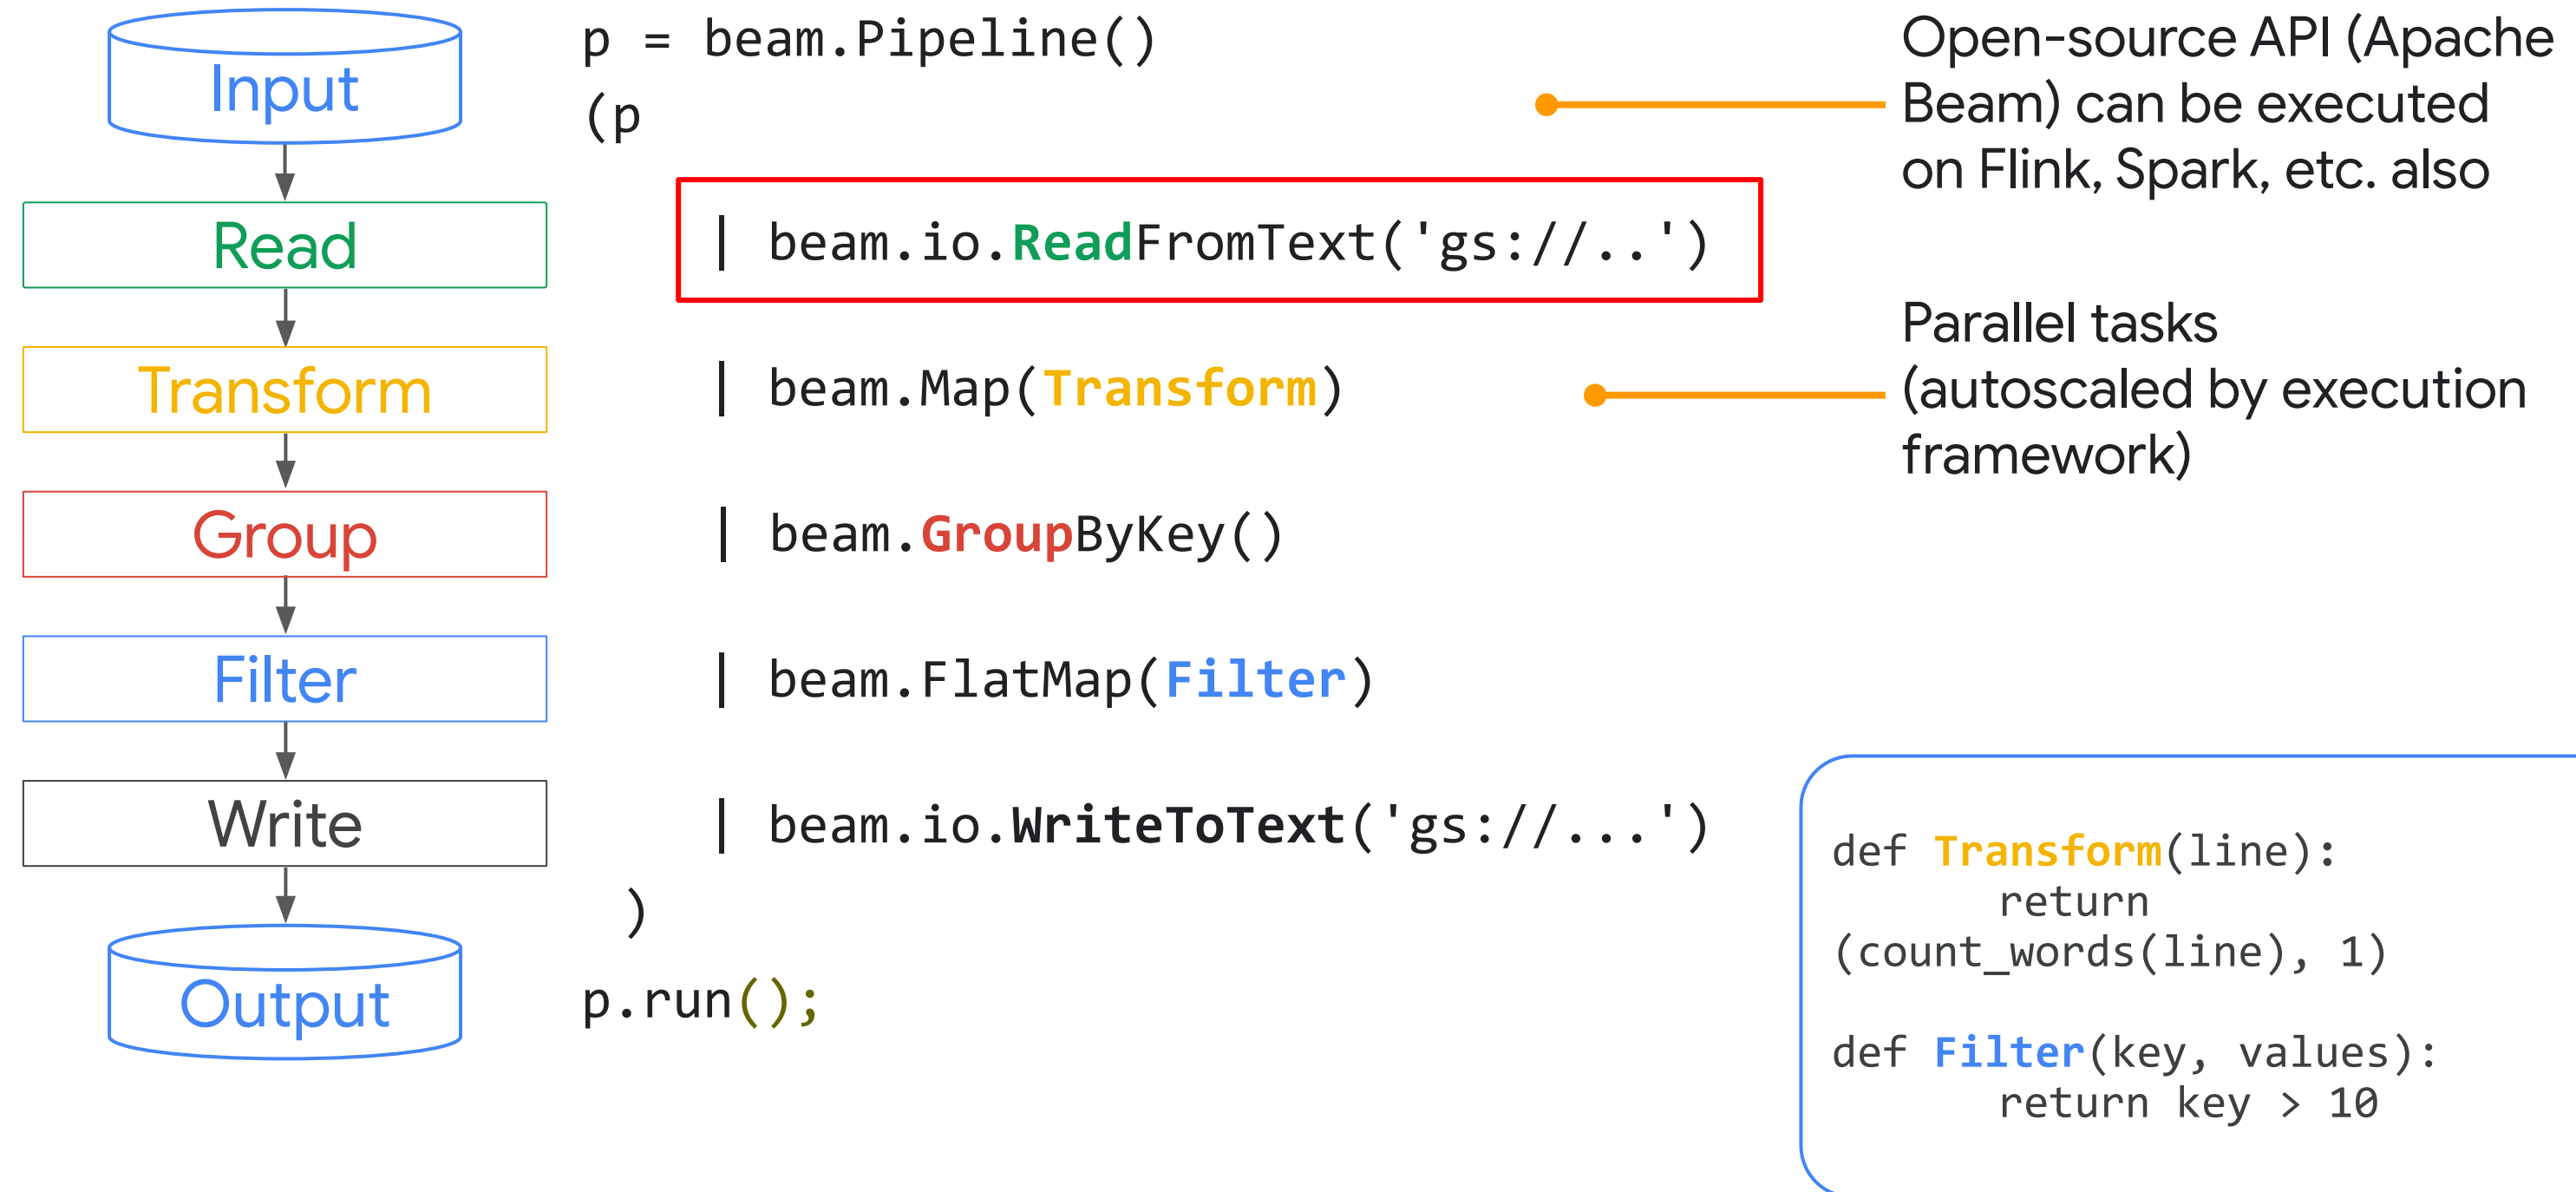
```
)  
p.run();
```

Open-source API (Apache Beam) can be executed on Flink, Spark, etc. also

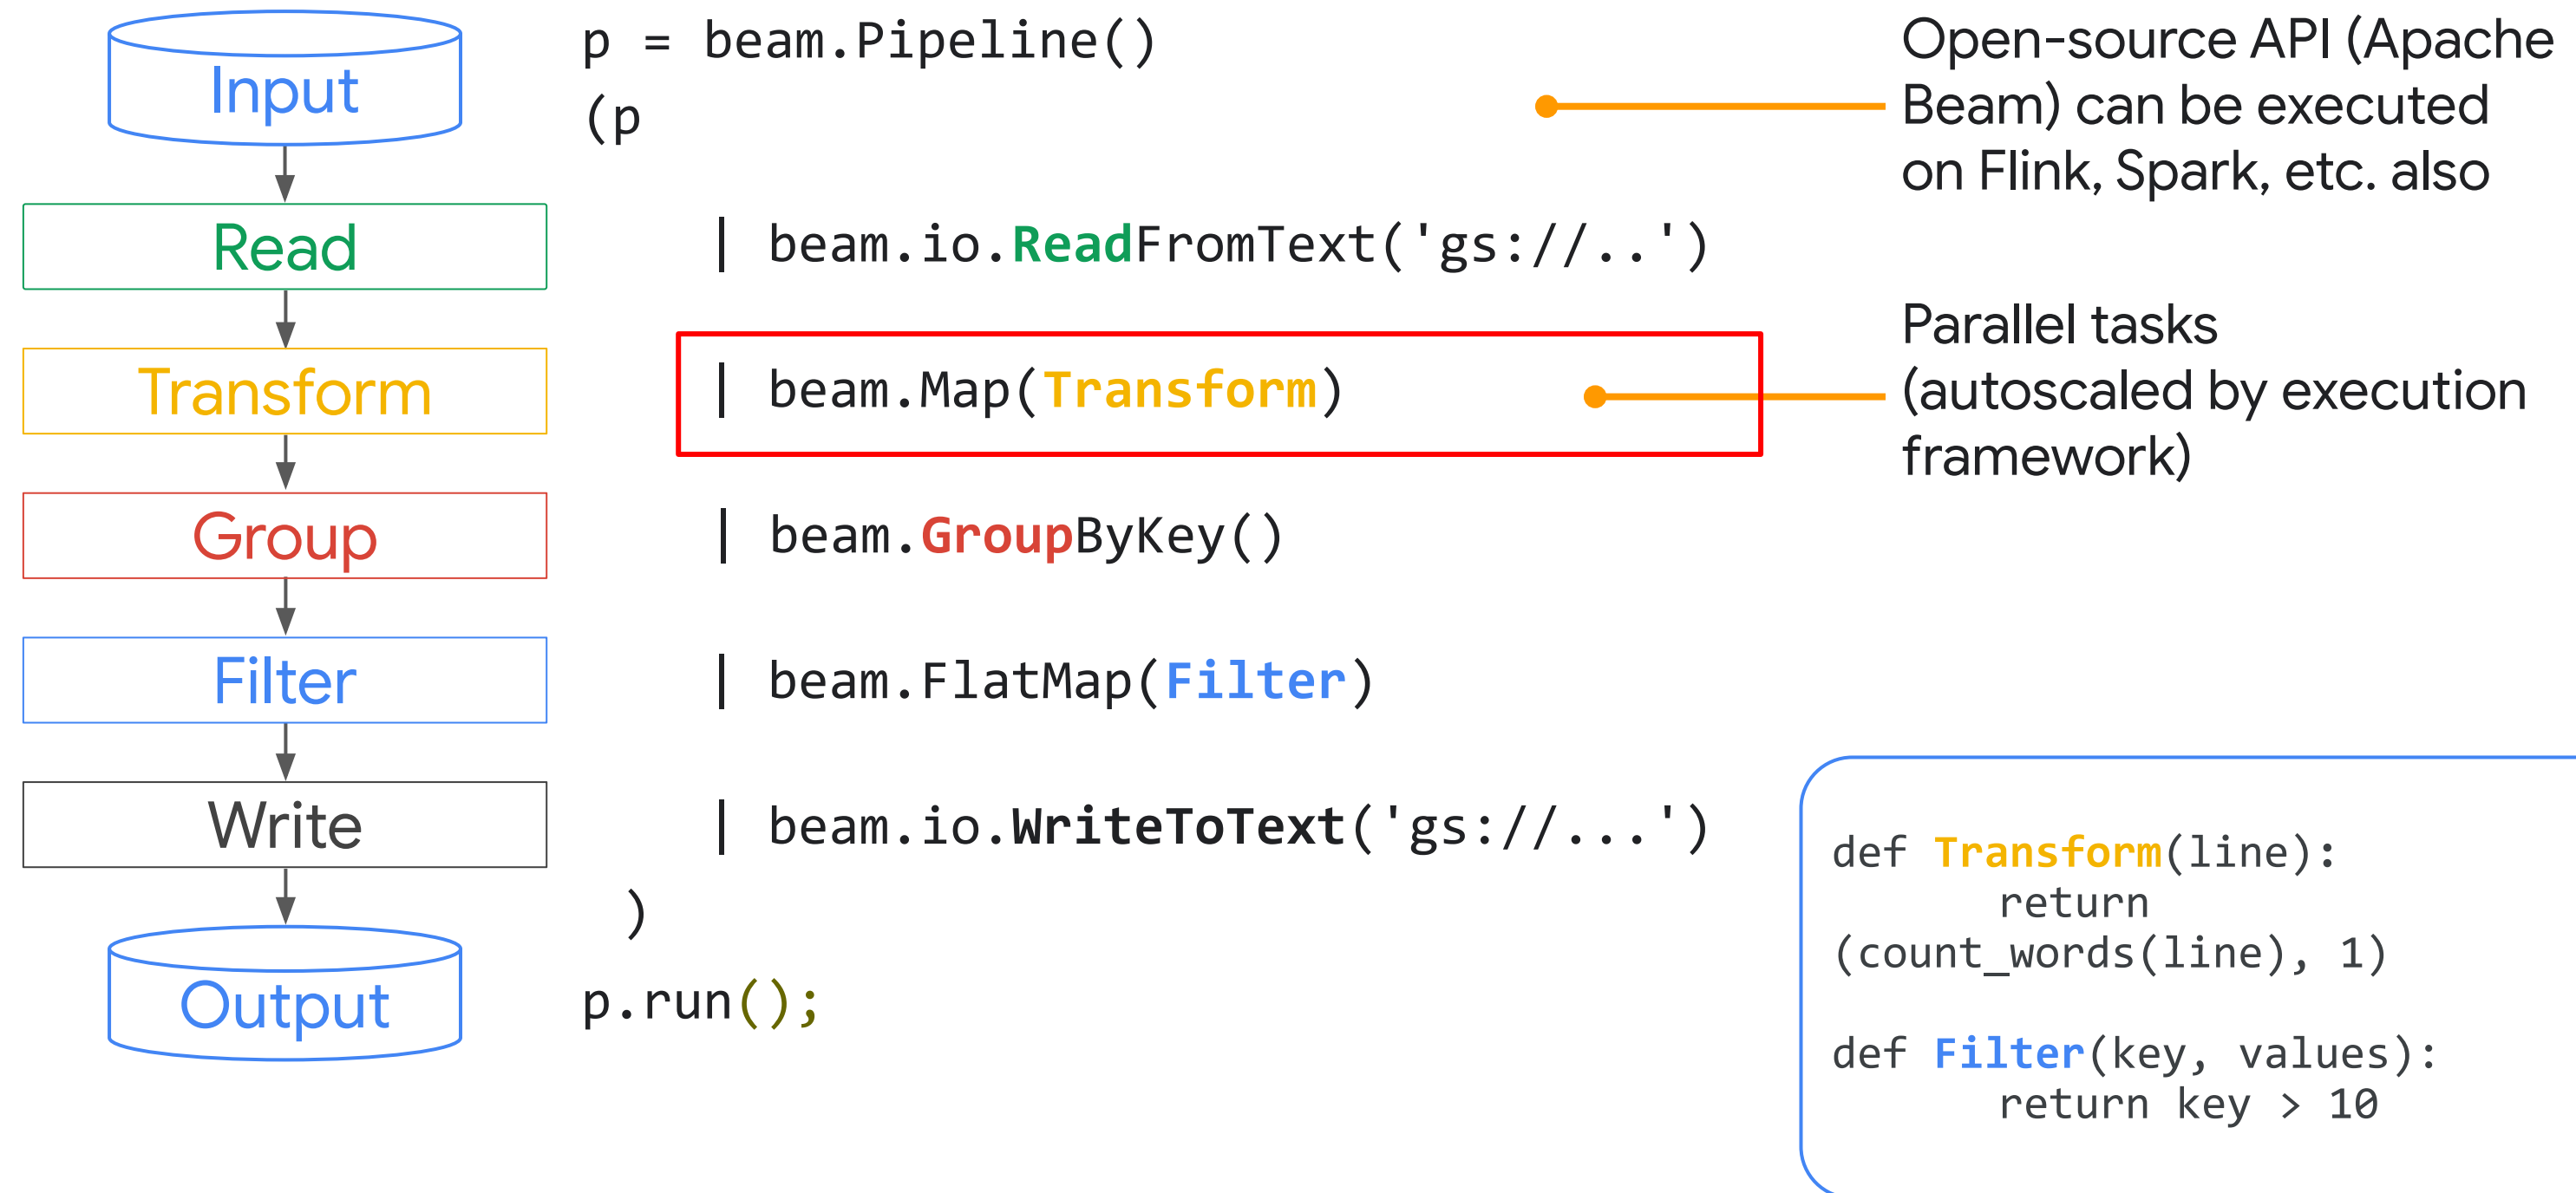
Parallel tasks (autoscaled by execution framework)

```
def Transform(line):  
    return  
    (count_words(line), 1)  
  
def Filter(key, values):  
    return key > 10
```

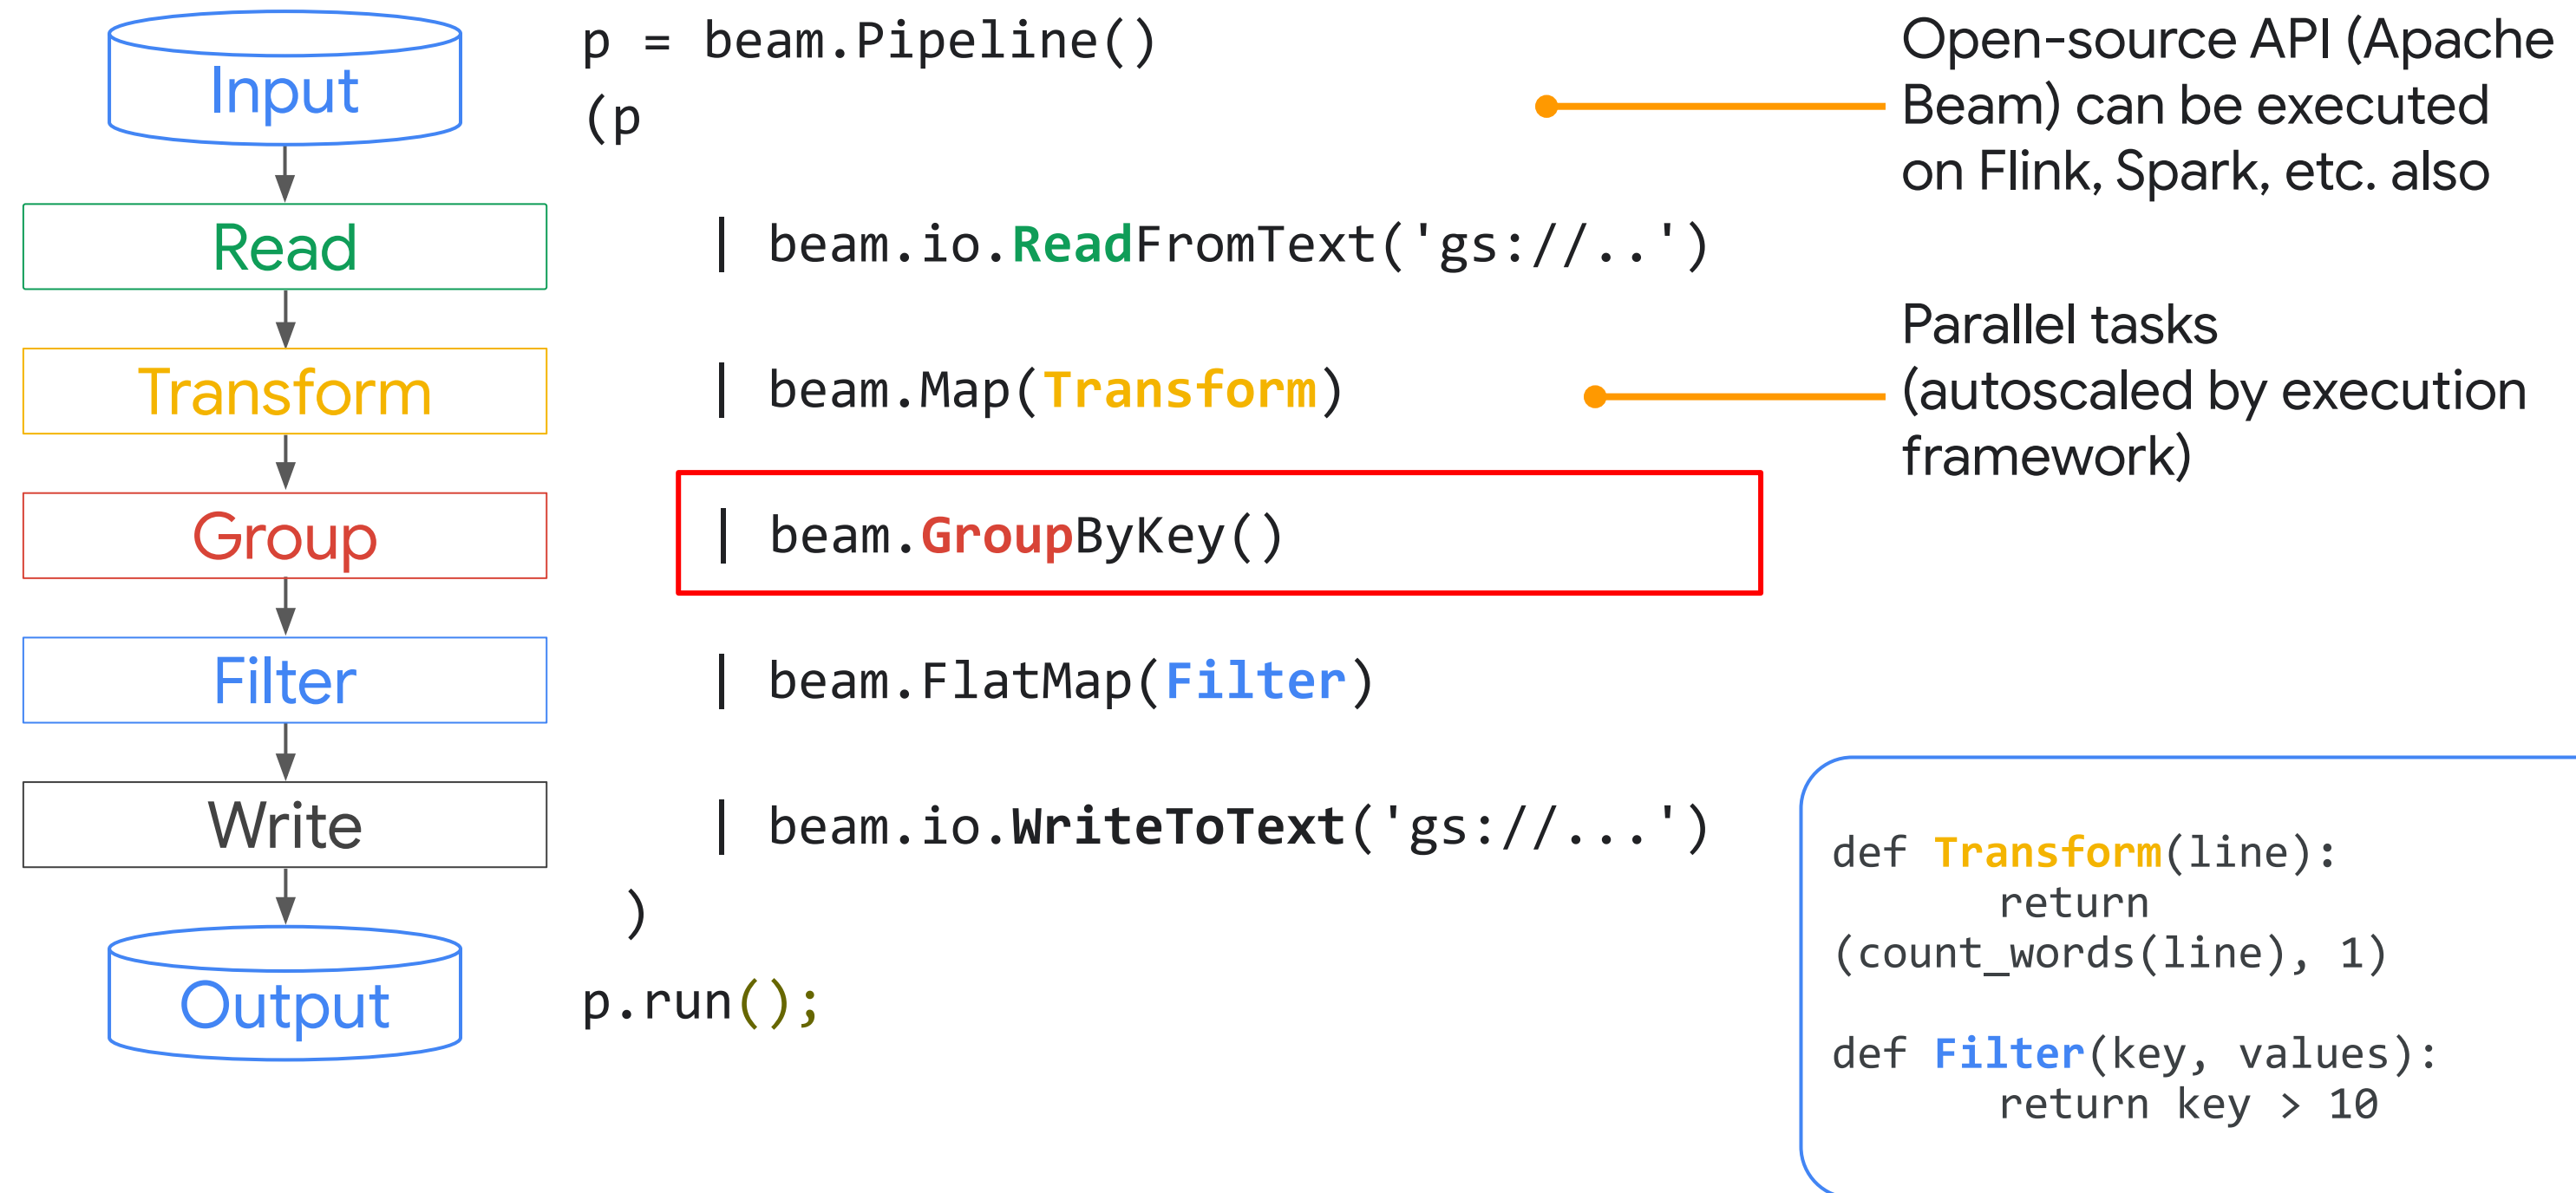
Open-source API, Google infrastructure



Open-source API, Google infrastructure



Open-source API, Google infrastructure



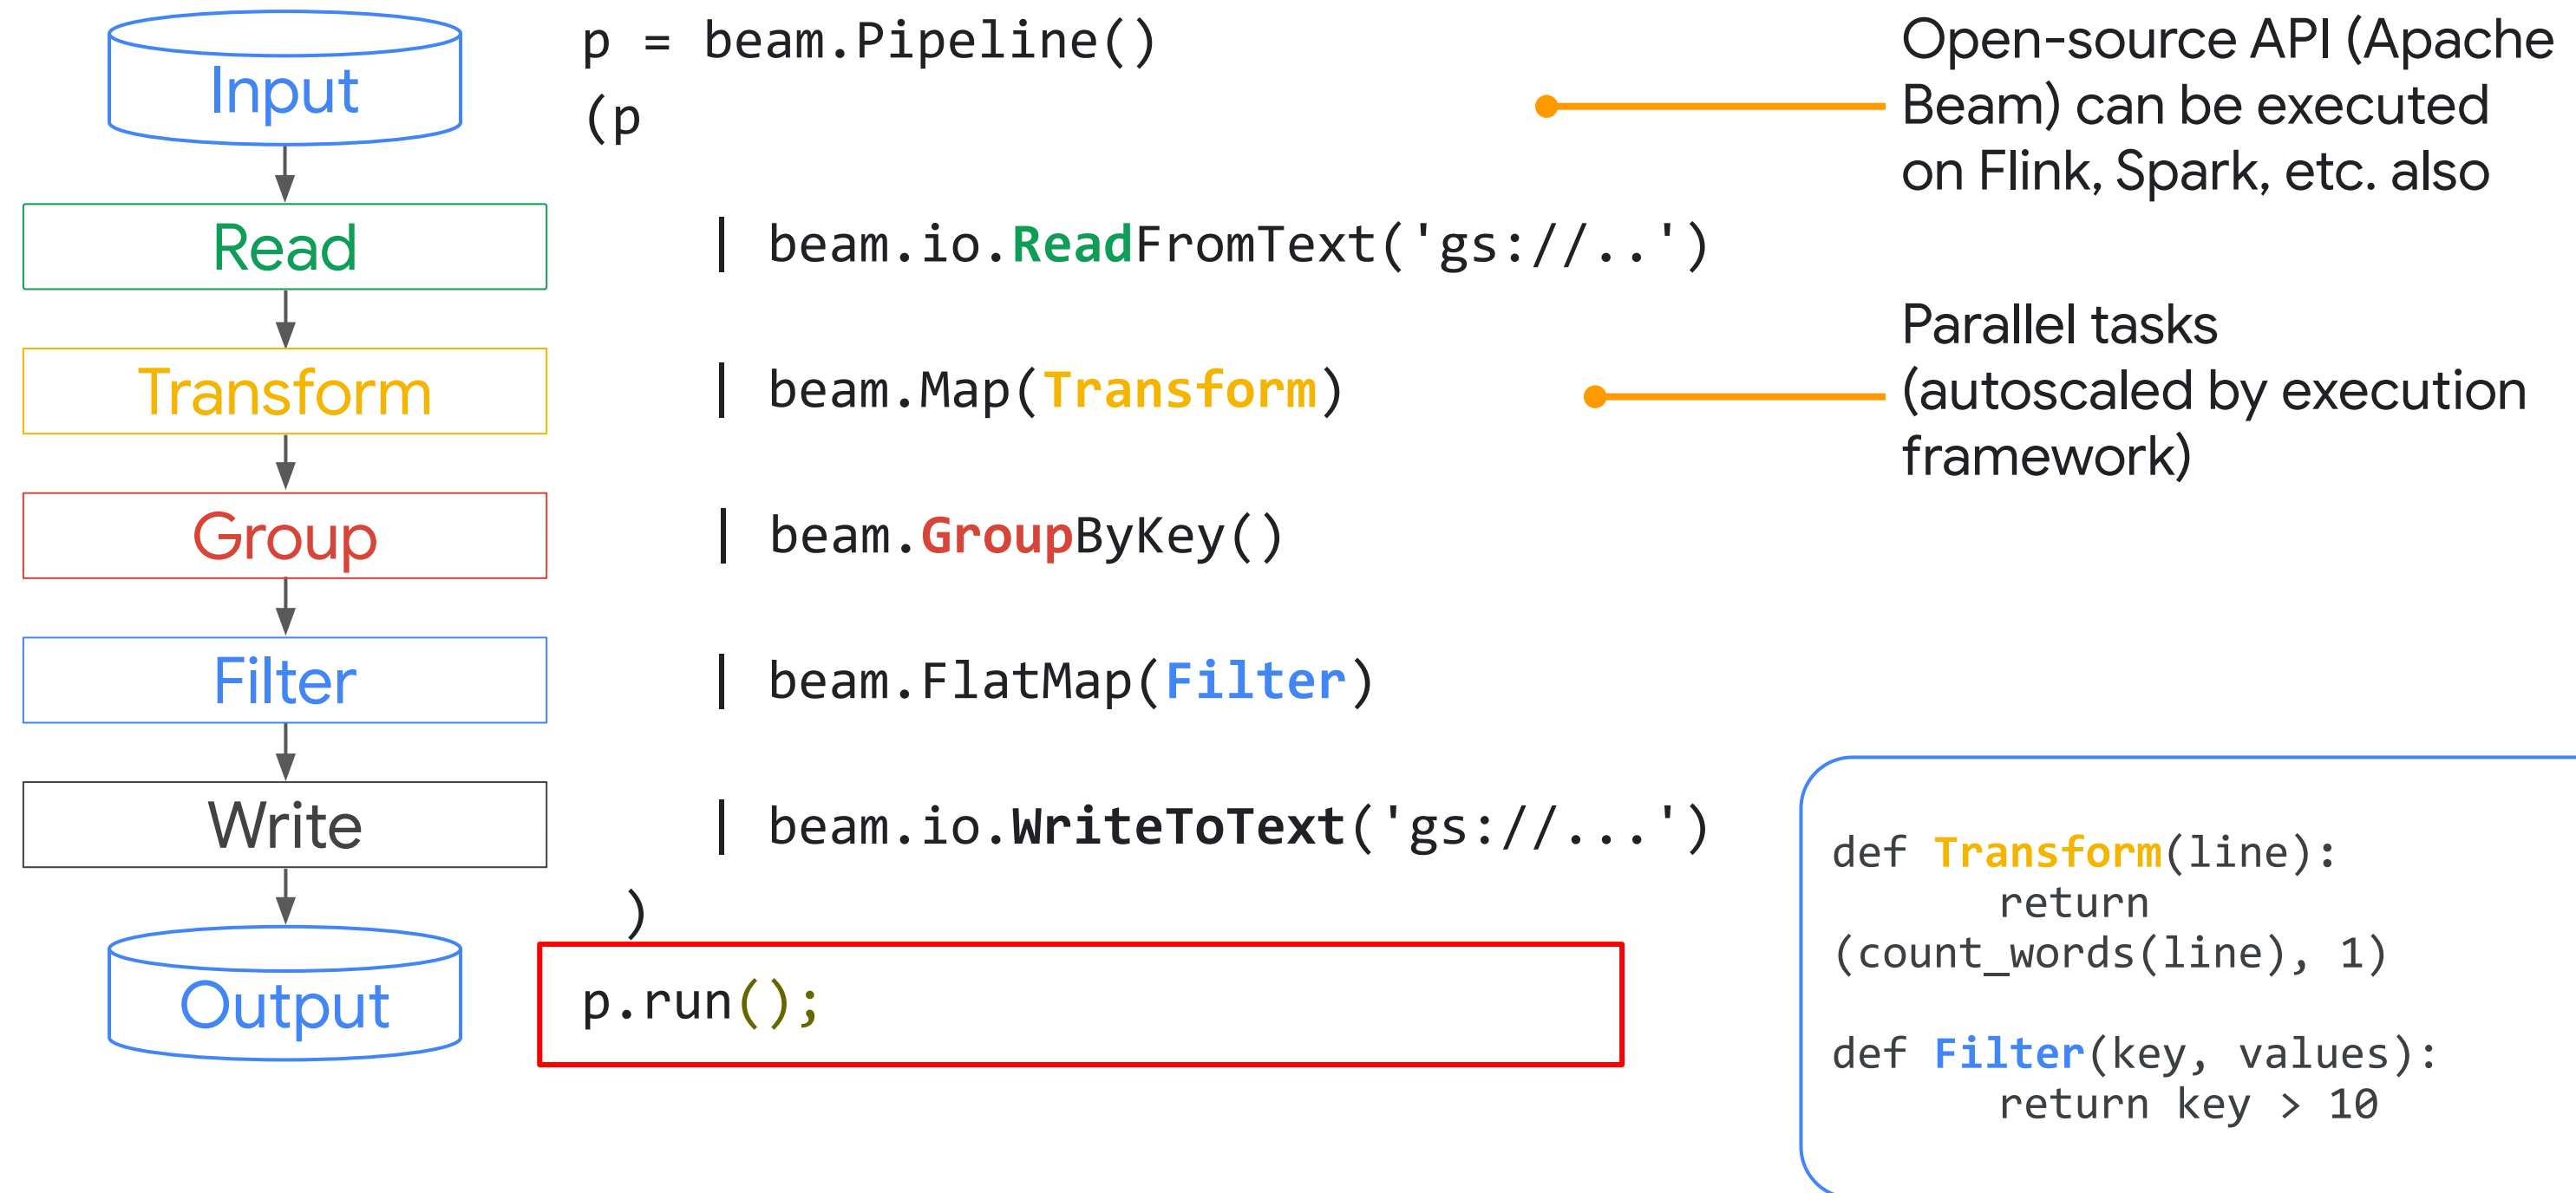
Open-source API, Google infrastructure



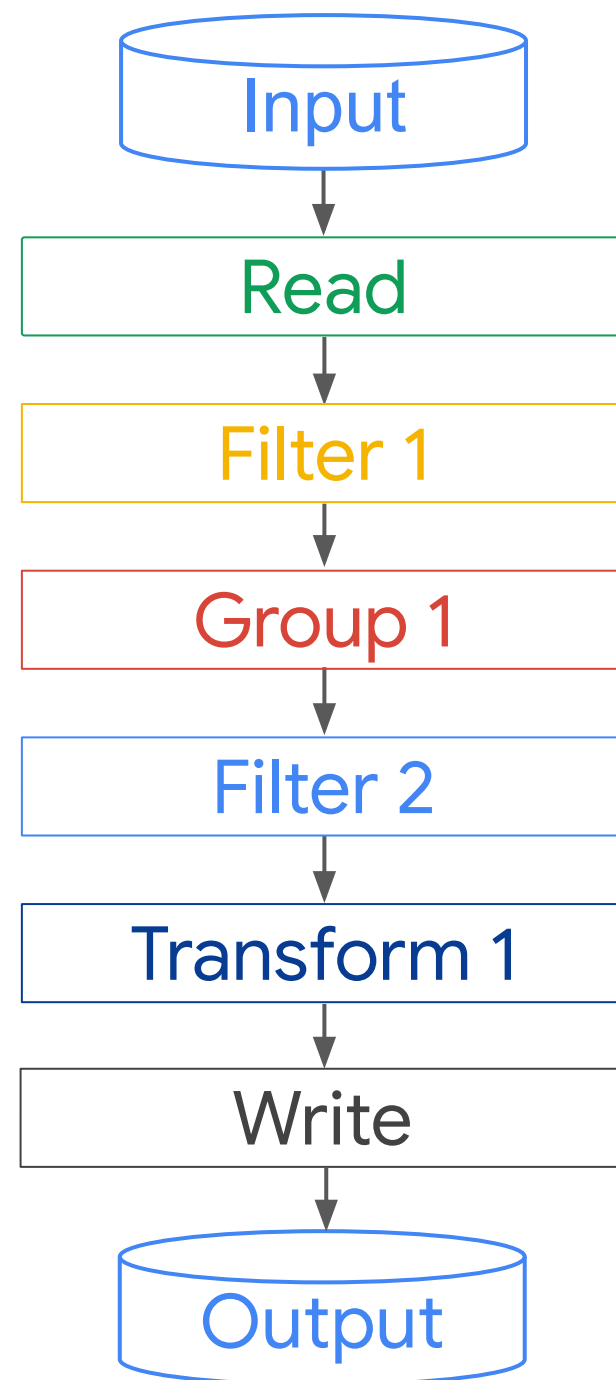
Open-source API, Google infrastructure



Open-source API, Google infrastructure



Open-source API, Google infrastructure



```
Pipeline p = Pipeline.create();
```

```
p
```

```
.apply(TextIO.read().from("gs://..."))
```

```
.apply(ParDo.of(new Filter1()))
```

```
.apply(new Group1())
```

```
.apply(ParDo.of(new Filter2()))
```

```
.apply(new Transform1())
```

```
.apply(TextIO.write().to("gs://..."));
```


```
p.run();
```

Open-source API (Apache Beam) can be executed on Flink, Spark, etc. also

Parallel task (autoscaled by execution framework)

```
class Filter1 extends DoFn<...> {  
    public void  
    processElement(ProcessContext c) {  
        ... = c.element();  
        ...  
        c.output(...);  
    }  
}
```

Open-source API, Google infrastructure



```
Pipeline p = Pipeline.create();
```

p

```
.apply(TextIO.read().from("gs://..."))  
.apply(ParDo.of(new Filter1()))  
.apply(new Group1())  
.apply(ParDo.of(new Filter2()))  
.apply(new Transform1())  
.apply(TextIO.write().to("gs://..."));
```


```
p.run();
```

Open-source API (Apache Beam) can be executed on Flink, Spark, etc. also

Parallel task (autoscaled by execution framework)

```
class Filter1 extends DoFn<...> {  
    public void  
    processElement(ProcessContext c) {  
        ... = c.element();  
        ...  
        c.output(...);  
    }  
}
```

Open-source API, Google infrastructure



```
Pipeline p = Pipeline.create();
```

p

```
.apply(TextIO.read().from("gs://..."))
```

```
.apply(ParDo.of(new Filter1()))
```

```
.apply(new Group1())
```

```
.apply(ParDo.of(new Filter2()))
```

```
.apply(new Transform1())
```

```
.apply(TextIO.write().to("gs://..."));
```

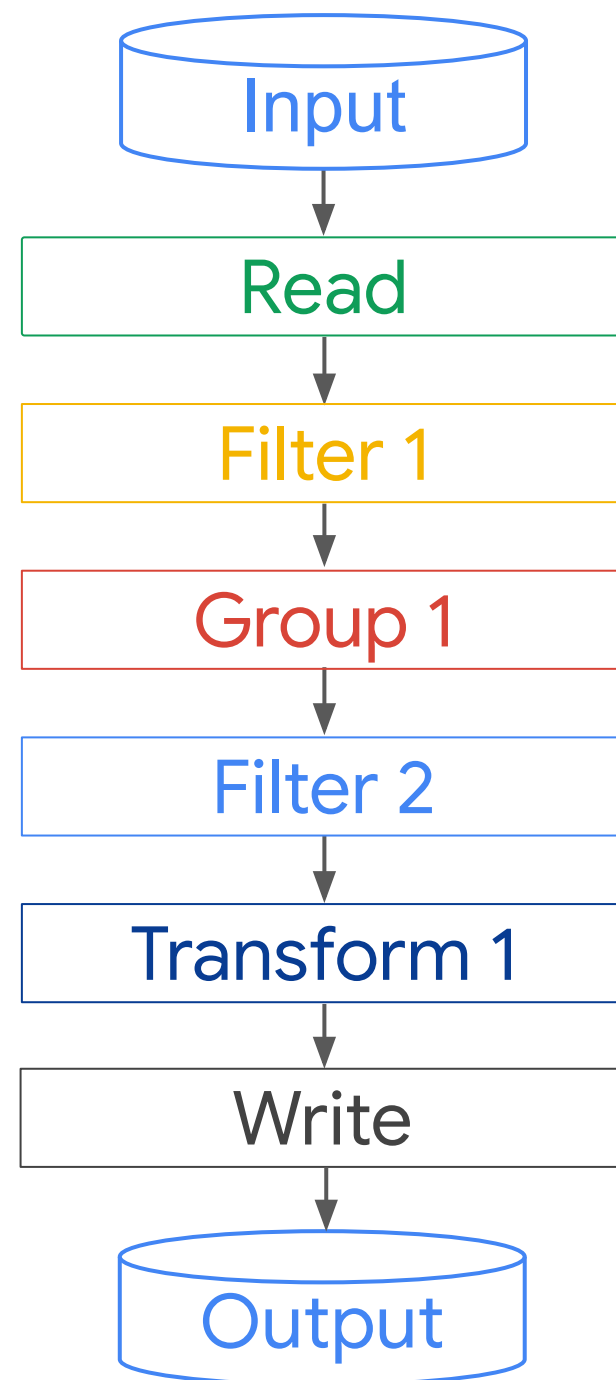
```
p.run();
```

Open-source API (Apache Beam) can be executed on Flink, Spark, etc. also

Parallel task (autoscaled by execution framework)

```
class Filter1 extends DoFn<...> {  
    public void  
    processElement(ProcessContext c) {  
        ... = c.element();  
        ...  
        c.output(...);  
    }  
}
```

Open-source API, Google infrastructure



```
Pipeline p = Pipeline.create();
```

p

```
.apply(TextIO.read().from("gs://..."))
```

```
.apply(ParDo.of(new Filter1()))
```

```
.apply(new Group1())
```

```
.apply(ParDo.of(new Filter2()))
```

```
.apply(new Transform1())
```

```
.apply(TextIO.write().to("gs://..."));
```


```
p.run();
```

Open-source API (Apache Beam) can be executed on Flink, Spark, etc. also

Parallel task (autoscaled by execution framework)

```
class Filter1 extends DoFn<...> {
    public void
    processElement(ProcessContext c) {
        ... = c.element();
        ...
        c.output(...);
    }
}
```

Open-source API, Google infrastructure



```
Pipeline p = Pipeline.create();
```

p

```
.apply(TextIO.read().from("gs://..."))
```

```
.apply(ParDo.of(new Filter1()))
```

```
.apply(new Group1())
```

```
.apply(ParDo.of(new Filter2()))
```

```
.apply(new Transform1())
```

```
.apply(TextIO.write().to("gs://..."));
```


```
p.run();
```

Open-source API (Apache Beam) can be executed on Flink, Spark, etc. also

Parallel task (autoscaled by execution framework)


```
class Filter1 extends DoFn<...> {  
    public void  
    processElement(ProcessContext c) {  
        ... = c.element();  
        ...  
        c.output(...);  
    }  
}
```


The code is the same between real-time and batch (Java)




```
p = beam.Pipeline()
(p
  | beam.io.ReadStringsFromPubSub('project/topic')
  | beam.WindowInto(SlidingWindows(60))
  | beam.Map(Transform)
  | beam.GroupByKey()
  | beam.FlatMap(Filter)
  | beam.io.WriteToBigQuery(table)
)
p.run()
```

The code is the same between real-time and batch (Java)



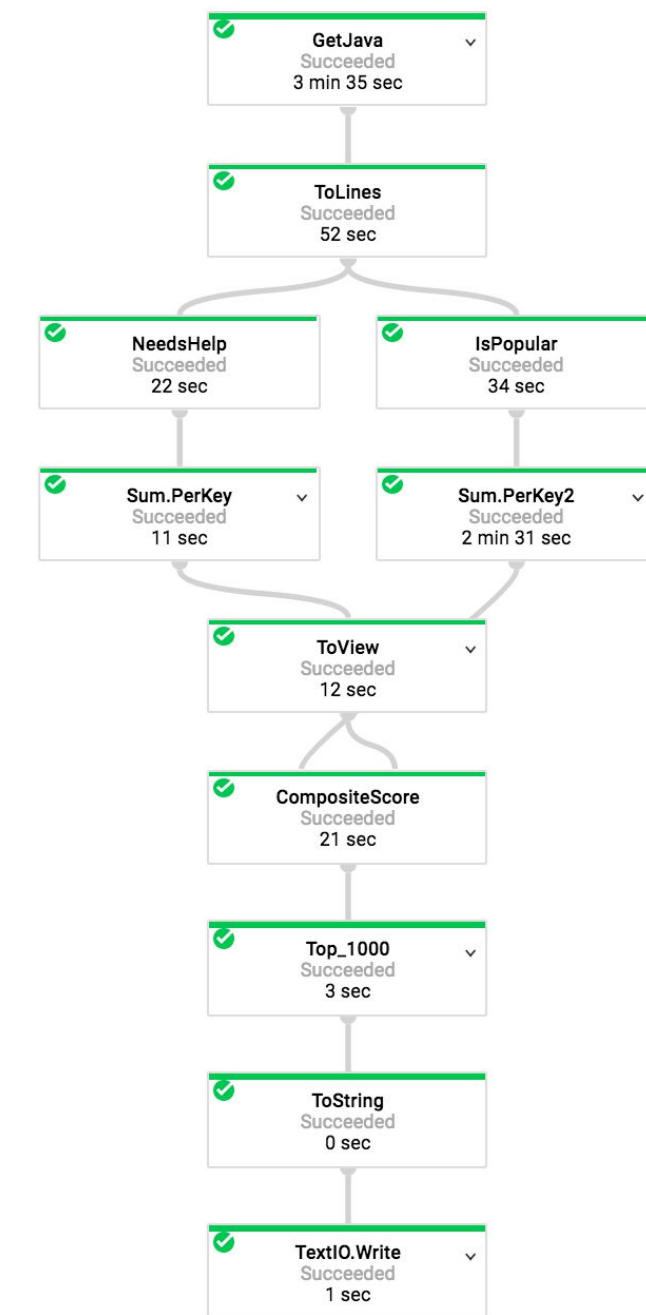
```
p = beam.Pipeline()
(p
  | beam.io.ReadStringsFromPubSub('project/topic')
  | beam.WindowInto(SlidingWindows(60))
  | beam.Map(Transform)
  | beam.GroupByKey()
  | beam.FlatMap(Filter)
  | beam.io.WriteToBigQuery(table)
)
p.run()
```

The code is the same between real-time and batch (Java)



```
p = beam.Pipeline()  
(p  
| beam.io.ReadStringsFromPubSub('project/topic')  
| beam.WindowInto(SlidingWindows(60))  
| beam.Map(Transform)  
| beam.GroupByKey()  
| beam.FlatMap(Filter)  
| beam.io.WriteToBigQuery(table)  
)  
p.run()
```

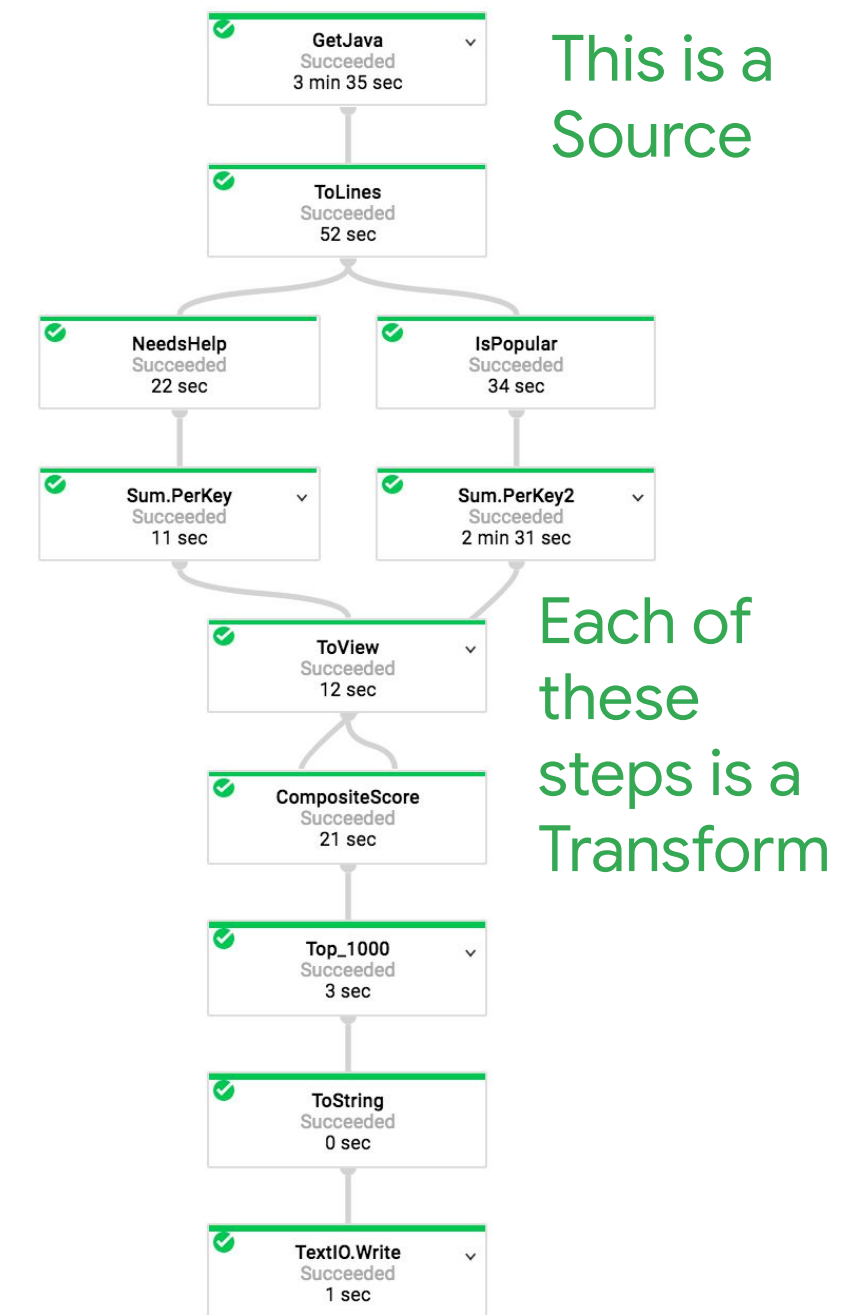
Dataflow terms and concepts



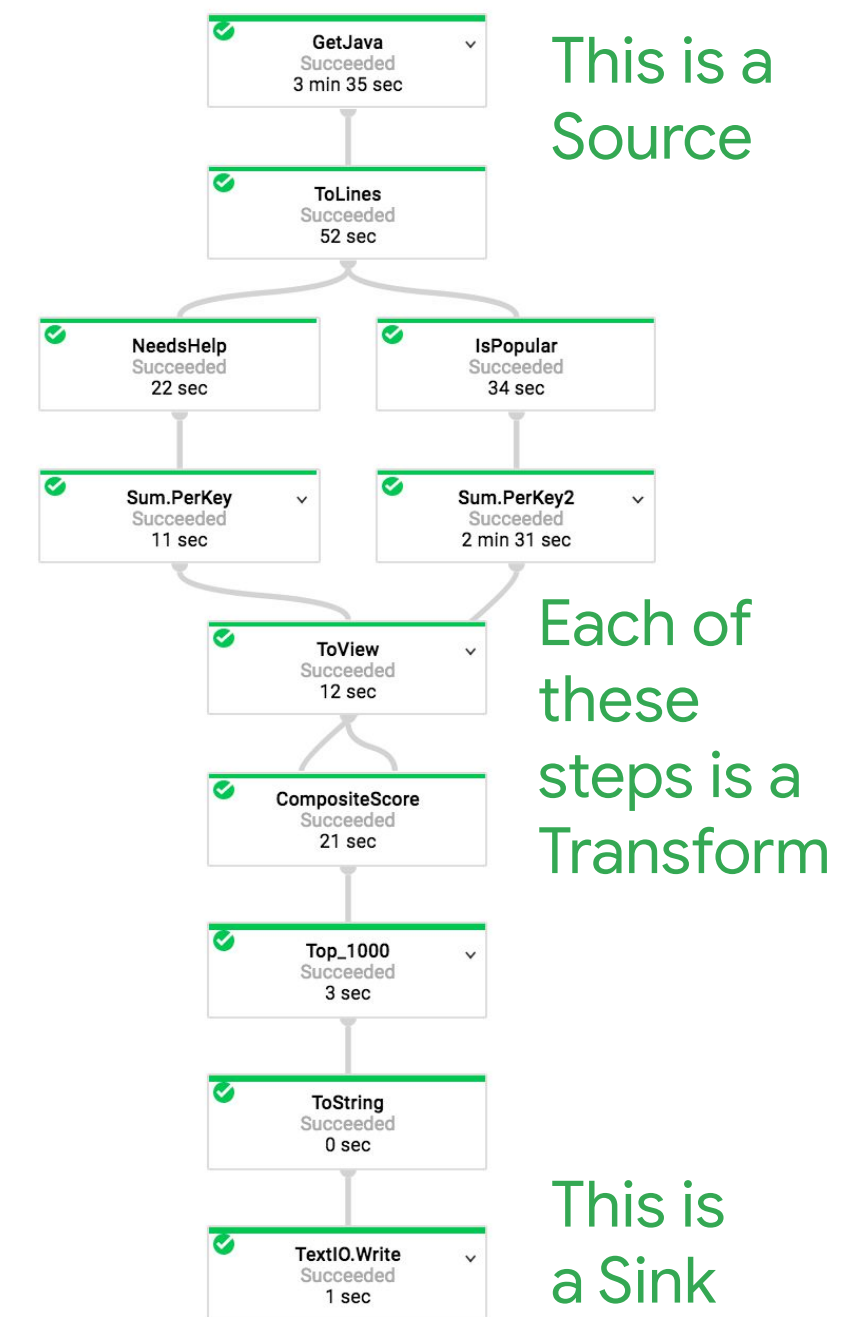
Dataflow terms and concepts



Dataflow terms and concepts




Dataflow terms and concepts



Dataflow terms and concepts

The Pipeline is executed on the cloud by a Runner; each step is elastically scaled

```
def packageHelp(record, keyword):  
    count=0  
    package_name=''  
    if record is not None:  
        lines=record.split('\n')  
        for line in lines:  
            if line.startswith(keyword):  
                package_name=line  
                if 'FIXME' in line or 'TODO' in line:  
                    count+=1  
        packages = (getPackages(package_name))  
        for p in packages:  
            yield (p, count)
```



Cloud Dataflow

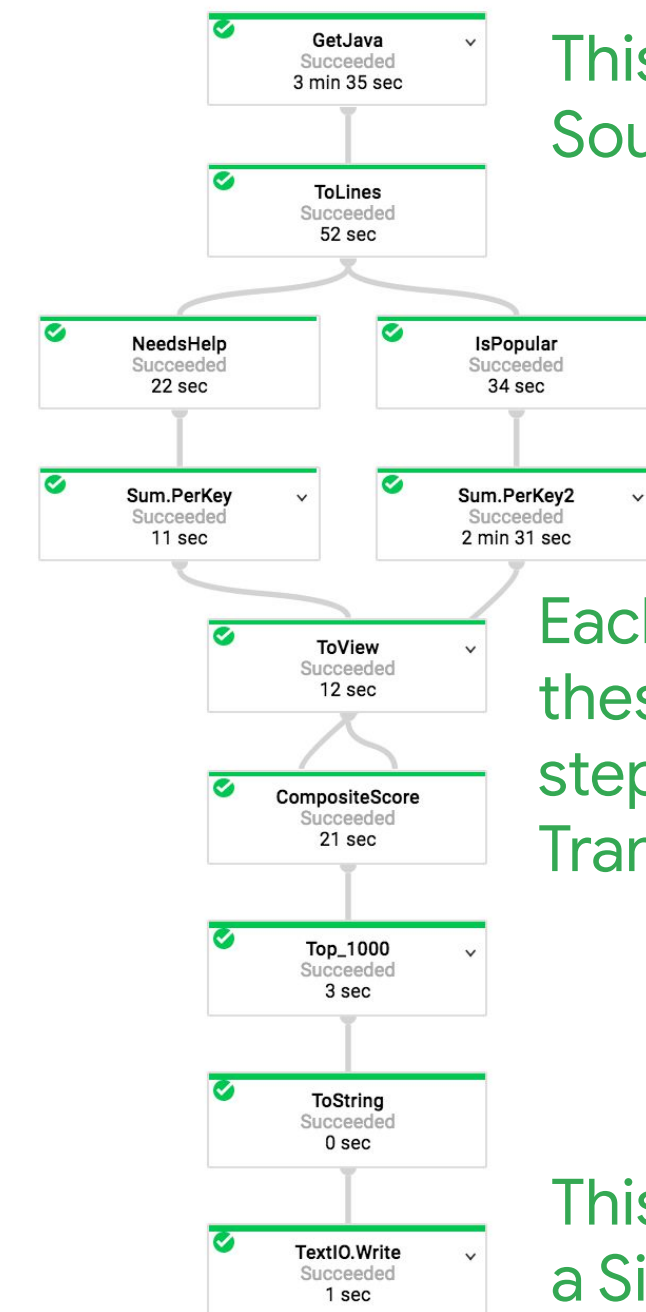


BigQuery



Cloud Storage

Together,
they form
a Pipeline



This is a
Source

Each of
these
steps is a
Transform

This is
a Sink

Dataflow terms and concepts


The Pipeline is executed on the cloud by a Runner; each step is elastically scaled

```
def packageHelp(record, keyword):  
    count=0  
    package_name=''  
    if record is not None:  
        lines=record.split('\n')  
        for line in lines:  
            if line.startswith(keyword):  
                package_name=line  
                if 'FIXME' in line or 'TODO' in line:  
                    count+=1  
        packages = (getPackages(package_name))  
        for p in packages:  
            yield (p, count)
```

Each Transform of the Pipeline is applied on a PCollection; the result of apply() is another PCollection



BigQuery

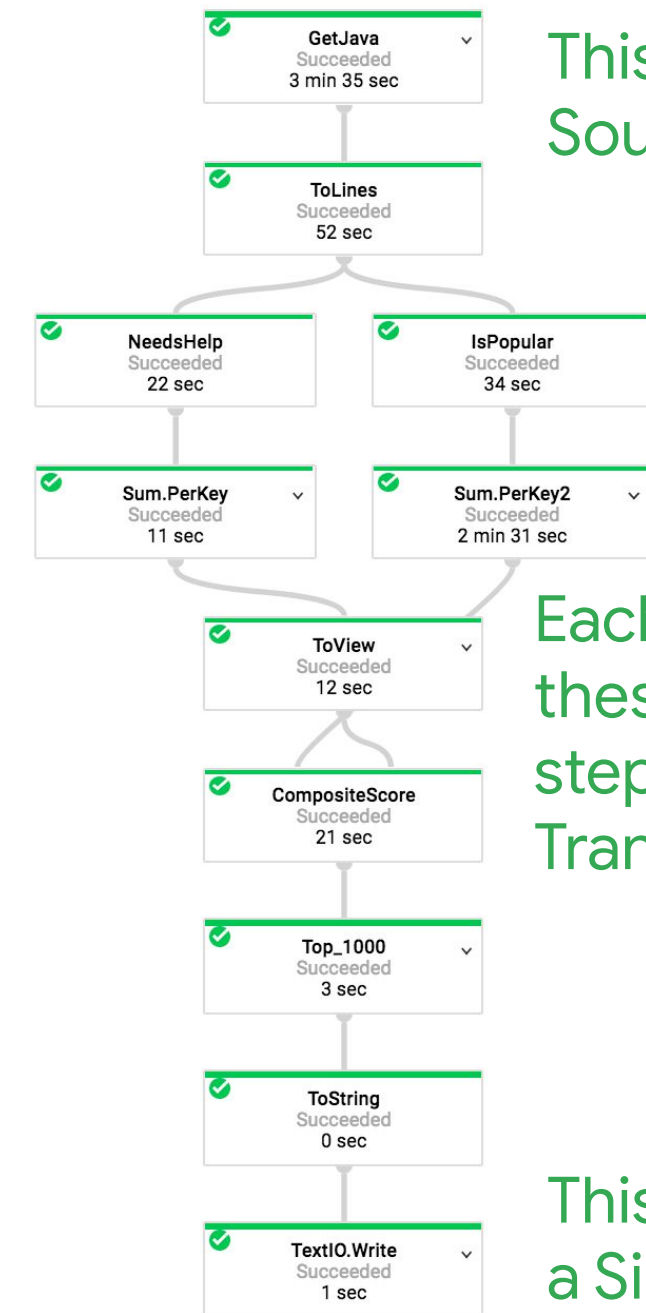


Cloud Dataflow



Cloud Storage

Together, they form a Pipeline



This is a Source

Each of these steps is a Transform

This is a Sink

A Pipeline is a directed graph of steps

Read in data, transform it, write out

Can branch, merge, use if-then statements, etc.

Pythonic syntax

```
import apache_beam as beam
if __name__ == '__main__':

    # create a pipeline parameterized by commandline flags
    p = beam.Pipeline(argv=sys.argv)

    (p
     | 'Read' >> beam.io.ReadFromText('gs://...') # read input
     | 'CountWords' >> beam.FlatMap(lambda line: count_words(line))
     | 'Write' >> beam.io.WriteToText('gs://...') # write output
    )

    p.run() # run the pipeline
```

A Pipeline is a directed graph of steps

Read in data, transform it, write out

Can branch, merge, use if-then statements, etc.

```
import org.apache.beam.sdk.Pipeline; // etc.

public static void main(String[] args) {
    // Create a pipeline parameterized by commandline flags.
    Pipeline p = Pipeline.create(PipelineOptionsFactory.fromArgs(args));

    p.apply(TextIO.read().from("gs://...")) // Read input.
      .apply(new CountWords())              // Do some processing.
      .apply(TextIO.write().to("gs://...")); // Write output.

    // Run the pipeline.
    p.run();
}
```

Python API conceptually similar

Read in data, transform it, write out

Pythonic syntax

```
import apache_beam as beam

if __name__ == '__main__':
    # create a pipeline parameterized by commandline flags
    p = beam.Pipeline(argv=sys.argv)

    (p
     | beam.io.ReadFromText('gs://...') # read input
     | beam.FlatMap(lambda line: count_words(line)) # do some processing
     | beam.io.WriteToText('gs://...') # write output
    )

    p.run() # run the pipeline
```

Apply Transform to PCollection

Data in a pipeline are represented by PCollection

Supports parallel processing

Not an in-memory collection; can be unbounded

```
lines = p | ...
```

Apply Transform to PCollection; returns PCollection

```
sizes = lines | 'Length' >> beam.Map(lambda line: len(line) )
```

Apply Transform to PCollection

Data in a pipeline are represented by PCollection

Supports parallel processing

Not an in-memory collection; can be unbounded

```
PCollection<String> lines = p.apply(...) //
```

Apply Transform to PCollection

Apply Transform to PCollection; returns PCollection

```
PCollection<Integer> sizes =  
    lines.apply("Length", ParDo.of(new DoFn<String, Integer>() {  
        @ProcessElement  
        public void processElement(ProcessContext c) throws  
Exception {  
            String line =  
c.element();  
            c.output(line.length());  
        }  
    })))
```

Apply Transform to PCollection (Python)

Data in a pipeline are represented by PCollection

Supports parallel processing

Not an in-memory collection; can be unbounded

```
lines = p | ...
```

Apply Transform to PCollection; returns PCollection

```
sizes = lines | 'Length' >> beam.Map(lambda line: len(line) )
```


Ingesting data into a pipeline (Python)

Read data from file system, GCS or BigQuery

Text formats return String

```
lines = beam.io.ReadFromText('gs://.../input-*.csv.gz')
```

BigQuery returns a TableRow

```
rows = beam.io.Read(beam.io.BigQuerySource(query='SELECT x, y, z ' \
                                             'FROM [project:dataset.tablename]', project='PROJECT'))
```

Ingesting data into a pipeline (Java)

Read data from file system, GCS, BigQuery, Pub/Sub

Text formats return String

```
PCollection<String> lines = p.apply(TextIO.read().from("gs://.../input-*.csv.gz"));
```

```
PCollection<String> lines = p.apply(PubsubIO.readStrings().fromTopic(topic));
```

BigQuery returns a TableRow

```
String javaQuery = "SELECT x, y, z FROM [project:dataset.tablename]";  
PCollection<TableRow> javaContent = p.apply(BigQueryIO.read().fromQuery(javaQuery))
```

Can write data out to same formats (Python)

Write data to file system, GCS or BigQuery

```
beam.io.WriteToText(file_path_prefix='/data/output', file_name_suffix='.txt')
```

Can prevent sharding of output (do only if it is small)

```
beam.io.WriteToText(file_path_prefix='/data/output',  
file_name_suffix='.txt', num_shards = 1)
```

The output must be a PCollection of Strings before writing out

Can write data out to same formats (Java)

Write data to file system, GCS, BigQuery, Pub/Sub

```
lines.apply(TextIO.write().to("/data/output").withSuffix(".txt"))
```

Can prevent sharding of output (do only if it is small)

```
.apply(TextIO.write().to("/data/output").withSuffix(".csv").withoutSharding()  
)
```

May have to transform PCollection<Integer>, etc. to PCollection<String> before writing out

Executing pipeline (Java)

Simply running main() runs pipeline locally

```
java -classpath ... com...
```

```
mvn compile -e exec:java -Dexec.mainClass=$MAIN
```

To run on cloud, submit job to Dataflow

```
mvn compile -e exec:java \  
    -Dexec.mainClass=$MAIN \  
    -Dexec.args="--project=$PROJECT \  
    --stagingLocation=gs://$BUCKET/staging/ \  
    --tempLocation=gs://$BUCKET/staging/ \  
    --runner=DataflowRunner"
```

Executing pipeline (Python)

Simply running main() runs pipeline locally

```
python ./grep.py
```

To run on cloud, specify cloud parameters, and submit the job to Dataflow

```
python ./grep.py \  
  --project=$PROJECT \  
  --job_name=myjob \  
  --staging_location=gs://$BUCKET/staging/ \  
  --temp_location=gs://$BUCKET/staging/ \  
  --runner=DataflowRunner
```

Executing pipeline (Python)

Simply running main() runs pipeline locally

```
python ./grep.py
```

To run on cloud, specify cloud parameters

```
python ./grep.py \  
  --project=$PROJECT \  
  --job_name=myjob \  
  --staging_location=gs://$BUCKET/staging/ \  
  --temp_location=gs://$BUCKET/staging/ \  
  --runner=DataflowRunner
```

Lab

A simple Dataflow pipeline

Carl Osipov

Lab: A simple Dataflow pipeline (Java/Python)

In this lab, you will learn how to:

1. Set up a Dataflow project
2. Write a simple pipeline
3. Execute the pipeline on the local machine
4. Execute the pipeline on the cloud

```
p //  
    .apply("GetJava", TextIO.  
    .apply("Grep", ParDo.of(  
        @ProcessElement  
        public void proces  
            String li  
            if (line.c  
                c  
            }  
        }  
    }  
    ))) //  
    .apply(TextIO.write().to(
```



MapReduce approach splits Big Data so that each compute node processes data local to it

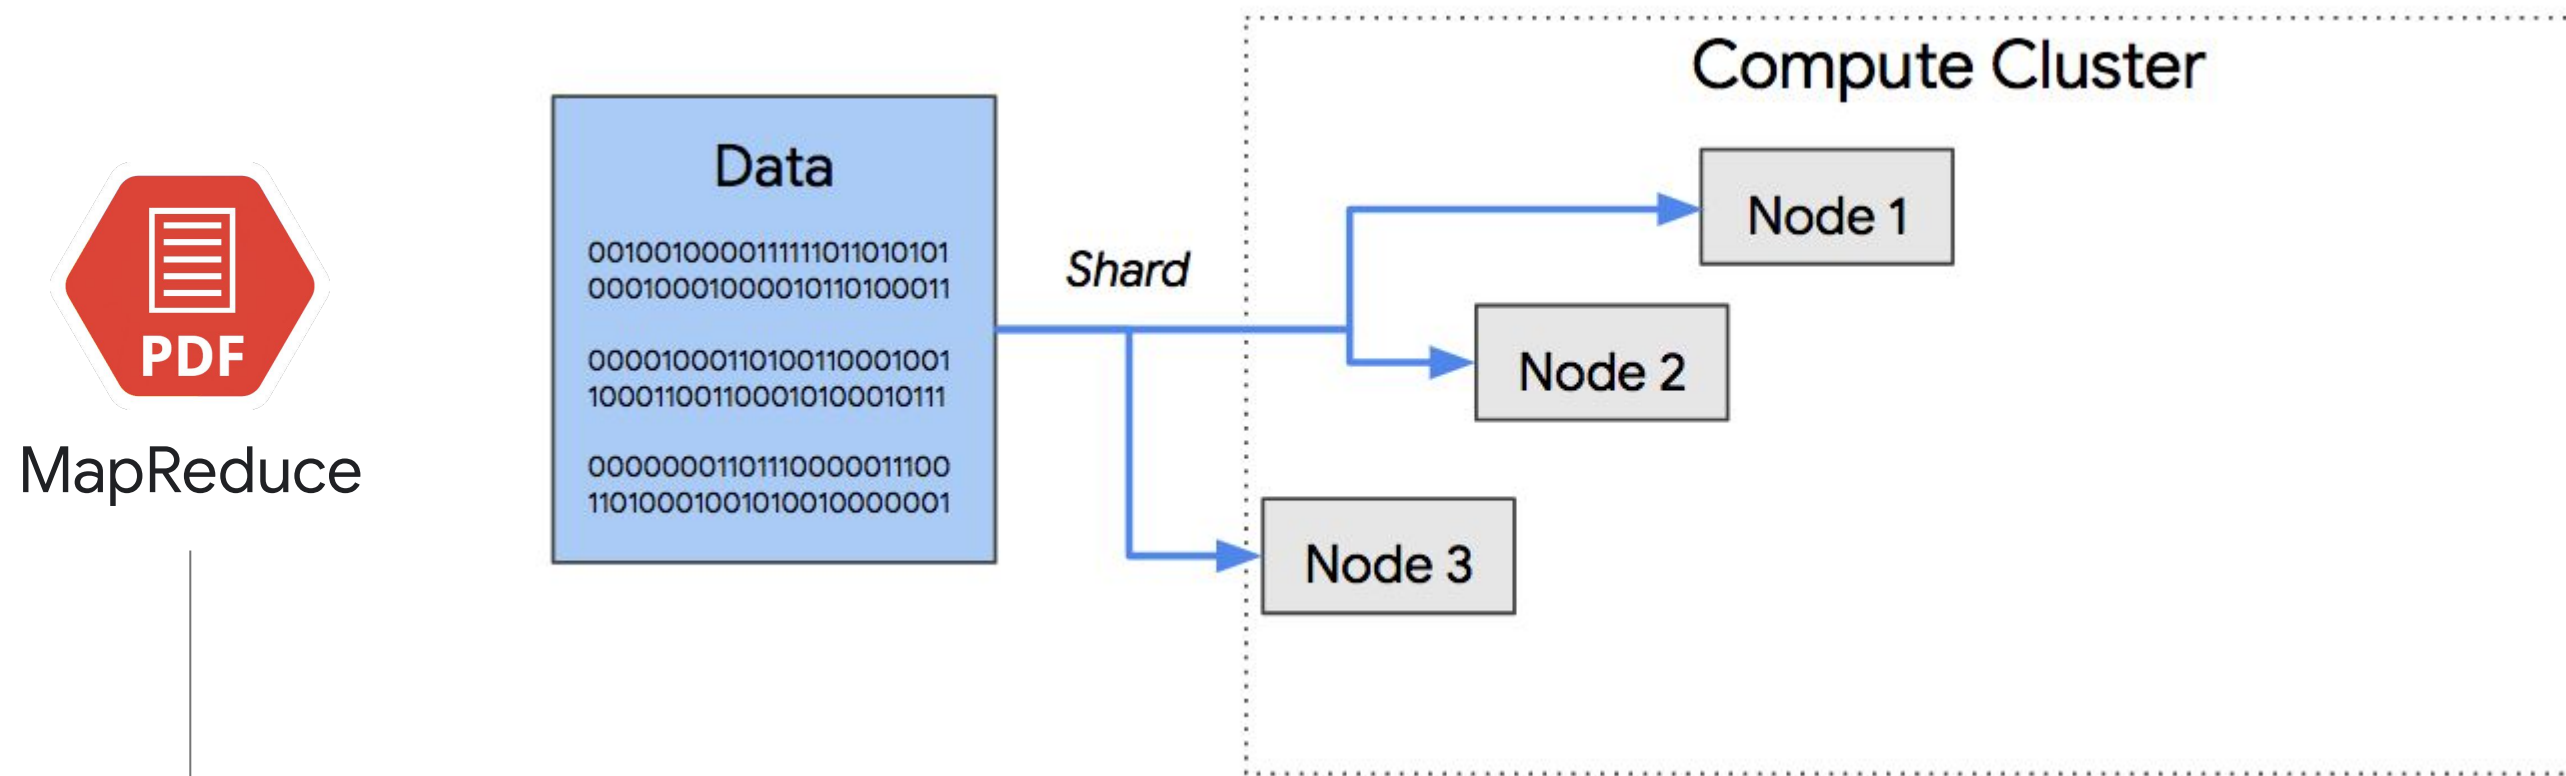
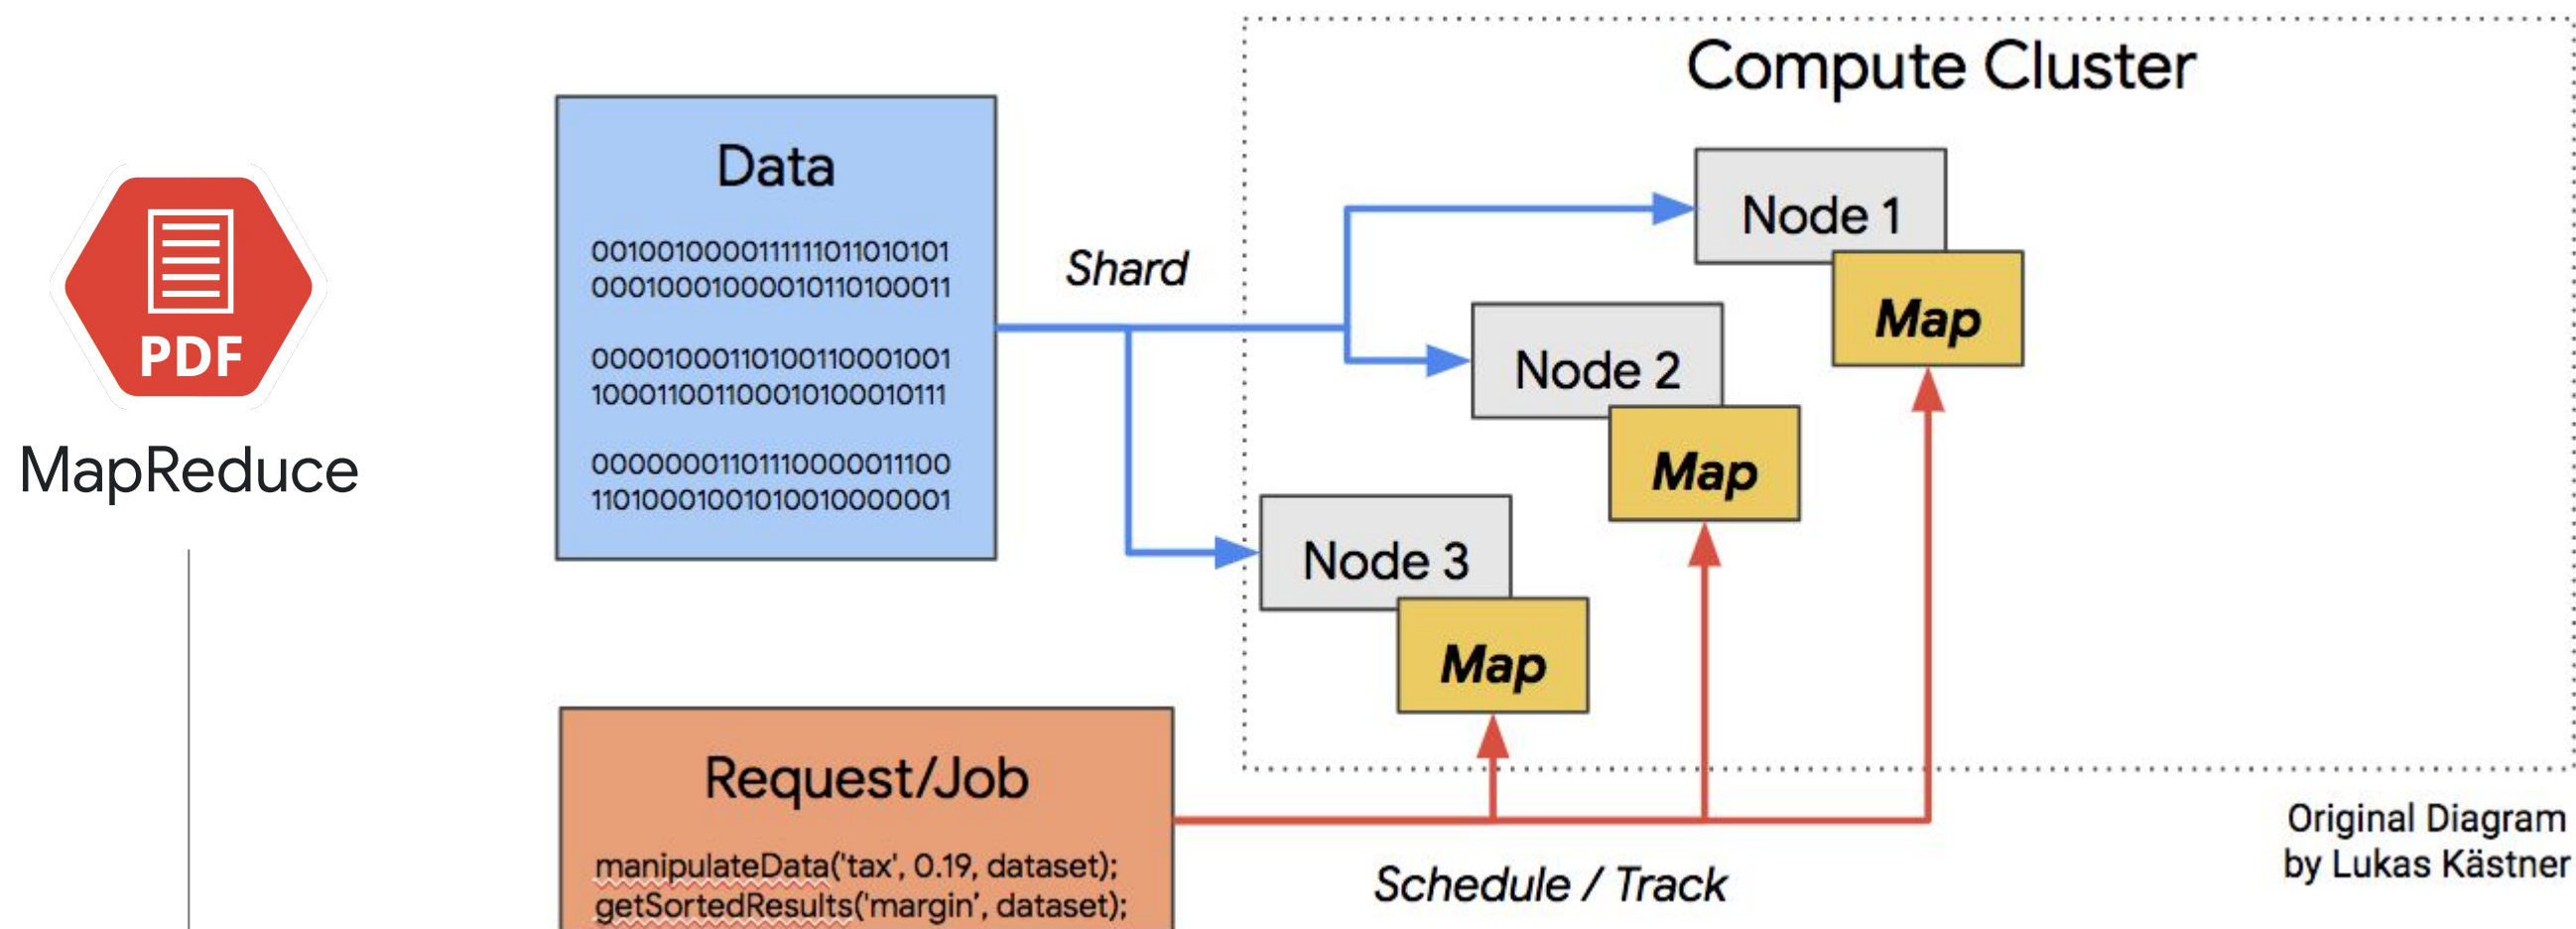


Diagram by Lukas Kästner

MapReduce approach splits Big Data so that each compute node processes data local to it



2002

2004

2006

2008

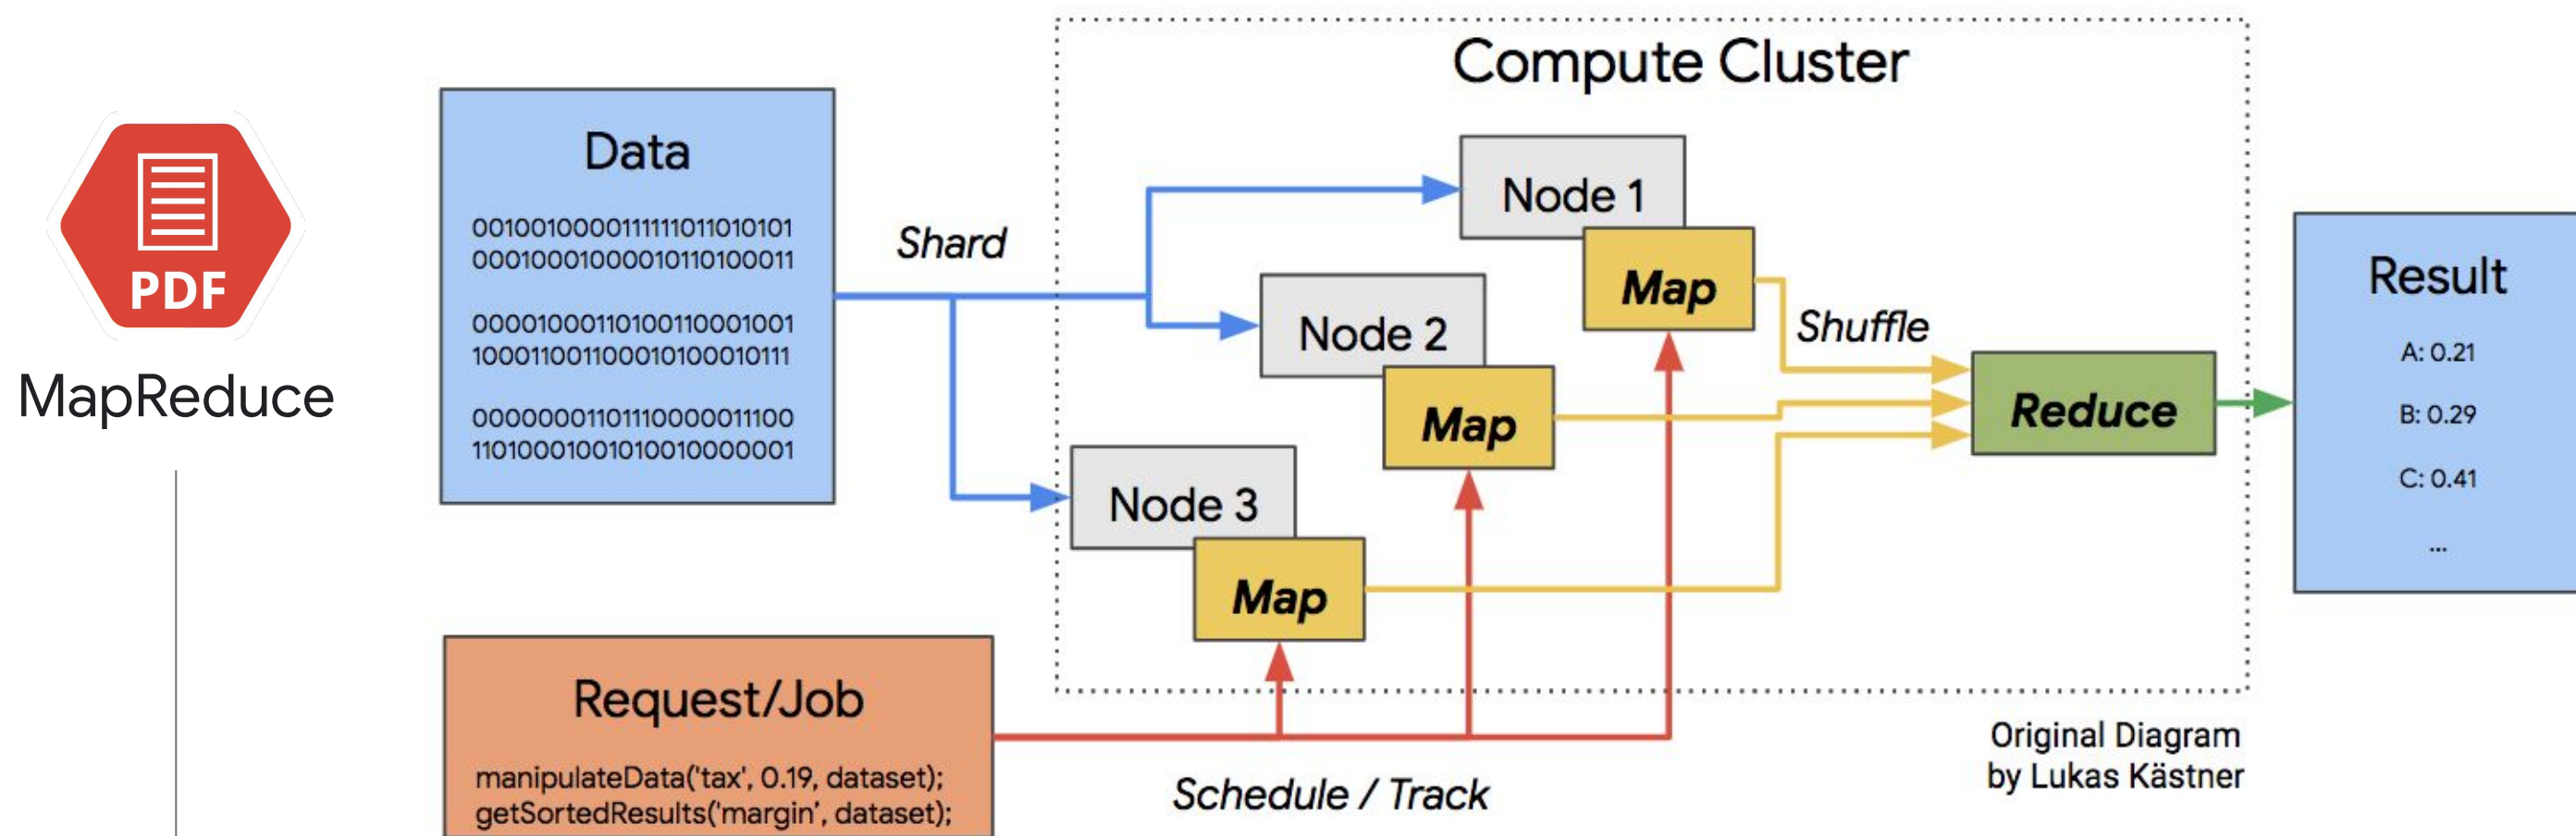
2010

2012

2014

2016

MapReduce approach splits Big Data so that each compute node processes data local to it



2002

2004

2006

2008

2010

2012

2014

2016

ParDo allows for parallel processing

ParDo acts on one item at a time (like a Map in MapReduce)

- Multiple instances of class on many machines

- Should not contain any state

Useful for:

- Filtering (choosing which inputs to emit)

- Extracting parts of an input (e.g., fields of TableRow)

- Converting one Java type to another

- Calculating values from different parts of inputs

Python: Map vs. FlatMap

Use Map for 1:1 relationship between input & output

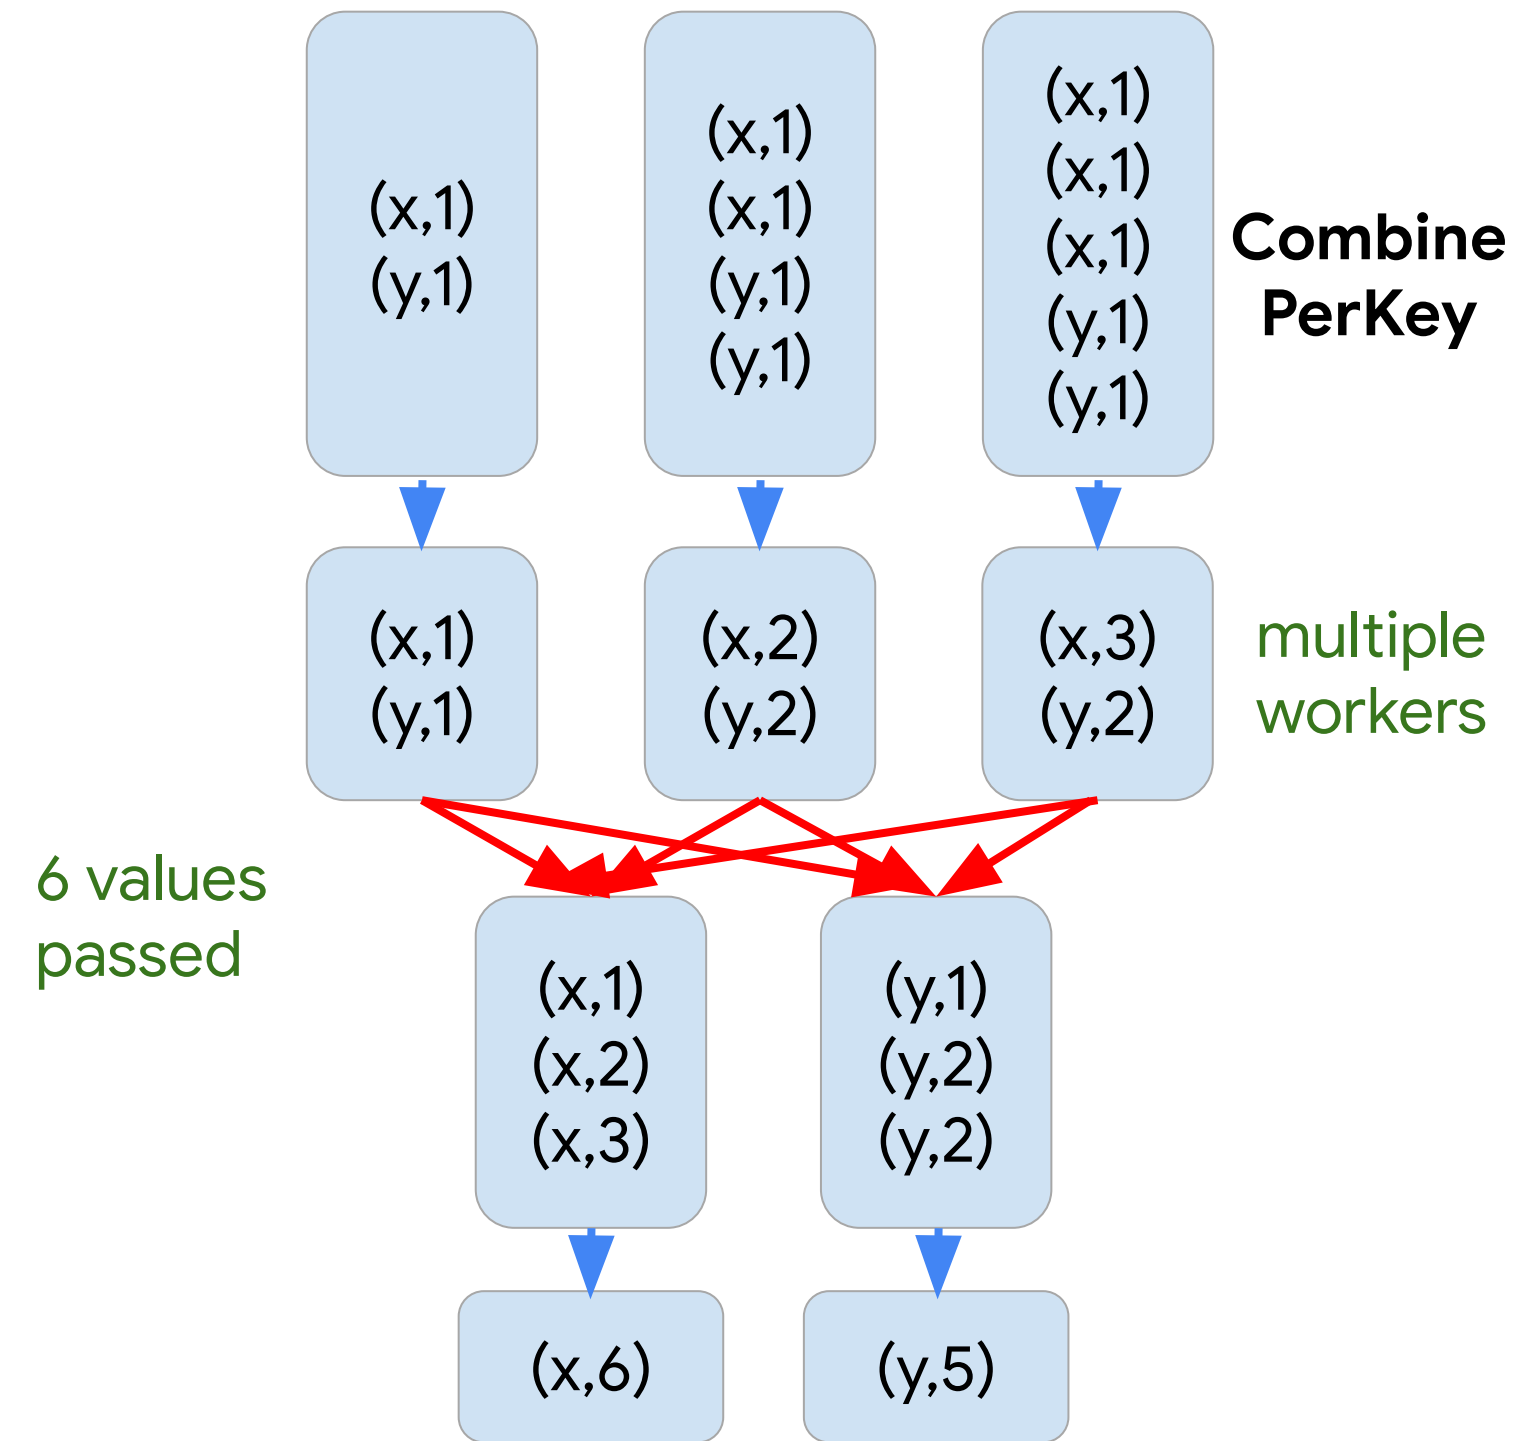
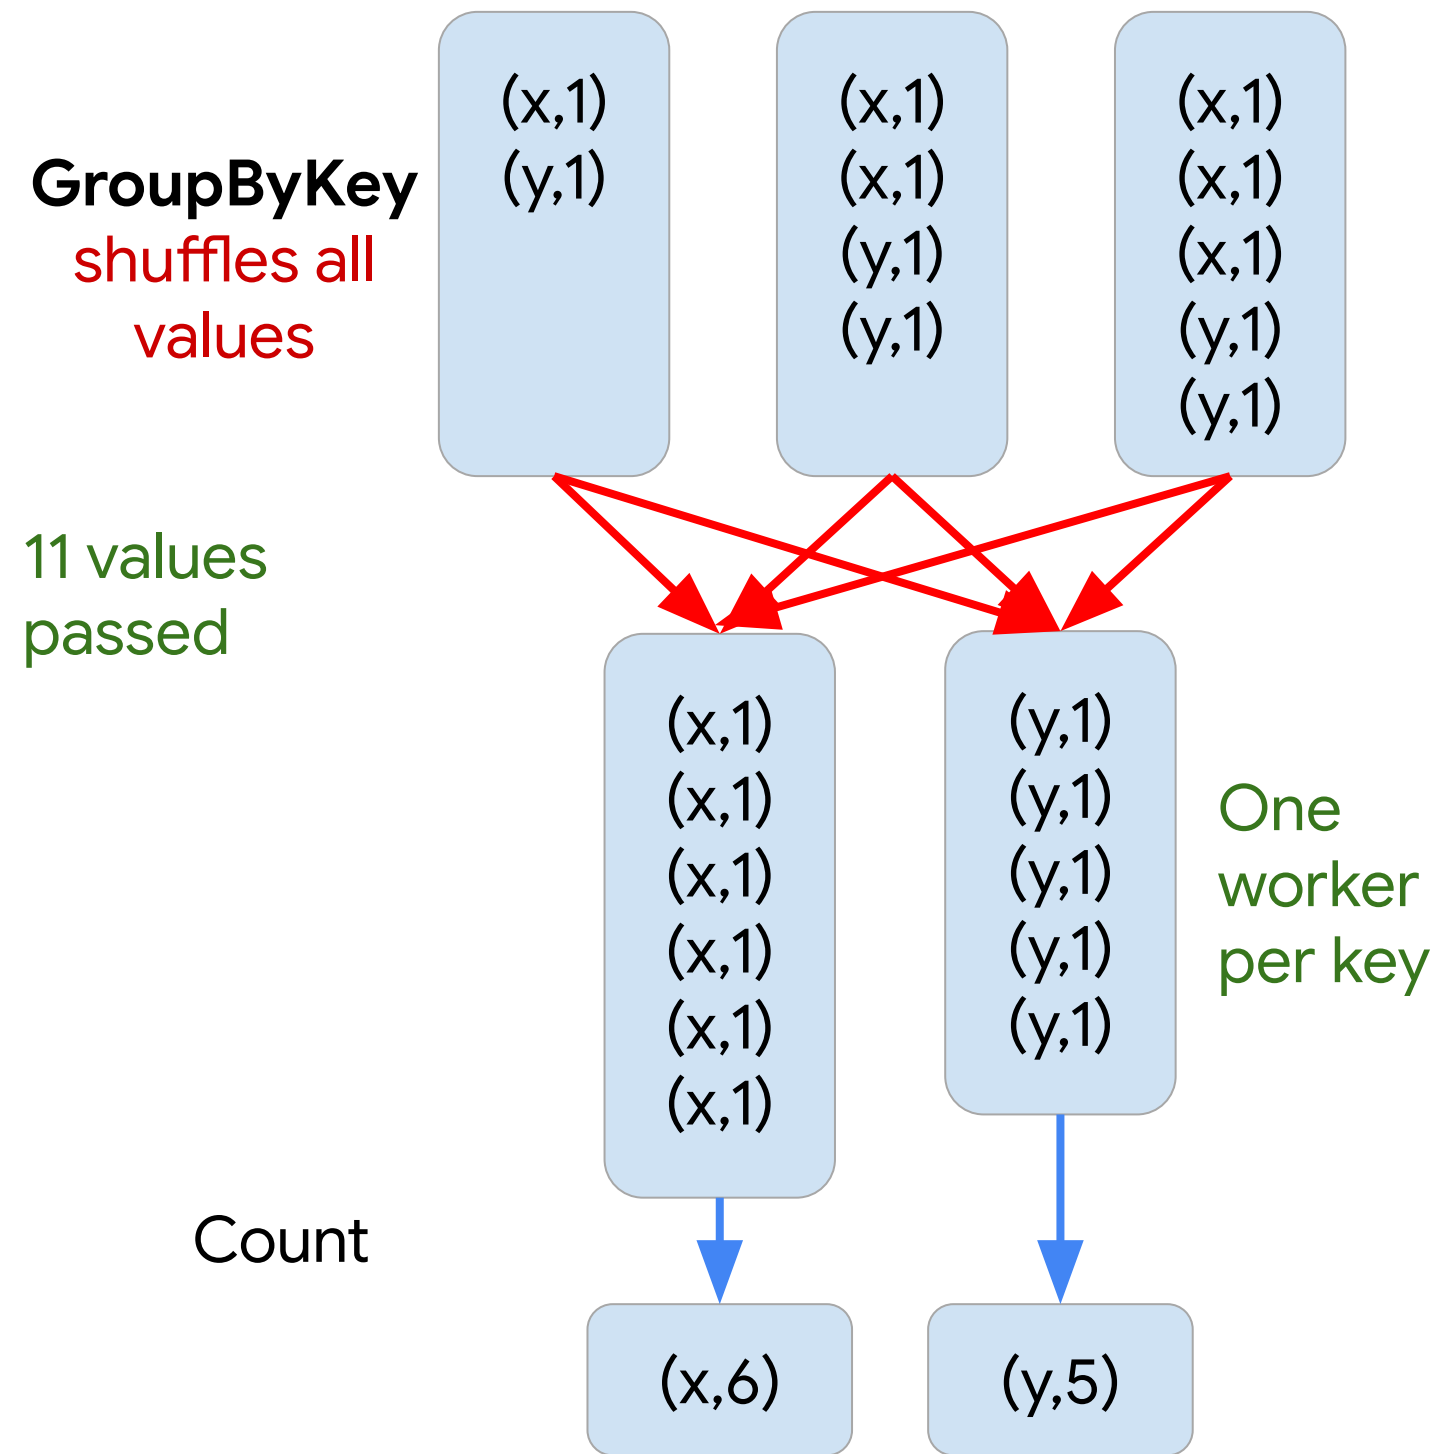
```
'WordLengths' >> beam.Map( lambda word: (word, len(word)) )
```

FlatMap for non 1:1 relationships, usually with generator

```
def vowels(word):  
    for ch in word:  
        if ch in ['a','e','i','o','u']:  
            yield ch  
  
'WordVowels' >> beam.FlatMap( lambda word: vowels(word) )
```

Java: Use apply(ParDo) for both cases

Reduce with GroupByKey or Combine.PerKey



GroupBy operation is akin to shuffle

In Dataflow, shuffle explicitly with a GroupByKey

Create a Key-Value pair in a ParDo

Then group by the key

```
cityAndZipcodes = p
    | beam.Map(lambda address: (address[1], address[3]) )
                        CITY      ZIPCODE
    | beam.GroupByKey()
```


Combine.PerKey lets you aggregate

Can be applied to a PCollection of values:

```
totalAmount = salesAmounts | Combine.globally(sum)
```

And also to a grouped Key-Value pair:

```
totalSalesPerPerson = salesRecords | Combine.perKey(sum)
```

Many built-in functions: Sum, Mean, etc.

Lab

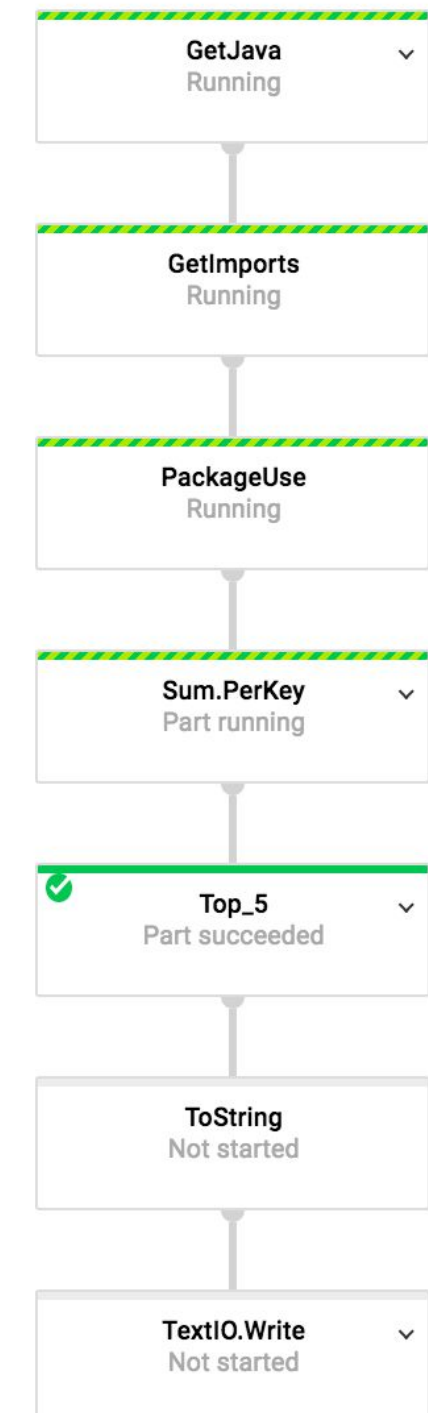
MapReduce in Dataflow

Carl Osipov

Lab: MapReduce in Dataflow (Java/Python)

In this lab, you will learn how to:

- Specify and use command-line options
- Carry out Map and Reduce operations



Use windows to specify how to aggregate unbounded collections

```
| 'Window' >> beam.WindowInto(  
    SlidingWindows(size = 120,  
                    period = 30))
```

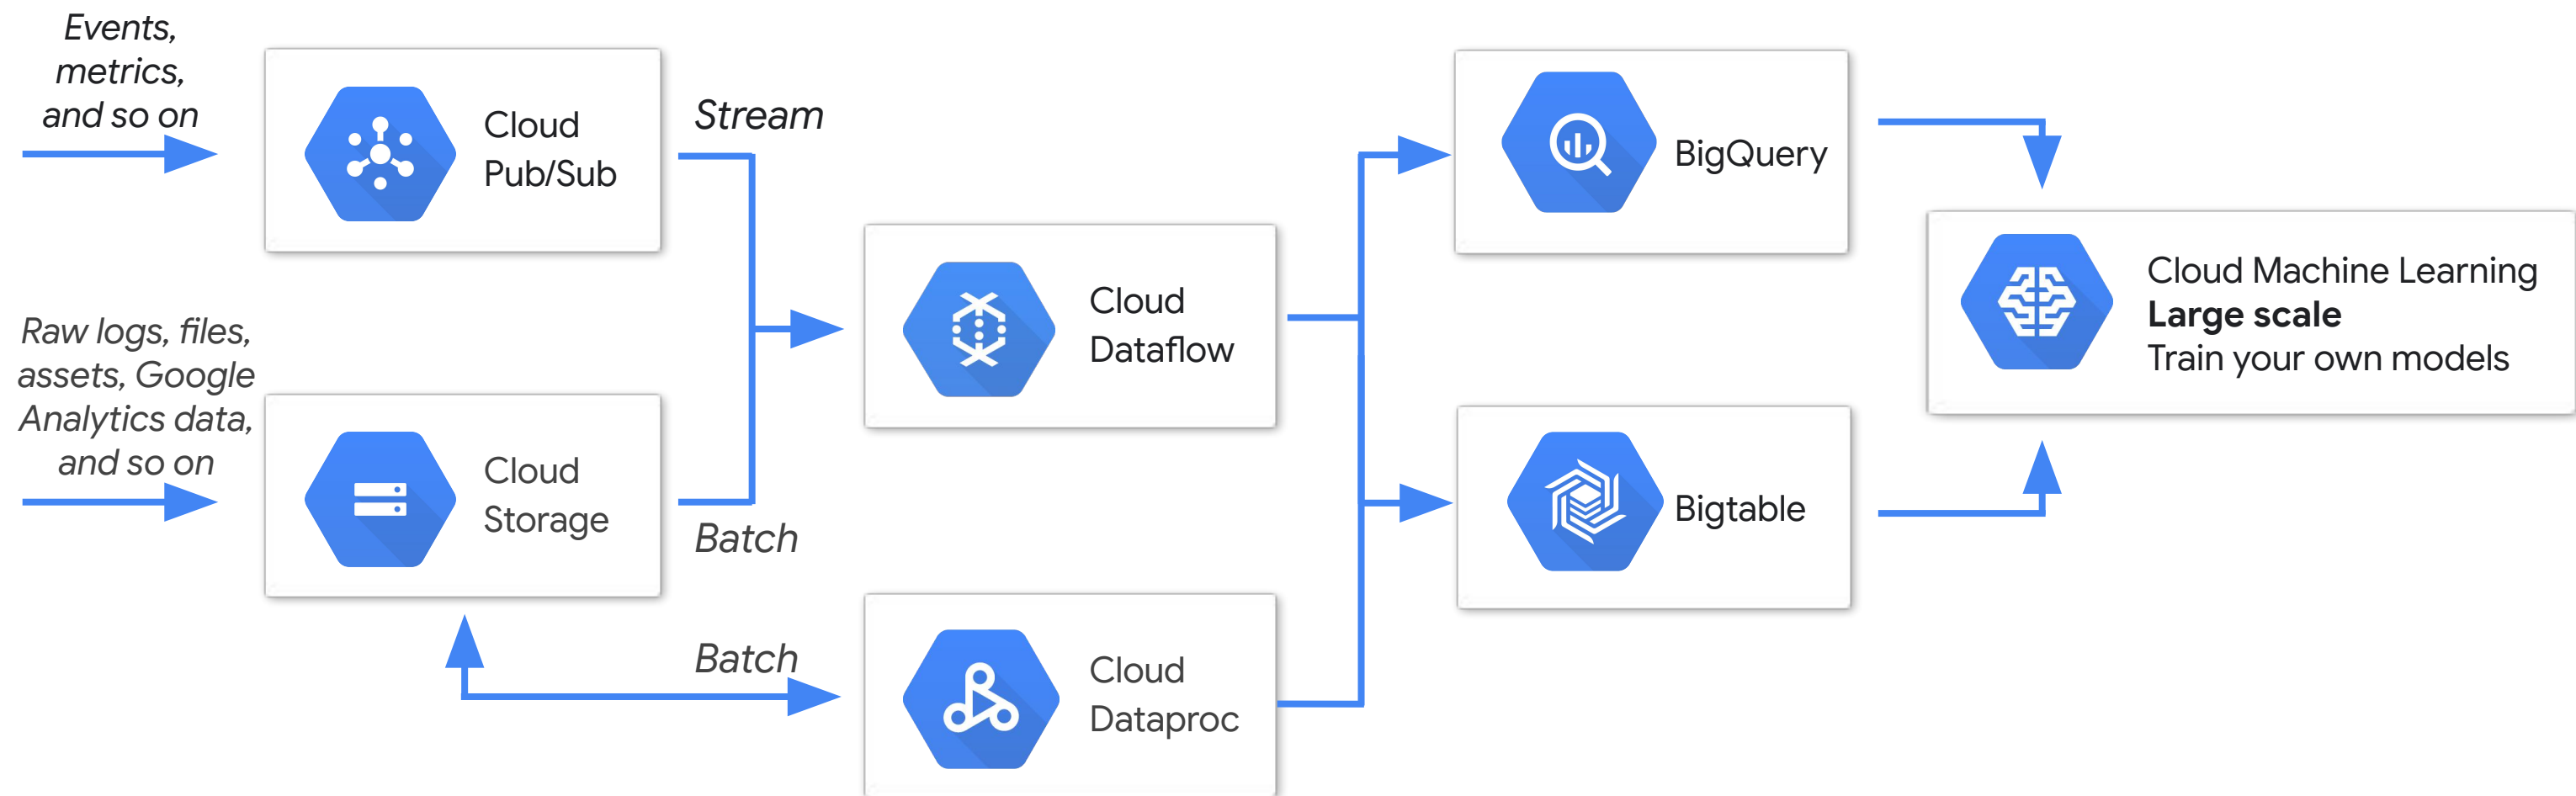
Subsequent Groups, aggregations, etc.
are computed only within time window

Use windows to specify how to aggregate unbounded collections

```
.apply("window", Window.into(SlidingWindows//  
    .of(Duration.standardMinutes(2))//  
    .every(Duration.standardSeconds(30)))) //
```

Subsequent Groups, aggregations, etc.
are computed only within time window

In the Google Cloud reference architecture, preprocessing is repeatable between training and prediction because of Dataflow






Preprocessing with Cloud Dataprep


Carl Osipov

Exploring and knowing your data is essential




Explore and Visualize
Common Values

Exploring and knowing your data is essential




Analyze Key
Statistics
(min, max, avg, std
dev)

Exploring and knowing your data is essential



Explore the
distributions

Exploring and knowing your data is essential



Collaborate with
Domain Experts

There are two general approaches to designing
preprocessing

There are two general approaches to designing preprocessing

1. Explore in Cloud Datalab
2. Write code in BigQuery / Dataflow / TensorFlow to transform data



Cloud
Dataflow






Cloud
Datalab





BigQuery

There are two general approaches to designing preprocessing

1. Explore in Cloud Datalab
2. Write code in BigQuery / Dataflow / TensorFlow to transform data






1. Explore in Cloud Dataprep
2. Design Recipe in UI to Preprocess Data
3. Apply generated Dataflow transformations to all data
4. Reuse Dataflow transformation in real-time pipeline





There are two general approaches to designing preprocessing

1. Explore in Cloud Datalab
2. Write code in BigQuery / Dataflow / TensorFlow to transform data




1. Explore in Cloud Dataprep
2. Design Recipe in UI to Preprocess Data
3. Apply generated Dataflow transformations to all data
4. Reuse Dataflow transformation in real-time pipeline



Example of exploring in Datalab: is there something wrong?

```
ax = sns.regplot(x="trip_distance", y="fare_amount", ci=None, truncate=True, data=trips)
```



Scatterplots and
in-memory pandas
dataframes don't scale


... is there a better
way?

Best to aggregate in BigQuery and plot in Datalab

```
query="""
SELECT
  departure_delay,
  COUNT(1) AS num_flights,
  APPROX_QUANTILES(arrival_delay, 10) AS arrival_delay_deciles
FROM
  `bigquery-samples.airline_ontime_data.flights`
GROUP BY
  departure_delay
HAVING
  num_flights > 100
ORDER BY
  departure_delay ASC
"""

import google.datalab.bigquery as bq
df = bq.Query(query).execute().result().to_dataframe()
df.head()
```

```
without_extremes = df.drop(['0%', '100%'], 1)
without_extremes.plot(x='departure_delay', xlim=(-30,50), ylim=(-50,50));
```

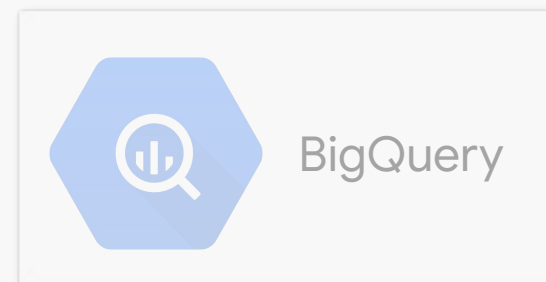




Write DataFlow code to do any transformations





Using Cloud Dataprep for exploring and preprocessing data

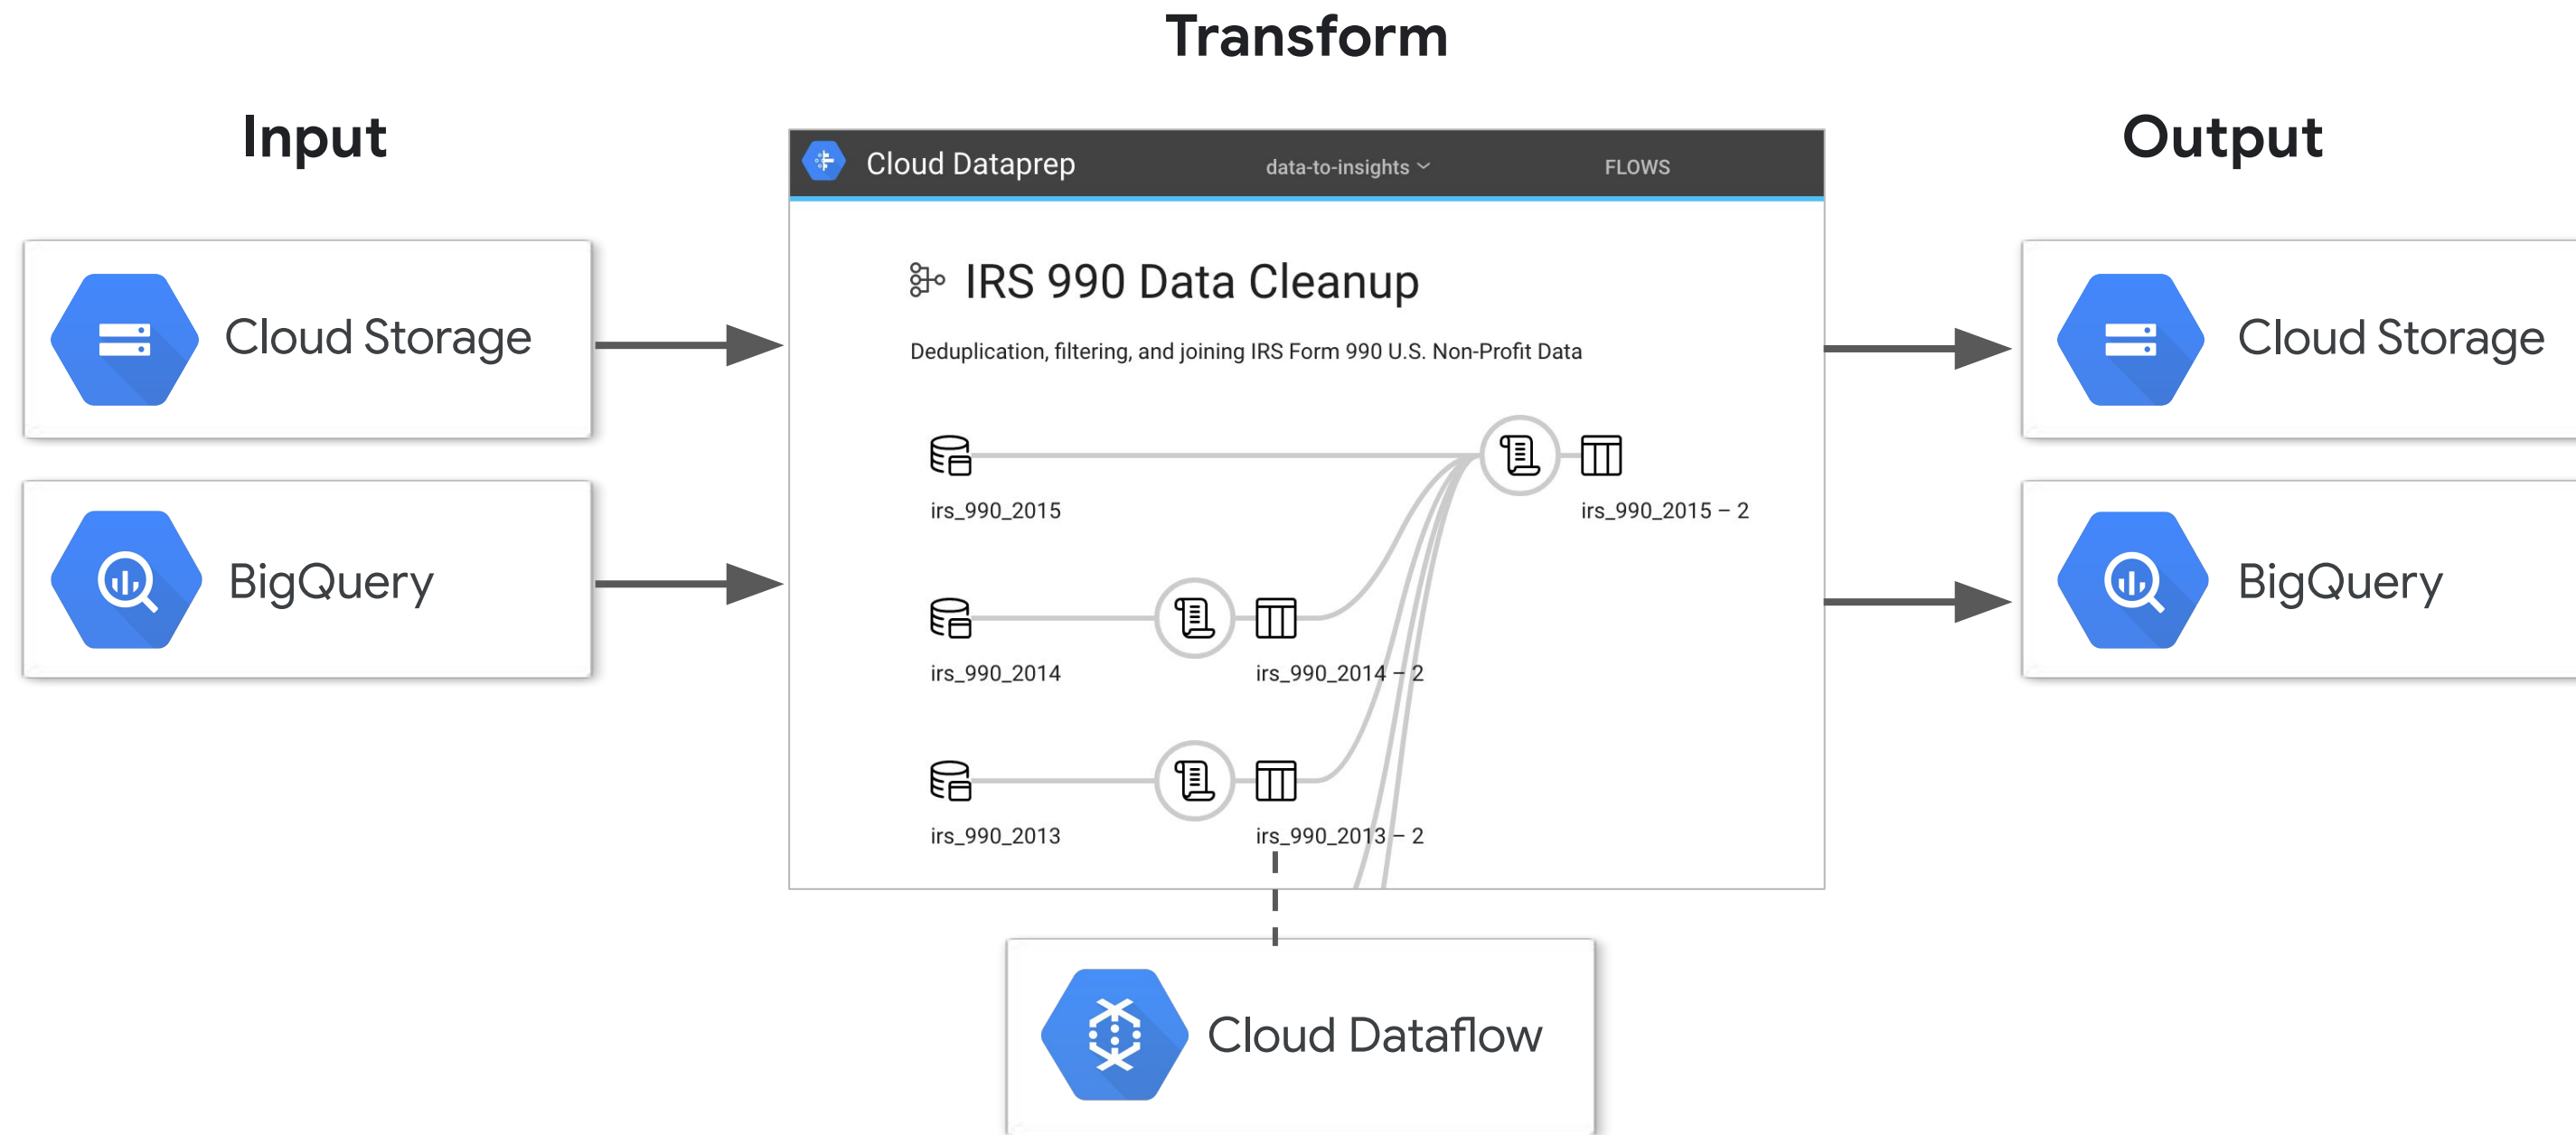
1. Explore in Cloud Datalab
2. Write code in BigQuery / Dataflow / TensorFlow to transform data



1. Explore in Cloud Dataprep
2. Design Recipe in UI to Preprocess Data
3. Apply generated Dataflow transformations to all data
4. Reuse Dataflow transformation in real-time pipeline

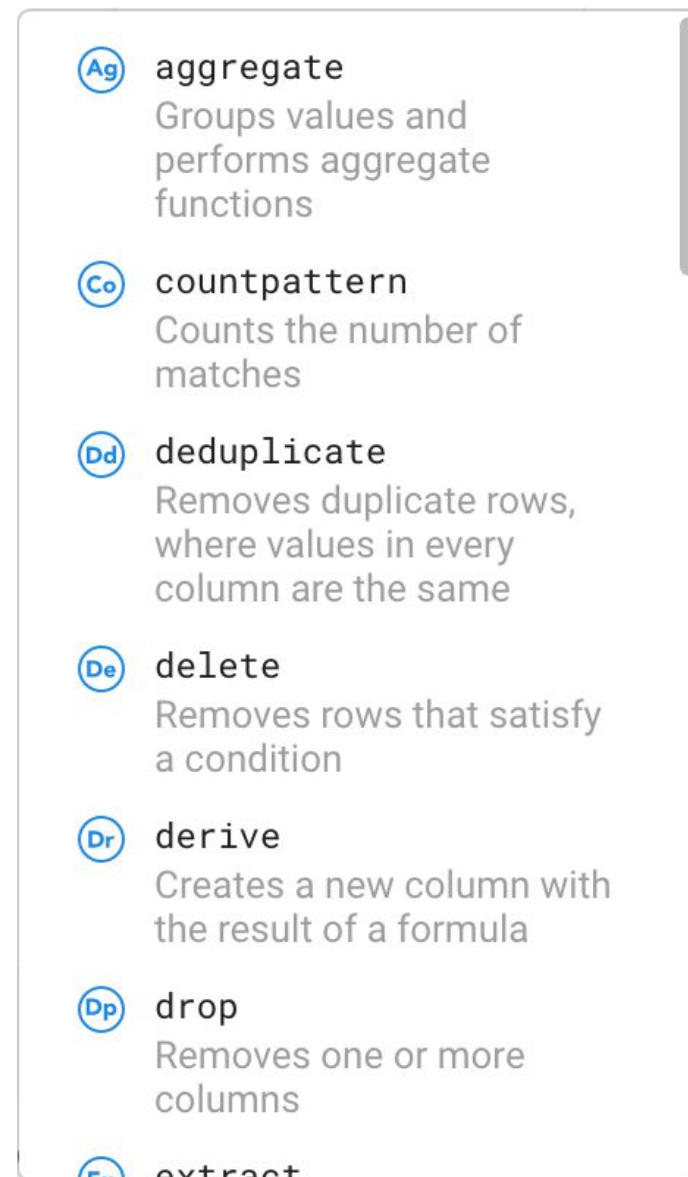


Cloud Dataprep supports the full preprocessing lifecycle



Cloud Dataprep wranglers write beam code in Dataflow

Build Recipes in Cloud Dataprep UI




Cloud Dataprep wranglers write beam code in Dataflow

Build Recipes in Cloud Dataprep UI

- Ag** aggregate
Groups values and performs aggregate functions
- Co** countpattern
Counts the number of matches
- Dd** deduplicate
Removes duplicate rows, where values in every column are the same
- De** delete
Removes rows that satisfy a condition
- Dr** derive
Creates a new column with the result of a formula
- Dp** drop
Removes one or more columns
- Ex** extract

Dataprep Converts Recipes to Beam




```
...  
.apply(Sum.integersPerKey()) //  
.apply("ToView", View.asMap());  
  
// packages in terms of use and which r  
javaContent //  
.apply("IsPopular", ParDo.of(ne  
    @ProcessElement  
    public void processElement(  
        String[] lines = c.ele  
        String[] packages = par  
        for (String packageName  
            c.output(KV.of(pack
```

Cloud Dataprep wranglers write beam code in Dataflow

Build Recipes in Cloud Dataprep UI

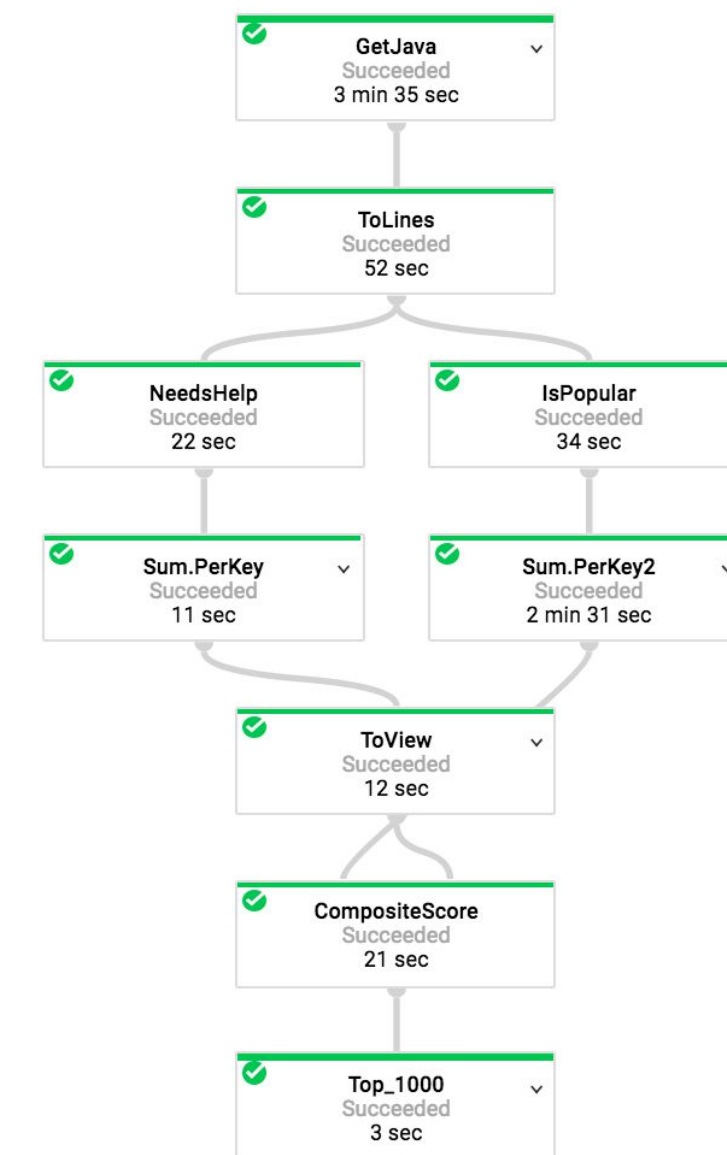
- Ag** aggregate
Groups values and performs aggregate functions
- Co** countpattern
Counts the number of matches
- Dd** deduplicate
Removes duplicate rows, where values in every column are the same
- De** delete
Removes rows that satisfy a condition
- Dr** derive
Creates a new column with the result of a formula
- Dp** drop
Removes one or more columns
- Ex** extract

Dataprep Converts Recipes to Beam



```
// ...  
.apply(Sum.integersPerKey()) //  
.apply("ToView", View.asMap());  
  
// packages in terms of use and which r  
javaContent //  
.apply("IsPopular", ParDo.of(ne  
    @ProcessElement  
    public void processElement(  
        String[] lines = c.ele  
        String[] packages = par  
        for (String packageName  
            c.output(KV.of(pack
```

Dataprep Runs a Dataflow Job



Wide array of transformation wranglers available

Data Ingestion (Upload, GCS, BigQuery)

Data Cleansing

Aggregations

Joins, Unions

Transformations

Type Conversions

Dataprep Wranglers

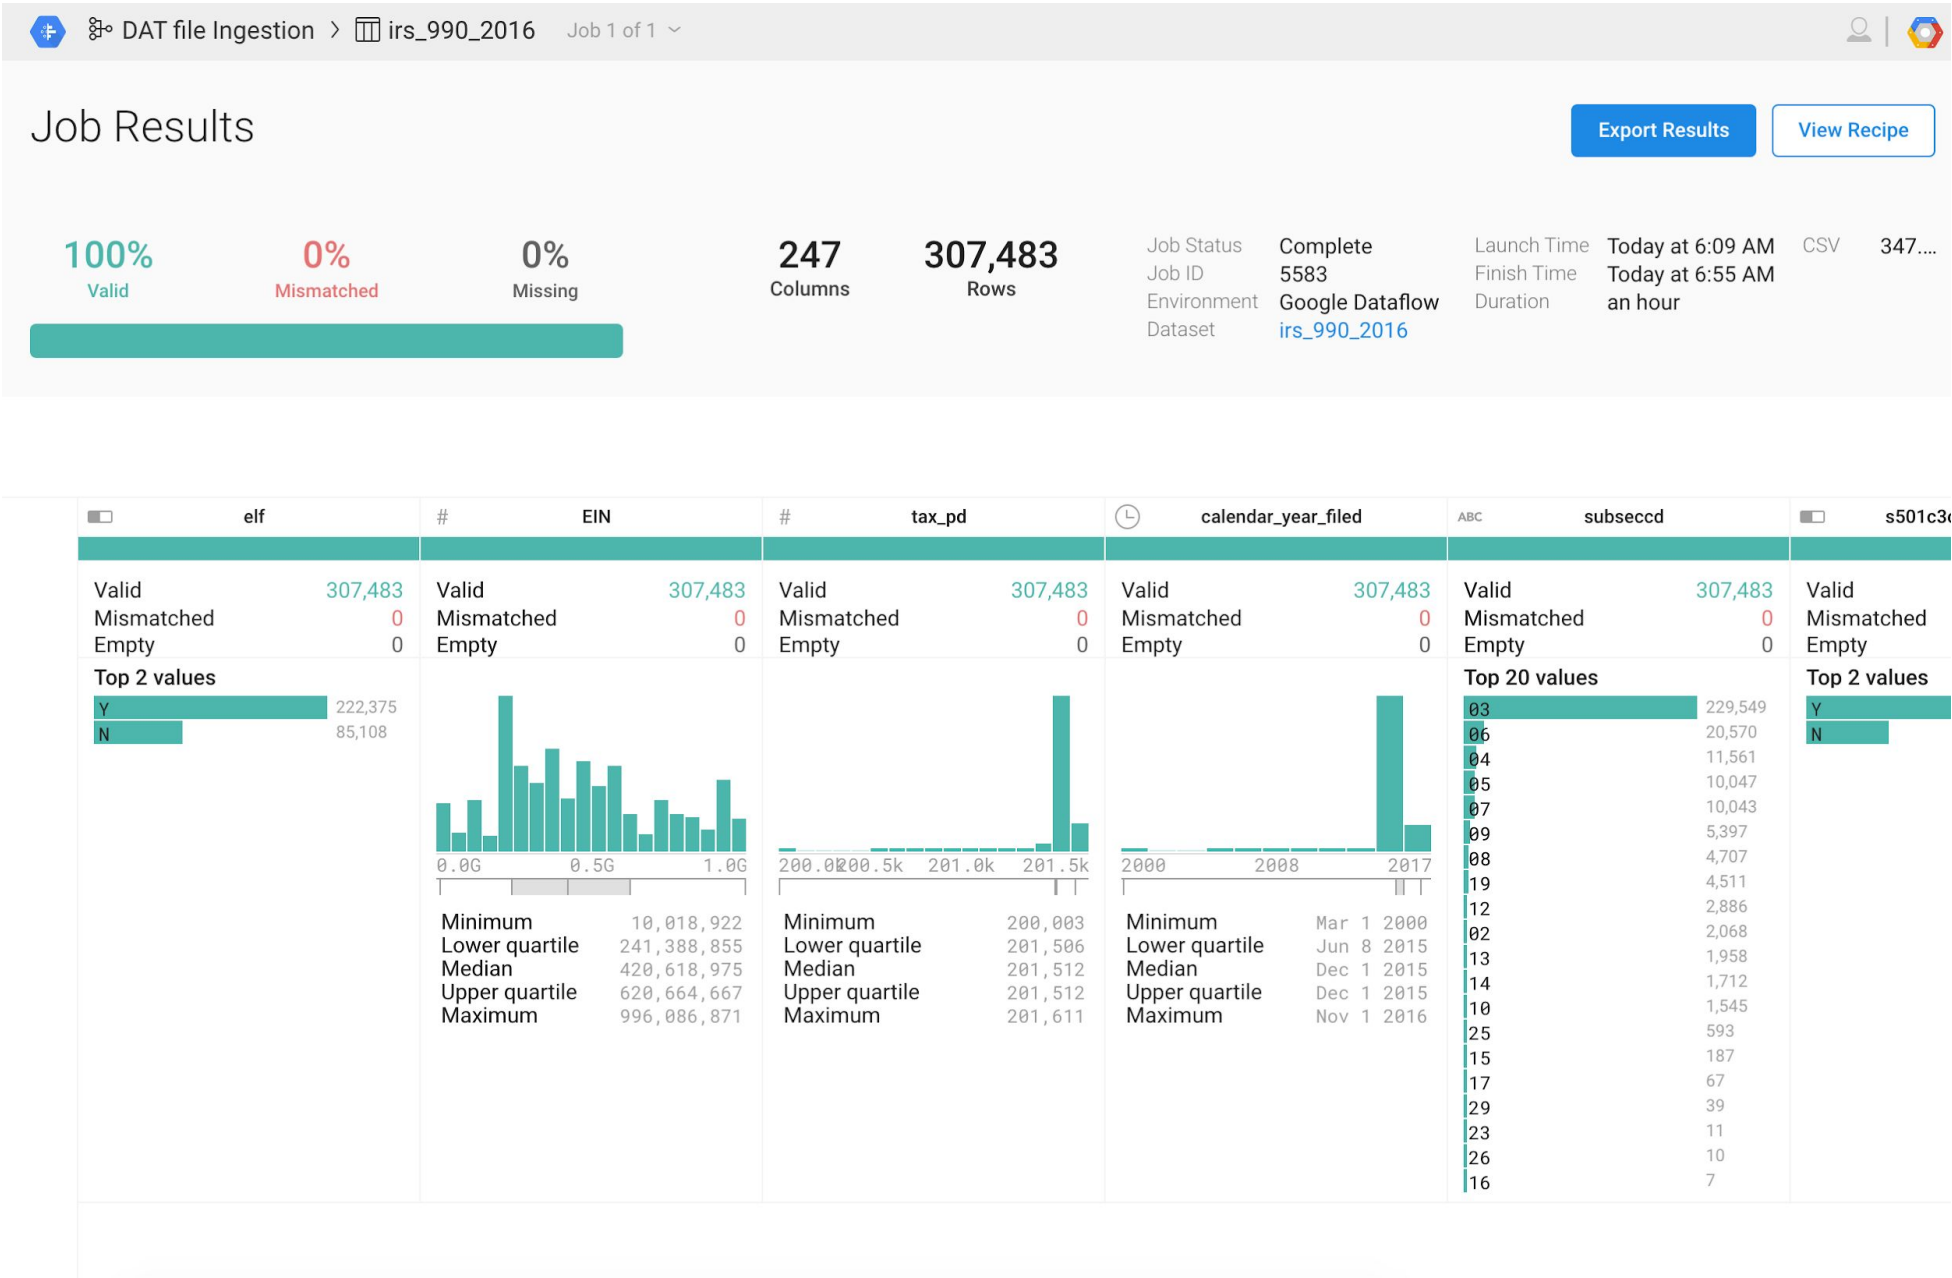
- Dr** **derive**
Creates a new column with the result of a formula
- Dp** **drop**
Removes one or more columns
- Ex** **extract**
Extracts matches into new columns
- Ek** **extractkv**
Extracts key-value pairs into a Map
- EI** **extractlist**
Extracts a list into an Array
- FI** **flatten**
Converts each element in an Array into a new row
- Jo** **join**
Adds additional columns from another dataset

Monitor Dataprep jobs and output results to BigQuery or GCS

Track completed and ongoing jobs

See the data quality metrics for transformed datasets

View histograms with summary statistics for each field



Lab

Computing Time-Windowed Features in Cloud Dataprep

Carl Osipov

Lab: Computing Time-Windowed Features in Cloud Dataprep

In this lab, you will learn how to:

- Build a new Flow using Cloud Dataprep

- Create and chain transformation steps with recipes

- Running Cloud Dataprep jobs



cloud.google.com



Feature Engineering

Carl Osipov



Preprocessing and feature
creation

Carl Osipov