



Feature Crosses

Feature Engineering

Machine Learning on Google Cloud Platform

Lak Lakshmanan

Learn how to...

Learn how to...

Recognize where feature crosses are
a powerful way to help machines
learn

Learn how to...

Recognize where feature crosses are a powerful way to help machines learn

Implement feature crosses in TensorFlow

Learn how to...

Recognize where feature crosses are a powerful way to help machines learn

Implement feature crosses in TensorFlow

Incorporate feature creation as part of your ML pipeline

Learn how to...

Recognize where feature crosses are a powerful way to help machines learn

Implement feature crosses in TensorFlow

Incorporate feature creation as part of your ML pipeline

Improve the taxifare model using feature crosses



Why feature crosses?

Lak Lakshmanan

Can you draw a line that separates these two classes?



Can you draw a line that separates these two classes?



Can you draw a line that separates these two classes?



Can you draw a line that separates these two classes?




Can you draw a line that separates these two classes?




Can you draw a line that separates these two classes?




Can you draw a line that separates these two classes?




This is a linear problem




How about this?
Is it a linear problem?




How about this?
Is it a linear problem?





How about this?
Is it a linear problem?





How about this?
Is it a linear problem?


























The feature cross provides a way to combine features in a linear model





Using non-linear inputs in a linear learner




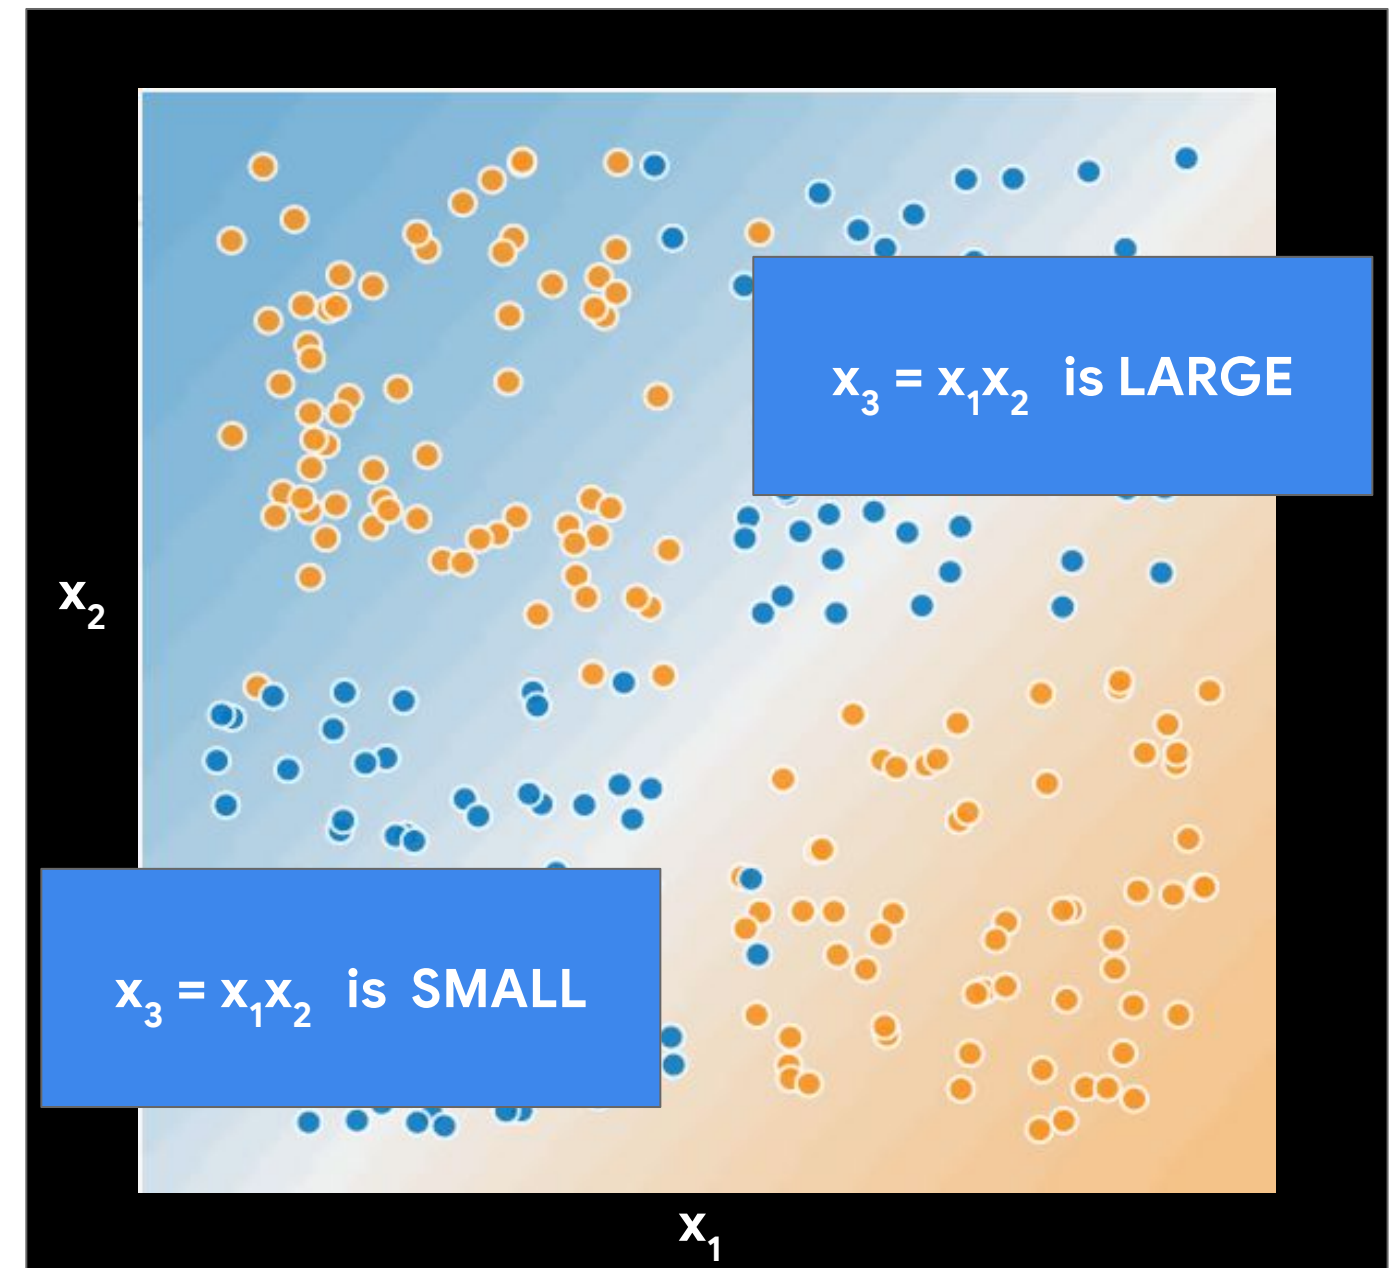
Using non-linear inputs in a linear learner



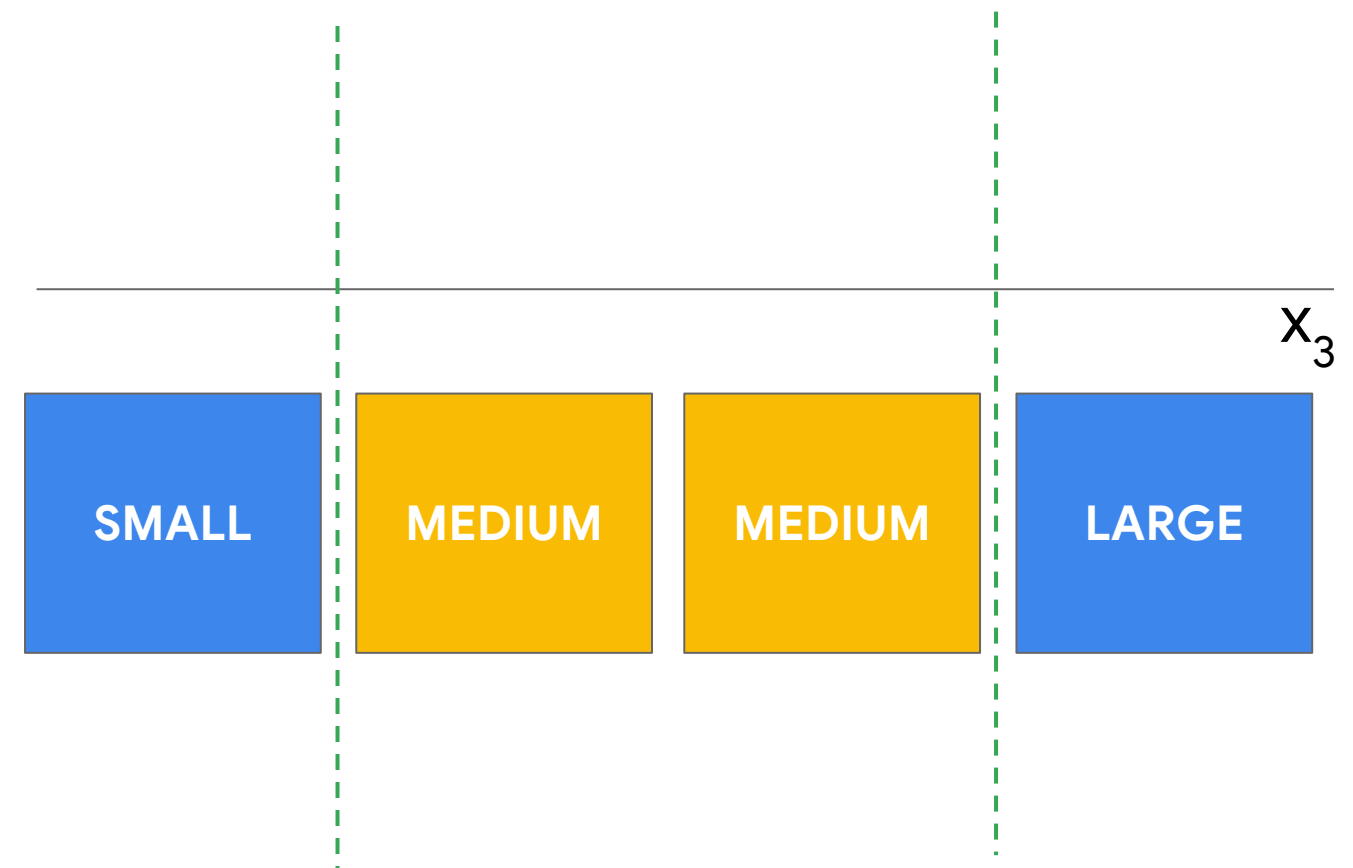










You now need two lines or
more free parameters




The white lines
discretize the space




Can a linear model
work for this problem?



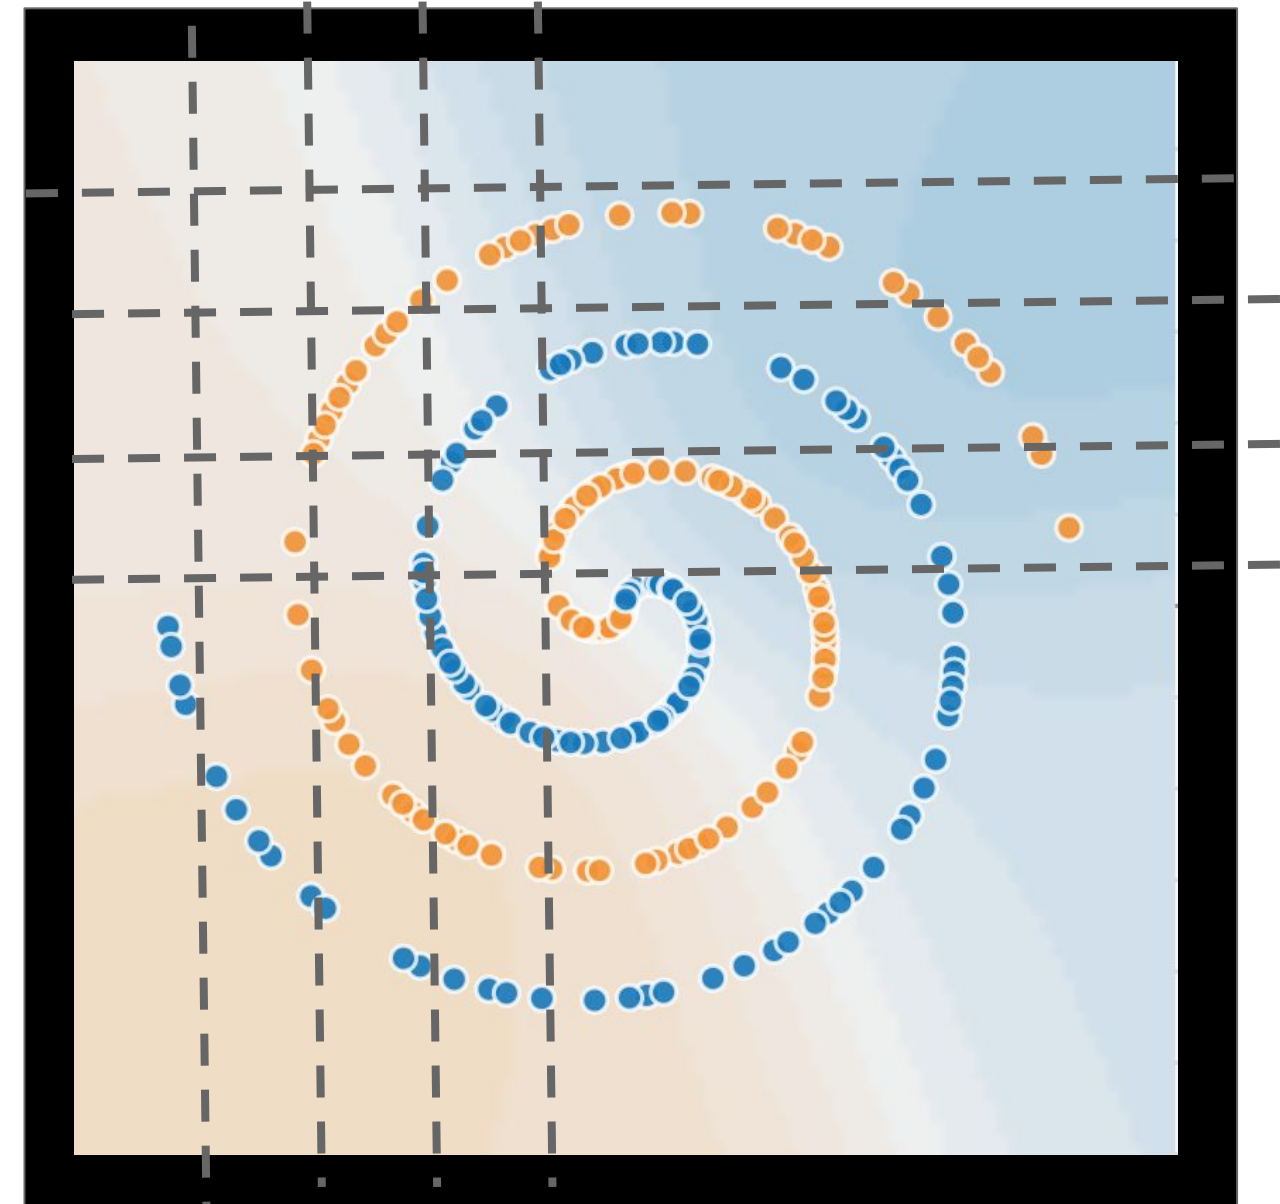
What if we discretize x_1
and x_2 and then multiply?




What if we discretize x_1
and x_2 and then multiply?




What if we discretize x_1
and x_2 and then multiply?



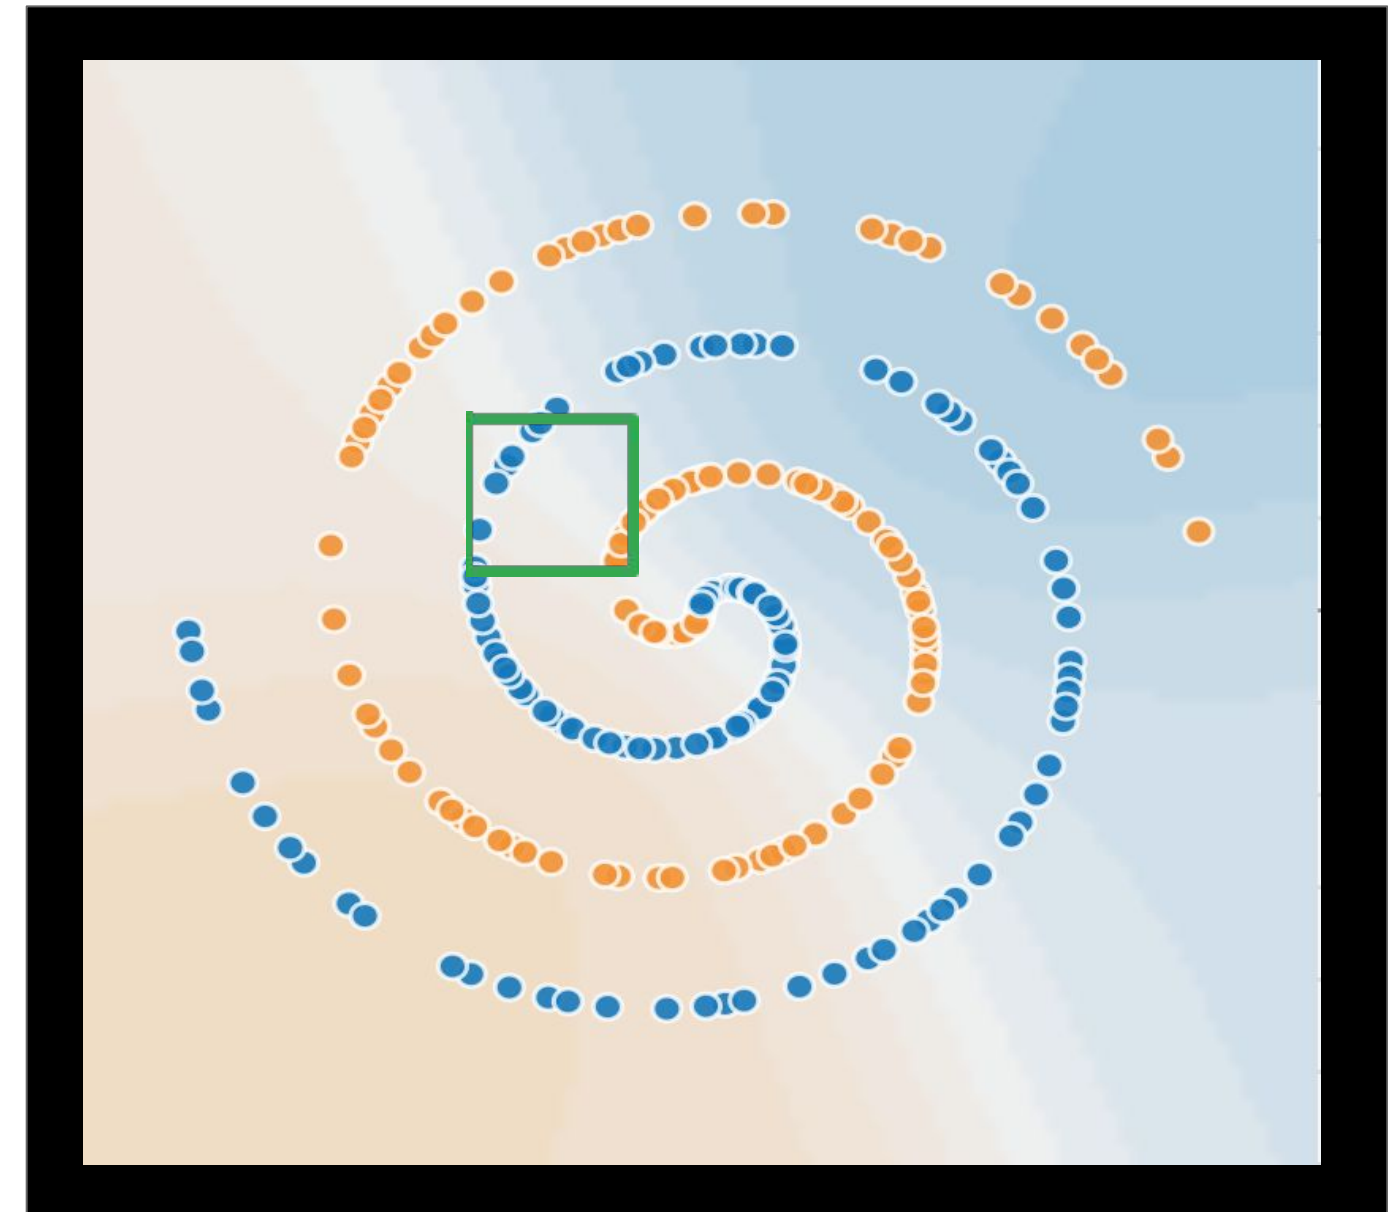
Dividing the input space
with two lines yields four
quadrants




Separate prediction per
grid cell



Separate prediction per
grid cell




Separate prediction per
grid cell



The weight of a cell is
essentially the prediction
for that cell


The weight of a cell is
essentially the prediction
for that cell




The weight of a cell is
essentially the prediction
for that cell




The weight of a cell is
essentially the prediction
for that cell




The weight of a cell is
essentially the prediction
for that cell



A feature cross memorizes
the input space



A feature cross memorizes
the input space



Feature crosses memorize!

Feature crosses memorize!

Goal of ML is generalization

Feature crosses memorize!

Goal of ML is generalization

Memorization works when you
have lots of data

Feature crosses memorize!

Goal of ML is generalization

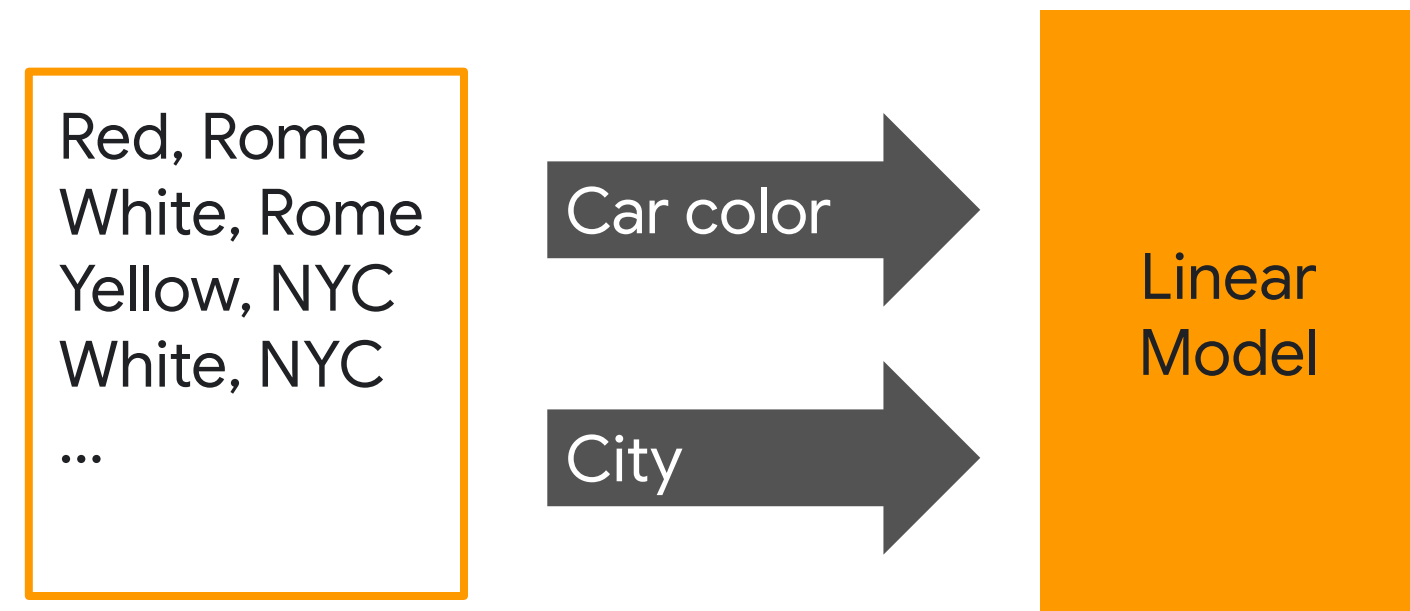
Memorization works when you
have lots of data

Feature crosses are powerful

Which of these cars is a taxi?



Assume that your input
data looks like this



The linear model has problems ...

Car color



The linear model has problems ...


Car color



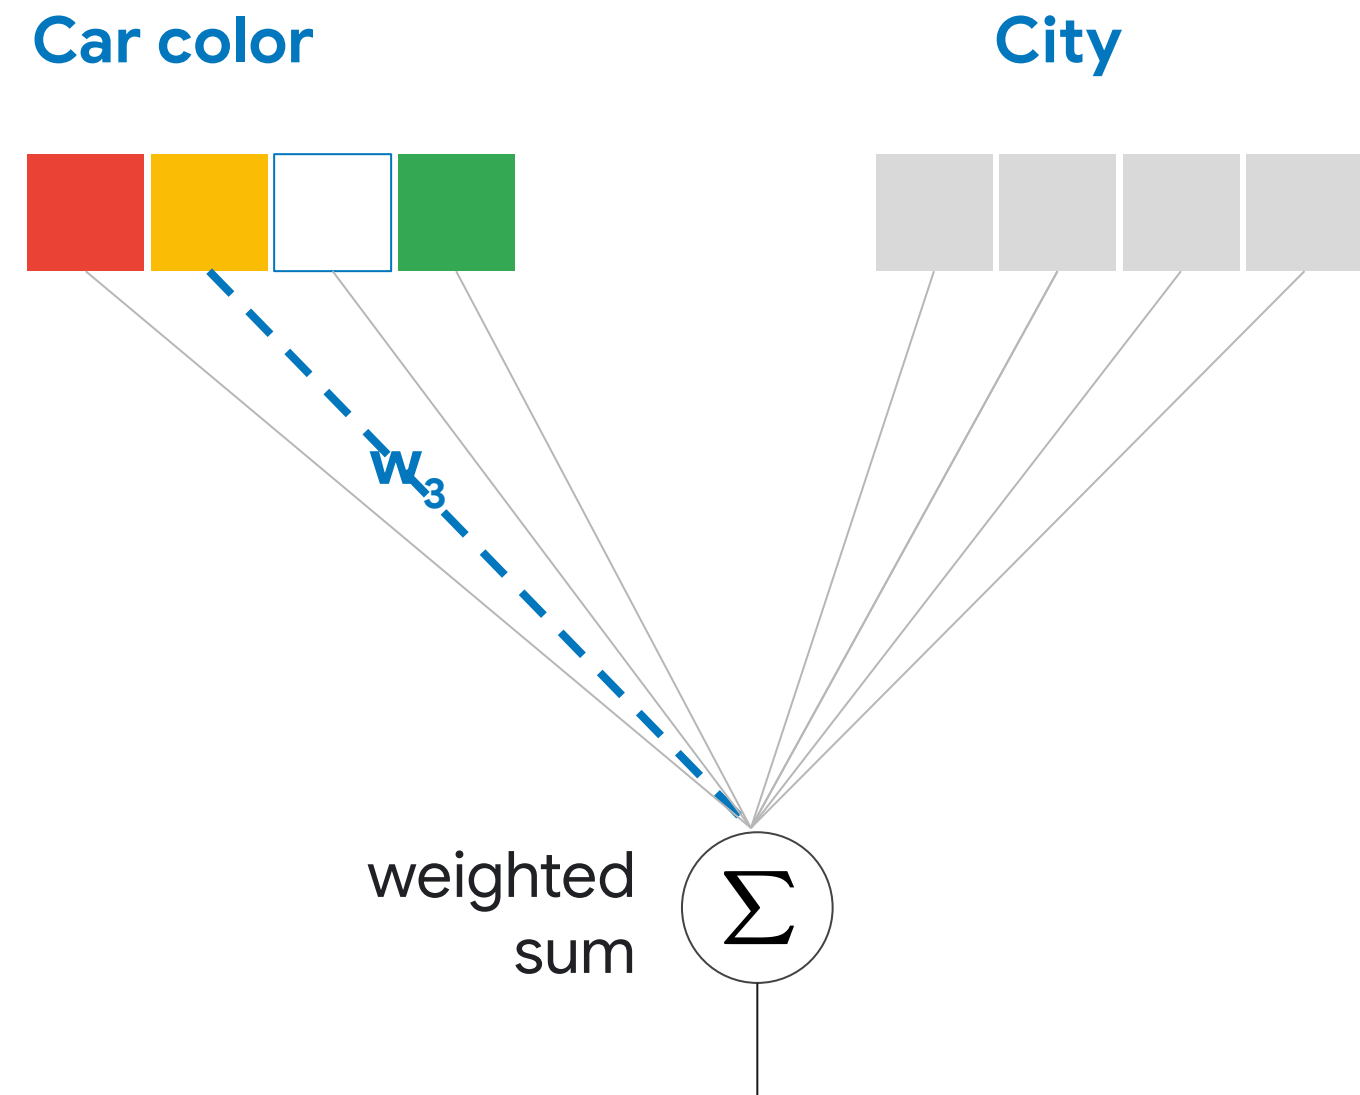
City




The linear model has problems ...



The linear model has problems ...



The linear model has problems ...



The feature cross has no problem

R = Rome
N = New York

Car color




City



$x_3 = \text{Color-City}$



weighted
sum



The feature cross has no problem

R = Rome
N = New York


Car color



City



weighted
sum



The feature cross has no problem

R = Rome
N = New York


Car color



City

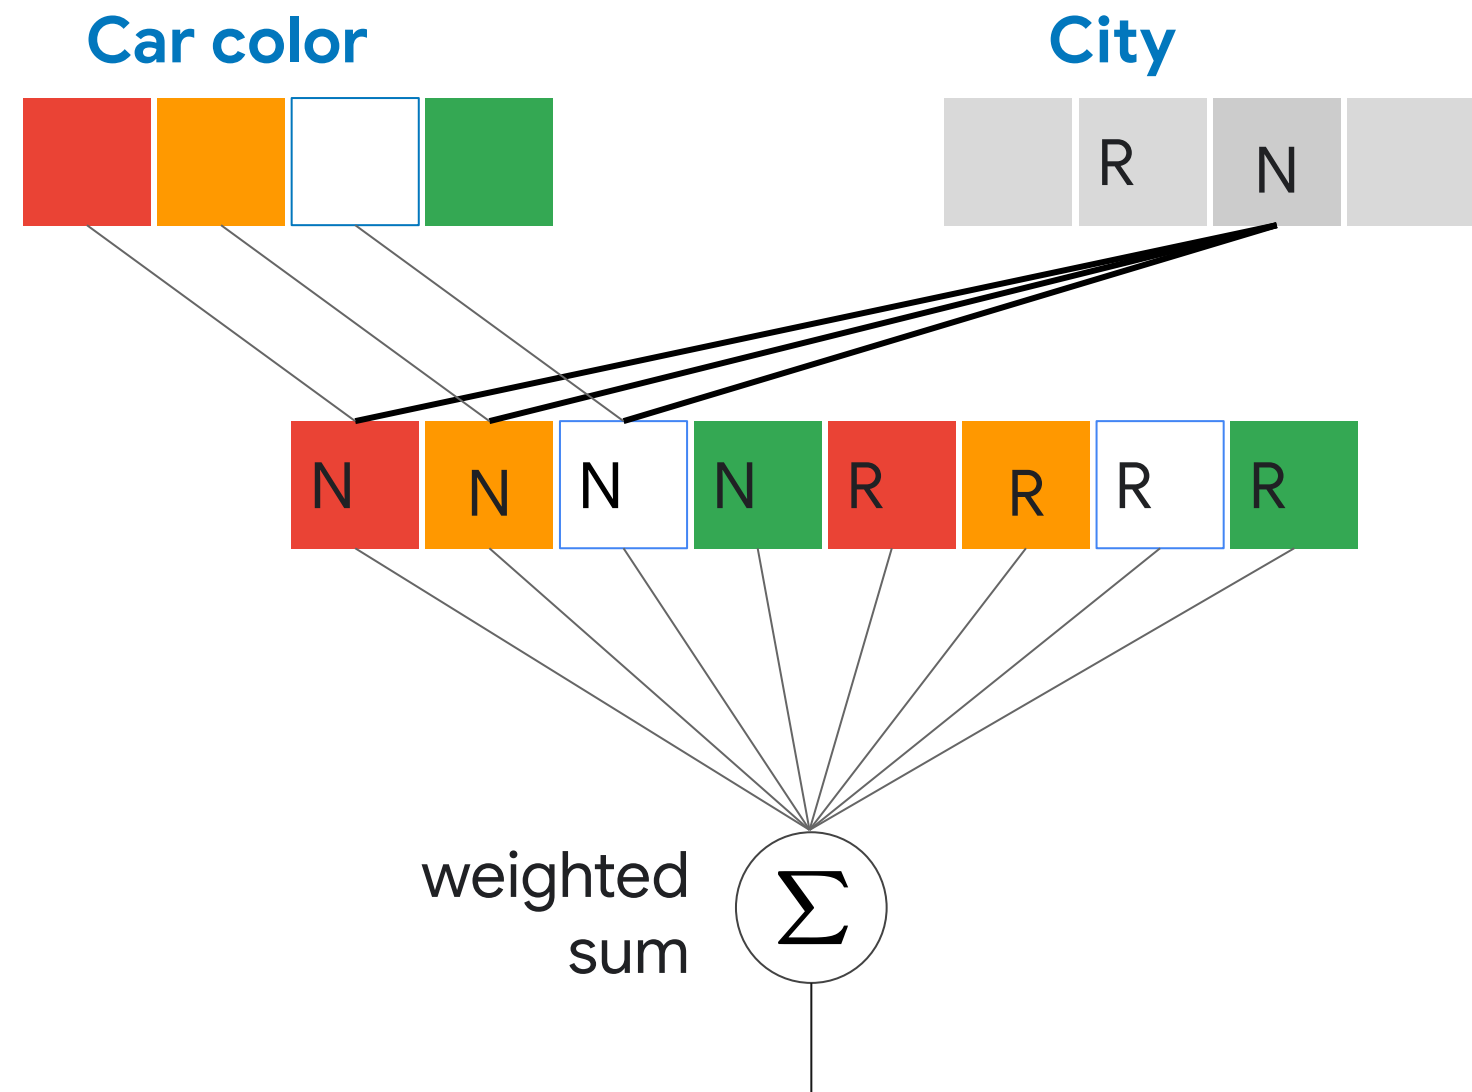


weighted
sum



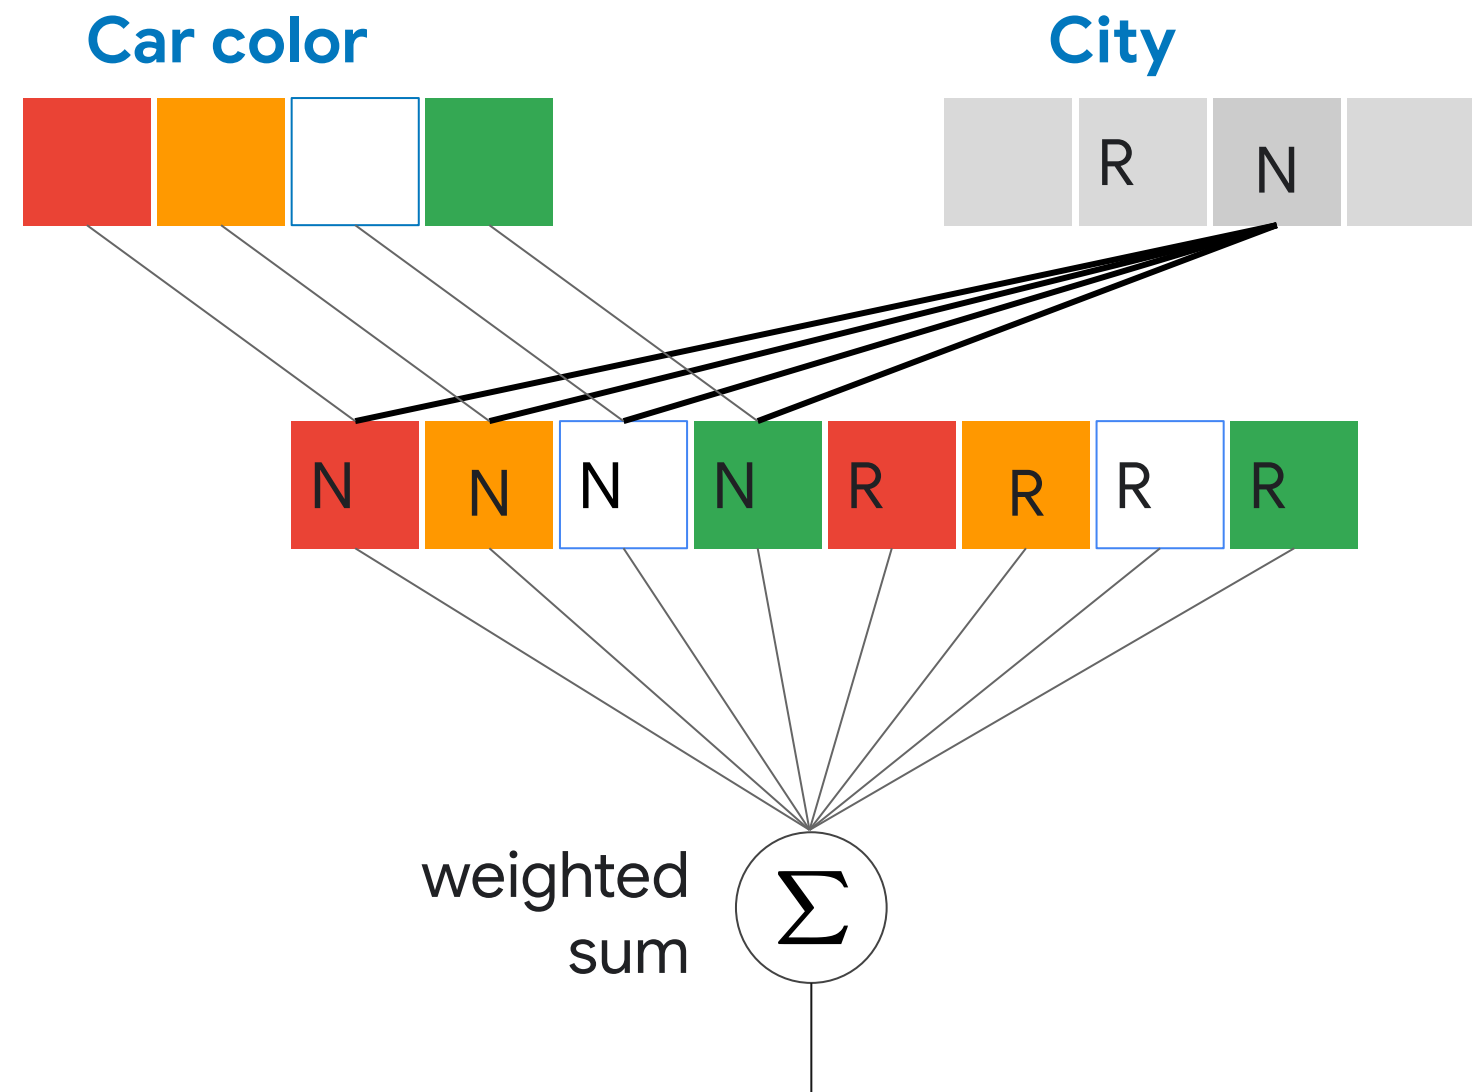
The feature cross has no problem

R = Rome
N = New York



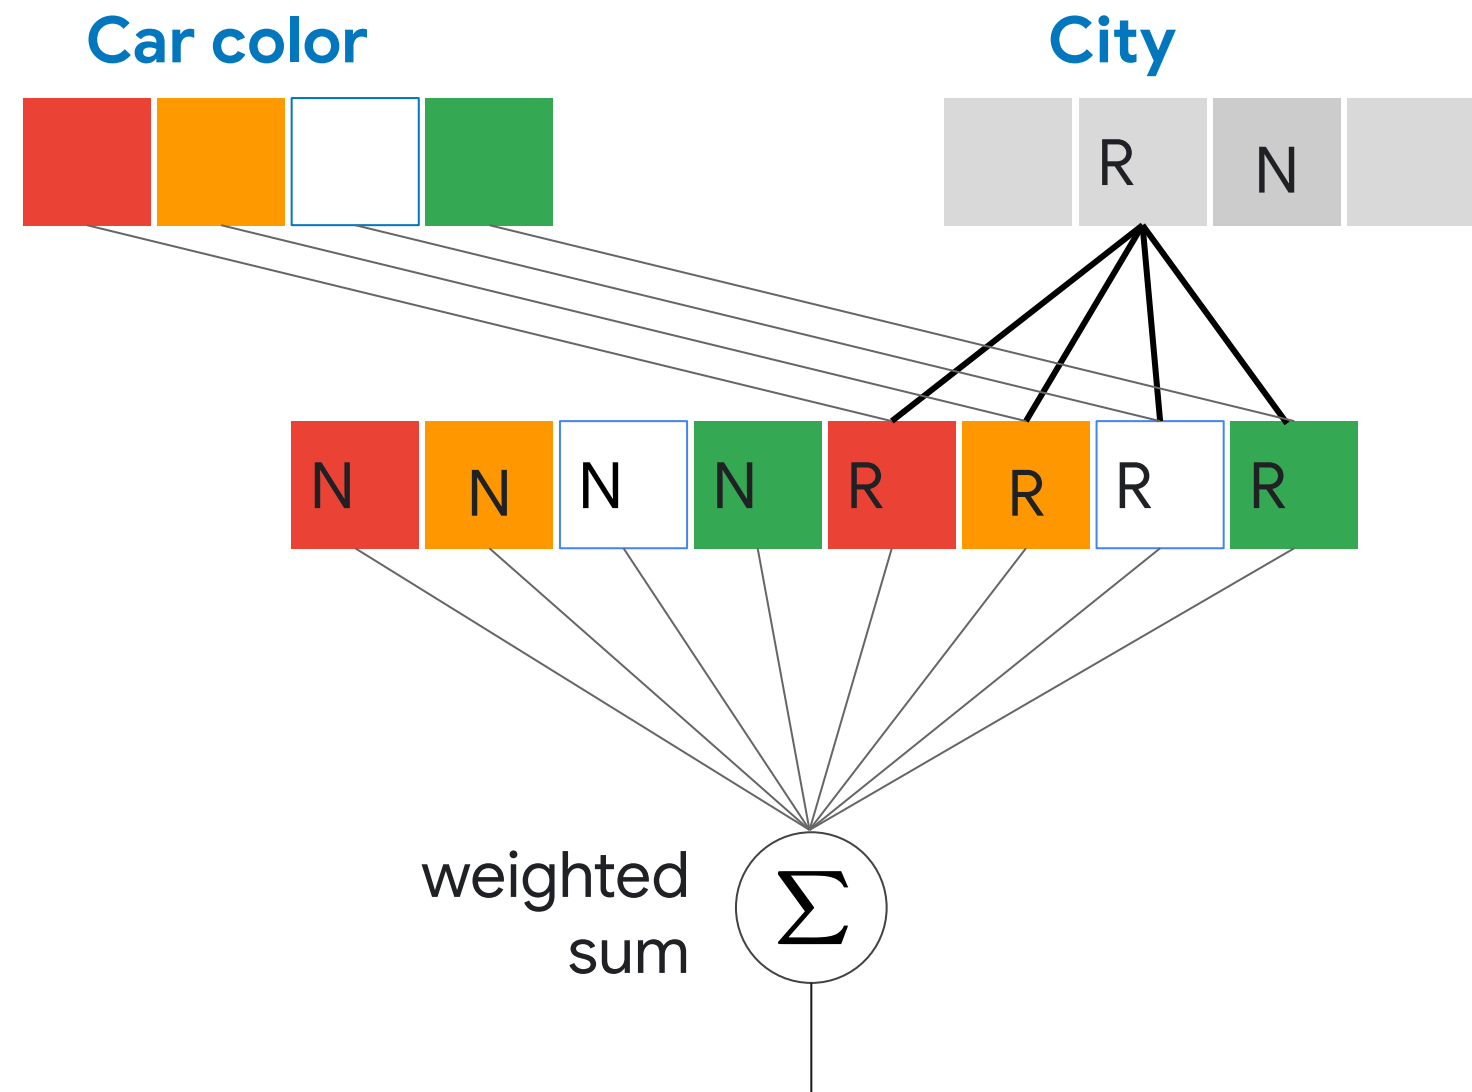
The feature cross has no problem

R = Rome
N = New York



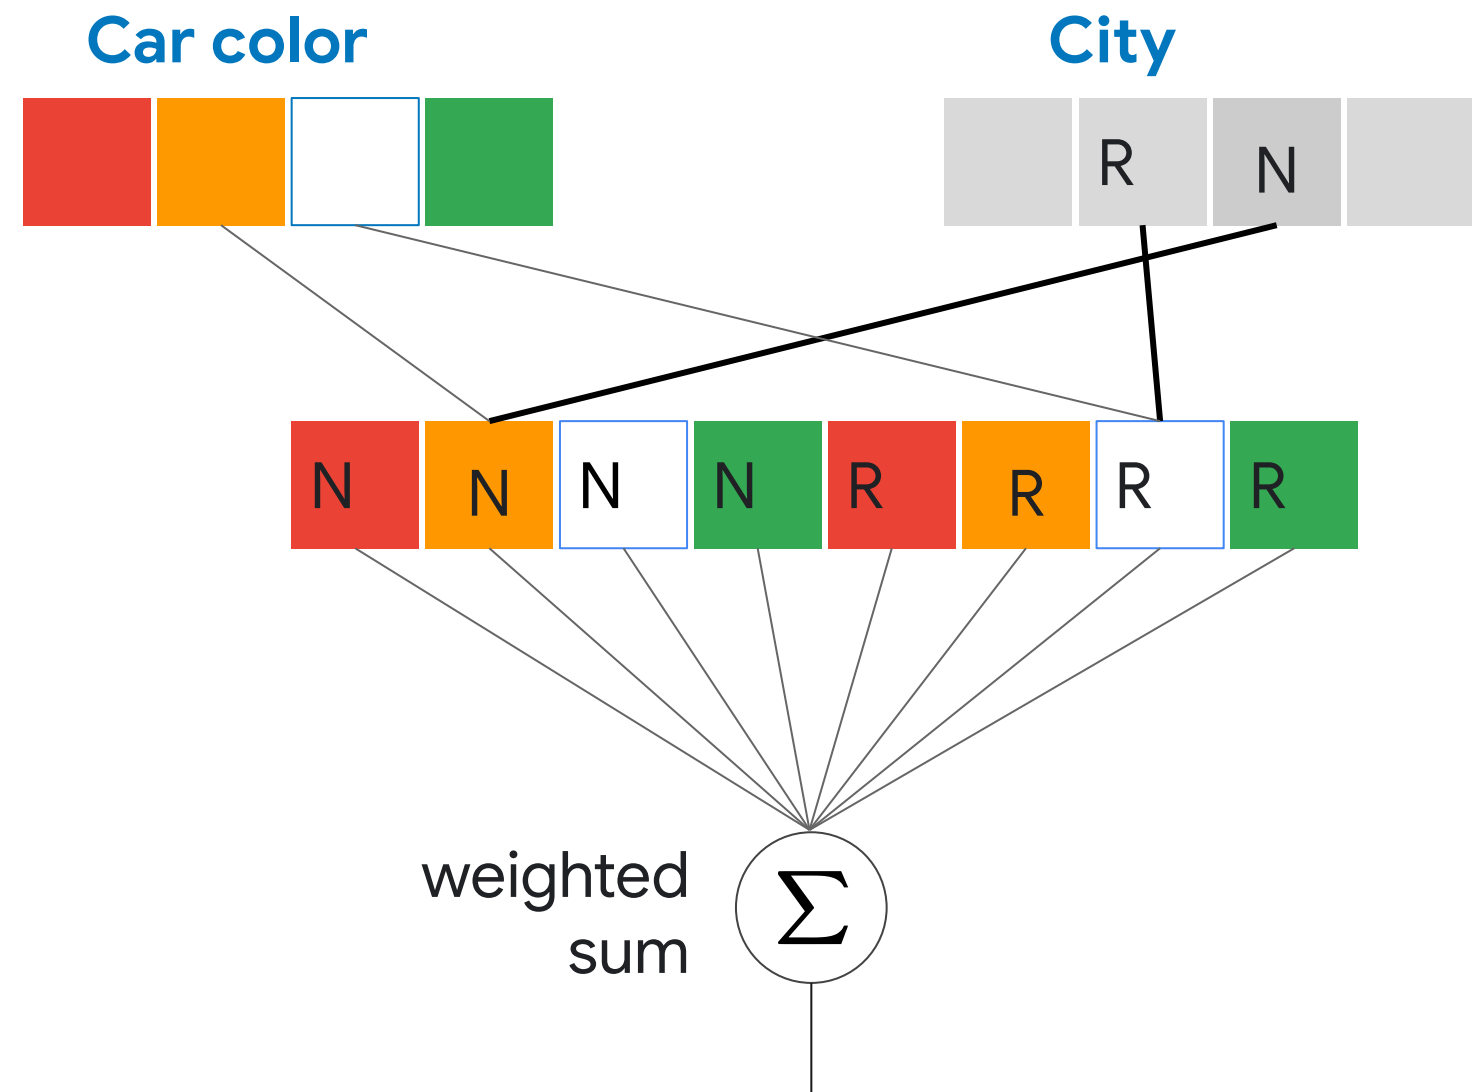
The feature cross has no problem

R = Rome
N = New York



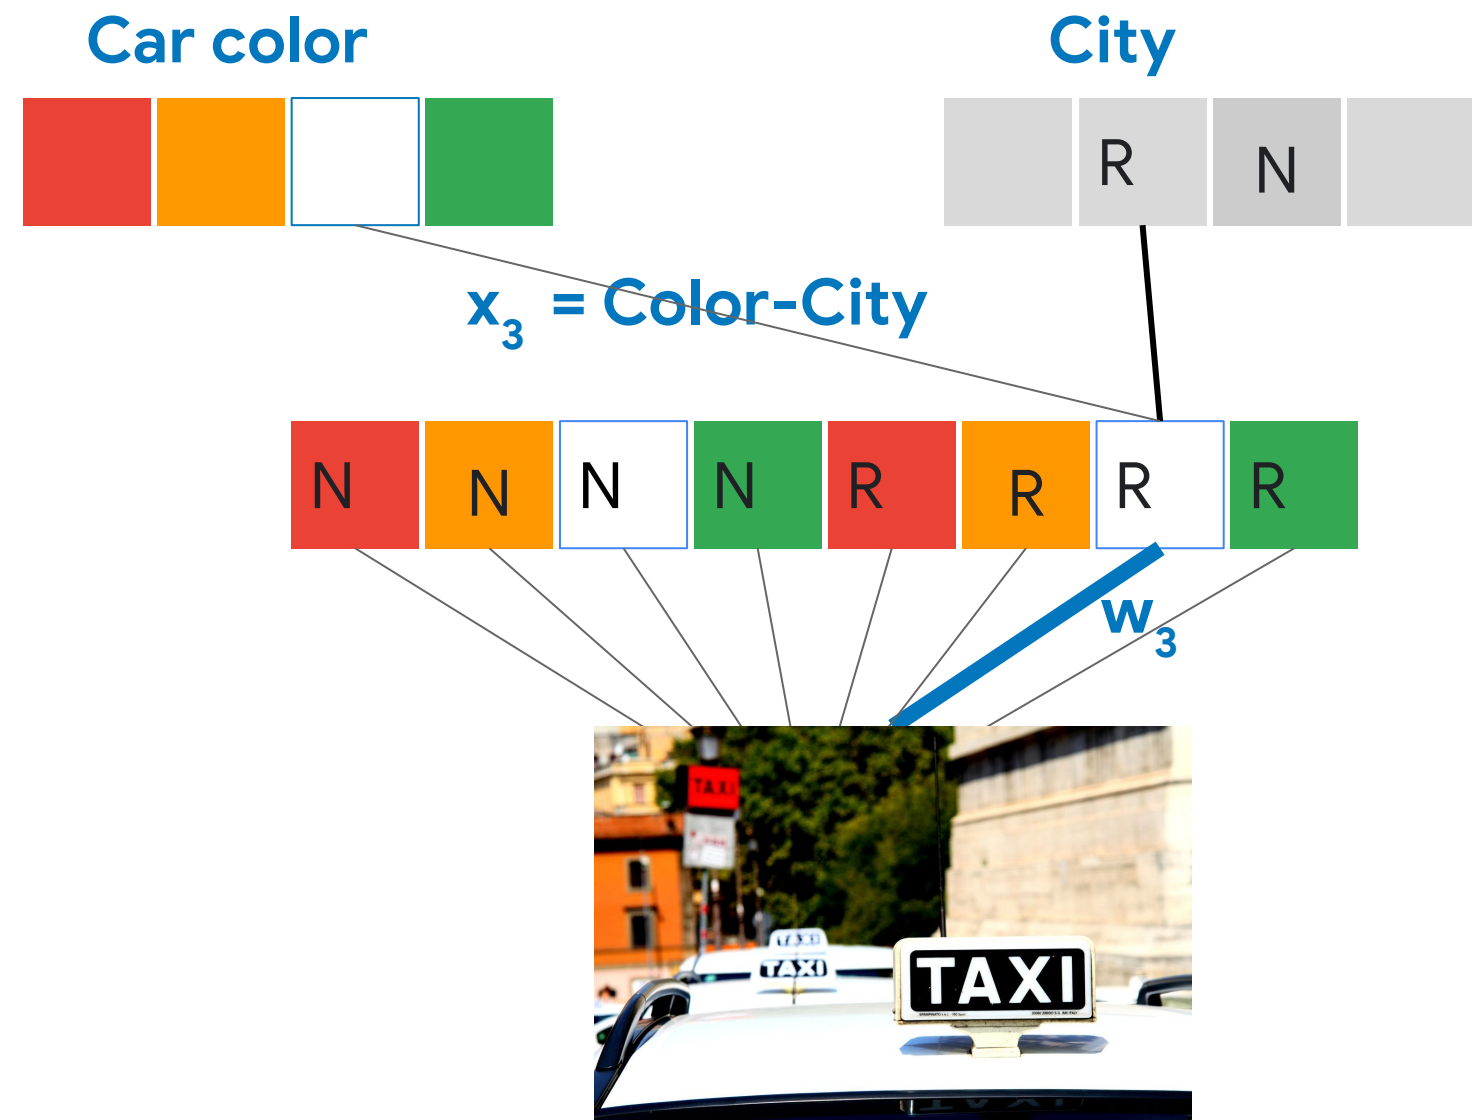
The feature cross has no problem

R = Rome
N = New York



The feature cross has no problem

R = Rome
N = New York



Feature Crosses bring a lot of power to linear models

Feature crosses + massive data is an efficient way for learning highly complex spaces

Feature Crosses bring a lot of power to linear models

Feature crosses + massive data is an efficient way for learning highly complex spaces

Feature crosses allow a linear model to memorize large datasets

Feature Crosses bring a lot of power to linear models

Feature crosses + massive data is an efficient way for learning highly complex spaces

Feature crosses allow a linear model to memorize large datasets

Optimizing linear models is a convex problem

Feature Crosses bring a lot of power to linear models

Feature crosses + massive data is an efficient way for learning highly complex spaces

Feature crosses allow a linear model to memorize large datasets

Optimizing linear models is a convex problem

Before TensorFlow, Google used massive scale learners

Feature Crosses bring a lot of power to linear models

Feature crosses + massive data is an efficient way for learning highly complex spaces

Feature crosses allow a linear model to memorize large datasets

Optimizing linear models is a convex problem

Before TensorFlow, Google used massive scale learners

Feature crosses, as a preprocessor, make neural networks converge a lot quicker

Lab

Use feature crosses to create
a good classifier

Lab: Use feature crosses to create a good classifier

<https://goo.gl/2NUCAF>

<https://goo.gl/ivd4x4>

What's the best performance you can get?


Which feature crosses help the most?

Does the model output surface look like a linear model?


Lab

Screencast


The model boundary
doesn't look linear!



The linear decision boundary gets transformed into the curve in the original coordinate space



Feature crosses combine discrete/categorical features



Cross hour_of_day with
day_of_week to predict
traffic

Cross hour_of_day with
day_of_week to predict
traffic

Hour of day



Day of week



Cross hour_of_day with
day_of_week to predict
traffic

Hour of day

24




Day of week

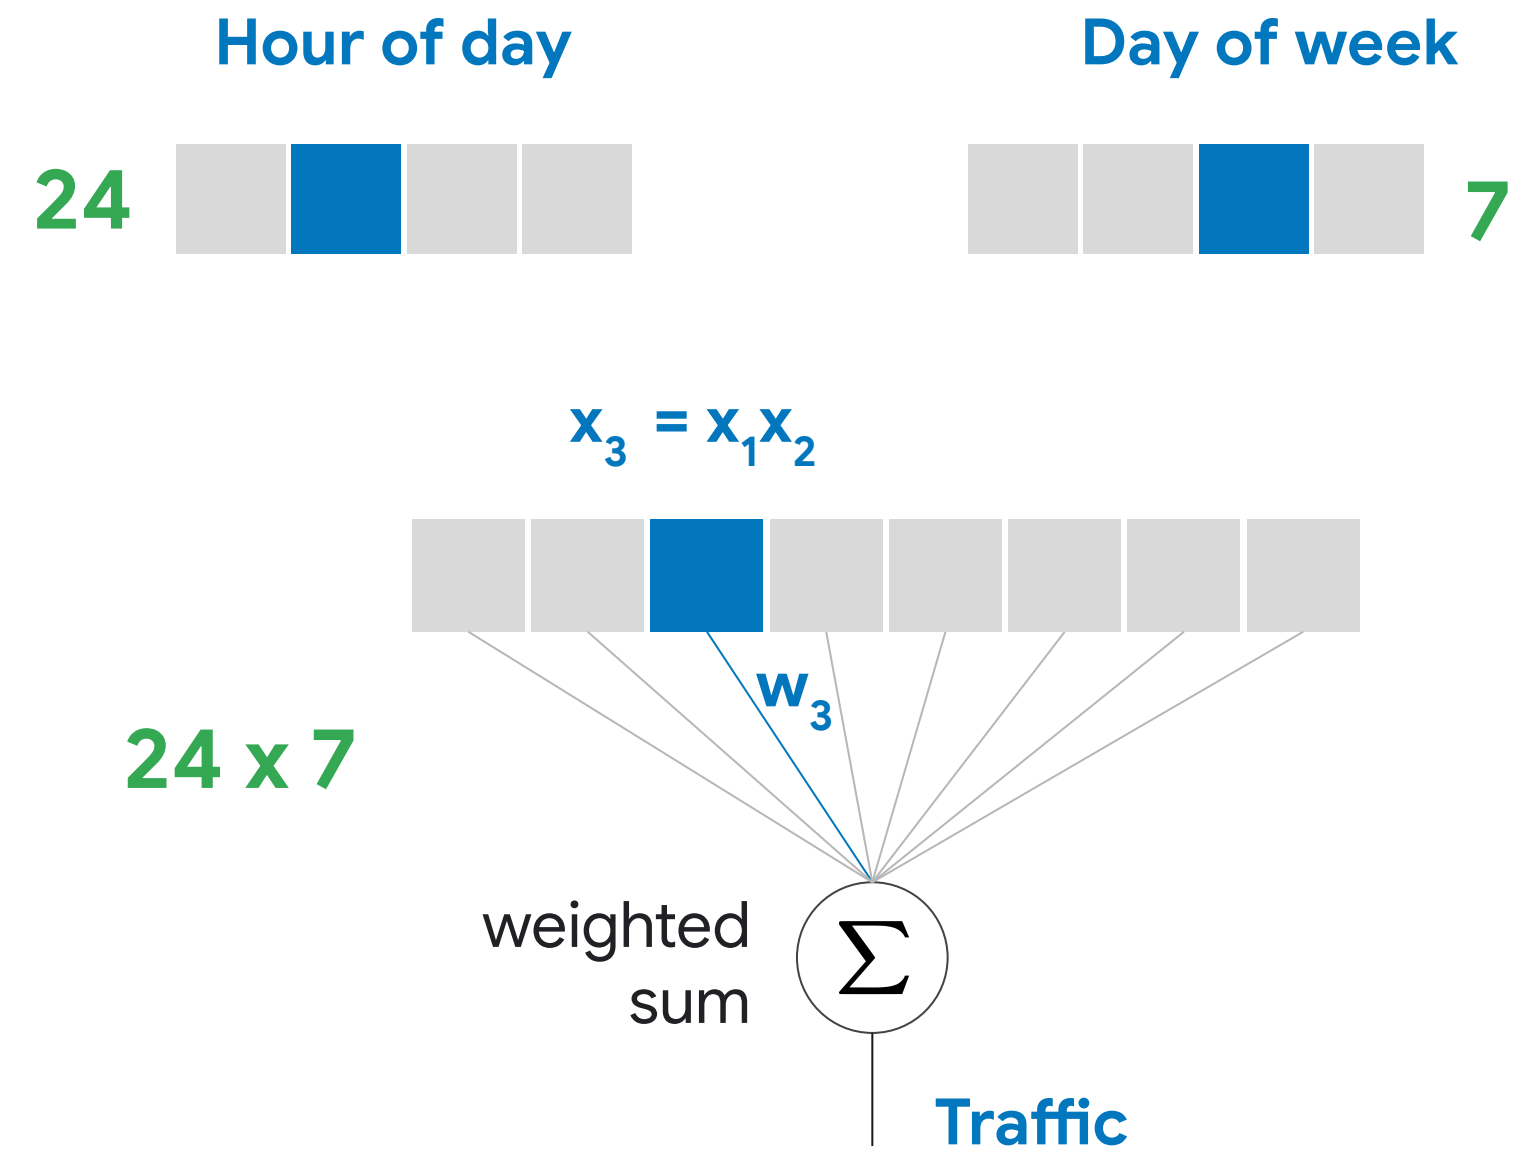
7



Cross hour_of_day with day_of_week to predict traffic



Cross hour_of_day with day_of_week to predict traffic



Feature Crosses lead to sparsity

Hour of day



Day of week




$$x_3 = x_1 x_2$$



w_3

weighted
sum



Traffic

Feature Crosses lead to sparsity

Hour of day



Day of week



$$x_3 = x_1 x_2$$




167 zeros

1 one

w_3

weighted
sum



Traffic

Quiz: Which of these is a good feature cross?

Different cities in California have markedly different housing prices. Suppose you must create a model to predict housing prices. Which of the following sets of features or feature crosses could learn city-specific relationships between house characteristic and housing price?

- a) Three separate binned features: [binned latitude], [binned longitude], [binned roomsPerPerson]
- b) Two feature crosses: [binned latitude X binned roomsPerPerson] and [binned longitude X binned roomsPerPerson]
- c) One feature cross: [binned latitude X binned longitude X binned roomsPerPerson]
- d) One feature cross: [latitude X longitude X roomsPerPerson]

Lab

Too much of a good thing

Lab: Too much of a good thing

<https://goo.gl/ofHCT>

Is the model behavior surprising?


What's the issue?

Try removing cross-product features. Does performance improve?


Lab

Screencast

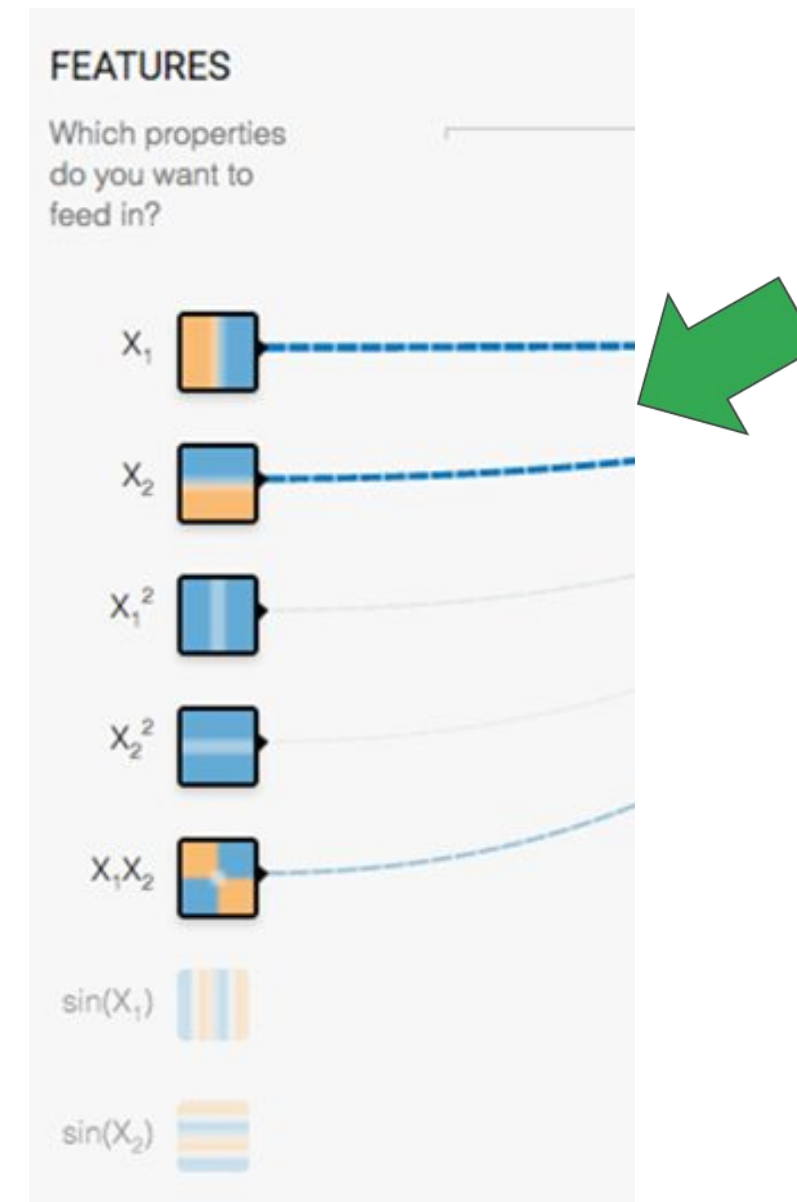
Lab: Too much of a good thing



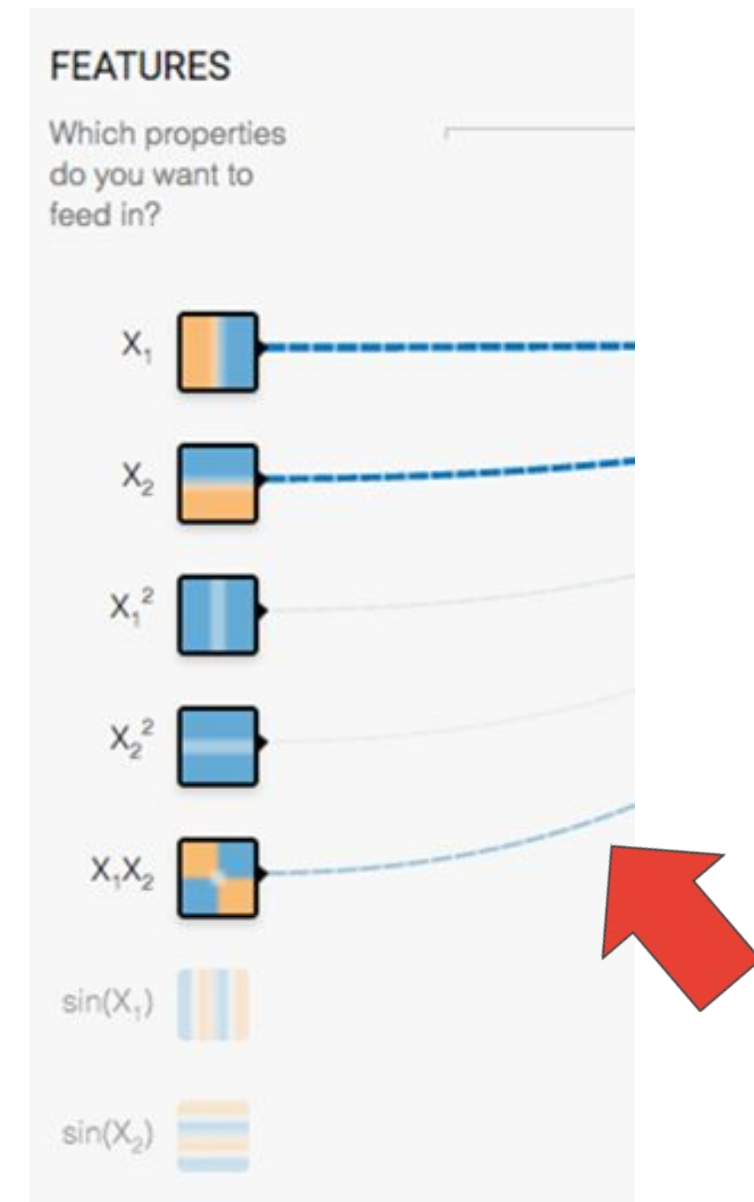
Lab: Too much of a good thing




Lab: Too much of a good thing



Lab: Too much of a good thing



After removing the feature
crosses ...





Implementing feature crosses


First Lastname

Creating feature crosses using TensorFlow

```
day_hr =  
tf.feature_column.crossed_column(  
    [dayofweek, hourofday],  
    24 * 7)
```

Creating feature crosses using TensorFlow

```
day_hr =  
tf.feature_column.crossed_column(  
    [dayofweek, hourofday],  
    24 * 7)
```



You can cross two
or more categorical
or bucketized
columns

Creating feature crosses using TensorFlow

```
day_hr =  
tf.feature_column.crossed_column(  
    [dayofweek, hourofday],  
    24 * 7)
```



This is the number
of hash buckets:
feature-cross %
hash_bucket_size

Choosing the number of hash buckets is an art, not a science

```
day_hr =  
tf.feature_column.crossed_column(  
    [dayofweek, hourofday],  
    24 * 7)
```

Choosing the number of hash buckets is an art, not a science

```
day_hr =  
tf.feature_column.crossed_column(  
    [dayofweek, hourofday],  
    6
```


Choosing the number of hash buckets is an art, not a science

```
day_hr =  
tf.feature_column.crossed_column(  
    [dayofweek, hourofday],  
    6
```

feature-cross % hash_bucket_size



3pm->15 Wed->3

3pm X Wed -> $15 + 3 \times 24 = 87$

Choosing the number of hash buckets is an art, not a science

```
day_hr =  
tf.feature_column.crossed_column(  
    [dayofweek, hourofday],  
    6
```

feature-cross % hash_bucket_size



3pm->15 Wed->3

3pm X Wed -> 15 + 3*24 = 87

The number of hash buckets controls sparsity and collisions

Small hash_buckets -> lots of collisions



The number of hash buckets controls sparsity and collisions

Small hash_buckets -> lots of collisions



High hash_buckets -> very sparse



The number of hash buckets controls sparsity and collisions

Small hash_buckets -> lots of collisions



High hash_buckets -> very sparse



Creating an embedding column from a feature cross

Hour of day



Day of week



$$x_3 = x_1 x_2$$



Creating an embedding column from a feature cross

Hour of day




Day of week



$$x_3 = x_1 x_2$$



weighted
sum



Traffic

Creating an embedding column from a feature cross

Hour of day



Day of week



boolean


$$x_3 = x_1 x_2$$



real-value

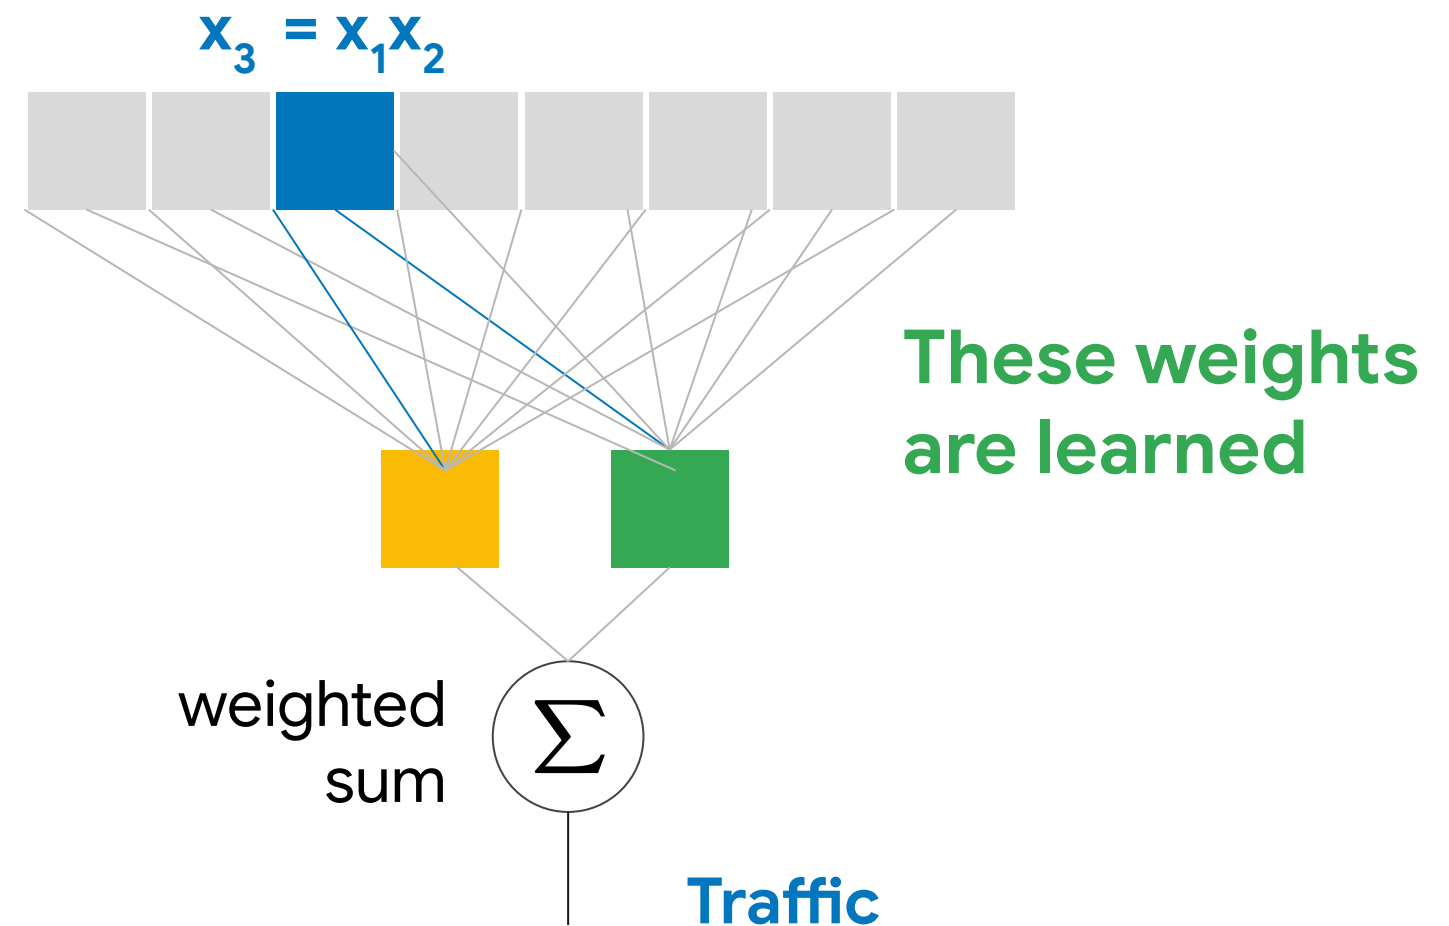


weighted
sum

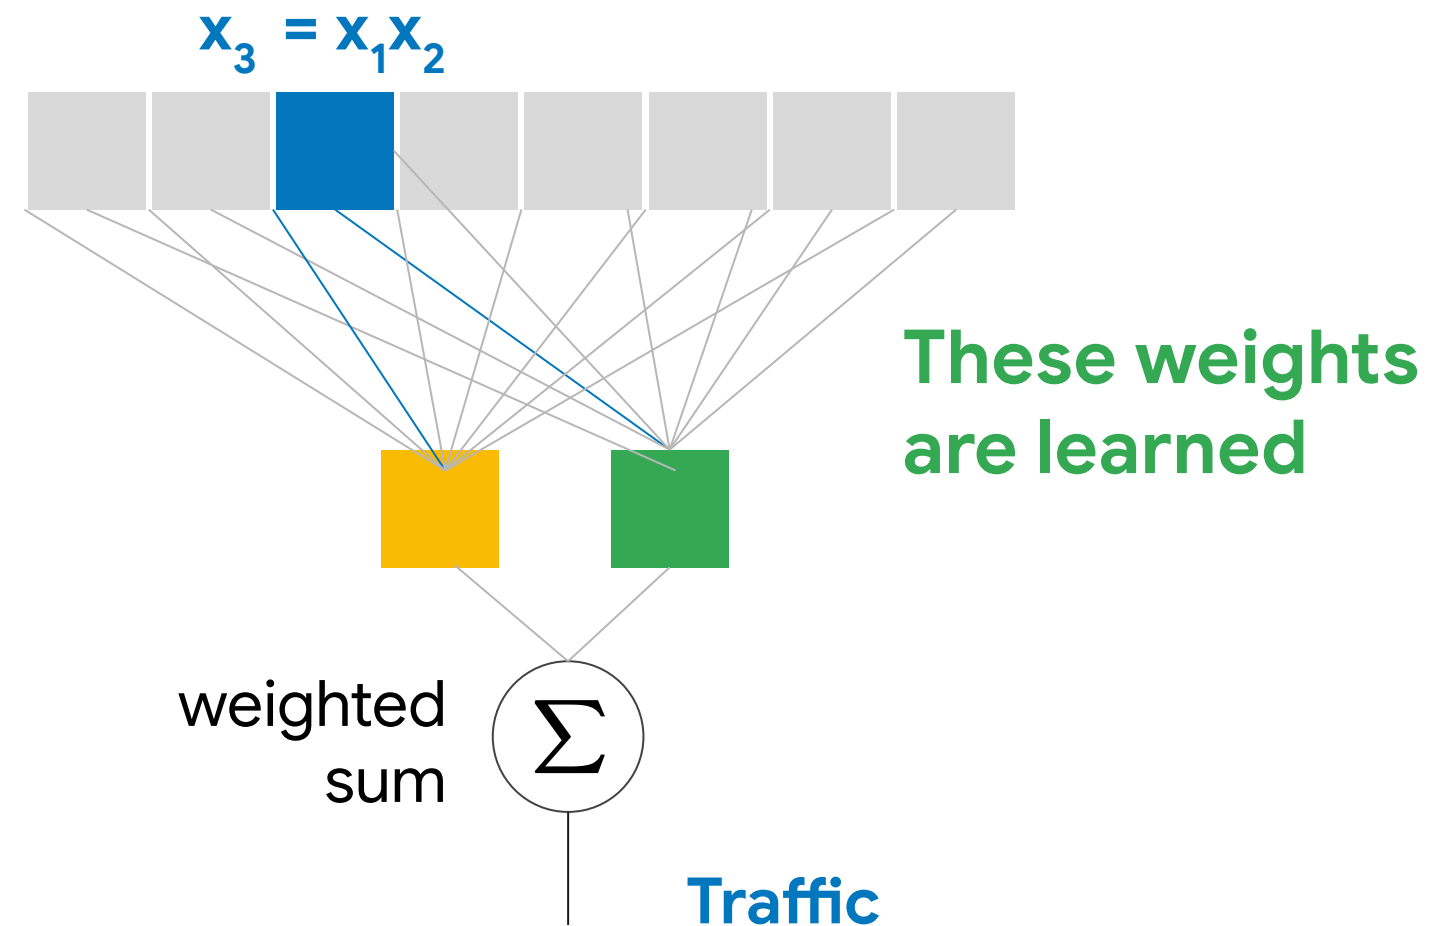


Traffic


The weights in the
embedding column are
learned from data





The weights in the
embedding column are
learned from data





The weights in the
embedding column are
learned from data





The model learns how to embed the feature cross in lower-dimensional space

		
8am Tue	0.8	0.7
9am Wed	0.7	0.9
11 am Tue	0.1	0.6
2 pm Wed	0.1	0.7
2 am Tue	0	0.1
2 am Wed	0	0.1



The model learns how to embed the feature cross in lower-dimensional space

		
8am Tue	0.8	0.7
9am Wed	0.7	0.9
11 am Tue	0.1	0.6
2 pm Wed	0.1	0.7
2 am Tue	0	0.1
2 am Wed	0	0.1
<div>similar</div>		



The model learns how to embed the feature cross in lower-dimensional space

		
8am Tue	0.8	0.7
9am Wed	0.7	0.9
11 am Tue	0.1	0.6
2 pm Wed	0.1	0.7
2 am Tue	0	0.1
2 am Wed	0	0.1
	<div>different</div>	<div>similar</div>

The model learns how to embed the feature cross in lower-dimensional space

		
8am Tue	0.8	0.7
9am Wed	0.7	0.9
11 am Tue	0.1	0.6
2 pm Wed	0.1	0.7
2 am Tue	0	0.1
2 am Wed	0	0.1
<div>similar</div>		

The model learns how to embed the feature cross in lower-dimensional space

		
8am Tue	0.8	0.7
9am Wed	0.7	0.9
11 am Tue	0.1	0.6
2 pm Wed	0.1	0.7
2 am Tue	0	0.1
2 am Wed	0	0.1

Embedding a feature cross in TensorFlow

```
import tf.feature_column as fc

day_hr = fc.crossed_column(
    [dayofweek, hourofday],
    24x7 )

day_hr_em = fc.embedding_column(
    day_hr,
    2,)
```

Transfer Learning of embeddings from similar ML models

```
import tf.feature_column as fc

day_hr = fc.crossed_column(
    [dayofweek, hourofday],
    24x7 )

day_hr_em = fc.embedding_column(
    day_hr,
    2,
    ckpt_to_load_from='london/*ckpt-1000*',
    tensor_name_in_ckpt='dayhr_embed',
    trainable=False
)
```


Where does the feature engineering code fit in?

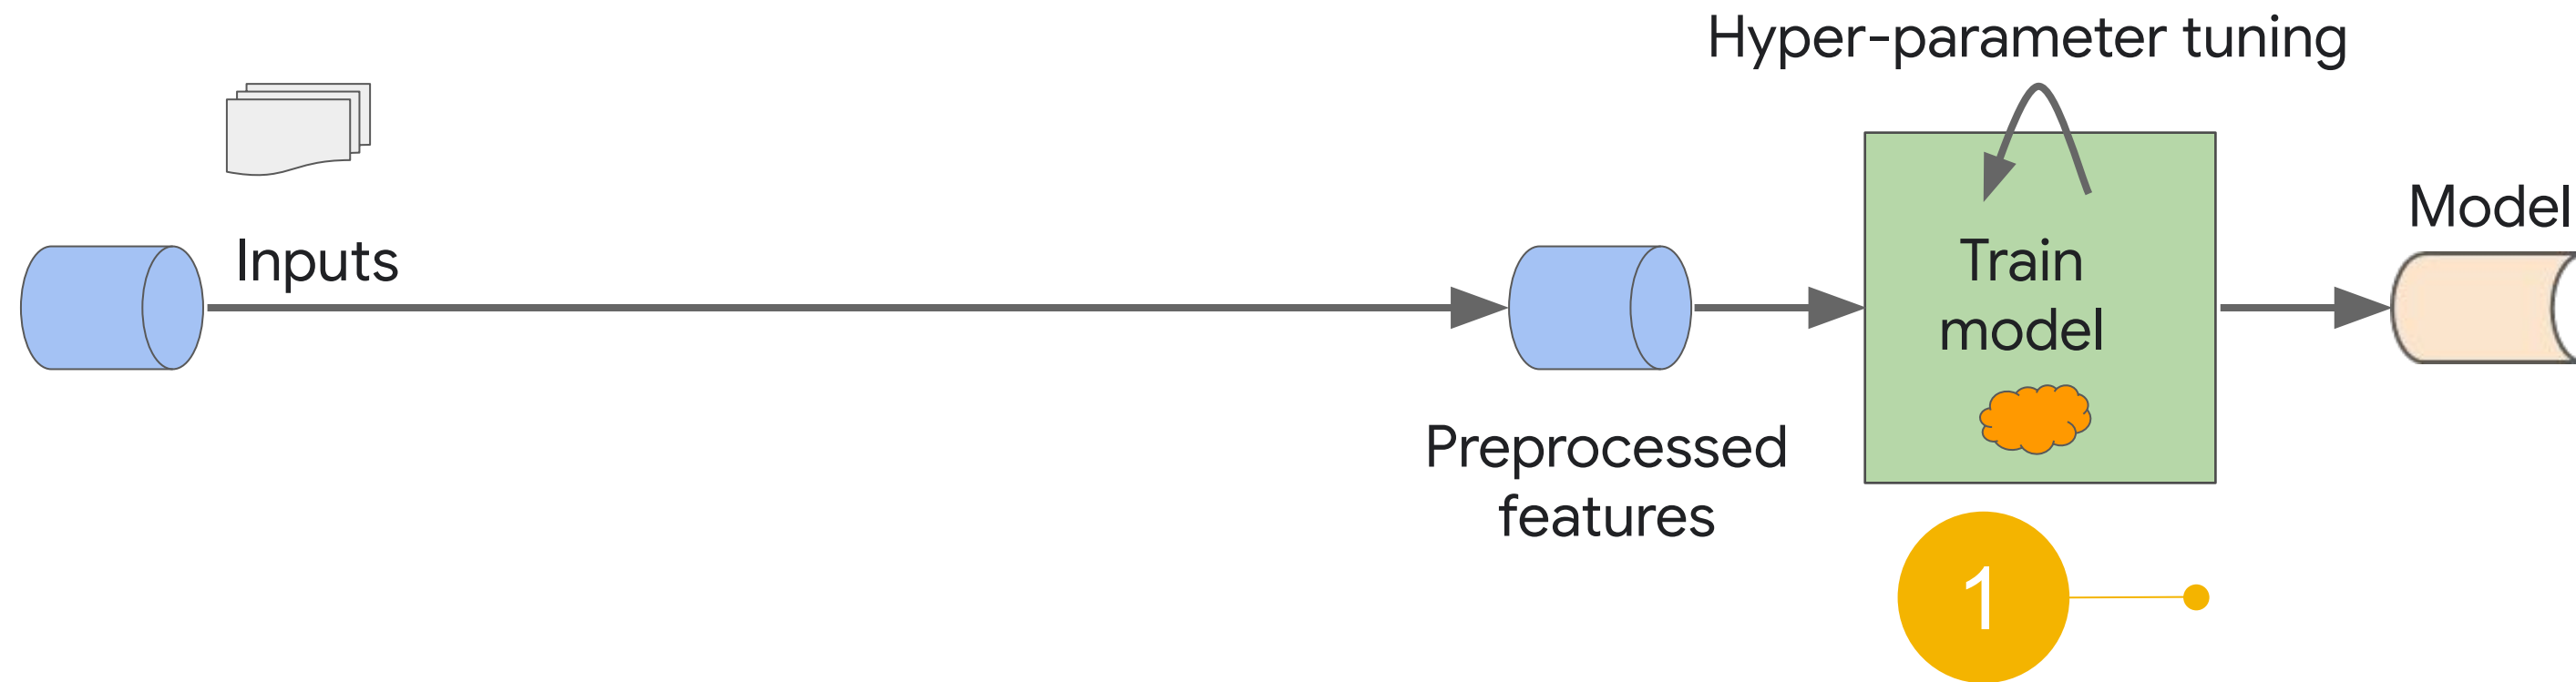
```
def train_input_fn(file_prefix):  
    ...  
    return features, labels  
  
featcols = [  
    fc.numeric_column("sq_footage"),  
    fc.categorical_column_with_vocabulary_list(  
        "type", ["house", "apt"])  
]  
  
model = tf.estimator.LinearRegressor(featcols)  
  
train_spec, eval_spec = ...  
  
model.train_and_evaluate(train_spec, ...)
```

Three possible places to do feature engineering



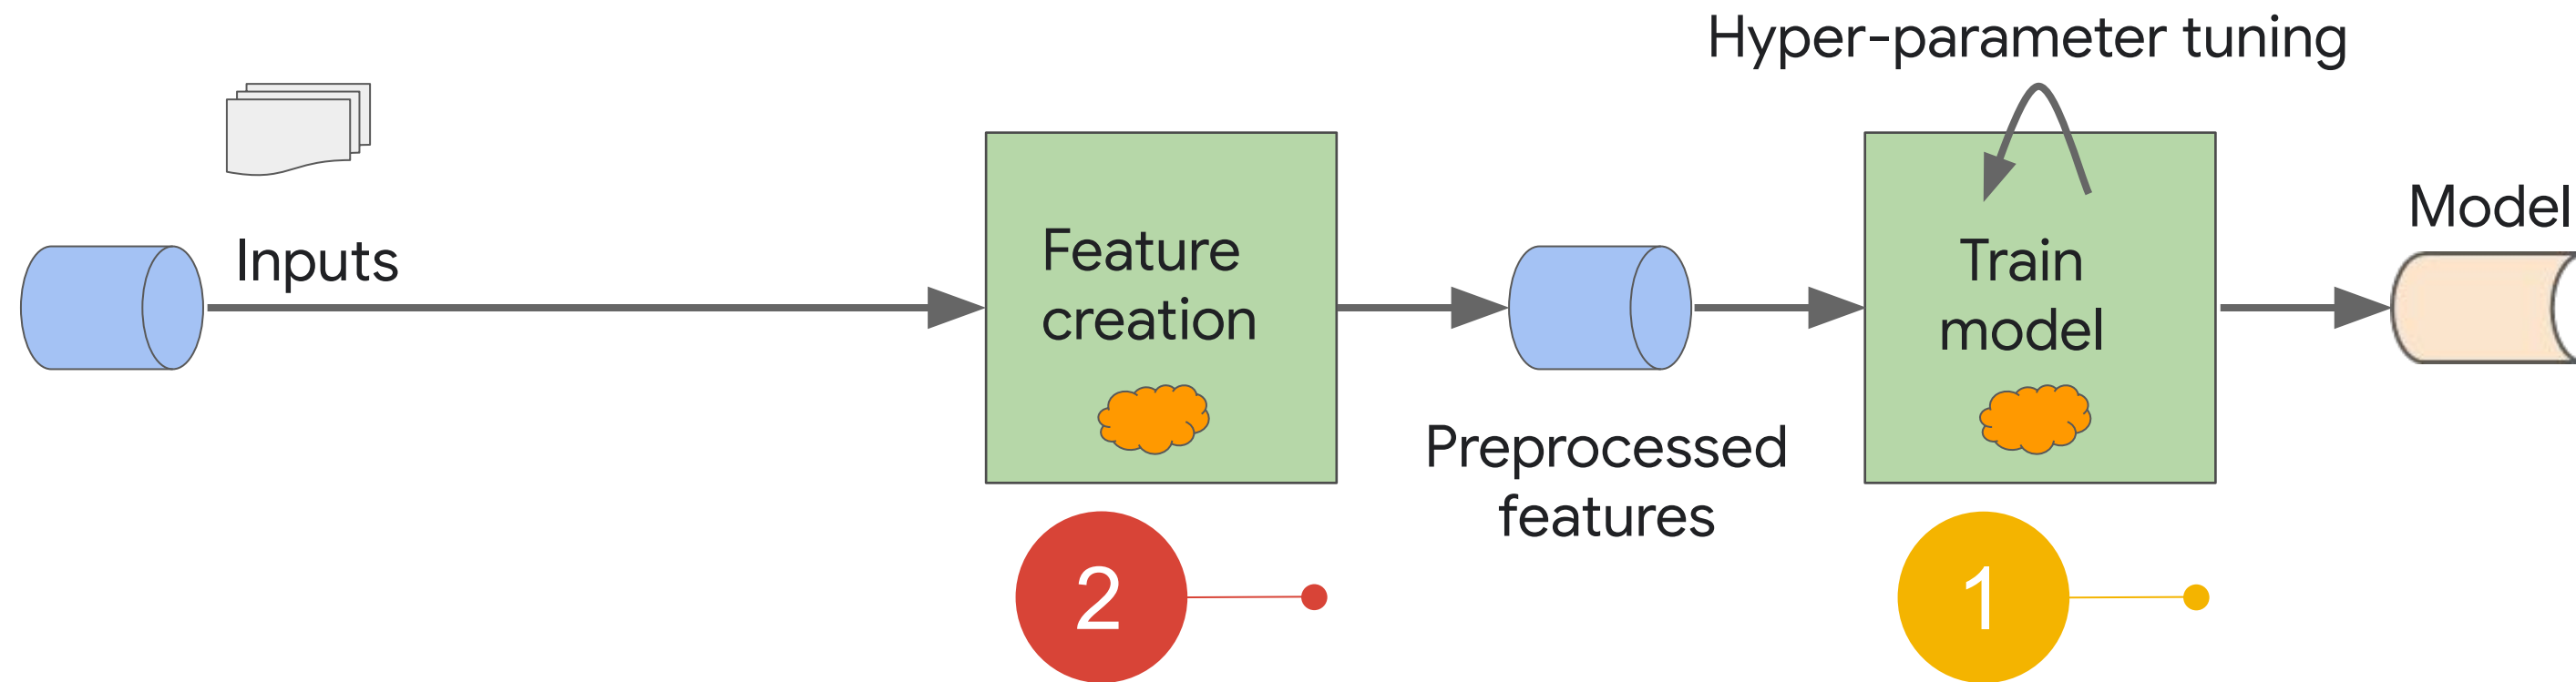
TensorFlow
feature_column
input_fn

Three possible places to do feature engineering



TensorFlow
feature_column
input_fn

Three possible places to do feature engineering



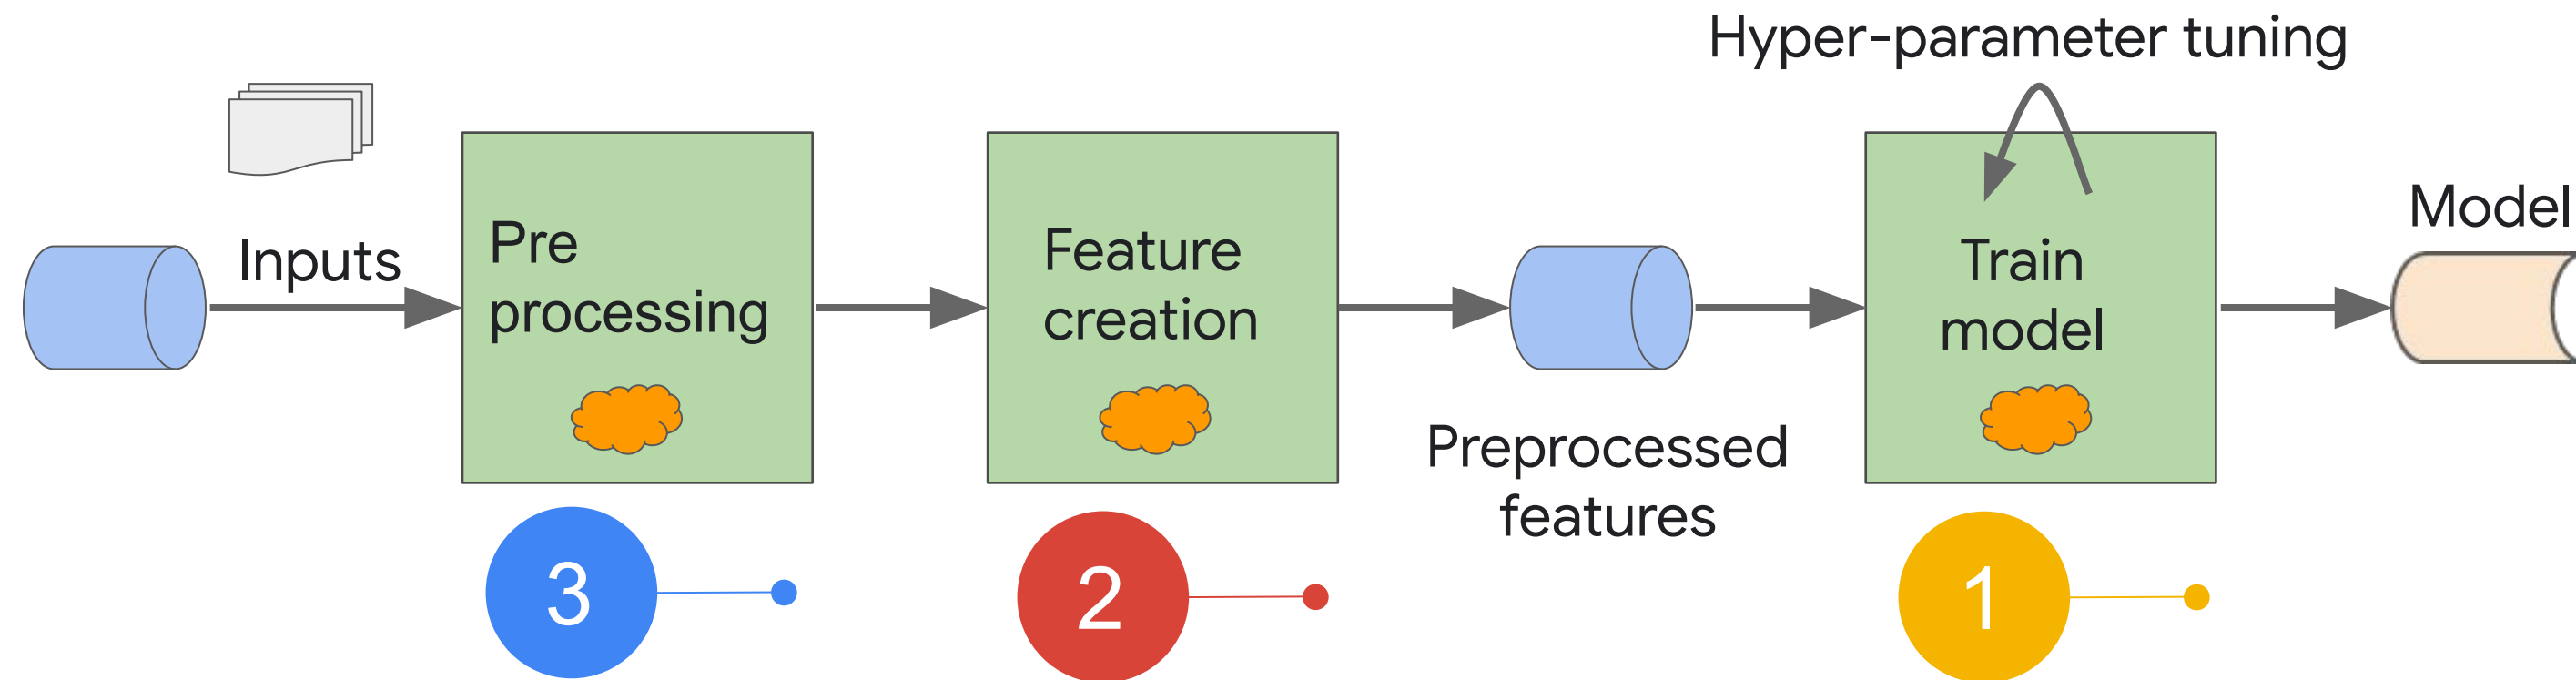
Dataflow

*If Dataflow is part
of your prediction
runtime also

TensorFlow

feature_column
input_fn

Three possible places to do feature engineering



Dataflow +
TensorFlow
(`tf.transform`)

Dataflow
**If Dataflow is part
of your prediction
runtime also*

TensorFlow
`feature_column`
`input_fn`

Some preprocessing can be done in `tf.feature_column`

1

```
def train_input_fn(file_prefix):
    ...
    return features, labels

featcols = [
    fc.numeric_column("sq_footage"),
    fc.categorical_column_with_vocabulary_list(
        "type", ["house", "apt"])
]

featcols.append(
    fc.bucketized_column(featcols[0],
        [500, 1000, 2500]))

model = tf.estimator.LinearRegressor(featcols)

train_spec, eval_spec = ...
model.train_and_evaluate(train_spec, ...)
```

Some preprocessing can be done in `tf.feature_column`

1

```
def train_input_fn(file_prefix):
    ...
    return features, labels

featcols = [
    fc.numeric_column("sq_footage"),
    fc.categorical_column_with_vocabulary_list(
        "type", ["house", "apt"])
]

featcols[0] = (
    fc.bucketized_column(featcols[0],
        [500, 1000, 2500]))

model = tf.estimator.LinearRegressor(featcols)

train_spec, eval_spec = ...
model.train_and_evaluate(train_spec, ...)
```

Powerful preprocessing can be done in TensorFlow

by creating a new feature column

1

```
latbuckets = np.linspace(38.0, 42.0, nbuckets).tolist()  
lonbuckets = np.linspace(-76.0, -72.0, nbuckets).tolist()
```

Preprocessing is part of the model graph,
so “automatic” at prediction time

Powerful preprocessing can be done in TensorFlow

by creating a new feature column

1

```
latbuckets = np.linspace(38.0, 42.0, nbuckets).tolist()
lonbuckets = np.linspace(-76.0, -72.0, nbuckets).tolist()

b_lat = fc.bucketized_column(house_lat, latbuckets)
b_lon = fc.bucketized_column(house_lon, lonbuckets)
```

Preprocessing is part of the model graph,
so “automatic” at prediction time

Powerful preprocessing can be done in TensorFlow

1 by creating a new feature column

```
latbuckets = np.linspace(38.0, 42.0, nbuckets).tolist()
lonbuckets = np.linspace(-76.0, -72.0, nbuckets).tolist()

b_lat = fc.bucketized_column(house_lat, latbuckets)
b_lon = fc.bucketized_column(house_lon, lonbuckets)

# feature cross and embed
loc = fc.crossed_column([b_lat, b_lon], nbuckets*nbuckets)
```

Preprocessing is part of the model graph,
so “automatic” at prediction time

Powerful preprocessing can be done in TensorFlow

by creating a new feature column

1

```
latbuckets = np.linspace(38.0, 42.0, nbuckets).tolist()
lonbuckets = np.linspace(-76.0, -72.0, nbuckets).tolist()

b_lat = fc.bucketized_column(house_lat, latbuckets)
b_lon = fc.bucketized_column(house_lon, lonbuckets)

# feature cross and embed
loc = fc.crossed_column([b_lat, b_lon], nbuckets*nbuckets)

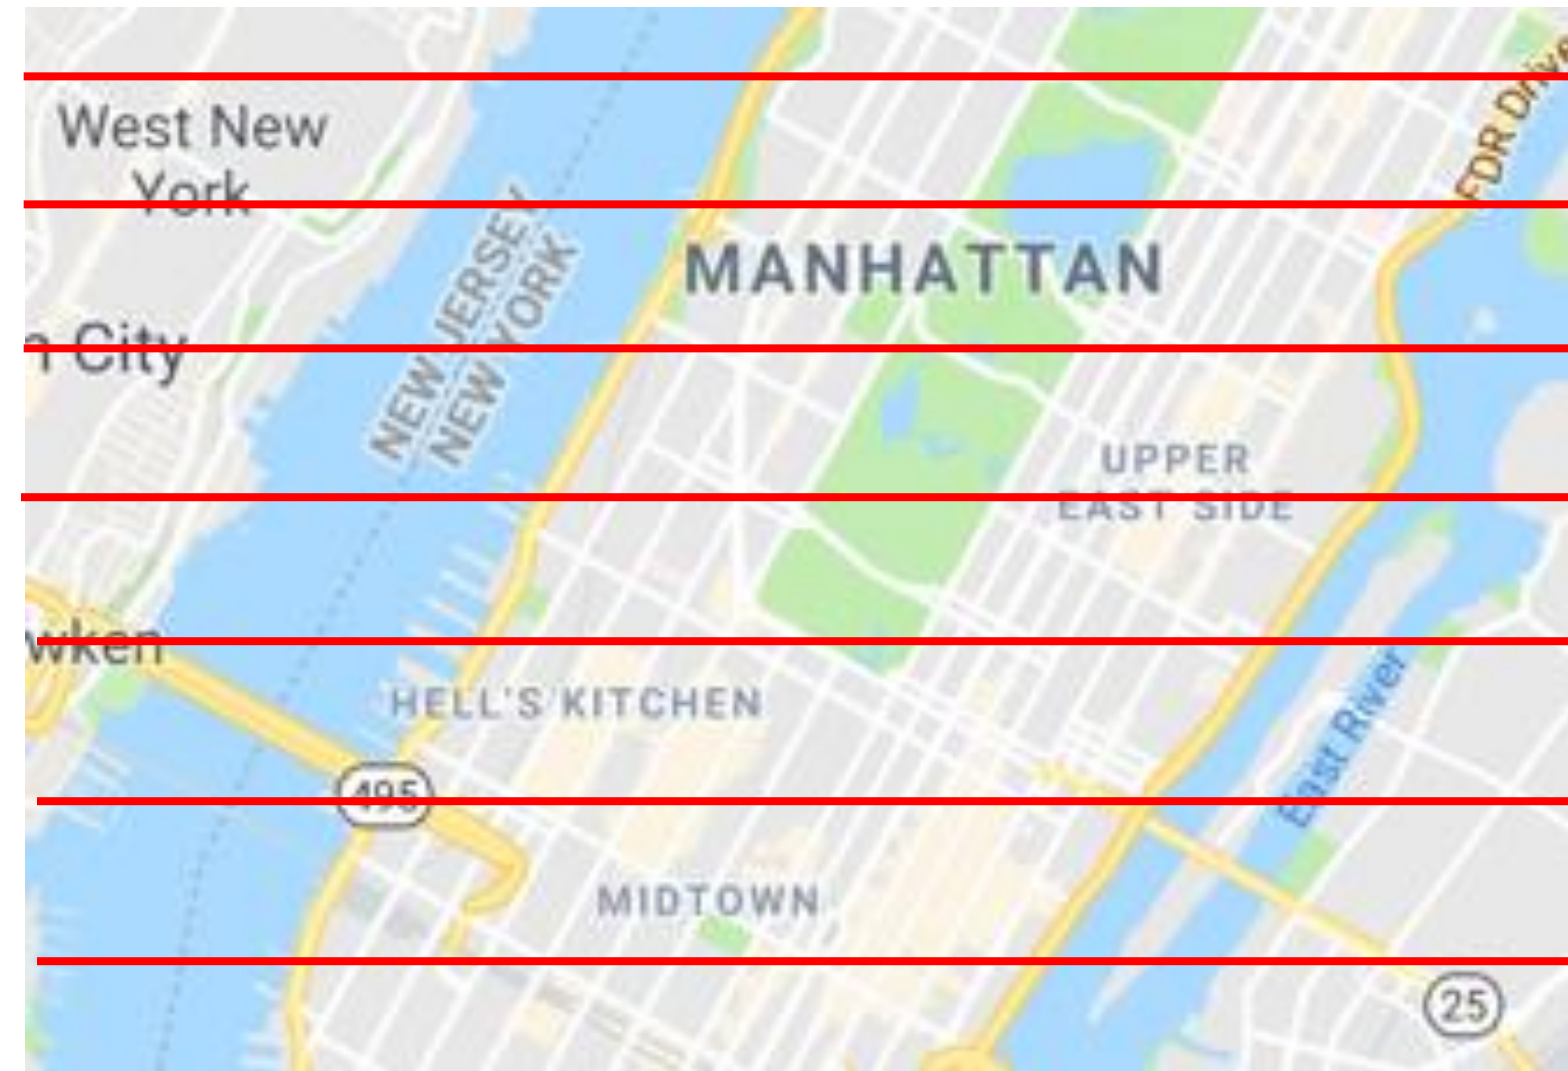
eloc = fc.embedding_column(loc, nbuckets//4)
```

Preprocessing is part of the model graph,
so “automatic” at prediction time

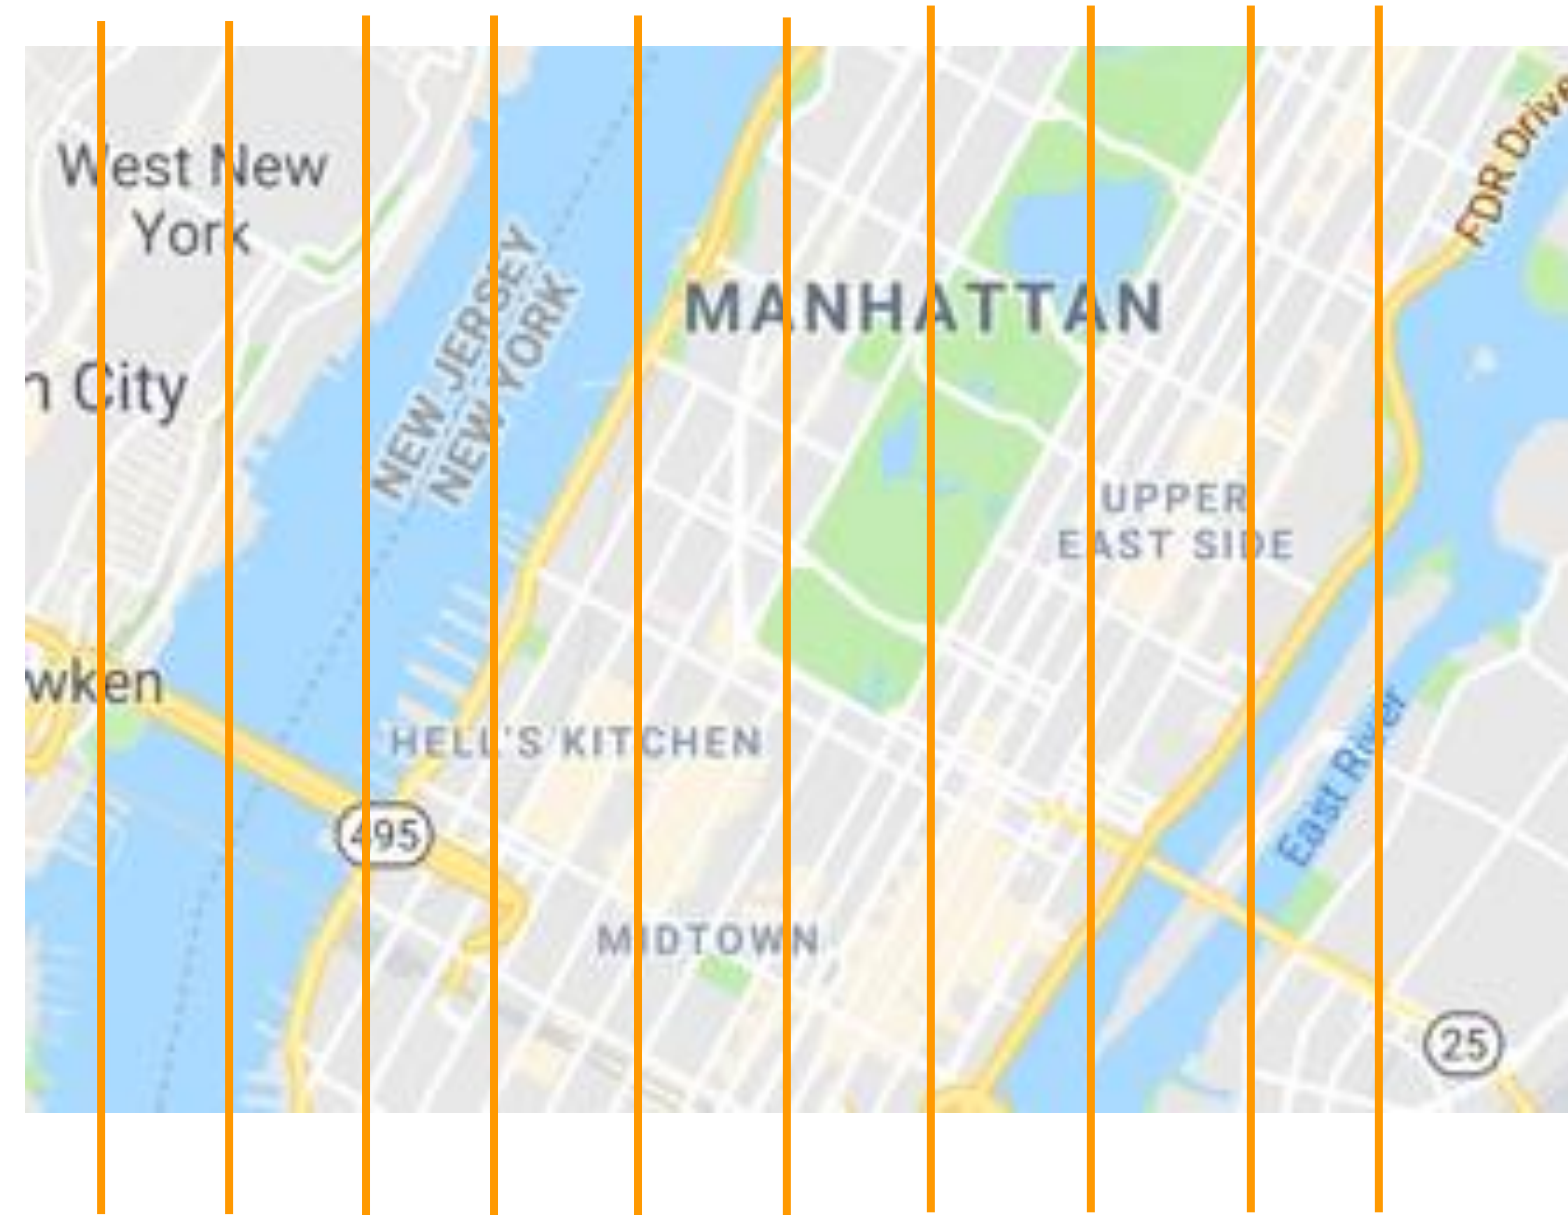
What feature crossing and embedding end up doing




What feature crossing and embedding end up doing




What feature crossing and embedding end up doing




What feature crossing and embedding end up doing



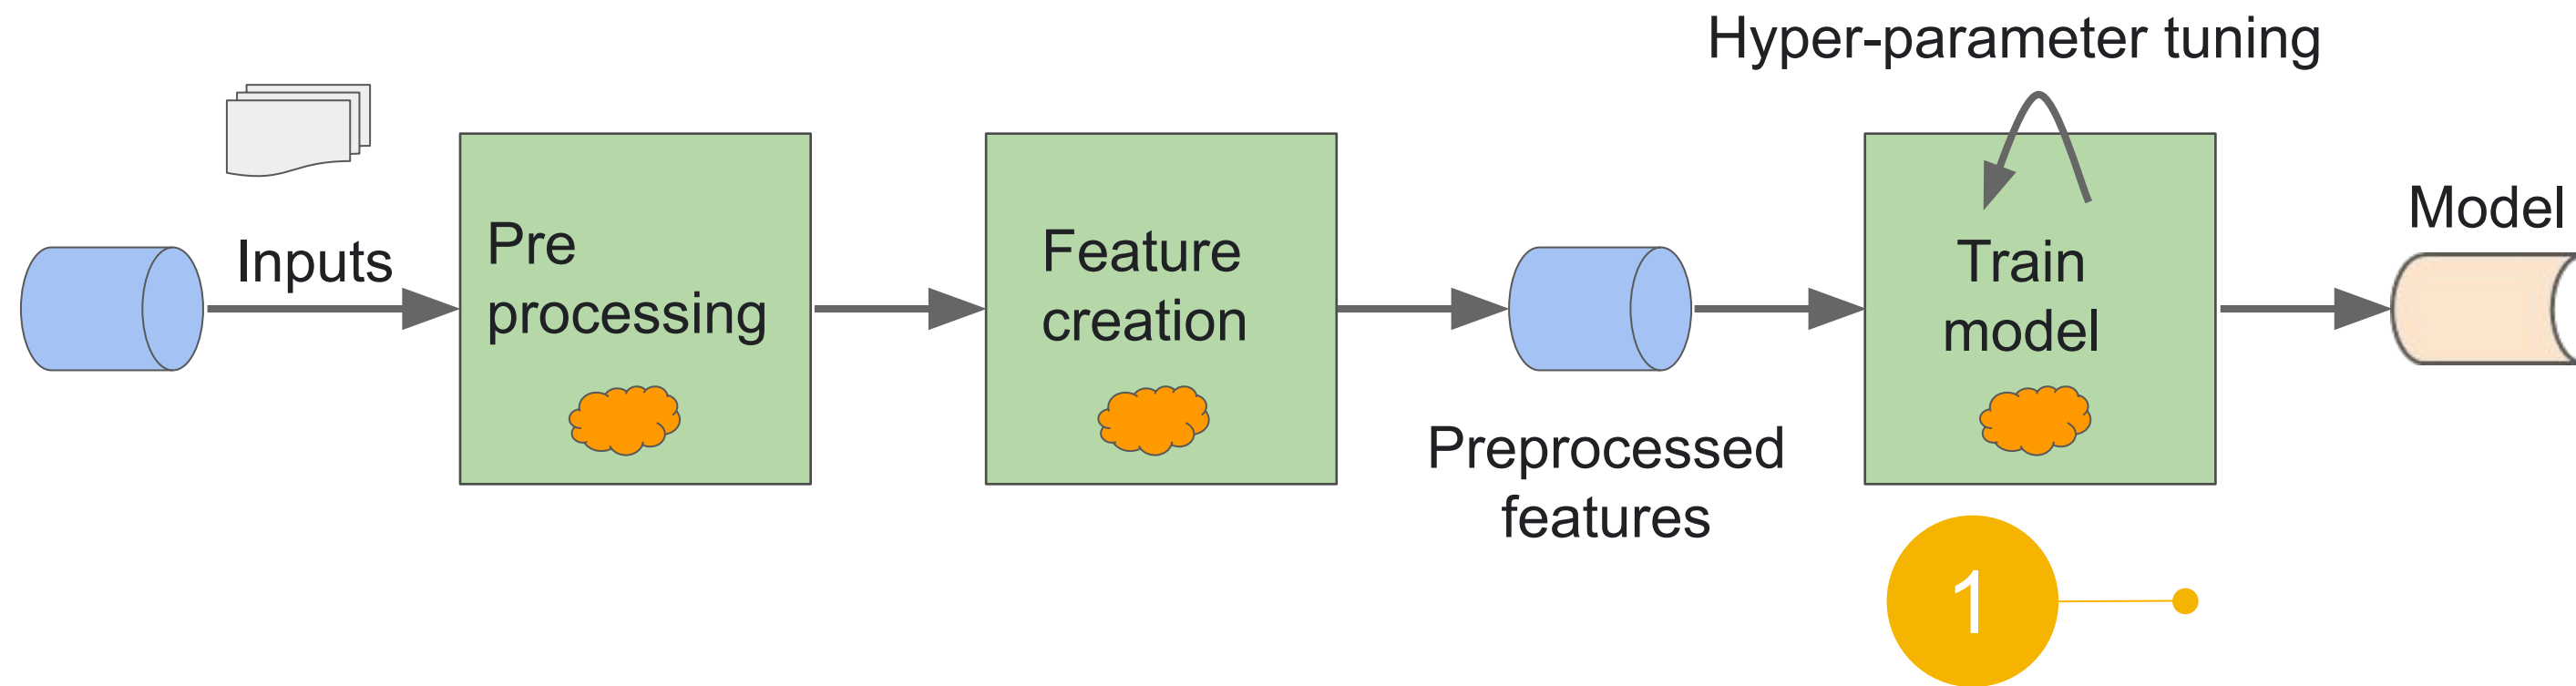
What feature crossing and embedding end up doing



What feature crossing and embedding end up doing



Three possible places to do feature engineering



Recall that the input function returns features and labels

1

```
def  
train_input_fn(file_prefix):  
    ...  
    return features, labels
```

What is the data type of
features?

Create new features from existing features in TensorFlow

```
def add_engineered(features):  
    lat1 = features['lat']  
    lat2 = features['metro_lat']  
    latdiff = lat1-lat2  
    ...  
    dist = tf.sqrt(latdiff*latdiff + londiff*londiff)  
    features['euclidean'] = dist  
    return features
```



Create new features from existing features in TensorFlow

```
def add_engineered(features):  
    lat1 = features['lat']  
    lat2 = features['metro_lat']  
    latdiff = lat1-lat2  
    ...  
    dist = tf.sqrt(latdiff*latdiff + londiff*londiff)  
    features['euclidean'] = dist  
    return features
```



Call the add_engineered method from all input functions

```
def add_engineered(features):  
    ...  
    features['euclidean'] = dist  
    return features
```

Call the `add_engineered` method from all input functions

```
def add_engineered(features):  
    ...  
    features['euclidean'] = dist  
    return features
```

```
def train_input_fn():  
    ...  
    features = ...  
    return add_engineered(features), label
```

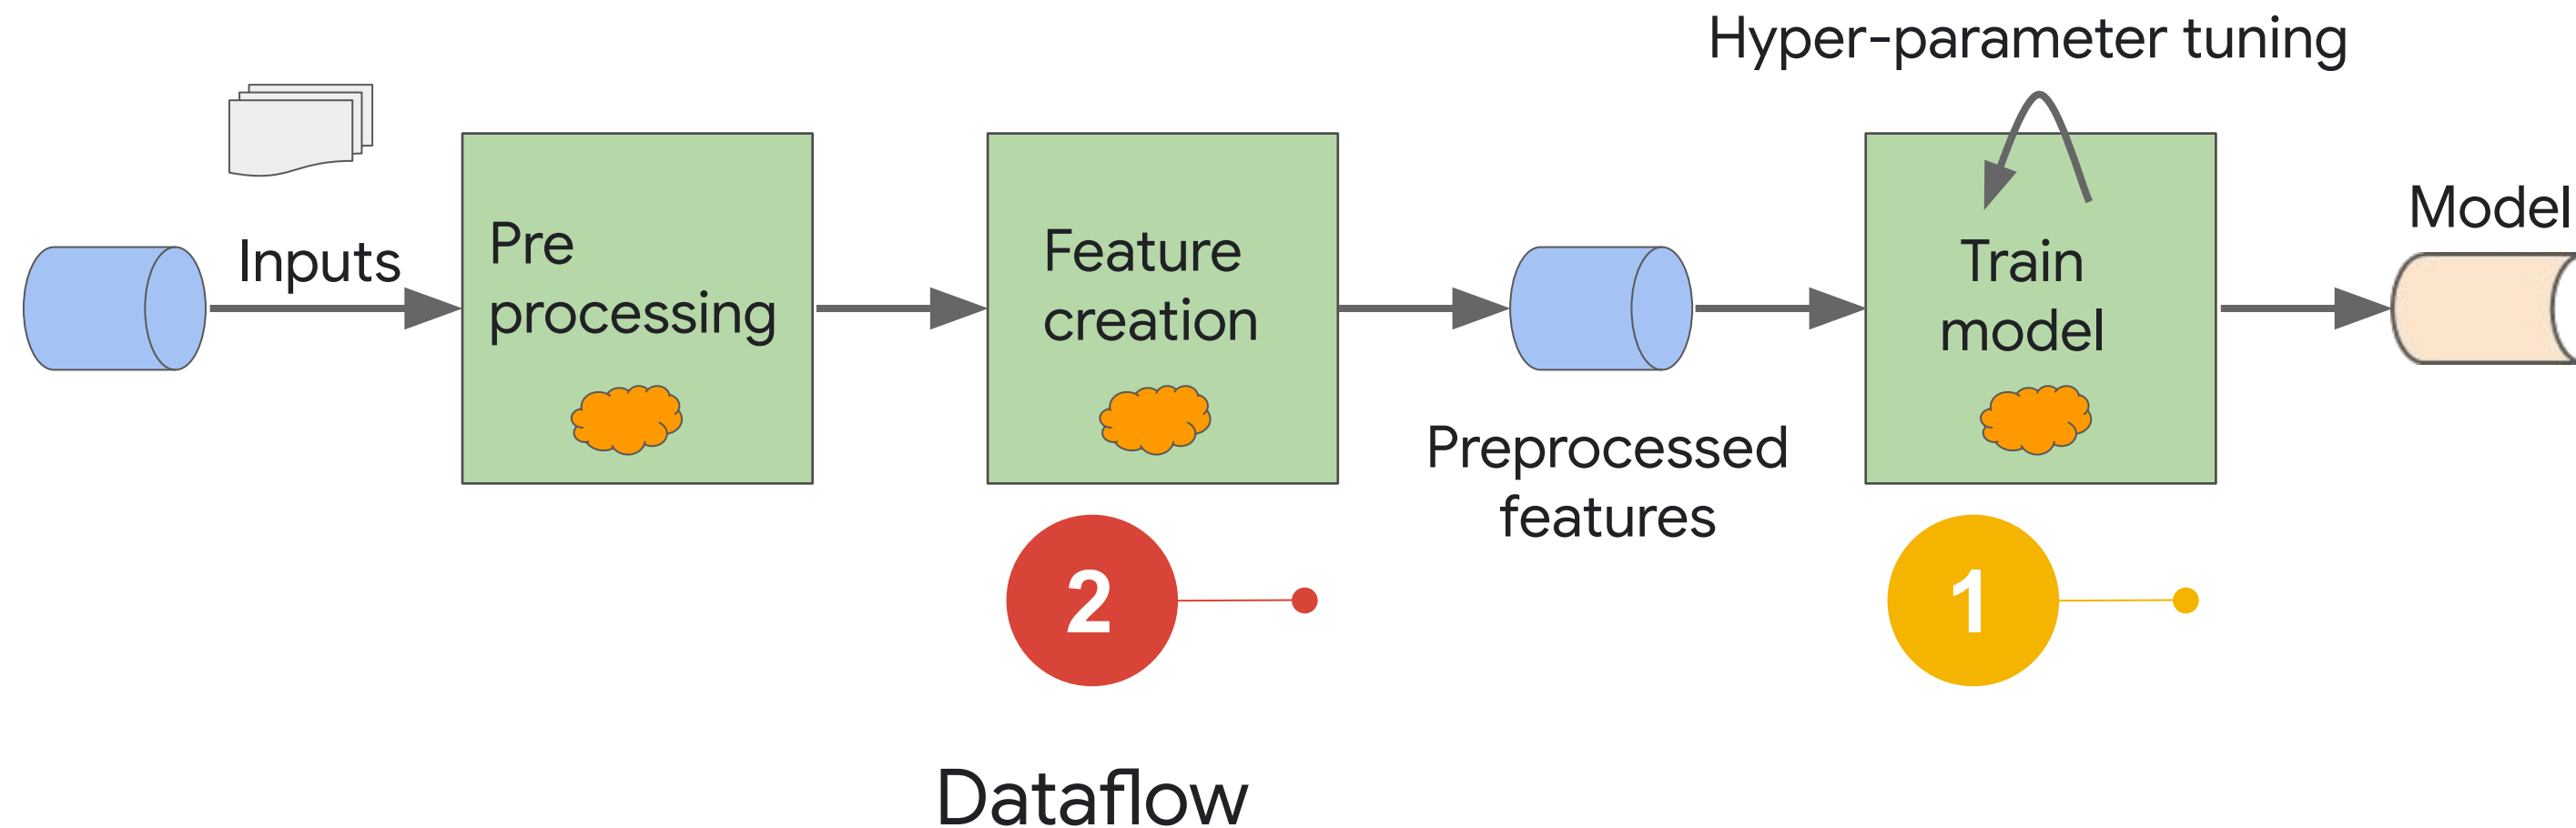
Call the `add_engineered` method from all input functions

```
def add_engineered(features):  
    ...  
    features['euclidean'] = dist  
    return features
```

```
def train_input_fn():  
    ...  
    features = ...  
    return add_engineered(features), label
```

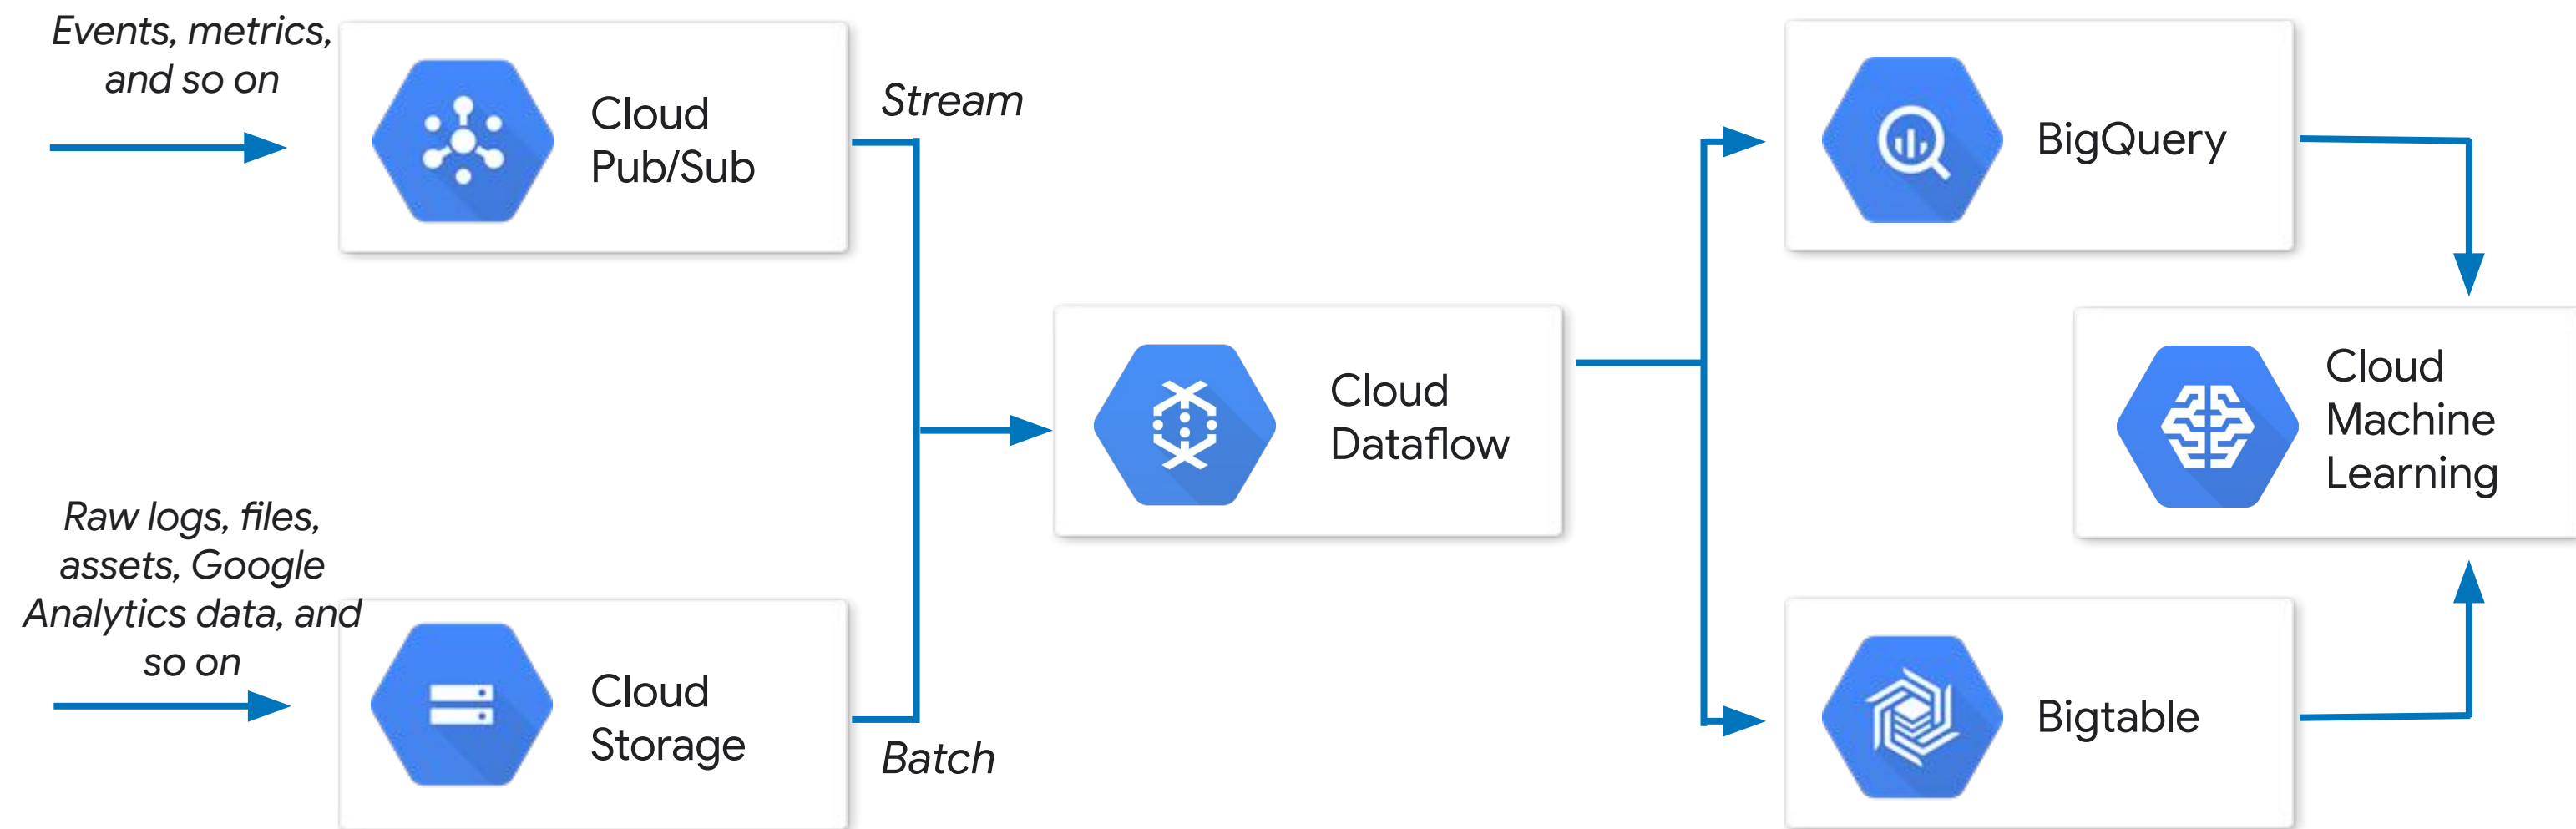
```
def serving_input_fn():  
    ...  
    return ServingInputReceiver(  
        add_engineered(features),  
        json_features_ph)
```

Three possible places to do feature engineering

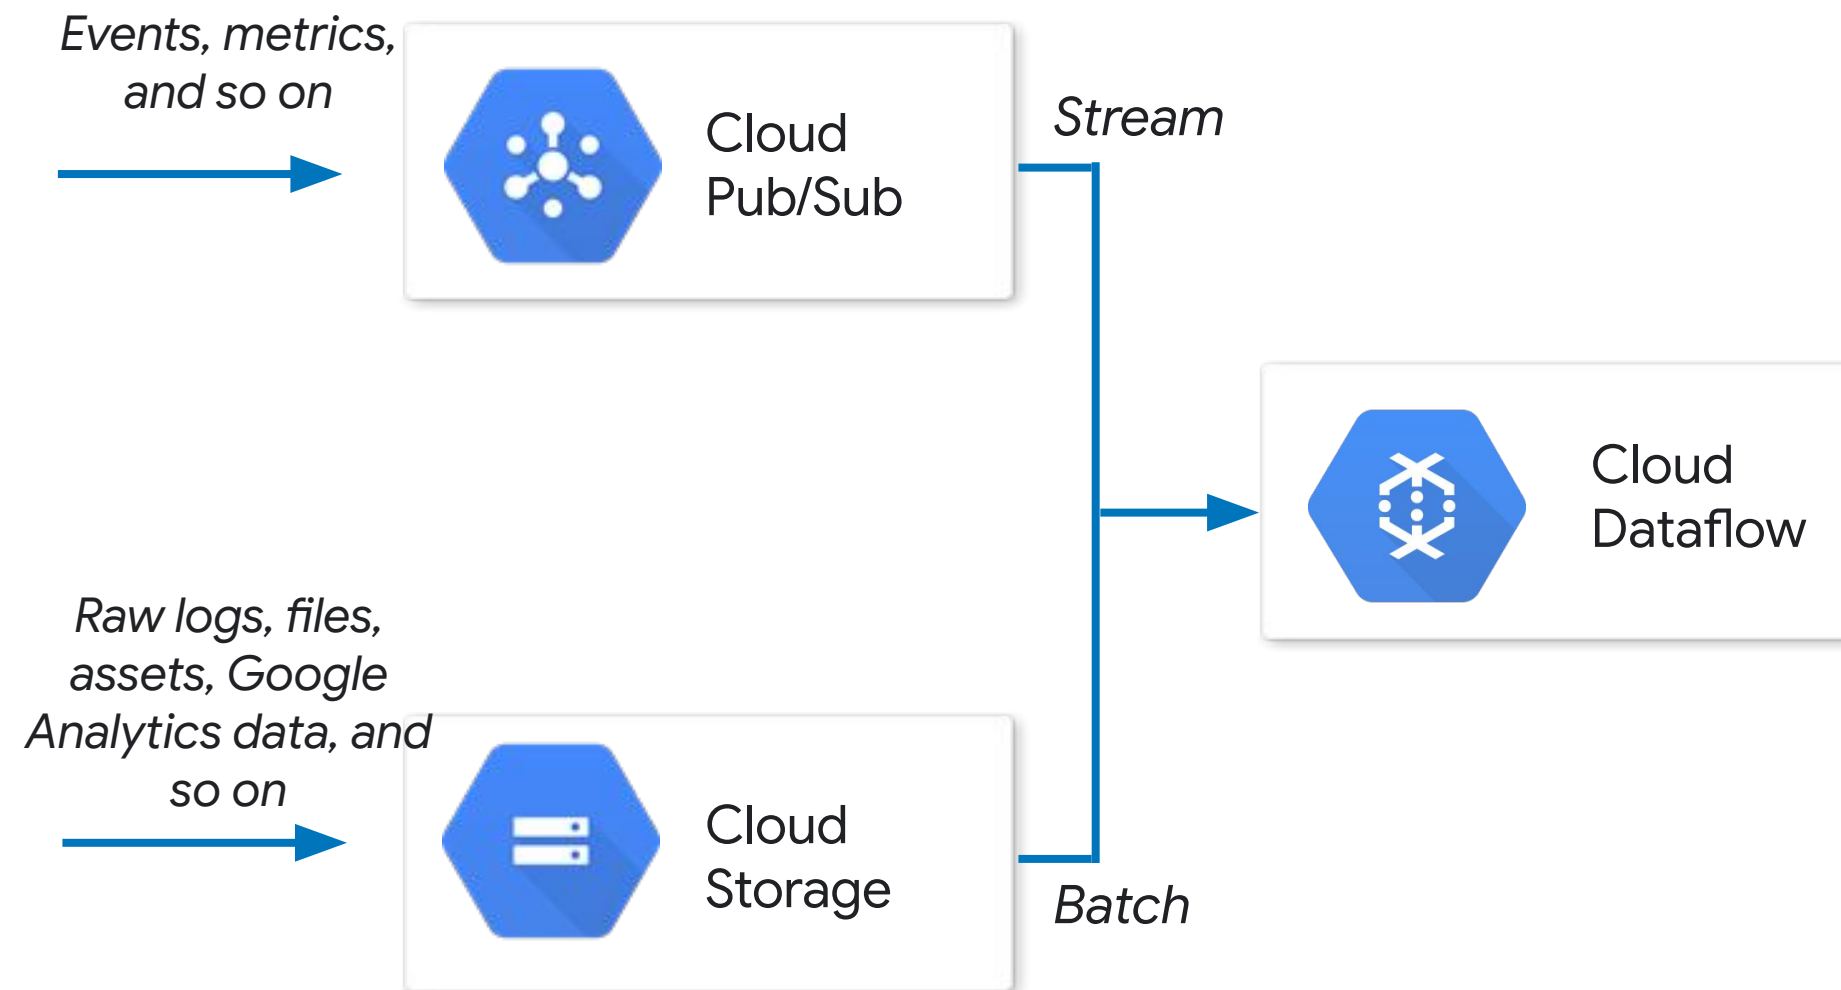


*If Dataflow is part of your prediction runtime also

Recall that the reference architecture for GCP involves Dataflow in both the training and prediction pipeline



Dataflow is ideal for time-windowed aggregations




Adding new features in Dataflow is like any other PTransform

2

```
train = pipeline
| beam.io.Read(beam.io.Read(
    beam.io.BigQuerySource(query=query)))
| beam.FlatMap(add_fields) #'pastHrCount'
| beam.io.Write(...)
```

```
predictions = pipeline
| beam.io.ReadStringsFromPubSub(...)
| beam.FlatMap(add_fields) #'pastHrCount'
| ...
```

tf.transform

Preprocessing for Machine Learning with tf.Transform

Wednesday, February 22, 2017

Posted by Kester Tong, David Soergel, and Gus Katsiapis, Software Engineers

When applying machine learning to real world datasets, a lot of effort is required to preprocess data into a format suitable for standard machine learning models, such as neural networks. This preprocessing takes a variety of forms, from converting between formats, to tokenizing and stemming text and forming vocabularies, to performing a variety of numerical operations such as normalization.

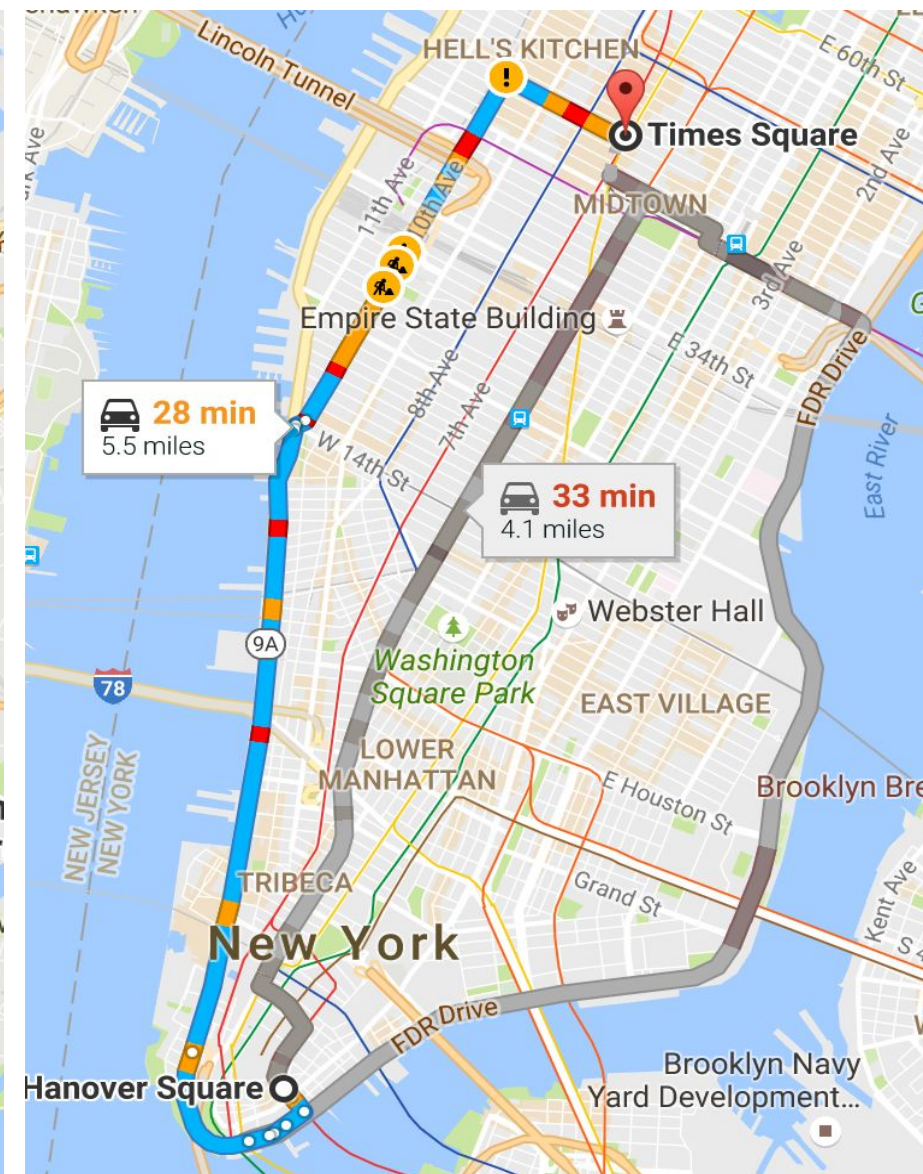
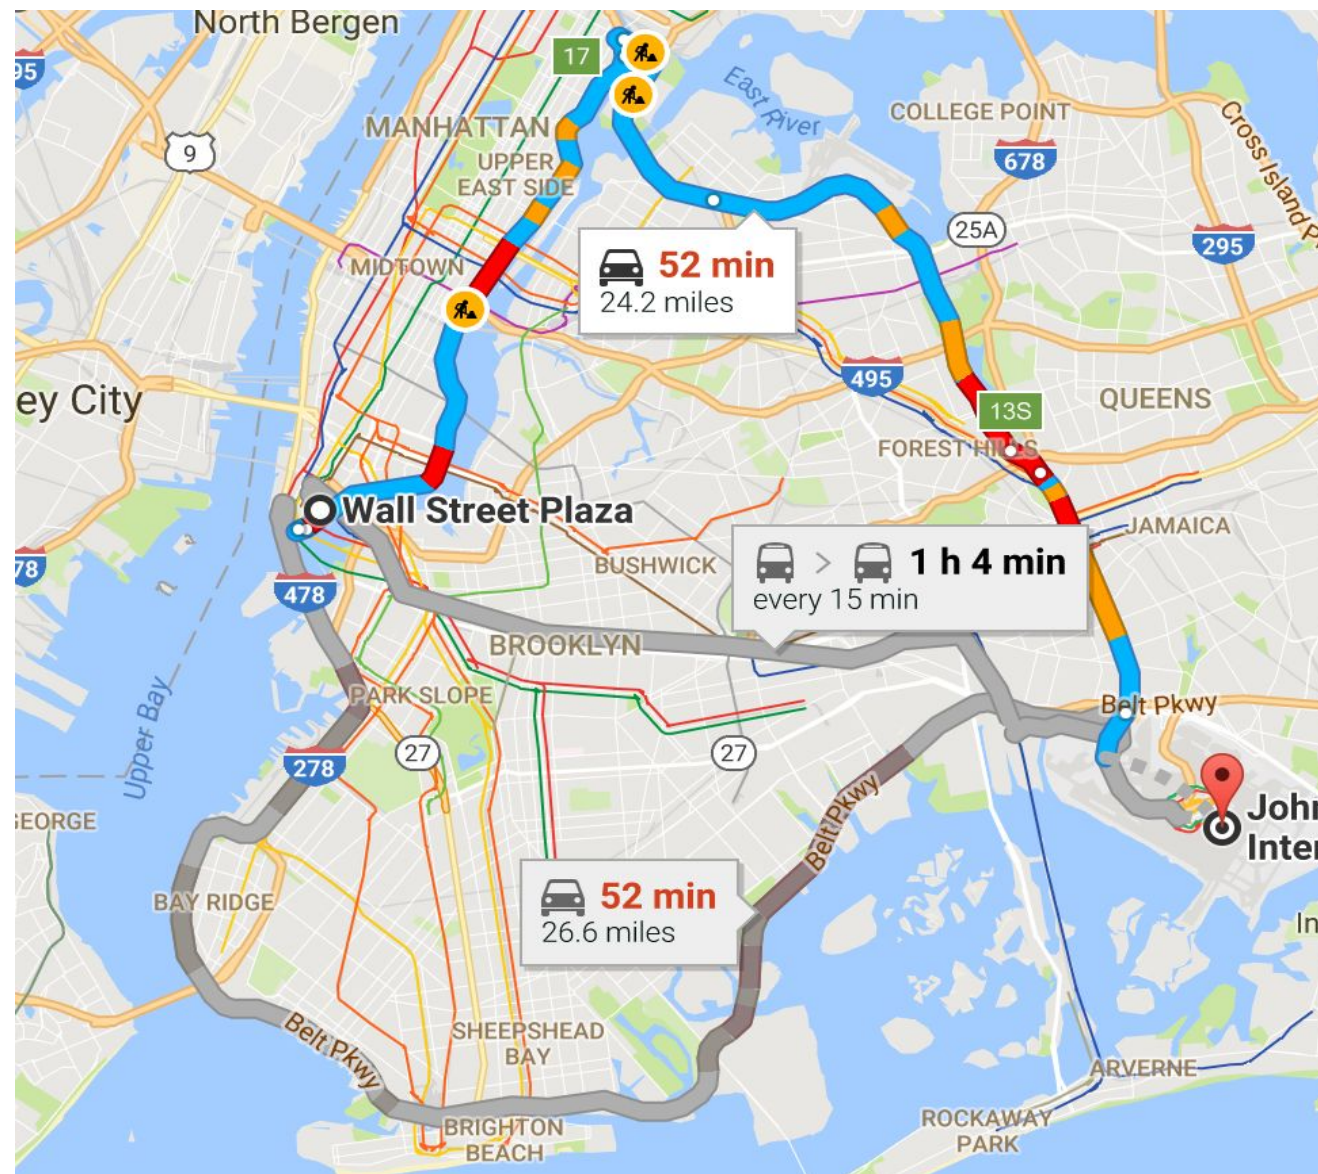
Today we are announcing [tf.Transform](#), a library for TensorFlow that allows users to define preprocessing pipelines and run these using large scale data processing frameworks, while also exporting the pipeline in a way that can be run as part of a TensorFlow graph. Users define a pipeline by composing modular Python functions, which tf.Transform then executes with [Apache Beam](#), a framework for large-scale, efficient, distributed data processing. Apache Beam pipelines can be run on [Google Cloud Dataflow](#) with planned support for running with [other frameworks](#). The TensorFlow graph exported by tf.Transform enables the preprocessing steps to be replicated when the trained model is used to make predictions, such as when serving the model with [Tensorflow Serving](#).

<https://research.googleblog.com/2017/02/preprocessing-for-machine-learning-with.html>

Lab

Improve ML model with
Feature Engineering

Goal: To estimate taxi fare



Taxi fares:

\$2.50 initial charge
+
50c per $\frac{1}{5}$ mile
(or)
50c per minute if stopped
+
Passenger pays tolls
+
Various special charges

http://www.nyc.gov/html/tlc/html/passenger/taxicab_rate.shtml

Lab: Improve ML model with Feature Engineering

In this lab, you will learn how to incorporate feature engineering into your pipeline.

- 1 Working with feature columns
- 2 Adding feature crosses in TensorFlow
- 3 Reading data from BigQuery
- 4 Creating datasets using Dataflow




Lab

Lab debrief


Screencast (Camtasia)

A question of ML Fairness ...


Is it fair to use feature crosses in the taxi-fare model?



Can the resolution of the feature cross of latitude & longitude amplify injustice?



Can the resolution of the feature cross of latitude & longitude amplify injustice?



cloud.google.com