



TensorFlow Transform

Lak Lakshmanan

Learn how to...

Learn how to...

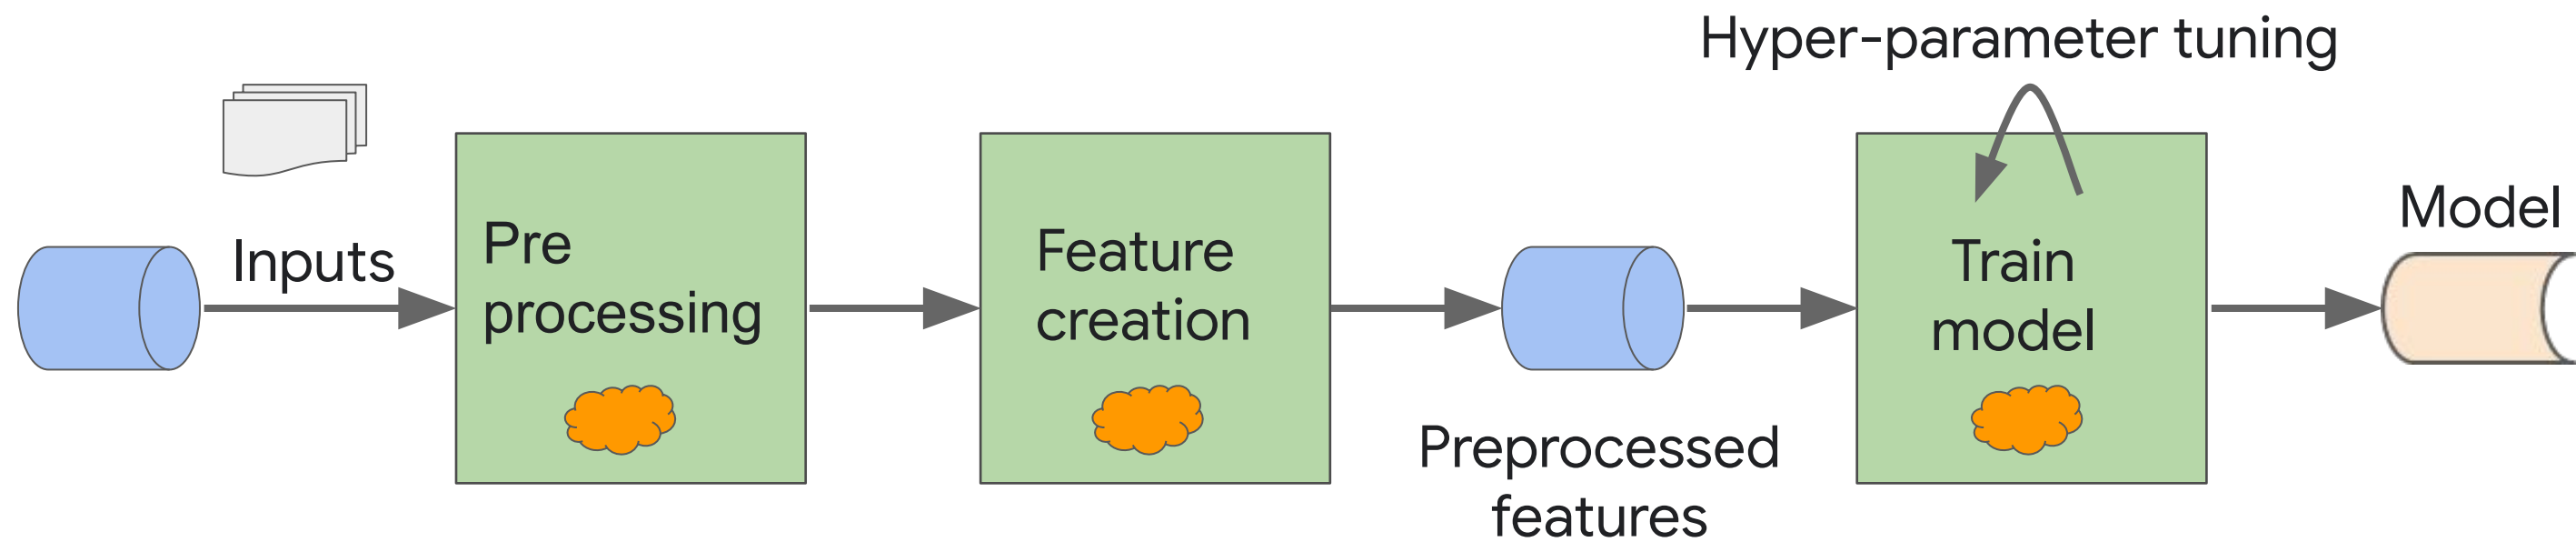
Implement feature preprocessing and
feature creation using `tf.transform`

Learn how to...

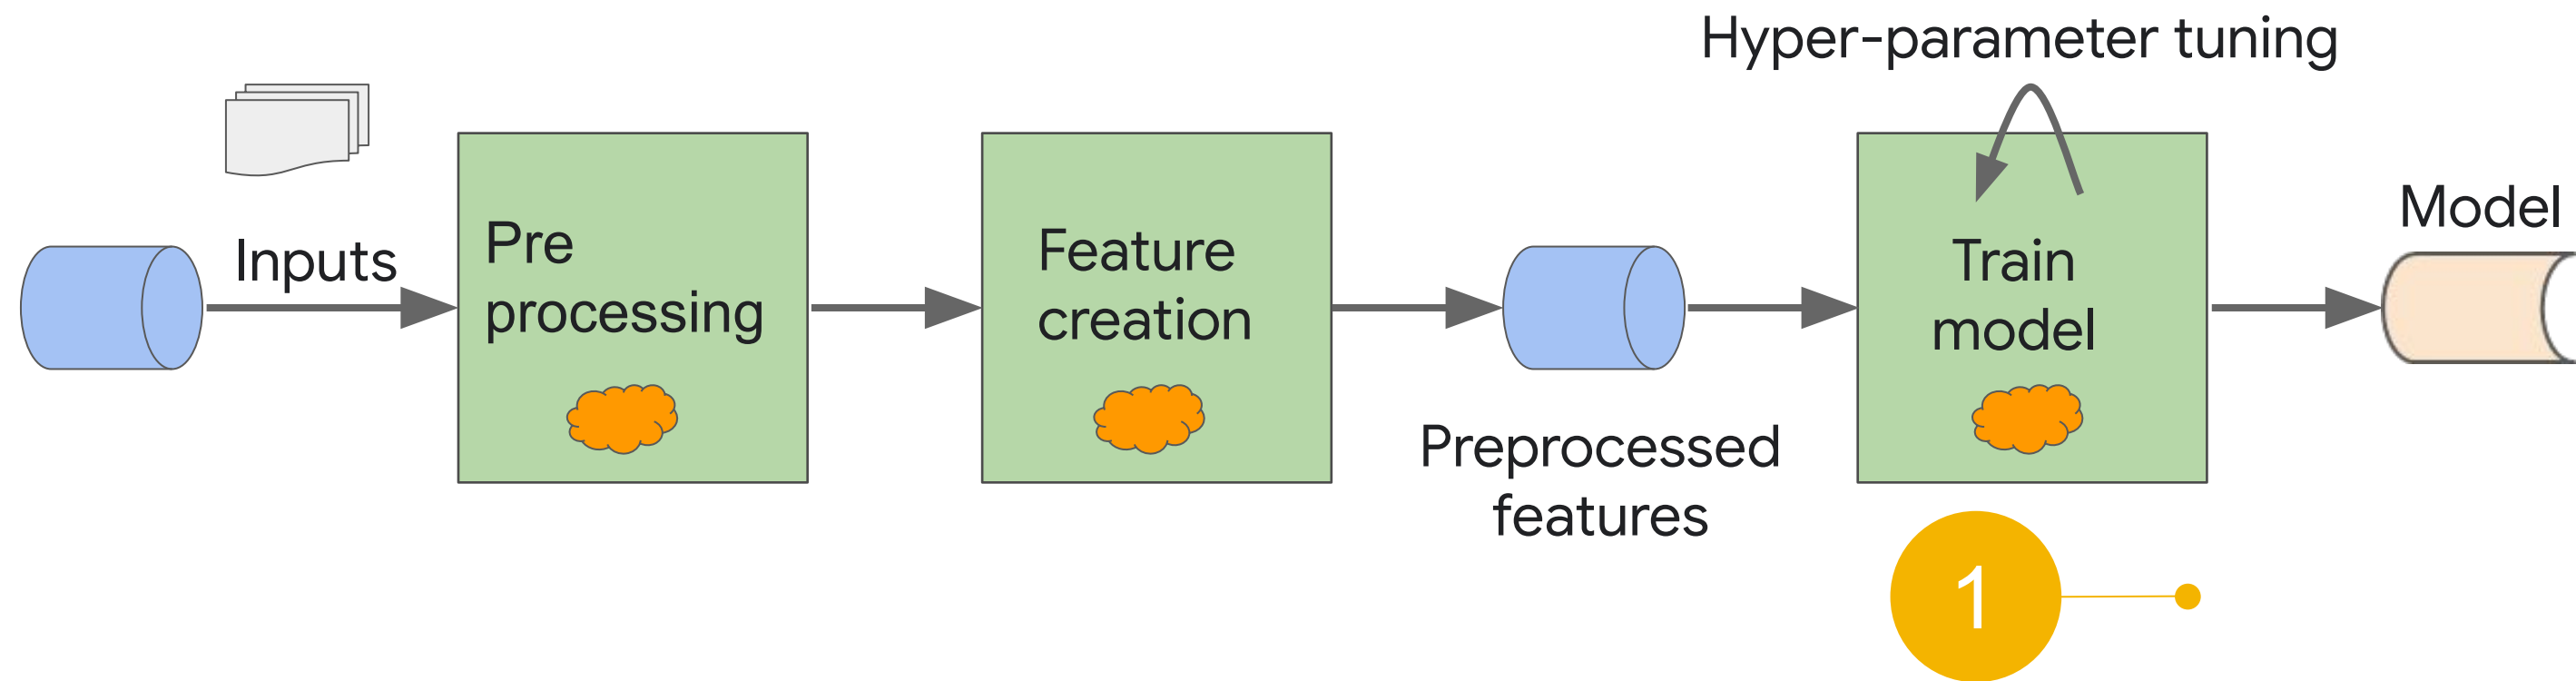
Implement feature preprocessing and feature creation using `tf.transform`

Carry out feature processing efficiently, at scale and on streaming data

Recall that there are three possible places to do feature engineering, each of which has its pros and cons



Recall that there are three possible places to do feature engineering, each of which has its pros and cons



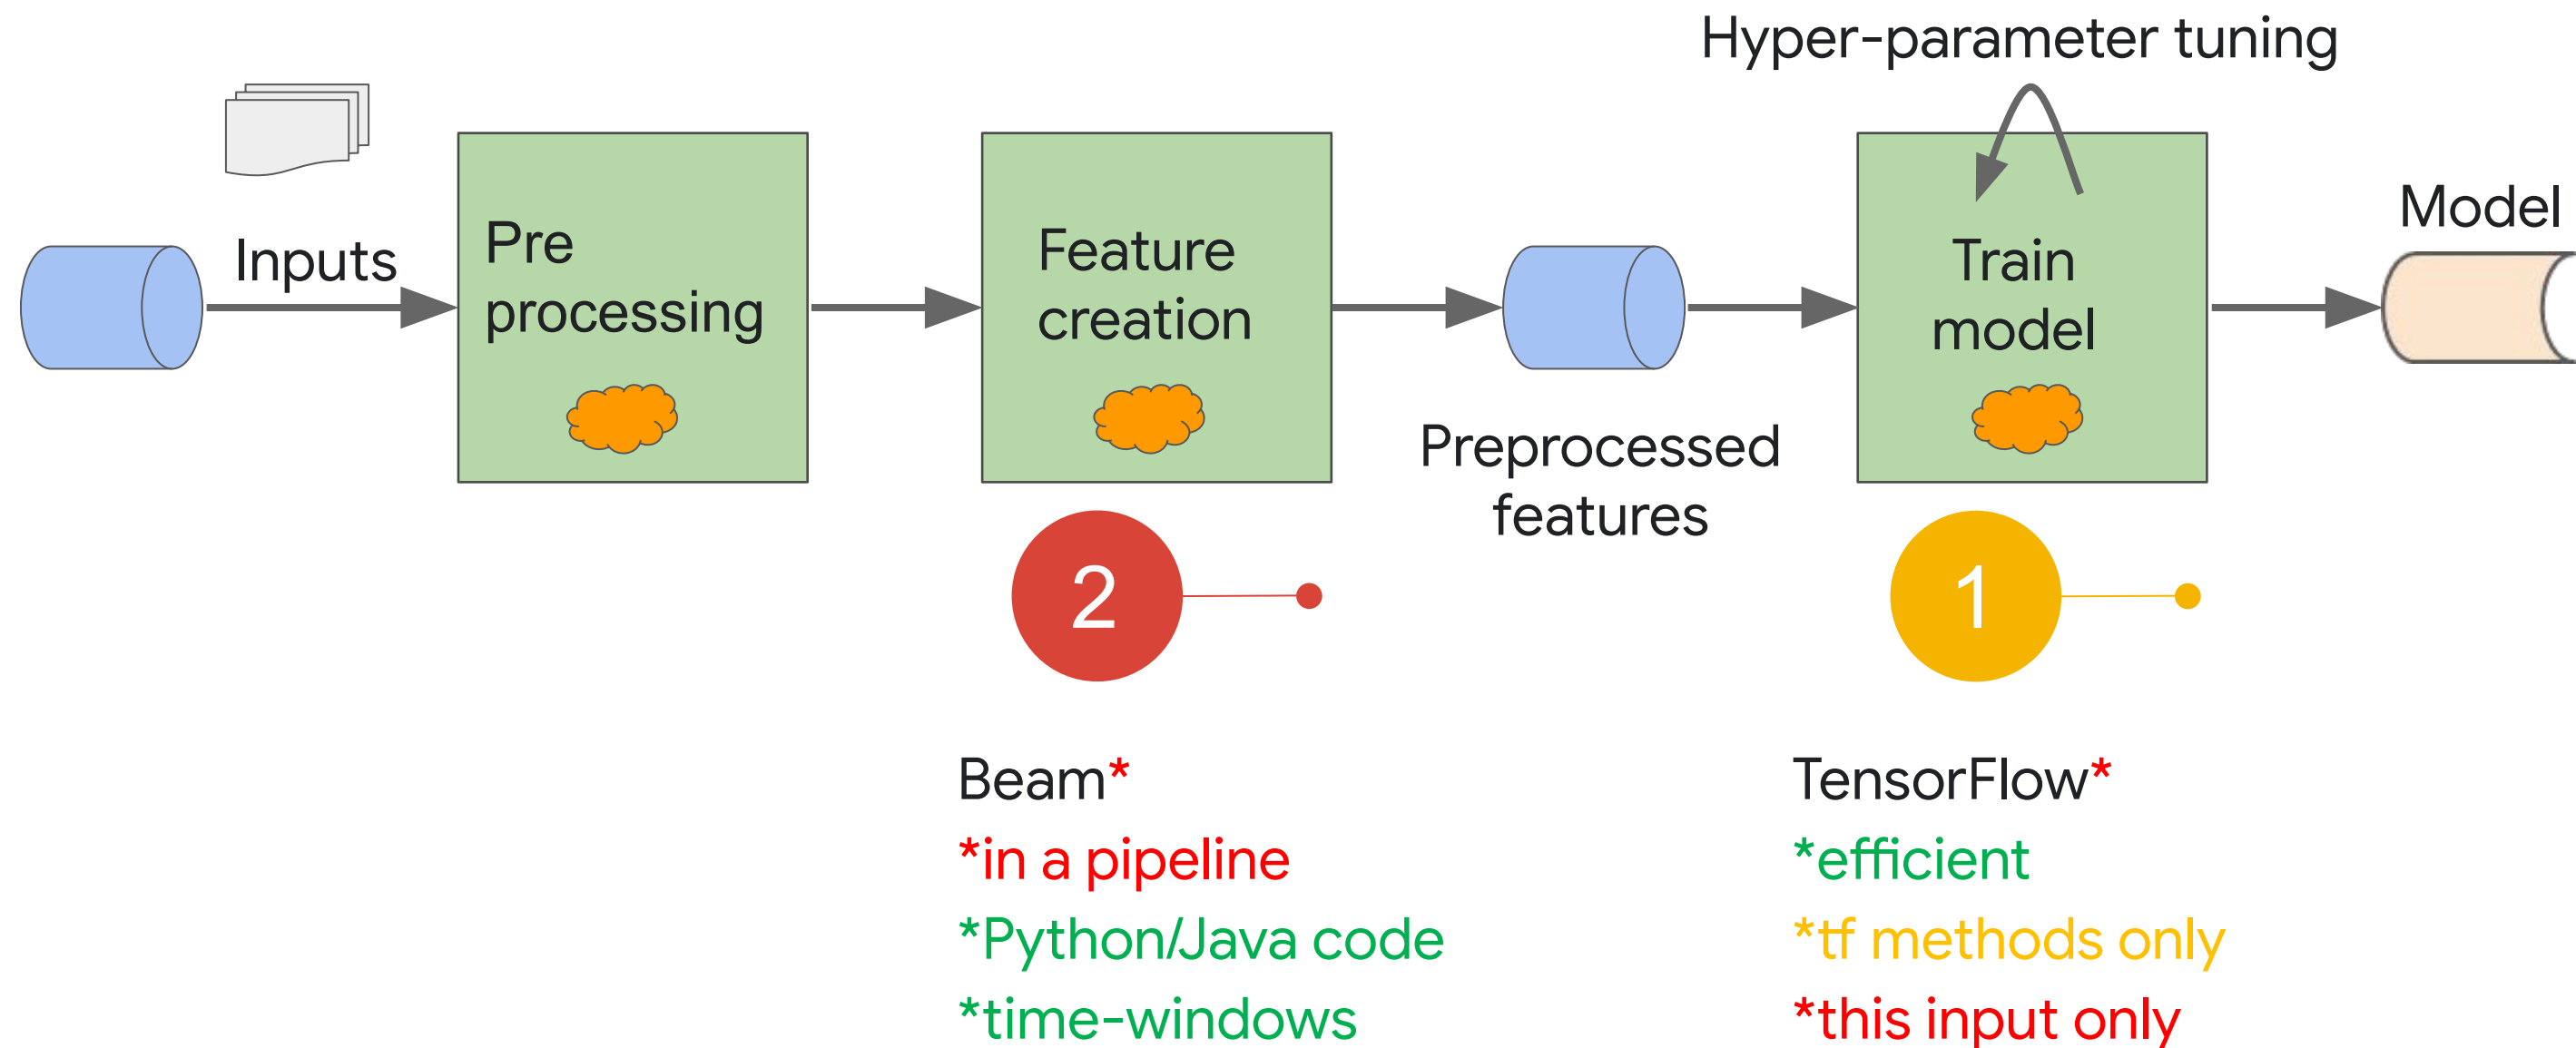
TensorFlow*

*efficient

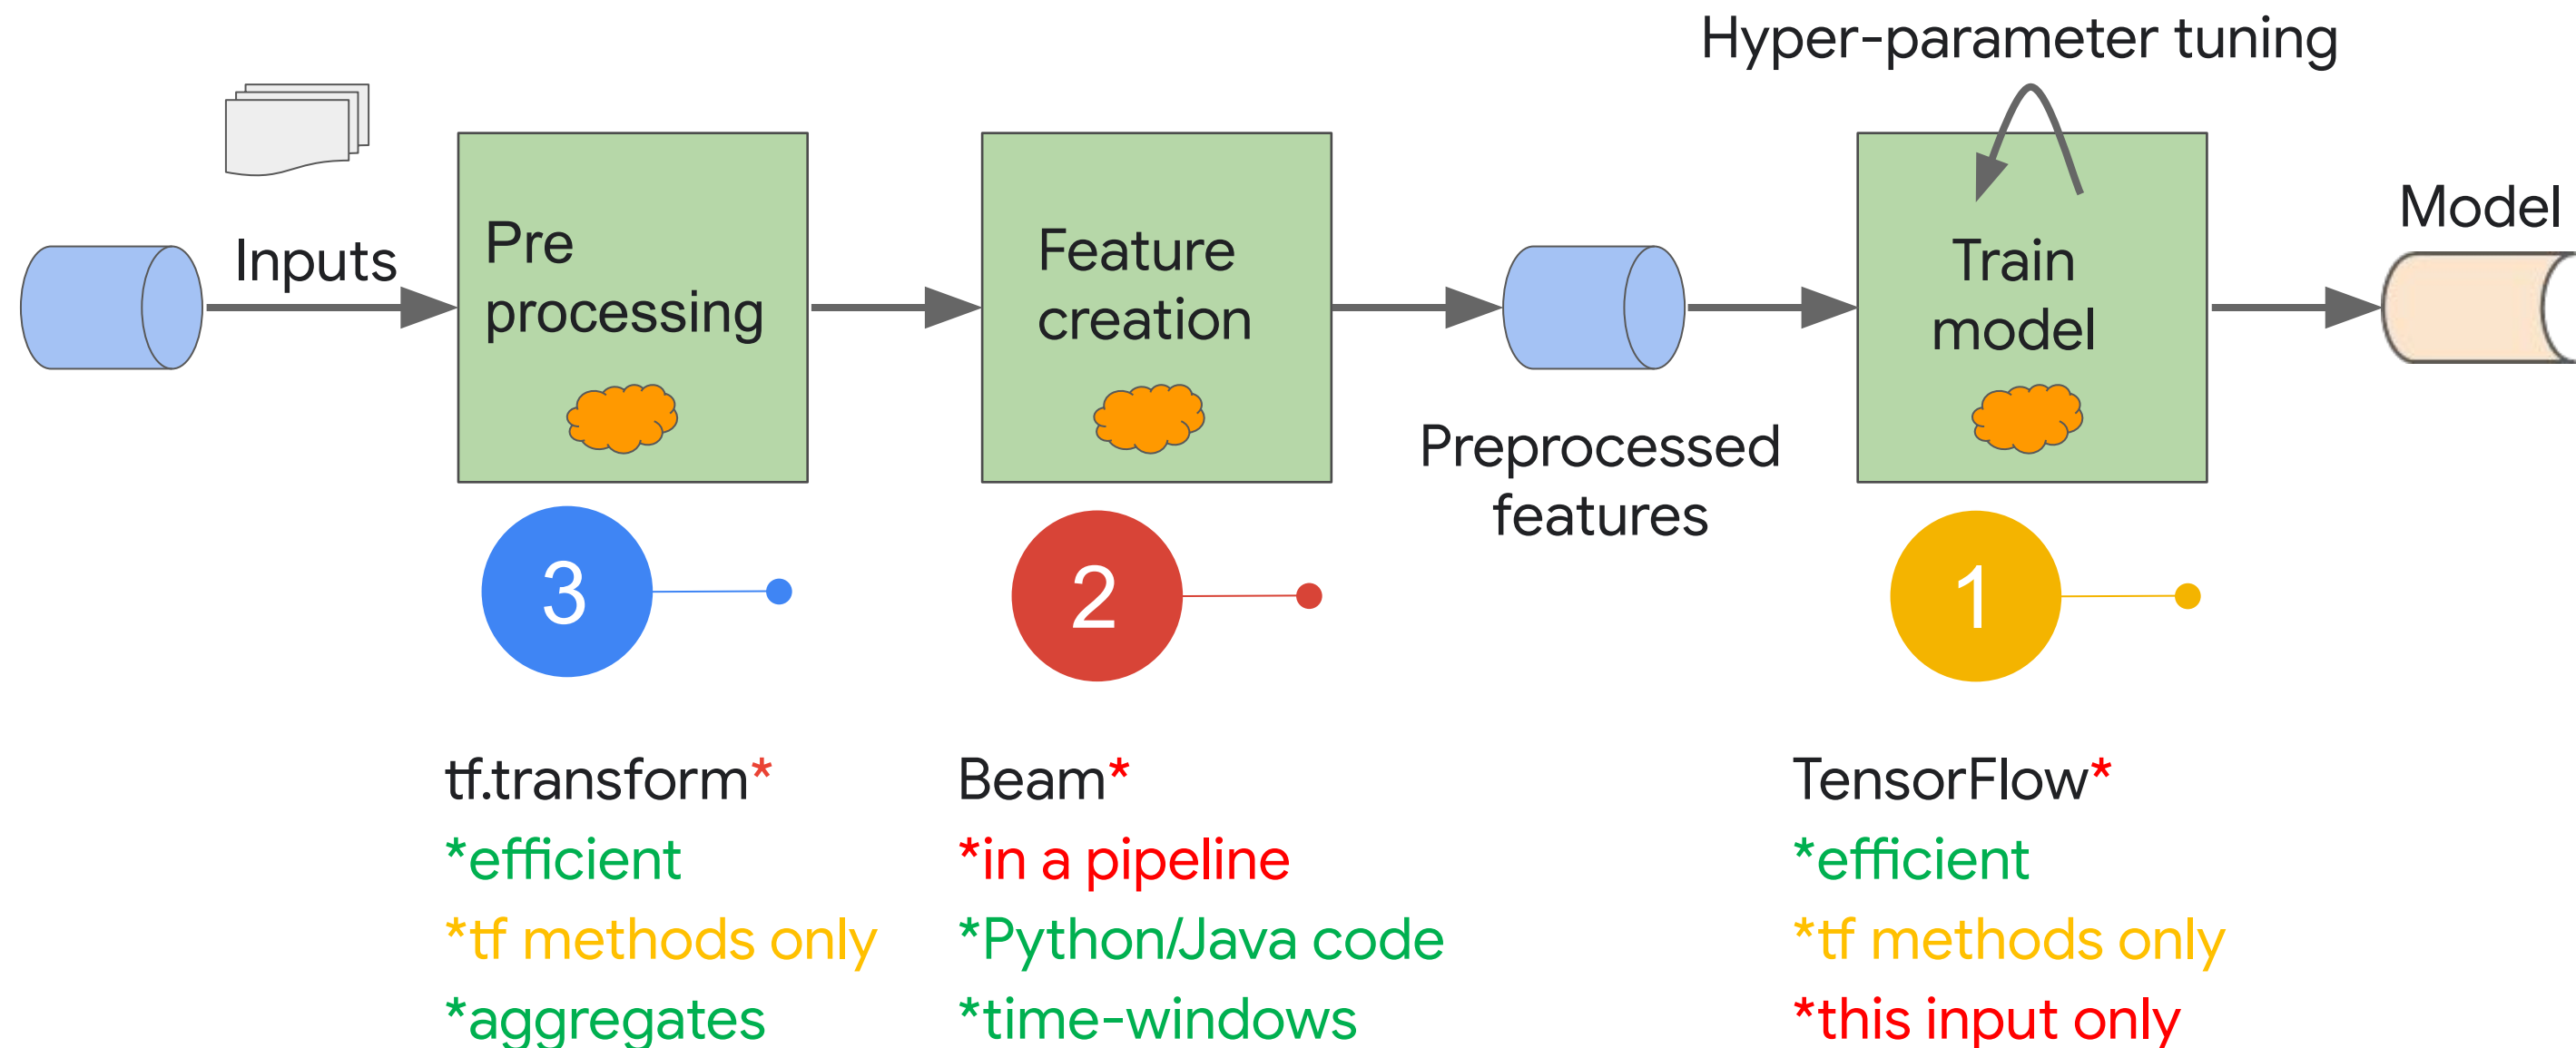
*tf methods only

*this input only

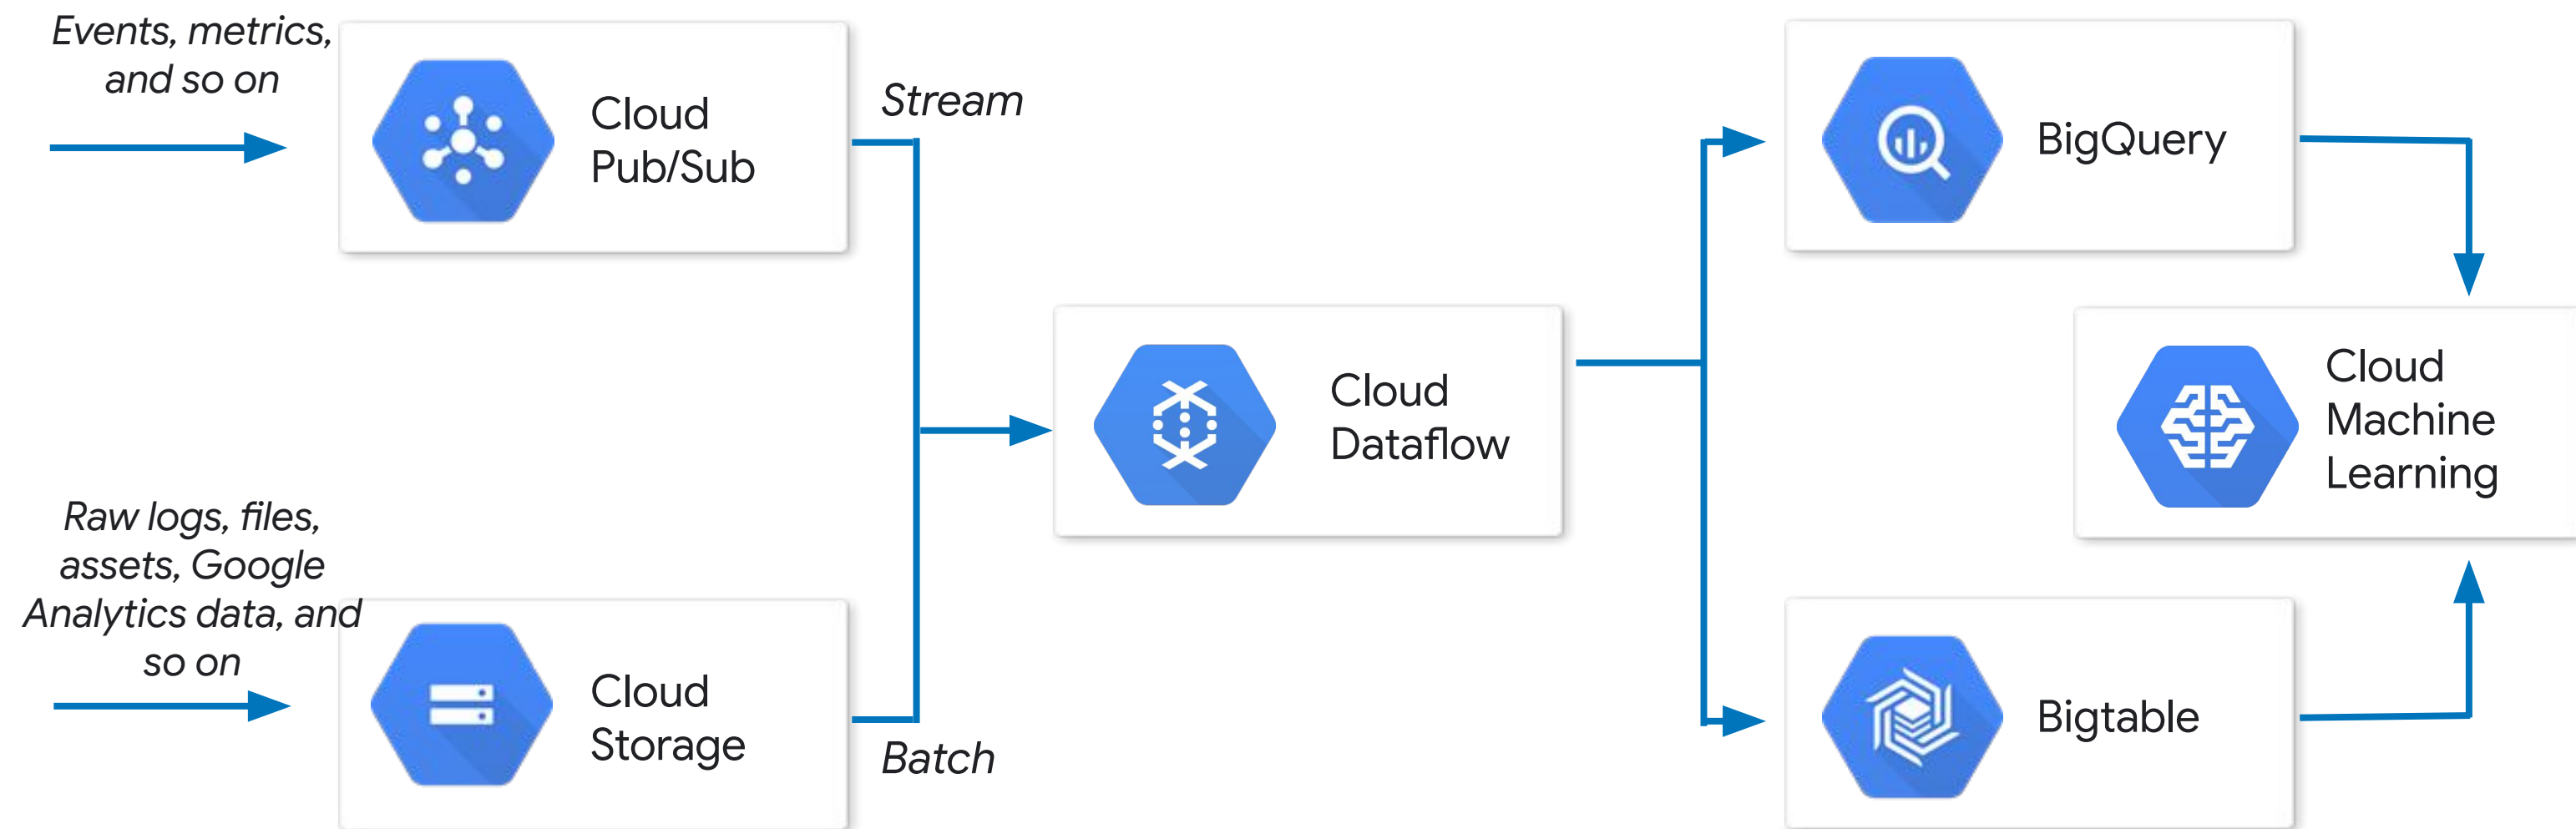
Recall that there are three possible places to do feature engineering, each of which has its pros and cons



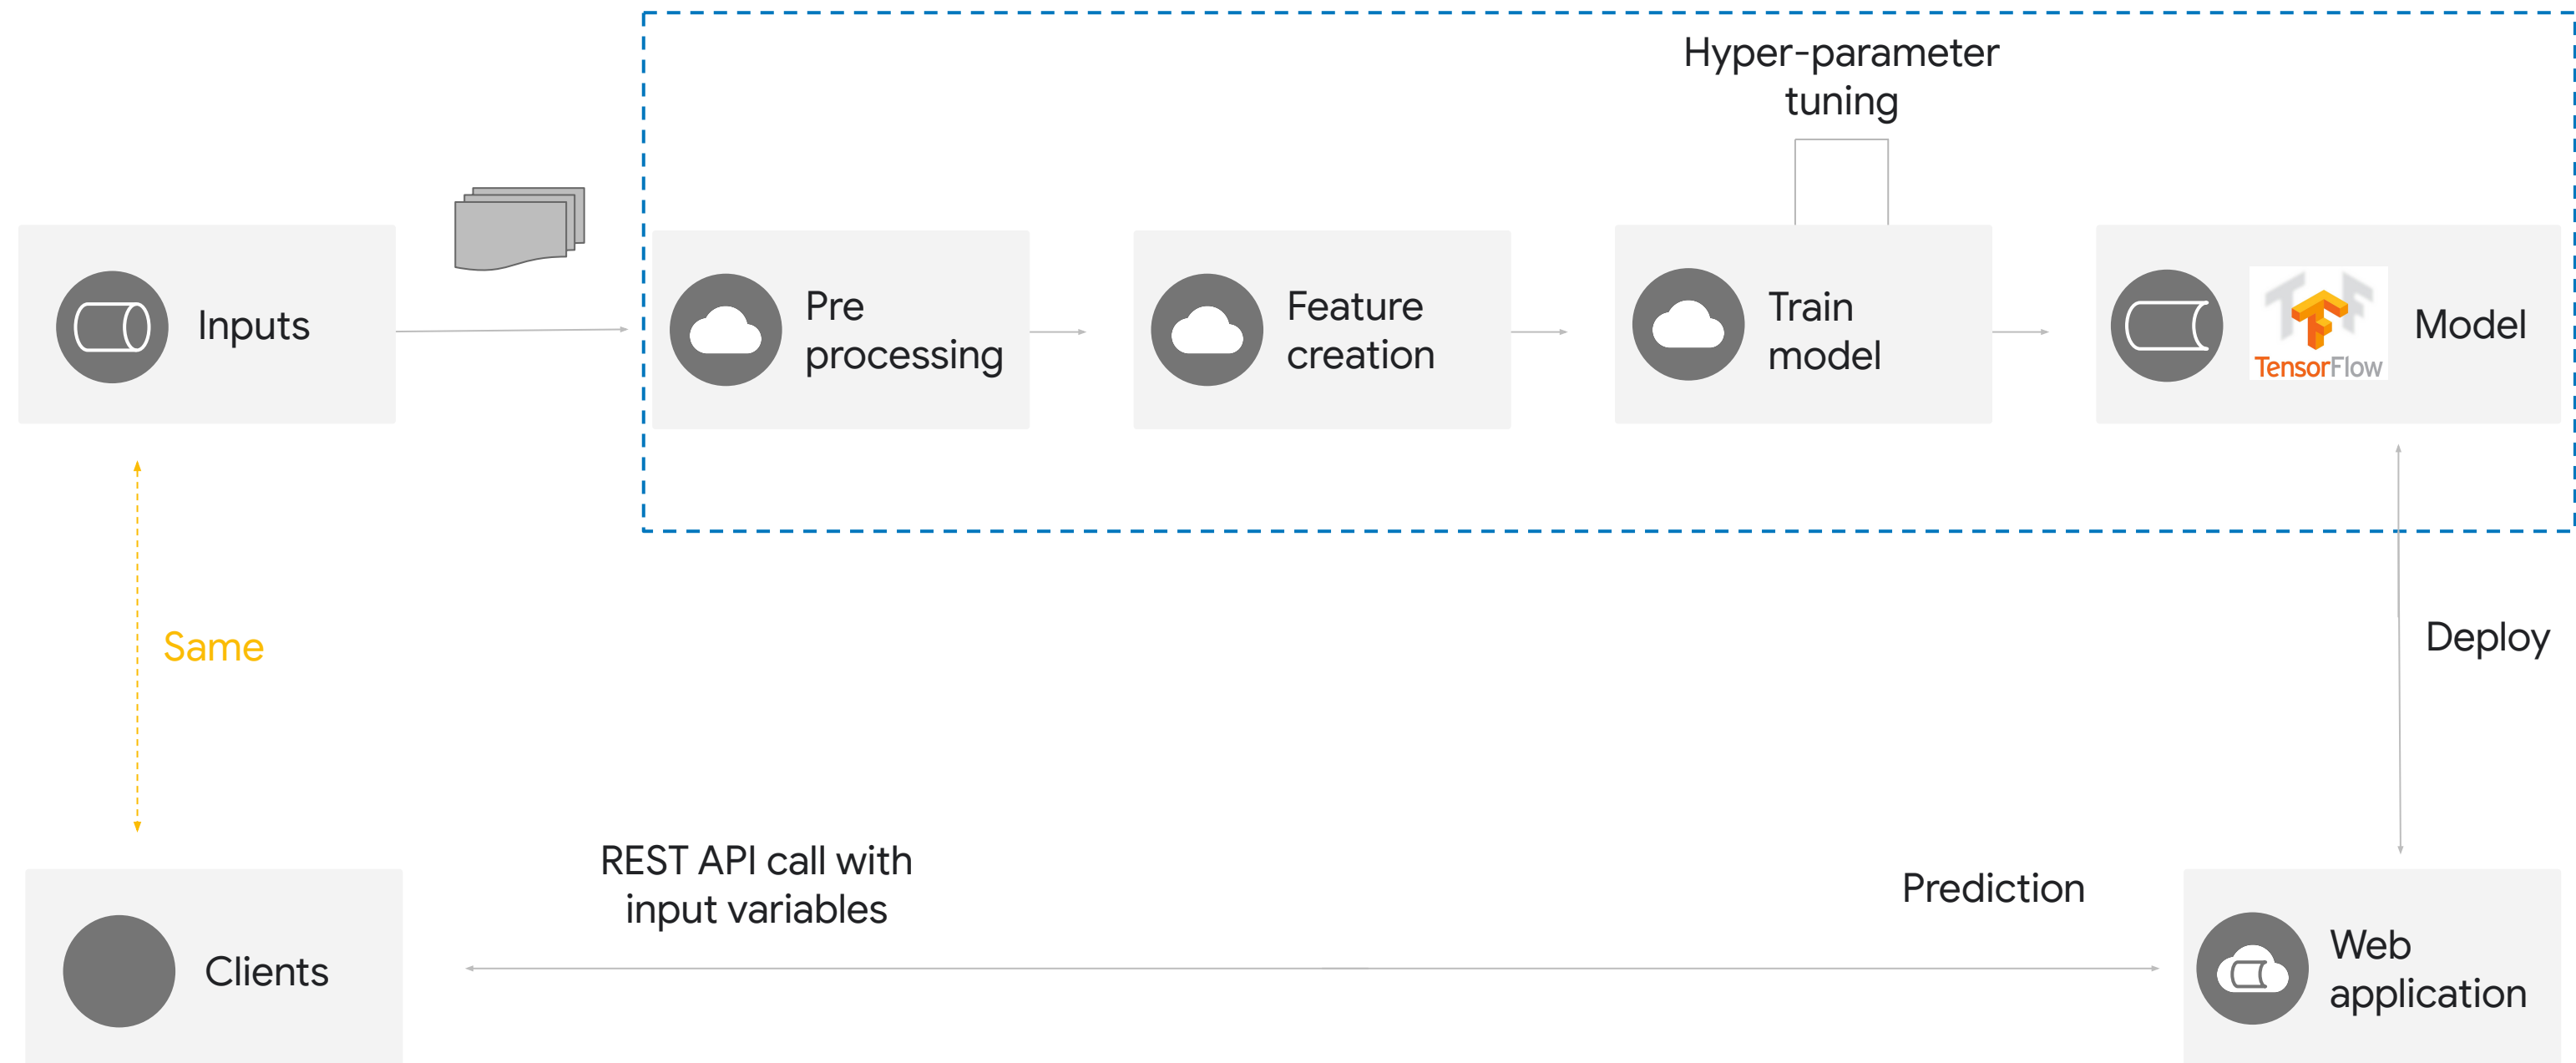
Recall that there are three possible places to do feature engineering, each of which has its pros and cons



Dataflow preprocessing works in the context of a pipeline



TensorFlow is good for on-demand, on-the-fly processing



tf.transform is a hybrid of
Beam and TensorFlow

tf.transform is a hybrid of Beam and TensorFlow

Find min/max value of
a numeric feature

tf.transform is a hybrid of Beam and TensorFlow

Find min/max value of
a numeric feature

Scale inputs by the
min & max

tf.transform is a hybrid of Beam and TensorFlow

Find min/max value of
a numeric feature

Scale inputs by the
min & max

Find all the unique values of a
categorical feature

tf.transform is a hybrid of Beam and TensorFlow

Find min/max value of
a numeric feature

Scale inputs by the
min & max

Find all the unique values of a
categorical feature

One-hot encode
inputs based on set of
unique values

tf.transform is a hybrid of Beam and TensorFlow

Find min/max value of
a numeric feature

Scale inputs by the
min & max

Find all the unique values of a
categorical feature

One-hot encode
inputs based on set of
unique values

Analyze

tf.transform is a hybrid of Beam and TensorFlow

Find min/max value of
a numeric feature

Scale inputs by the
min & max

Find all the unique values of a
categorical feature

One-hot encode
inputs based on set of
unique values

Analyze

Transform

tf.transform is a hybrid of Beam and TensorFlow

Find min/max value of
a numeric feature

Scale inputs by the
min & max

Find all the unique
values of a
categorical feature

One-hot encode
inputs based on set of
unique values

Analyze

Beam

Transform

TensorFlow

tf.transform provides two
PTransforms

tf.transform provides two
PTransforms

AnalyzeAndTransformDataset

tf.transform provides two PTransforms

AnalyzeAndTransformDataset

Executed in Beam to create the training
dataset

tf.transform provides two PTransforms

AnalyzeAndTransformDataset

Executed in Beam to create the training
dataset

TransformDataset

tf.transform provides two PTransforms

AnalyzeAndTransformDataset

Executed in Beam to create the training
dataset

TransformDataset

Executed in Beam to create the evaluation
dataset

tf.transform provides two PTransforms

AnalyzeAndTransformDataset

Executed in Beam to create the training
dataset

TransformDataset

Executed in Beam to create the evaluation
dataset

The underlying transformations are
executed in TensorFlow at prediction time

tf.transform has
two phases

tf.transform has
two phases

Analysis phase (compute min/max/vocab
etc. using Beam)

tf.transform has
two phases

Analysis phase (compute min/max/vocab
etc. using Beam)

Executed in Beam while creating training
dataset

tf.transform has two phases

Analysis phase (compute min/max/vocab
etc. using Beam)

Executed in Beam while creating training
dataset

Transform phase (scale/vocabulary etc.
using TensorFlow)

tf.transform has two phases

Analysis phase (compute min/max/vocab
etc. using Beam)

Executed in Beam while creating training
dataset

Transform phase (scale/vocabulary etc.
using TensorFlow)

Executed in TensorFlow during prediction

tf.transform has two phases

Analysis phase (compute min/max/vocab
etc. using Beam)

Executed in Beam while creating training
dataset

Transform phase (scale/vocabulary etc.
using TensorFlow)

Executed in TensorFlow during prediction

Executed in Beam to create
training/evaluation datasets

First, set up the schema of the training dataset

First, set up the schema of the training dataset

```
raw_data_schema = {  
    colname : dataset_schema.ColumnSchema(tf.string, ...)   
    for colname in 'dayofweek,key'.split(',')  
}
```


First, set up the schema of the training dataset

```
raw_data_schema = {                                TensorFlow type for input column
    colname : dataset_schema.ColumnSchema(tf.string, ...)
    for colname in 'dayofweek,key'.split(',')
}
```

First, set up the schema of the training dataset

```
raw_data_schema = {
    colname : dataset_schema.ColumnSchema(tf.string, ...)
    for colname in 'dayofweek,key'.split(',')
}
raw_data_schema.update({
    colname : dataset_schema.ColumnSchema(tf.float32, ...)
    for colname in 'fare_amount,pickuplon, ... ,dropofflat'.split(',')
})
```

First, set up the schema of the training dataset

```
raw_data_schema = {
    colname : dataset_schema.ColumnSchema(tf.string, ...)
    for colname in 'dayofweek,key'.split(',')
}
raw_data_schema.update({
    colname : dataset_schema.ColumnSchema(tf.float32, ...)
    for colname in 'fare_amount,pickuplon, ... ,dropofflat'.split(',')
})
raw_data_metadata =
    dataset_metadata.DatasetMetadata(dataset_schema.Schema(raw_data_schema))
```


First, set up the schema of the training dataset

```
raw_data_schema = {                                TensorFlow type for input column
    colname : dataset_schema.ColumnSchema(tf.string, ...)
    for colname in 'dayofweek,key'.split(',')
}
raw_data_schema.update({
    colname : dataset_schema.ColumnSchema(tf.float32, ...) float32
    for colname in 'fare_amount,pickuplon, ... ,dropofflat'.split(',')
})
raw_data_metadata = Use the schema to create metadata "template"
    dataset_metadata.DatasetMetadata(dataset_schema.Schema(raw_data_schema))
```

Next, run the analyze-and-transform PTransform on training dataset to get back preprocessed training data and the transform function

```
raw_data = (p
    | beam.io.Read(beam.io.BigQuerySource(query=myquery, use_standard_sql=True))
    | beam.Filter(is_valid))


transformed_dataset, transform_fn = ((raw_data, raw_data_metadata)
    | beam_impl.AnalyzeAndTransformDataset(preprocess))
```



Next, run the analyze-and-transform PTransform on training dataset to get back preprocessed training data and the transform function

```
raw_data = (p                                     1. Read in data as usual for Beam
    | beam.io.Read(beam.io.BigQuerySource(query=myquery, use_standard_sql=True))
    | beam.Filter(is_valid))


transformed_dataset, transform_fn = ((raw_data, raw_data_metadata)
    | beam_impl.AnalyzeAndTransformDataset(preprocess))
```



Next, run the analyze-and-transform PTransform on training dataset to get back preprocessed training data and the transform function

```
raw_data = (p                                     1. Read in data as usual for Beam
    | beam.io.Read(beam.io.BigQuerySource(query=myquery, use_standard_sql=True))
    | beam.Filter(is_valid)) 2. Filter out data that you don't want to train with


transformed_dataset, transform_fn = ((raw_data, raw_data_metadata)
    | beam_impl.AnalyzeAndTransformDataset(preprocess))
```



Next, run the analyze-and-transform PTransform on training dataset to get back preprocessed training data and the transform function

```
raw_data = (p                                     1. Read in data as usual for Beam
    | beam.io.Read(beam.io.BigQuerySource(query=myquery, use_standard_sql=True))
    | beam.Filter(is_valid)) 2. Filter out data that you don't want to train with

                                     3. Pass raw data + metadata template to AnalyzeAndTransformDataset
transformed_dataset, transform_fn = ((raw_data, raw_data_metadata)
    | beam_impl.AnalyzeAndTransformDataset(preprocess))
```



Next, run the analyze-and-transform PTransform on training dataset to get back preprocessed training data and the transform function

```
raw_data = (p
    | beam.io.Read(beam.io.BigQuerySource(query=myquery, use_standard_sql=True))
    | beam.Filter(is_valid))
transformed_dataset, transform_fn = ((raw_data, raw_data_metadata)
    | beam_impl.AnalyzeAndTransformDataset(preprocess))
```

1. Read in data as usual for Beam

2. Filter out data that you don't want to train with

3. Pass raw data + metadata template to AnalyzeAndTransformDataset

4. Get back transformed dataset and a reusable transform function

Write out the preprocessed training data into TFRecords, the most efficient format for TensorFlow

```
transformed_data |
    tfrecordio.WriteToTFRecord(
        os.path.join(OUTPUT_DIR, 'train'),

        coder=ExampleProtoCoder(
            transformed_metadata.schema)
    )
```

Write out the preprocessed training data into TFRecords, the most efficient format for TensorFlow

```
transformed_data |
    tfrecordio.WriteToTFRecord(
        os.path.join(OUTPUT_DIR, 'train'),
        The filenames will be like train-0003-of0015
        coder=ExampleProtoCoder(
            transformed_metadata.schema)
    )
```

Write out the preprocessed training data into TFRecords, the most efficient format for TensorFlow

```
transformed_data |
    tfrecordio.WriteToTFRecord(
        os.path.join(OUTPUT_DIR, 'train'),
        The filenames will be like train-0003-of0015
        coder=ExampleProtoCoder(
            transformed_metadata.schema)
    )
```

Note that we use the transformed metadata schema here

The preprocessing function
is restricted to TensorFlow
functions


```
transformed_dataset, transform_fn =  
((raw_data, raw_data_metadata)  
 |  
 beam_impl.AnalyzeAndTransformDataset(  
     preprocess))
```

The preprocessing function is restricted to TensorFlow functions

```
transformed_dataset, transform_fn =  
((raw_data, raw_data_metadata)  
 |  
 beam_impl.AnalyzeAndTransformDataset(  
     preprocess))
```

The things you do in `preprocess()` will get added to the TensorFlow graph, and be executed in TensorFlow during serving

The preprocessing function is restricted to functions you can call from TensorFlow graph



The preprocessing function is restricted to functions you can call from TensorFlow graph

```
def preprocess(inputs):
```


The preprocessing function is restricted to functions you can call from TensorFlow graph

```
def preprocess(inputs):
```

```
    result = {}    Create features from the input tensors and put into “results” dict
```

The preprocessing function is restricted to functions you can call from TensorFlow graph

```
def preprocess(inputs):  
  
    result = {}    Create features from the input tensors and put into “results” dict  
  
    result['fare_amount'] = inputs['fare_amount']    Pass through
```

The preprocessing function is restricted to functions you can call from TensorFlow graph

```
def preprocess(inputs):  
  
    result = {}    Create features from the input tensors and put into “results” dict  
  
    result['fare_amount'] = inputs['fare_amount']    Pass through  
  
    result['dayofweek'] = tft.string_to_int(inputs['dayofweek'])    vocabulary
```

The preprocessing function is restricted to functions you can call from TensorFlow graph

```
def preprocess(inputs):  
  
    result = {}    Create features from the input tensors and put into “results” dict  
  
    result['fare_amount'] = inputs['fare_amount']    Pass through  
  
    result['dayofweek'] = tft.string_to_int(inputs['dayofweek'])    vocabulary  
    ...  
    result['dropofflat'] = (tft.scale_to_0_1(inputs['dropofflat']))    scaling
```

The preprocessing function is restricted to functions you can call from TensorFlow graph

```
def preprocess(inputs):  
  
    result = {}  Create features from the input tensors and put into “results” dict  
  
    result['fare_amount'] = inputs['fare_amount']  Pass through  
  
    result['dayofweek'] = tft.string_to_int(inputs['dayofweek'])  vocabulary  
    ...  
    result['dropofflat'] = (tft.scale_to_0_1(inputs['dropofflat']))  scaling  
  
    result['passengers'] = tf.cast(inputs['passengers'], tf.float32)  Other TF fns
```

The preprocessing function is restricted to functions you can call from TensorFlow graph

```
def preprocess(inputs):  
  
    result = {}    Create features from the input tensors and put into “results” dict  
  
    result['fare_amount'] = inputs['fare_amount']    Pass through  
  
    result['dayofweek'] = tft.string_to_int(inputs['dayofweek'])    vocabulary  
    ...  
    result['dropofflat'] = (tft.scale_to_0_1(inputs['dropofflat']))    scaling  
  
    result['passengers'] = tf.cast(inputs['passengers'], tf.float32)    Other TF fns  
  
    return result
```

Analyze and Transform happens on the training dataset

```
transformed_dataset, transform_fn =  
((raw_data, raw_data_metadata)  
 |  
 beam_impl.AnalyzeAndTransformDataset(  
     preprocess))
```

Writing out the eval dataset is similar, except that we reuse the transform function computed from the training data


```
raw_test_data = (p
    | beam.io.Read(beam.io.BigQuerySource(...))
    | 'eval_filter' >> beam.Filter(is_valid))
transformed_test_dataset = (((raw_test_data, raw_data_metadata), transform_fn)
    | beam_impl.TransformDataset())
```


Writing out the eval dataset is similar, except that we reuse the transform function computed from the training data

```
raw_test_data = (p
    | beam.io.Read(beam.io.BigQuerySource(...))
    | 'eval_filter' >> beam.Filter(is_valid))
transformed_test_dataset = (((raw_test_data, raw_data_metadata), transform_fn)
    | beam_impl.TransformDataset())

transformed_test_data, _ = transformed_test_dataset
_ = transformed_test_data | tfrecordio.WriteToTFRecord(
    os.path.join(OUTPUT_DIR, 'eval'),
    coder=example_proto_coder.ExampleProtoCoder(
        transformed_metadata.schema))
```

For training and evaluation,
we created preprocessed
features using Beam



Created by
AnalyzeAndTransformDataset
Or by TransformDataset

For serving, we need to write out the transformation metadata

```
_ = transform_fn |
  transform_fn_io.WriteTransformFn(
    os.path.join(OUTPUT_DIR,
  'metadata'))
```

[Buckets](#) / [cloud-training-demos-ml](#) / [taxifare](#) / [preproc_tft](#) / metadata

| <input type="checkbox"/> | Name | Size | Type |
|--------------------------|---|------|--------|
| <input type="checkbox"/> |  rawdata_metadata/ | — | Folder |
| <input type="checkbox"/> |  transform_fn/ | — | Folder |
| <input type="checkbox"/> |  transformed_metadata/ | — | Folder |

Change input function to read preprocessed features

```
def read_dataset(args, mode):
    if mode == tf.estimator.ModeKeys.TRAIN:
        input_paths = args['train_data_paths']
    else:
        input_paths = args['eval_data_paths']      Reading transformed features
    transformed_metadata = metadata_io.read_metadata(
        os.path.join(args['metadata_path'], 'transformed_metadata'))
    return input_fn_maker.build_training_input_fn(
        metadata = transformed_metadata,
        file_pattern = (
            input_paths[0] if len(input_paths) == 1 else input_paths),
        ...)
```

Change input function to read preprocessed features

```
def read_dataset(args, mode):
    if mode == tf.estimator.ModeKeys.TRAIN:
        input_paths = args['train_data_paths']
    else:
        input_paths = args['eval_data_paths']
    transformed_metadata = metadata_io.read_metadata(
        os.path.join(args['metadata_path'], 'transformed_metadata'))
    return input_fn_maker.build_training_input_fn(
        metadata = transformed_metadata,
        file_pattern = (
            input_paths[0] if len(input_paths) == 1 else input_paths),
        ...)
```

Reading transformed features

Change input function to read preprocessed features

```
def read_dataset(args, mode):
    if mode == tf.estimator.ModeKeys.TRAIN:
        input_paths = args['train_data_paths']
    else:
        input_paths = args['eval_data_paths']      Reading transformed features
    transformed_metadata = metadata_io.read_metadata(
        os.path.join(args['metadata_path'], 'transformed_metadata'))
    return input_fn_maker.build_training_input_fn(
        metadata = transformed_metadata,
        file_pattern = (
            input_paths[0] if len(input_paths) == 1 else input_paths),
        ...)
```

The serving input function accepts the raw data

```
def make_serving_input_fn(args):  
    raw_metadata = metadata_io.read_metadata(  
        os.path.join(args['metadata_path'], 'rawdata_metadata'))  
    transform_savedmodel_dir = (  
        os.path.join(args['metadata_path'], 'transform_fn'))  
    return input_fn_maker.build_parsing_transforming_serving_input_receiver_fn(  
        raw_metadata,  
        transform_savedmodel_dir,  
        exclude_raw_keys = [LABEL_COLUMN])
```

The serving input function accepts the raw data

```
def make_serving_input_fn(args):  
    raw_metadata = metadata_io.read_metadata(  
        os.path.join(args['metadata_path'], 'rawdata_metadata'))  
    transform_savedmodel_dir = (  
        os.path.join(args['metadata_path'], 'transform_fn'))  
    return input_fn_maker.build_parsing_transforming_serving_input_receiver_fn(  
        raw_metadata,  
        transform_savedmodel_dir,  
        exclude_raw_keys = [LABEL_COLUMN])
```


The serving input function accepts the raw data


```
def make_serving_input_fn(args):  
    raw_metadata = metadata_io.read_metadata(  
        os.path.join(args['metadata_path'], 'rawdata_metadata'))  
    transform_savedmodel_dir = (  
        os.path.join(args['metadata_path'], 'transform_fn'))  
    return input_fn_maker.build_parsing_transforming_serving_input_receiver_fn(  
        raw_metadata,  
        transform_savedmodel_dir,  
        exclude_raw_keys = [LABEL_COLUMN])
```

The serving input function accepts the raw data

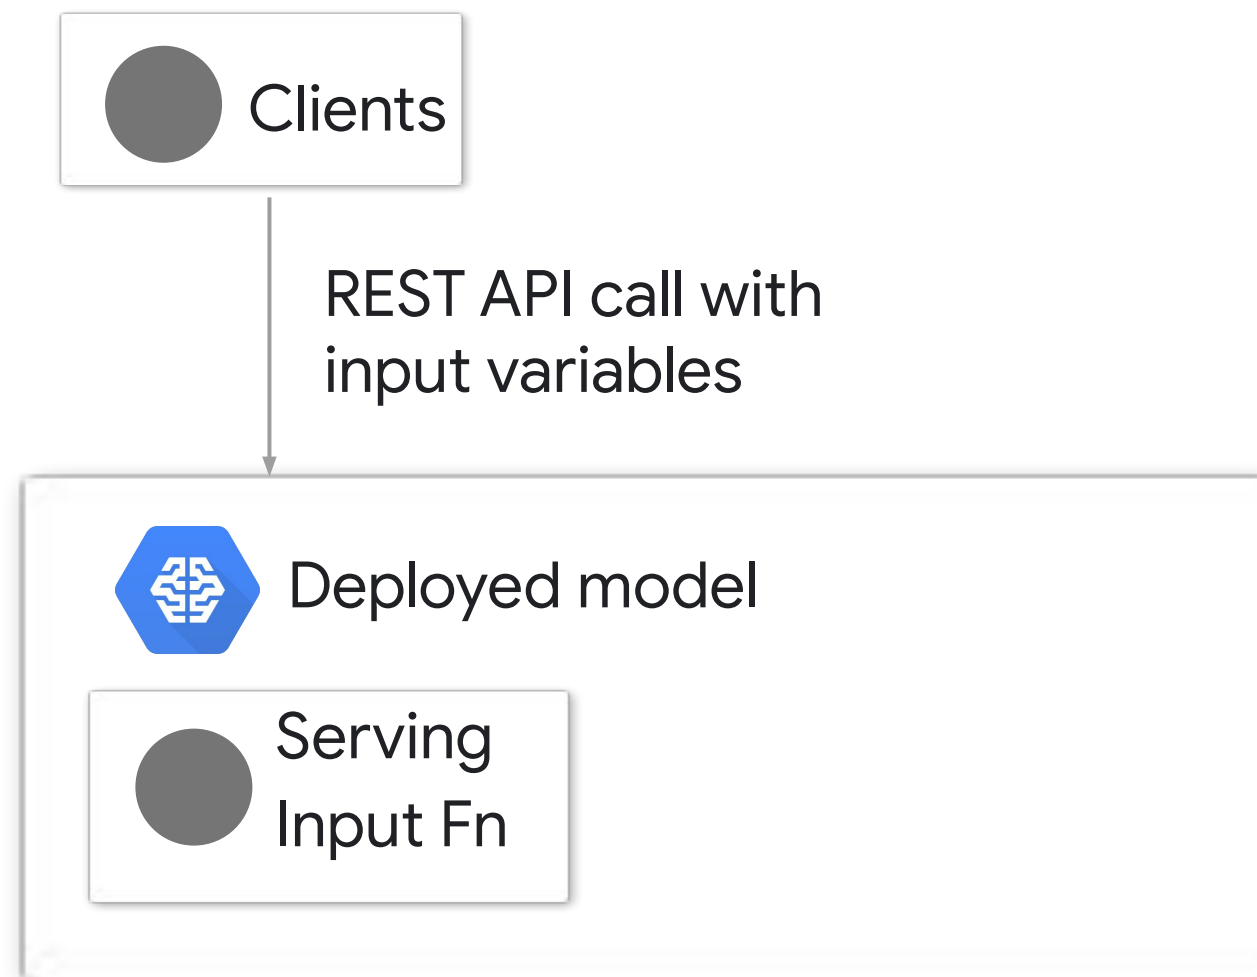
```
def make_serving_input_fn(args):  
    raw_metadata = metadata_io.read_metadata(  
        os.path.join(args['metadata_path'], 'rawdata_metadata'))  
    transform_savedmodel_dir = (  
        os.path.join(args['metadata_path'], 'transform_fn'))  
    return input_fn_maker.build_parsing_transforming_serving_input_receiver_fn(  
        raw_metadata,  
        transform_savedmodel_dir,  
        exclude_raw_keys = [LABEL_COLUMN])
```

The model graph includes
the preprocessing code

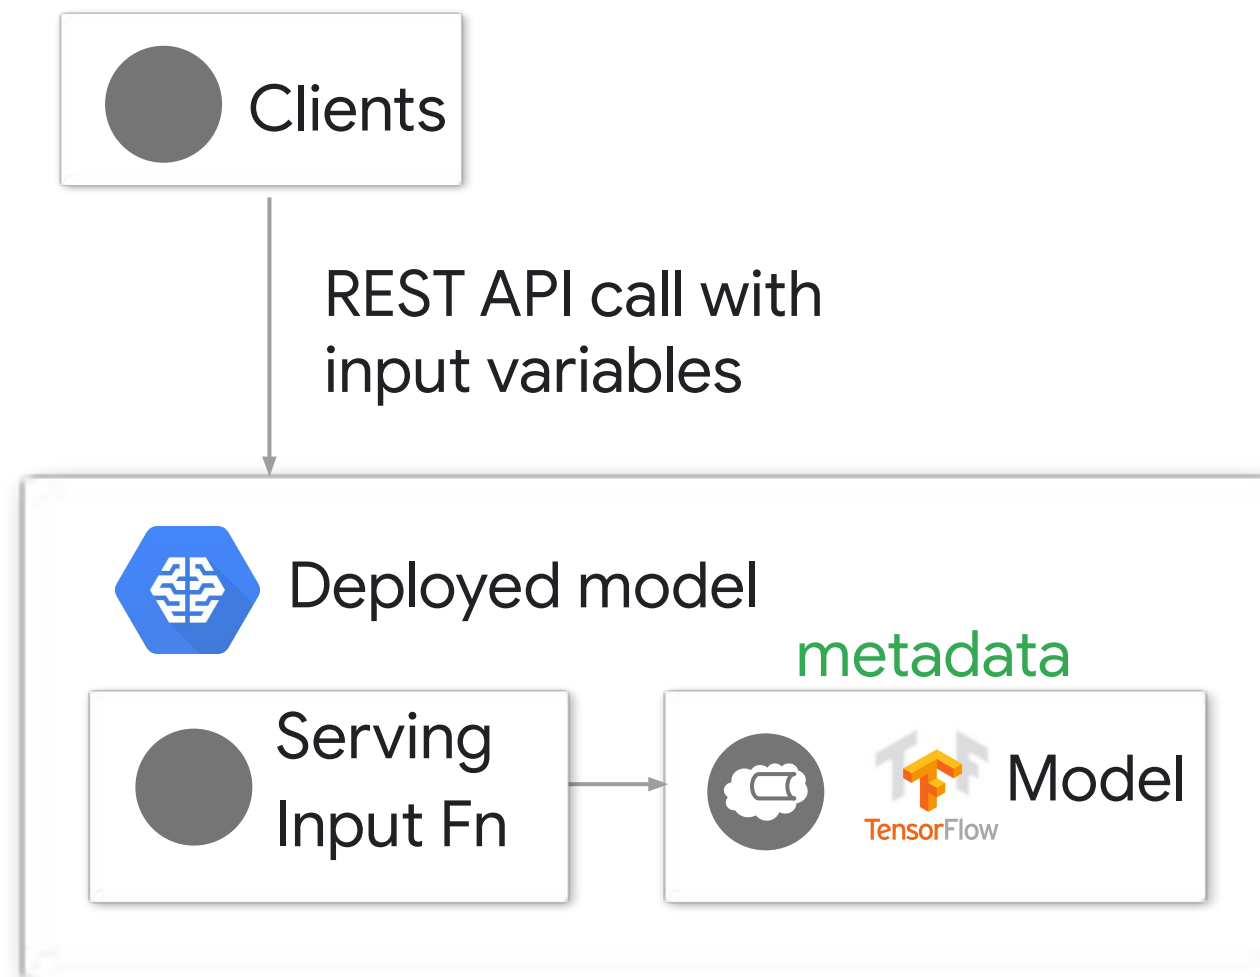
The model graph includes
the preprocessing code



The model graph includes
the preprocessing code



The model graph includes the preprocessing code



Lab

Exploring `tf.transform`

Look at `tftransform.ipynb`

cloud.google.com